



**SAIR**  
Spatial AI & Robotics Lab

# CSE 473/573-A

## W3: FILTERING & EDGE DETECTION

Chen Wang  
Spatial AI & Robotics Lab  
Department of Computer Science and Engineering

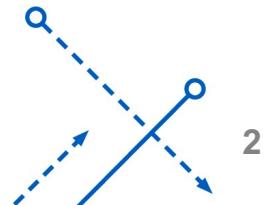
**UB** University at Buffalo The State University of New York



# Content

---

- Filtering
  - Linear filters, Correlation, Equivariance, Invariance
  - Gaussian Filter, Smoothing, Median filter, Sobel operator
- Template Matching
  - Cross-correlation, kernel cross-correlation
- Edge Detection
  - Image differentiation and gradient
  - Derivative theorem of convolution
  - Derivative of Gaussian filter, Laplacian of Gaussian
  - 2D edge detection filters
  - Canny edge detector, Hysteresis thresholding





# IMAGE PROCESSING

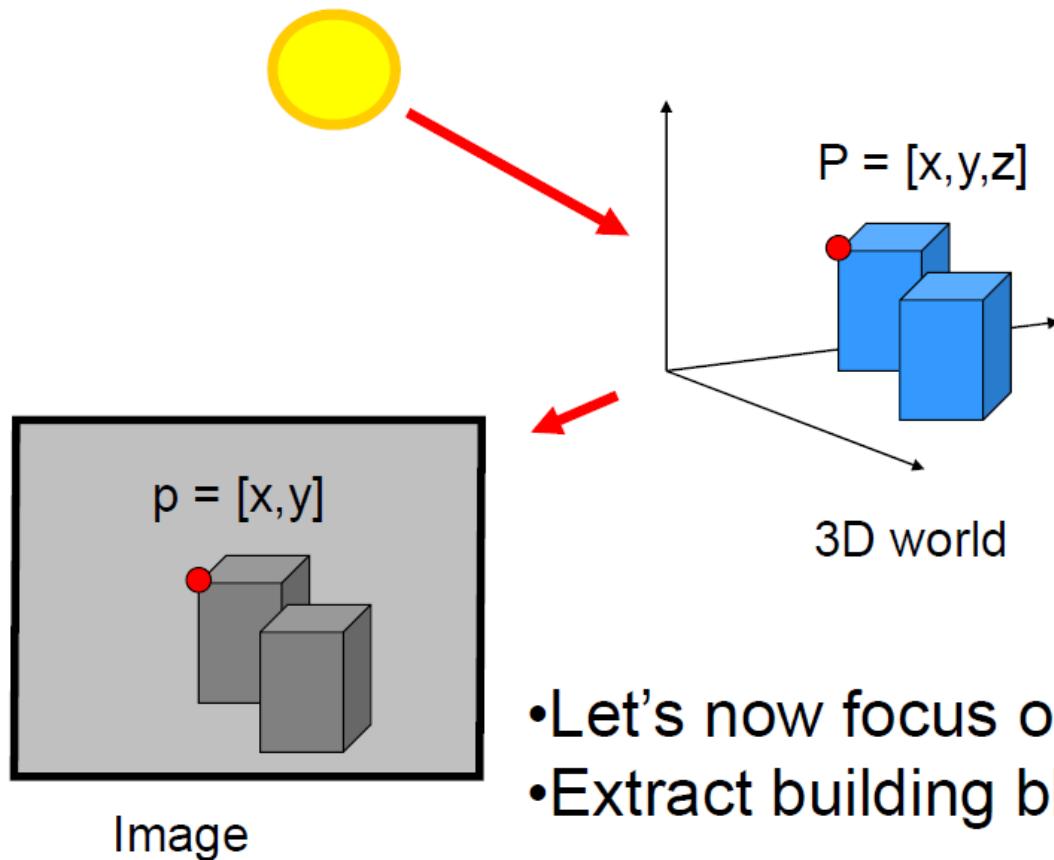
## Filtering



# Recap

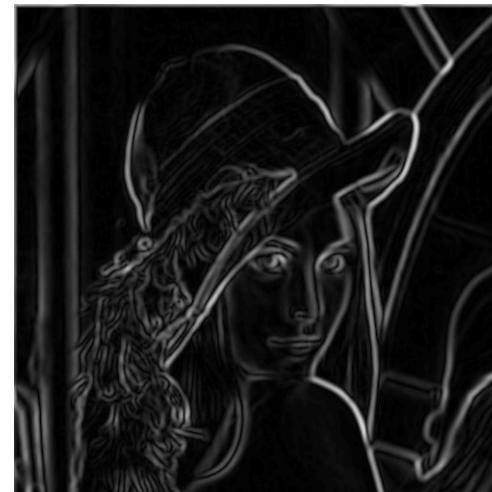
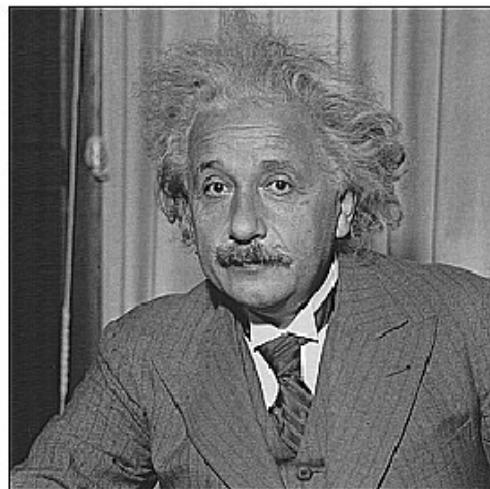
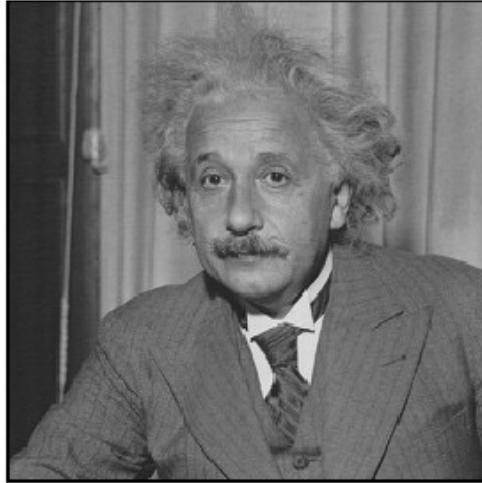
---

## From the 3D to 2D



- Let's now focus on 2D
- Extract building blocks

# Image Filtering



*Smooth/Sharpen Images...*

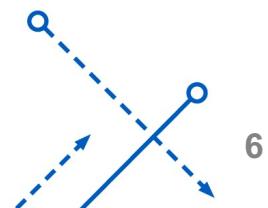
*Find edges...*

*Find Waldo...*

# Image filtering

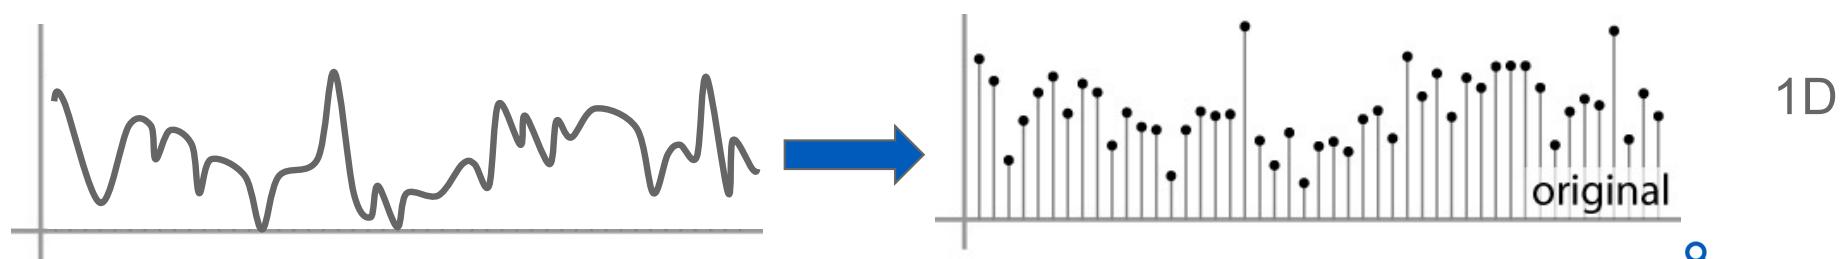
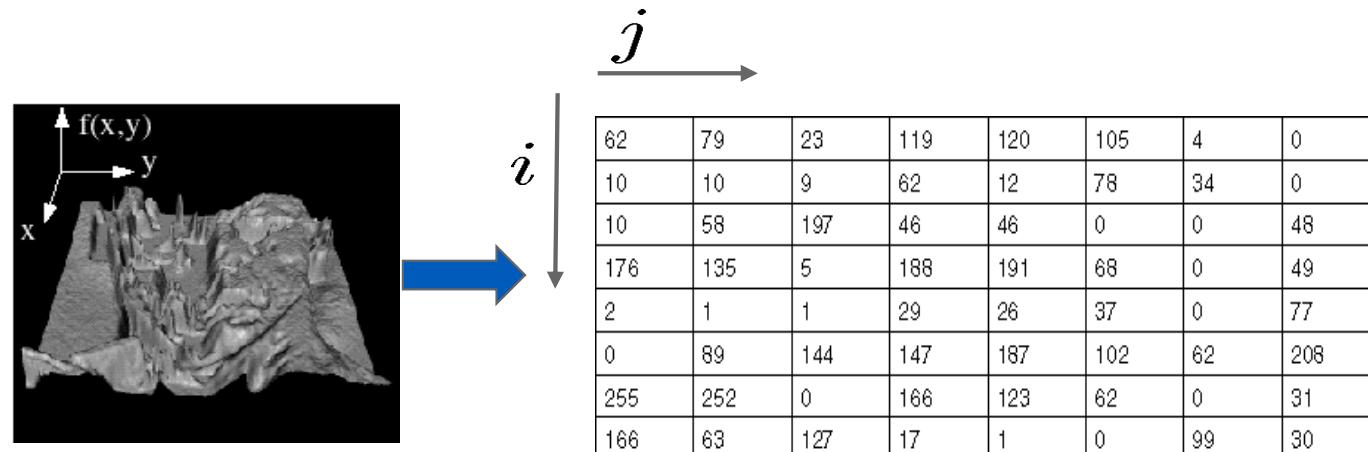
---

- Image filtering: compute a function of the local neighborhood at each position
- Really important!
  - Enhance images
    - Denoising, resize, increase contrast, etc.
  - Extract information from images
    - Texture, edges, distinctive points, etc.
  - Detect patterns
    - Template matching
    - Deep Convolutional Networks



# Digital images (Recap)

- In computer vision, we operate on **digital (discrete)** images:
  - **Sample** the 2D space on a regular grid
  - **Quantize** each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.



Adapted from S. Seitz

# Images as functions (Recap)

---

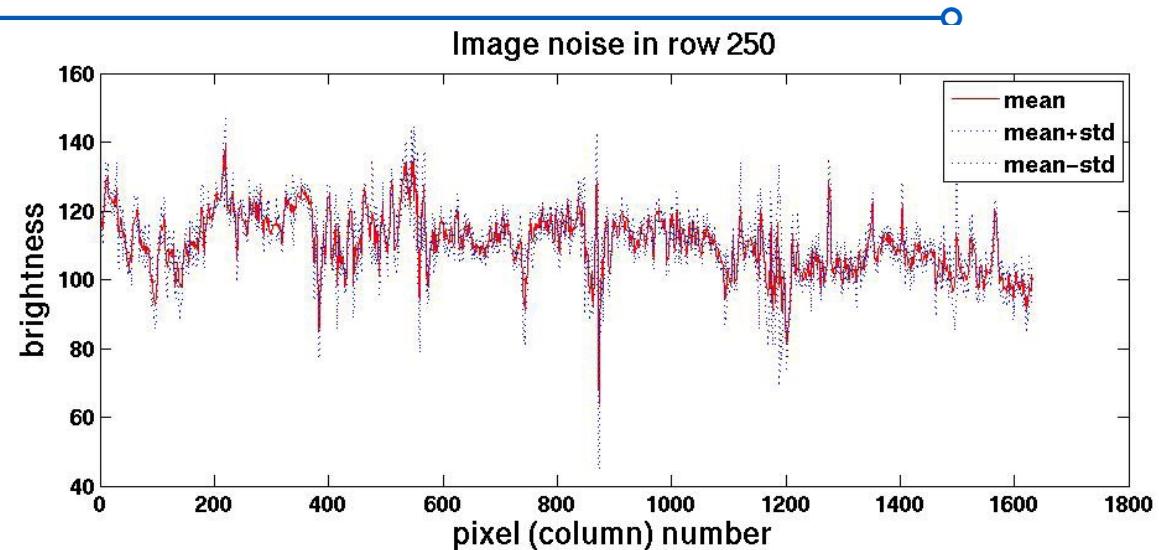
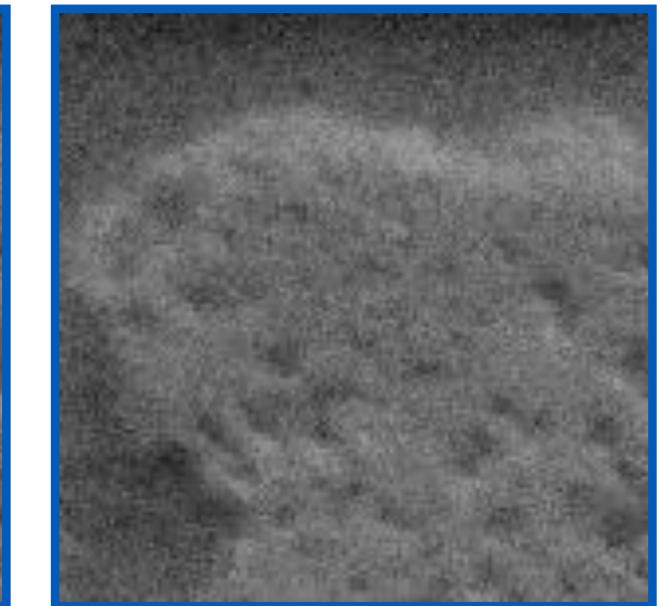
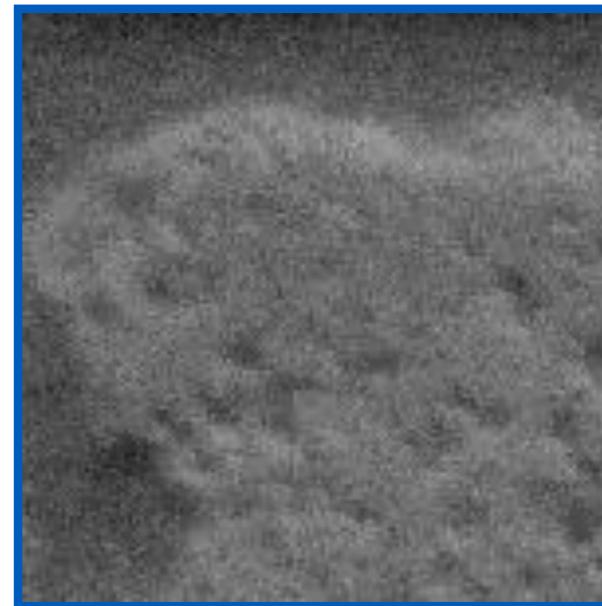
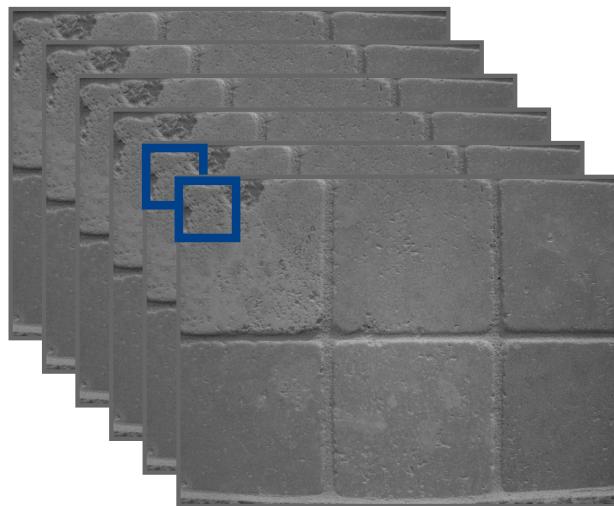
- Take an image as a function,  $f$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$ :
  - $f(x, y)$  gives the intensity at position  $(x, y)$
  - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
    - $f: [a, b] \times [c, d] \rightarrow [0, 255]$
    - Important: we often convert  $[0, 255]$  to **[0, 1.0]**.
- A color image is three functions pasted together, a “vector-valued” function

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$



# How can we do noise reduction?

- We can measure **noise** in multiple images of the same static scene.
- How could we reduce the noise, i.e., give an estimate of the true intensities?



# Common types of noise

- **Impulse noise:** random occurrences of white pixels
- **Salt and pepper noise:** random occurrences of black and white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise



Gaussian noise

Source: S. Seitz



# Gaussian noise

- Additive Noise

$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

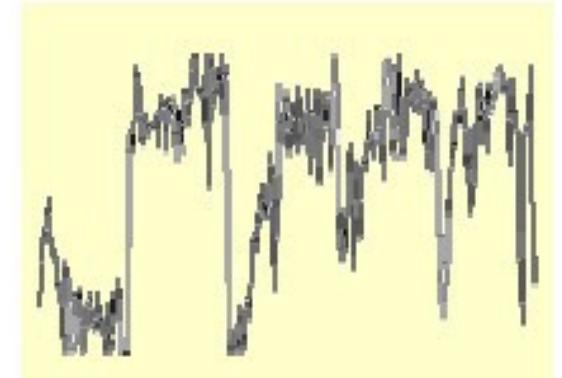
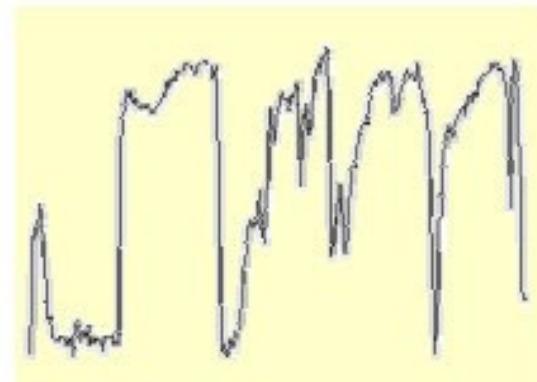
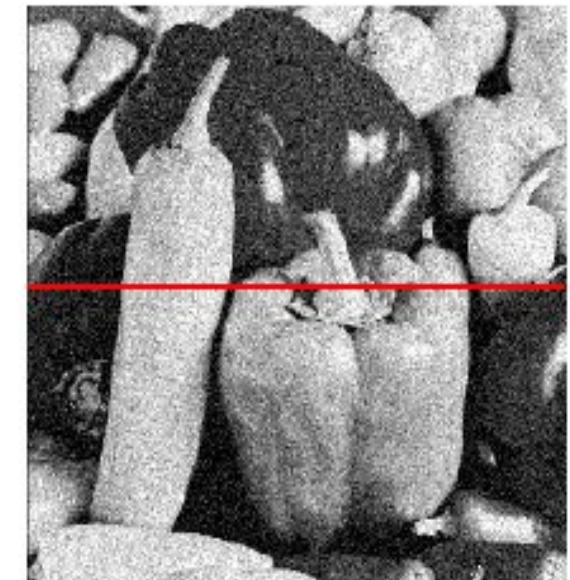
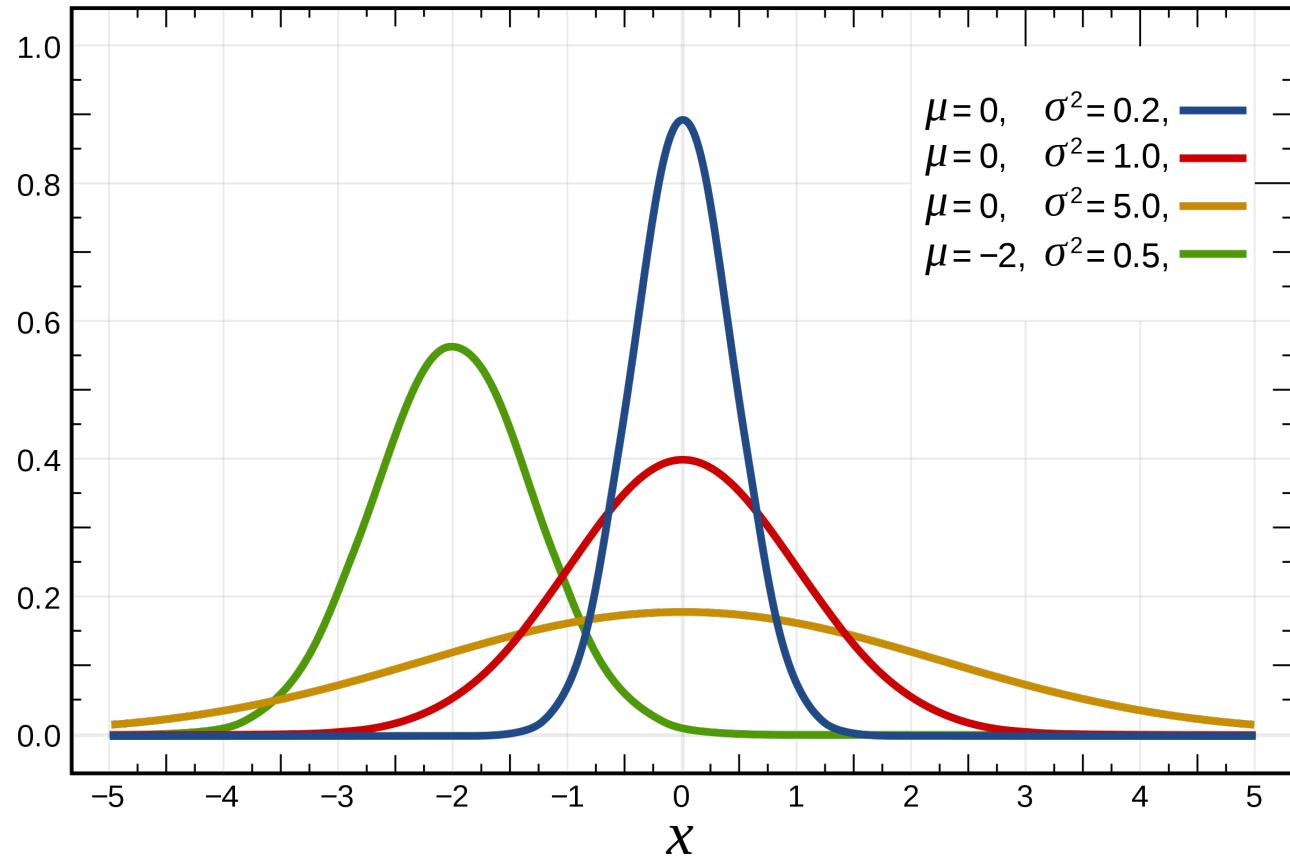


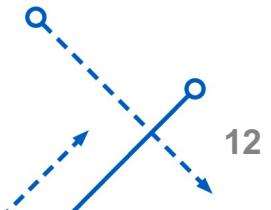
Fig: M. Hebert

# PDF of Gaussian distribution

- Probability density function



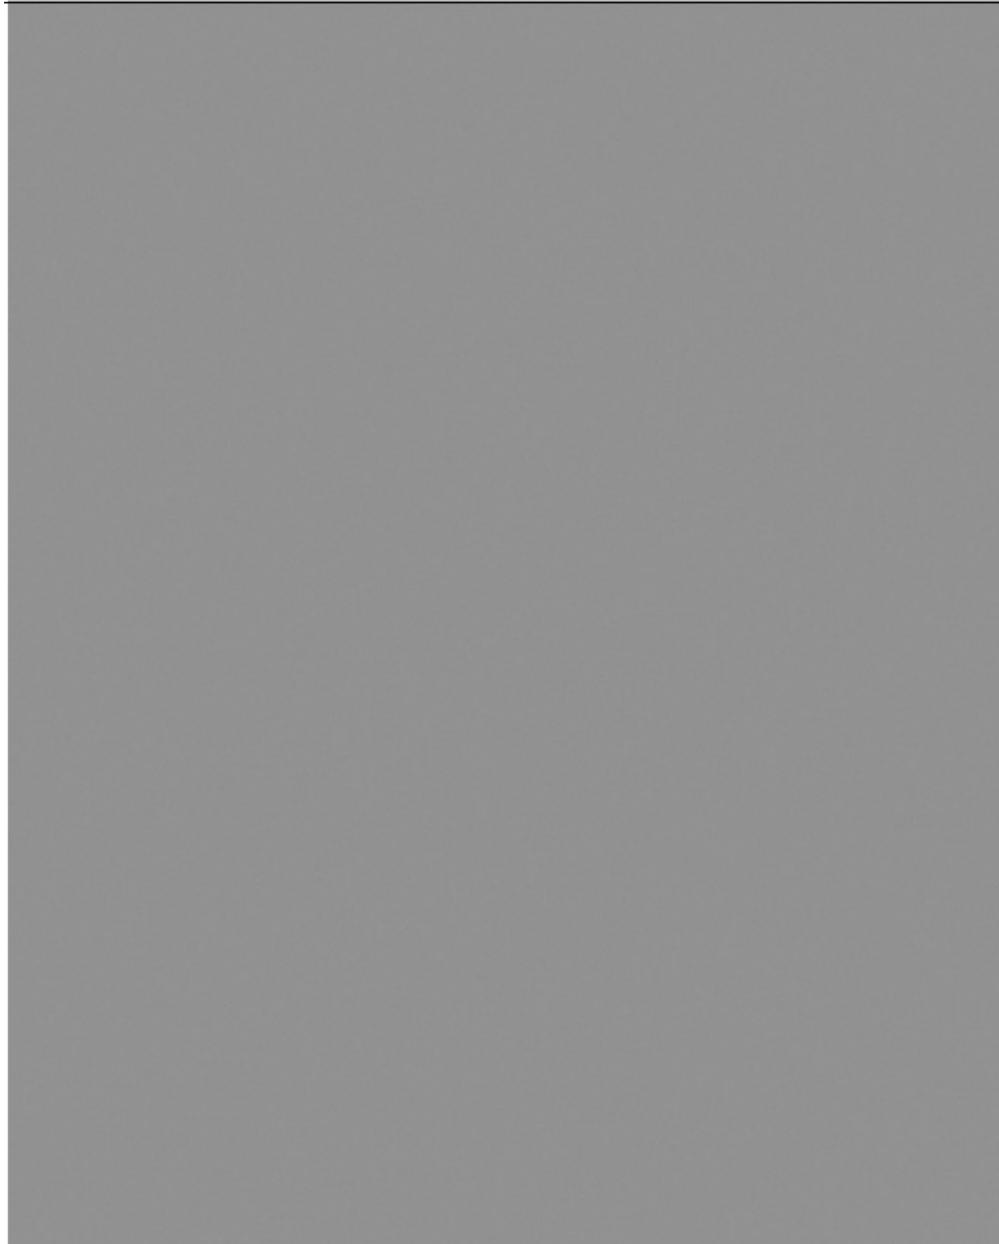
$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$



# Gaussian noise

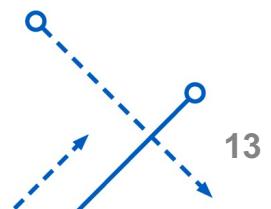
---

$\sigma = 1$



Effect of  $\sigma$  on Gaussian noise:

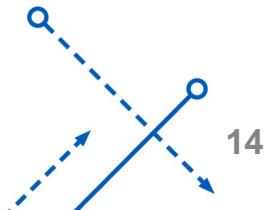
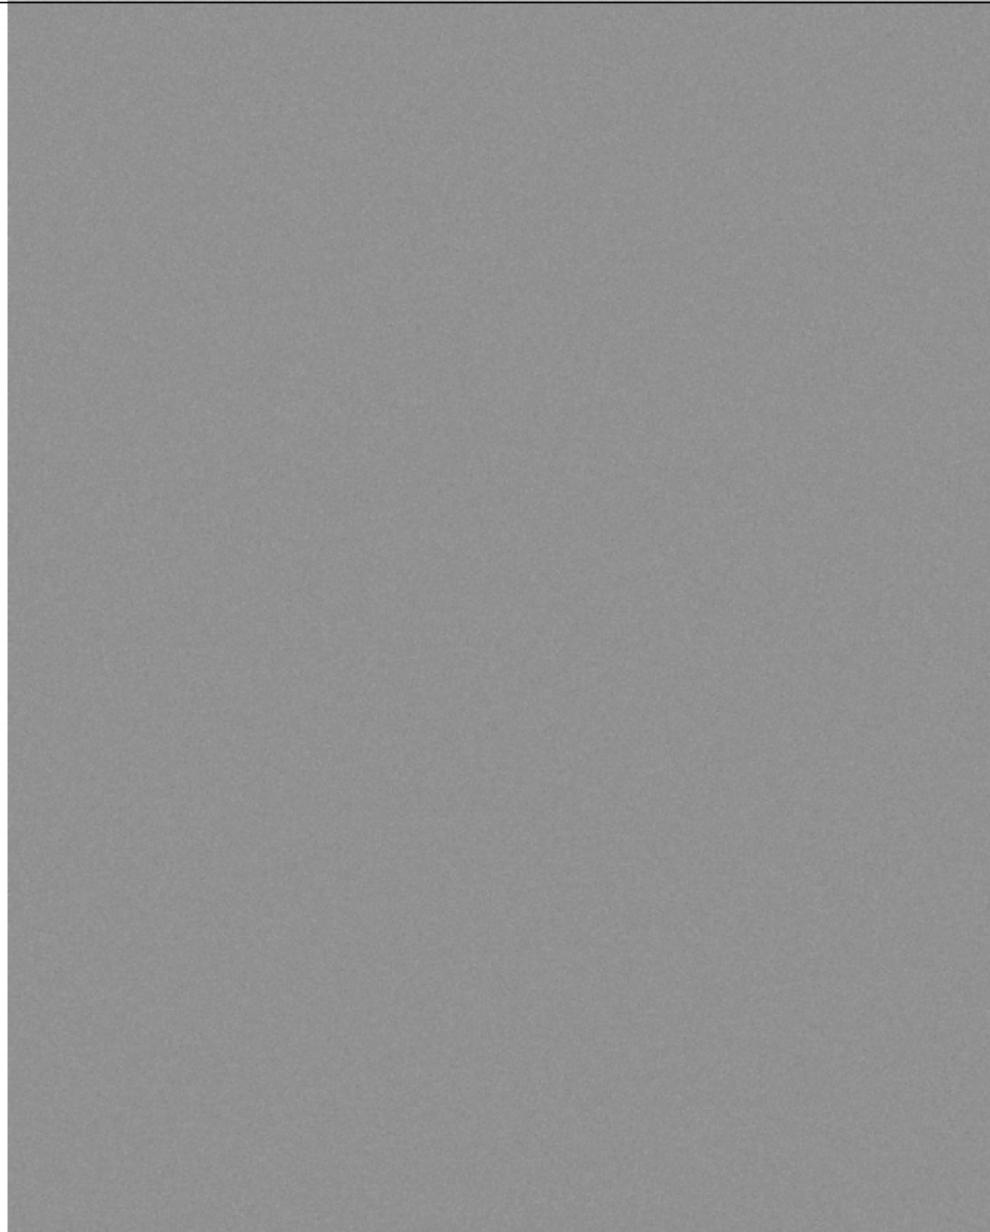
Image shows the noise values themselves.



# Gaussian noise

---

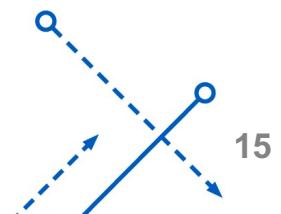
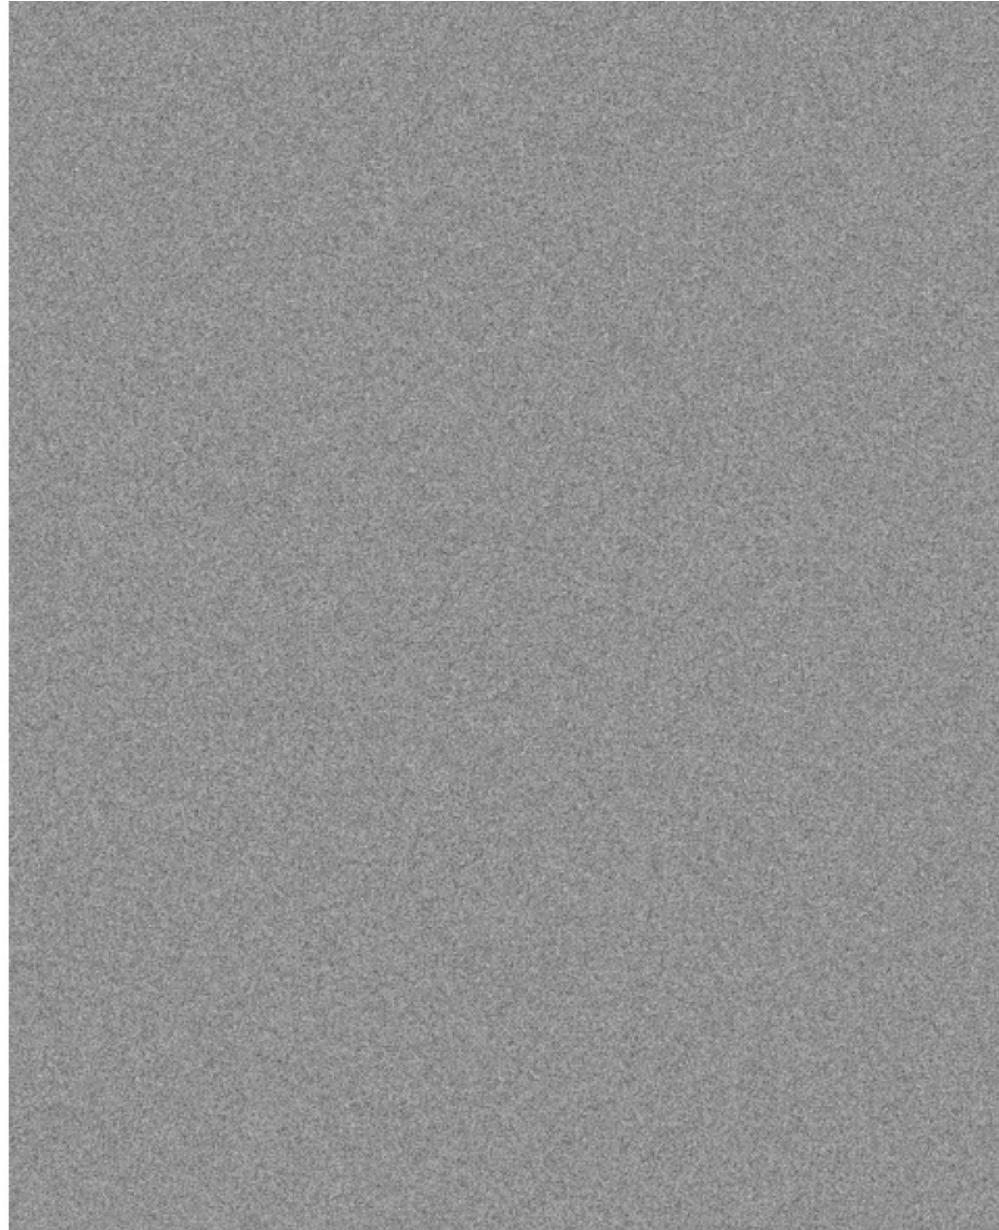
$\sigma = 4$



# Gaussian noise

---

$\sigma = 16$



# Gaussian noise

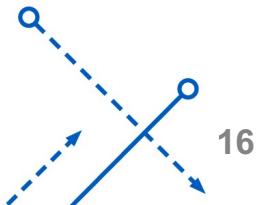
---

$\sigma = 1$



Effect of sigma on Gaussian noise:

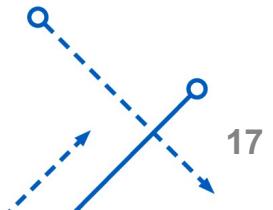
This shows the noise values added to the raw intensities of an image.



# Gaussian noise

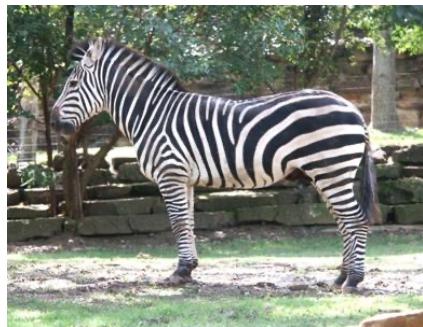
---

$\sigma = 16$

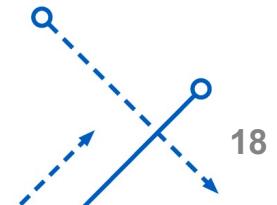


# Pixel neighborhoods are important.

- Q: What happens if we reshuffle all pixels?

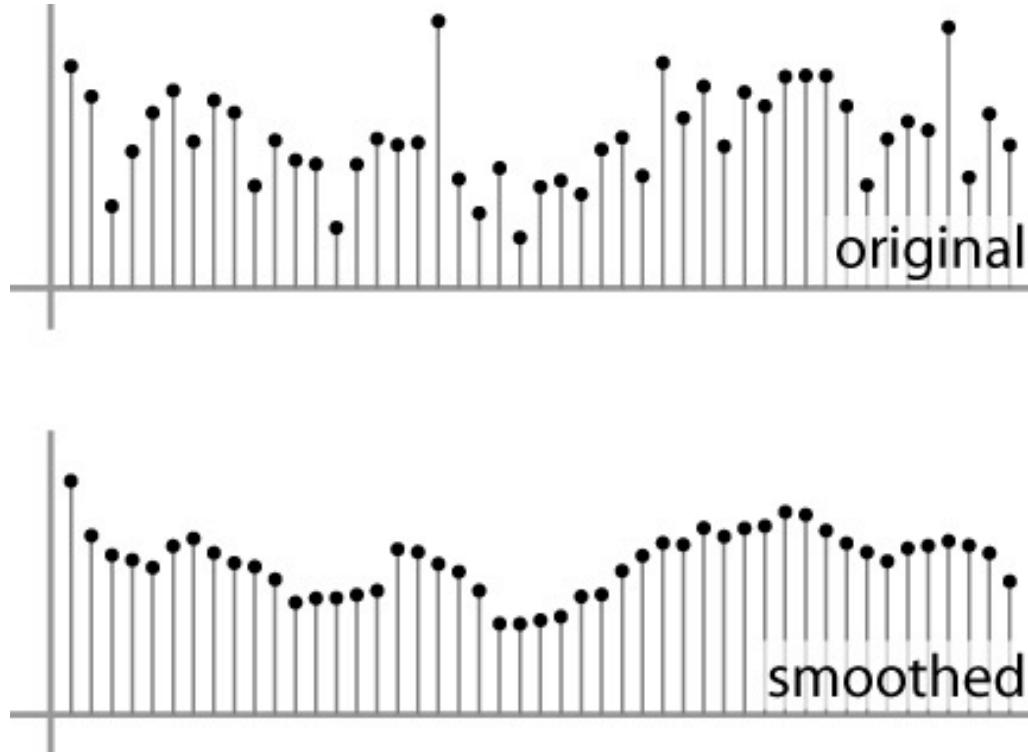


- A: Its histogram won't change.  
Point-wise processing unaffected.
- Can we use neighborhoods to remove image noise?



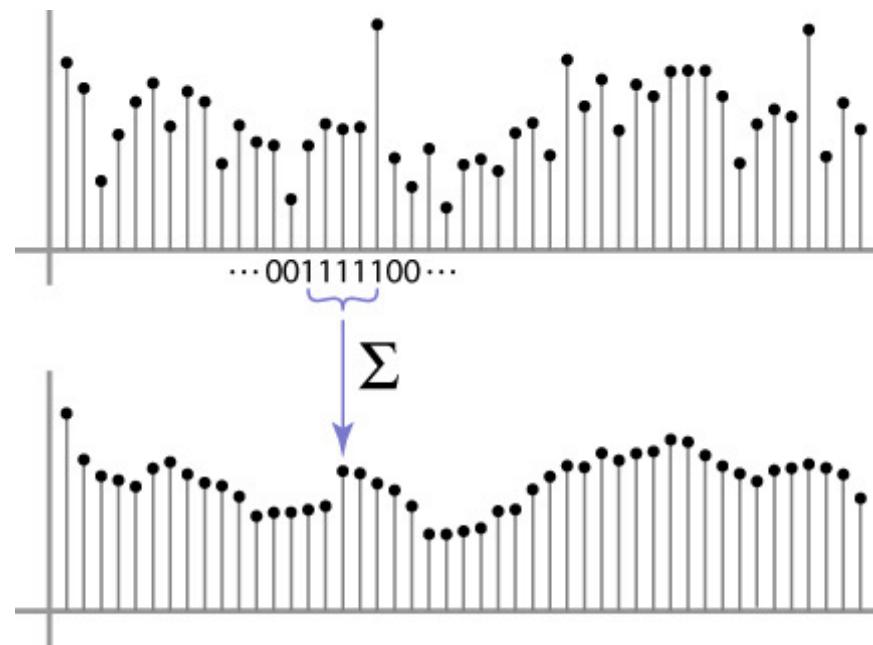
# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
  - Expect pixels to be like their neighbors
  - Expect noise processes to be independent from pixel to pixel.
  - Moving average in 1D:



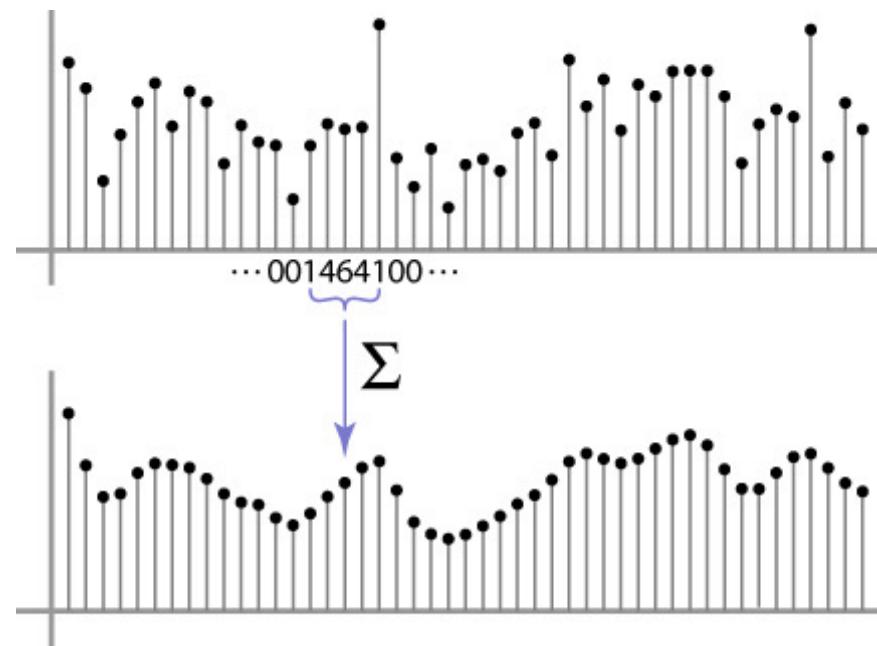
# Weighted Moving Average

- We can add weights to moving average
- *Weights*  $[1, 1, 1, 1, 1] / 5$



# Weighted Moving Average

- Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



# Example: Box Filter

---

$$f[\cdot, \cdot]$$

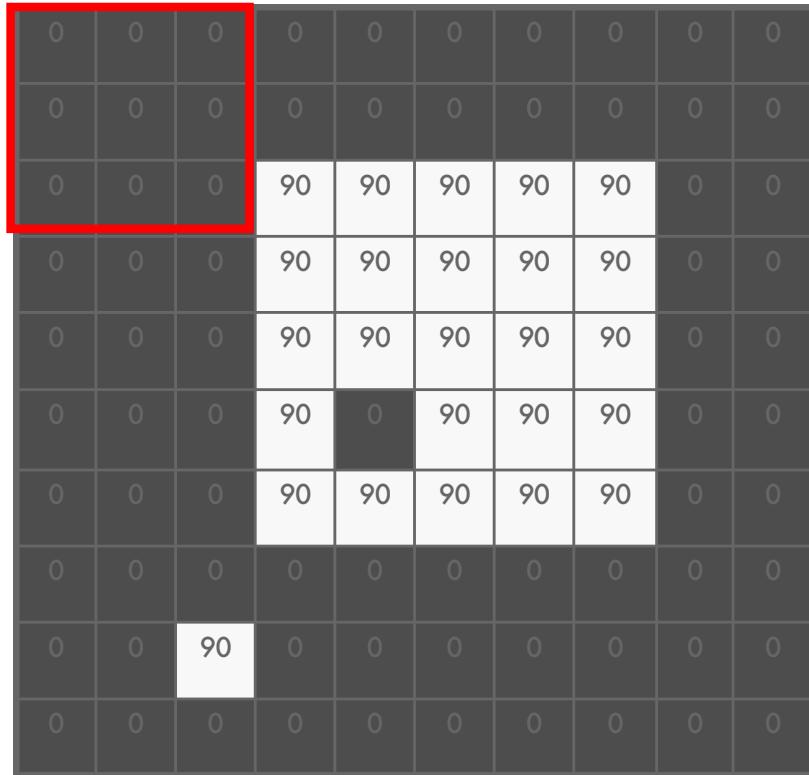
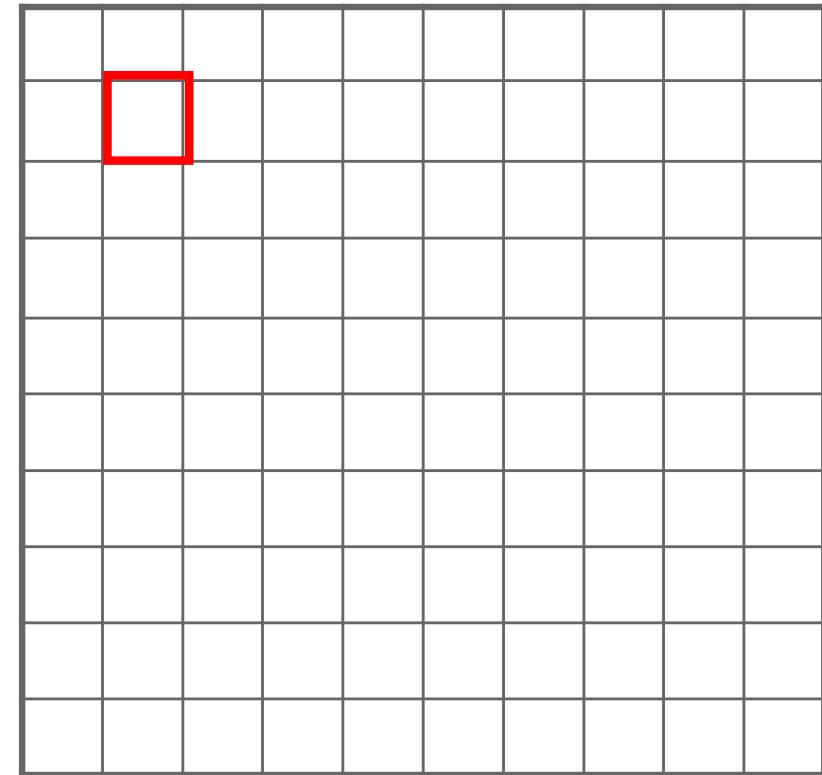
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot]$$

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

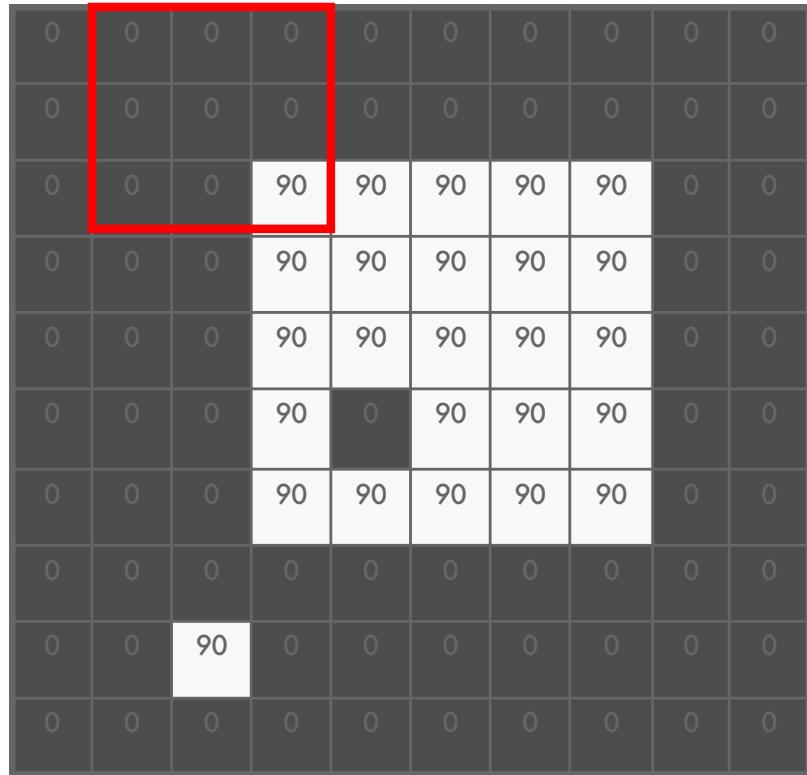
# Box Filter

 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

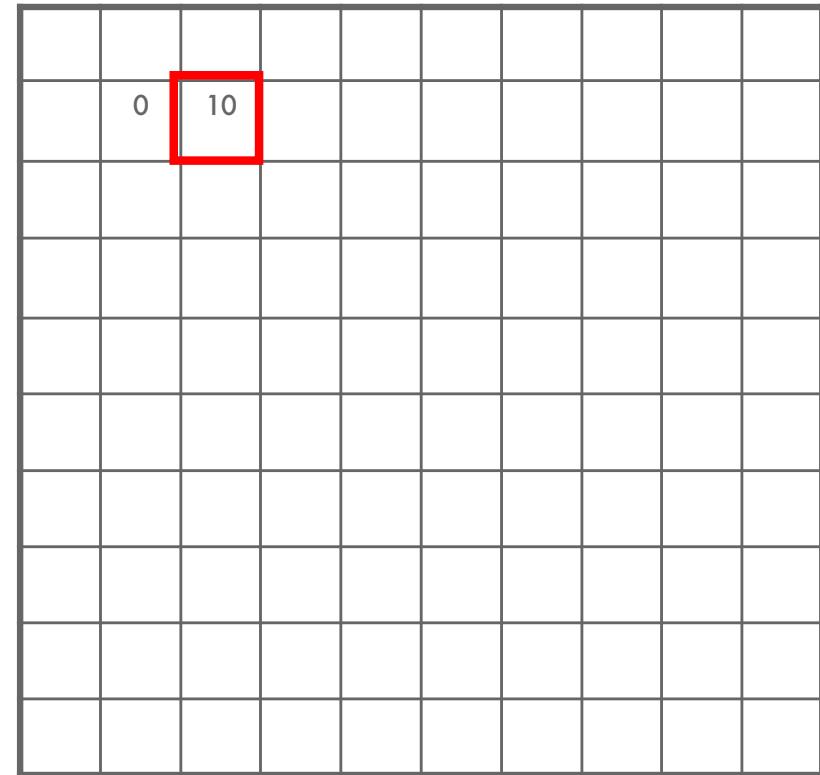
$$\begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

# Box Filter

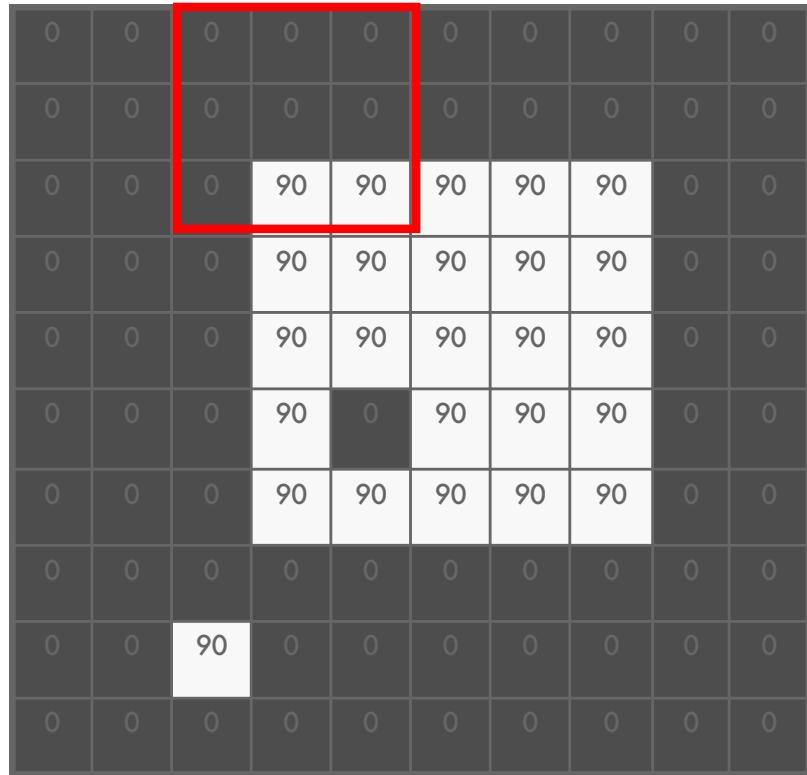
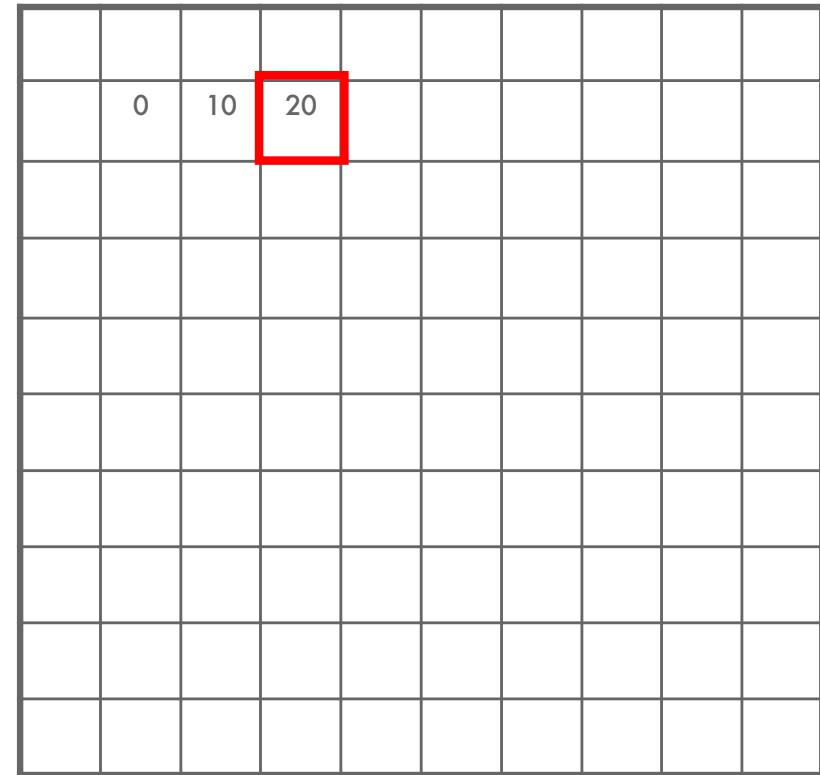
 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

1	1	1
1	1	1
1	1	1



$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

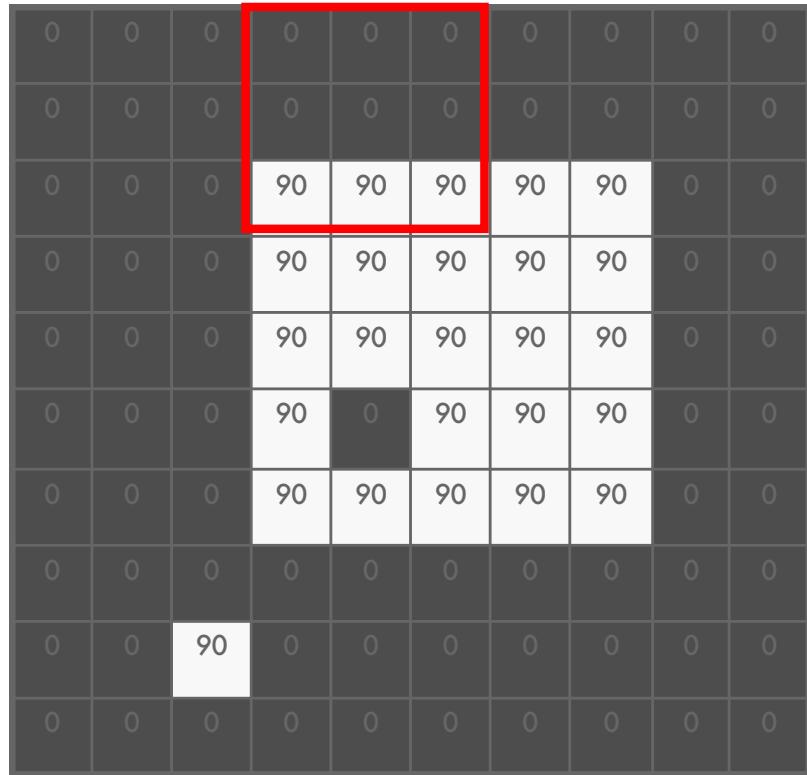
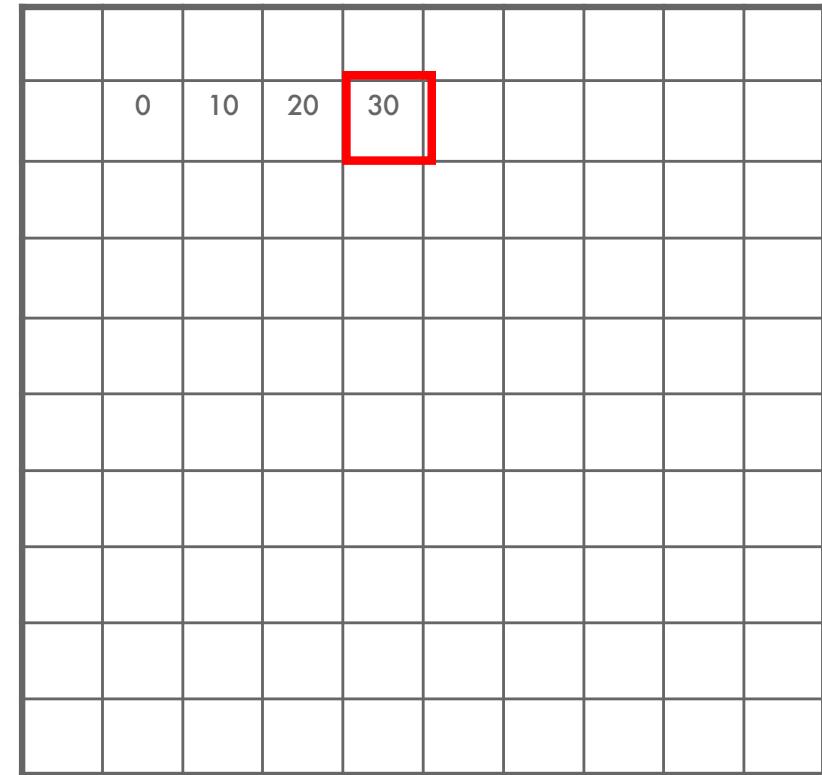
# Box Filter

 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

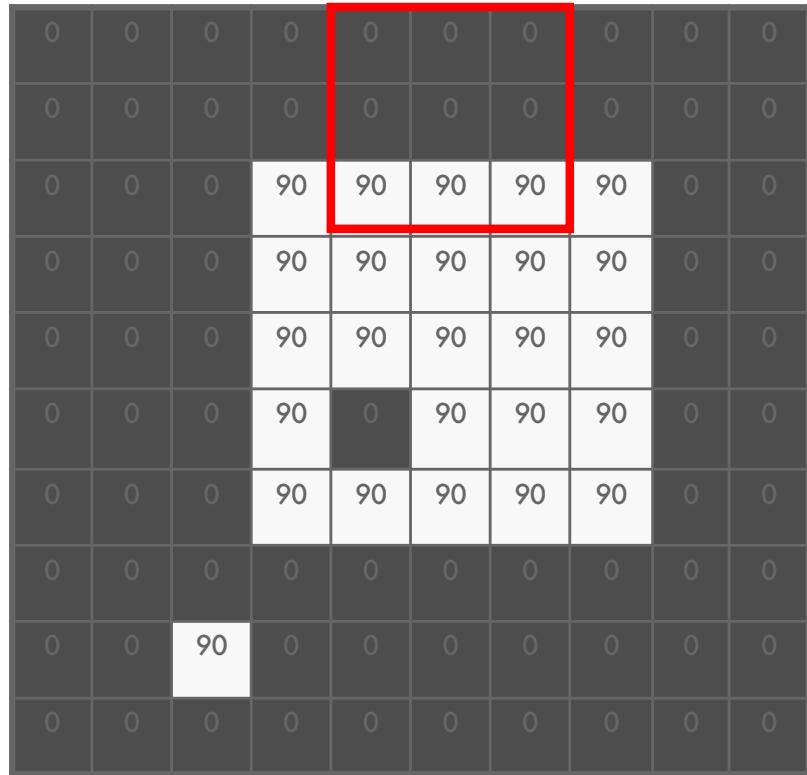
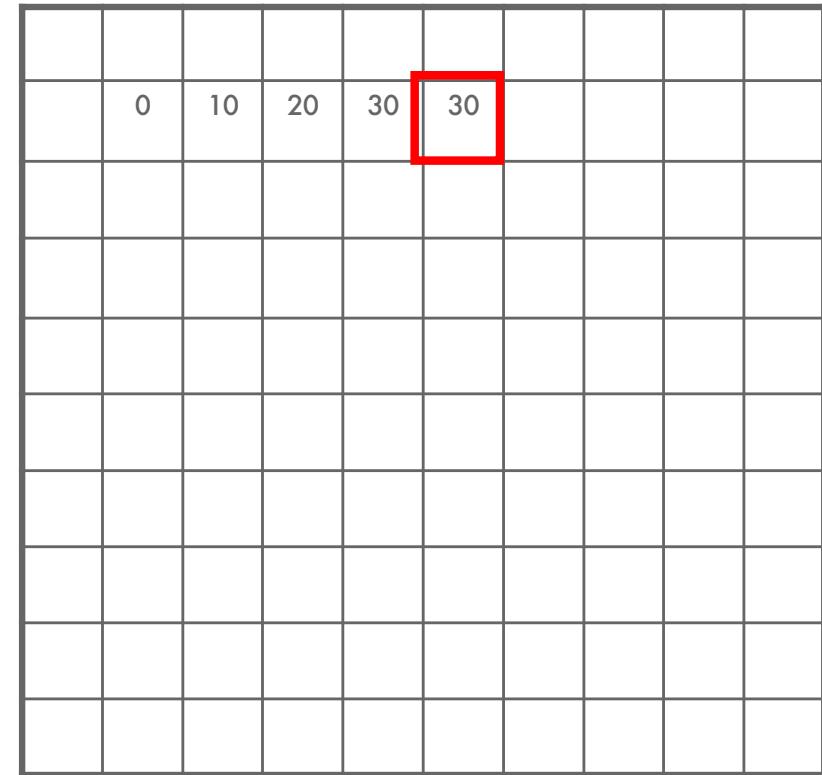
# Box Filter

 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

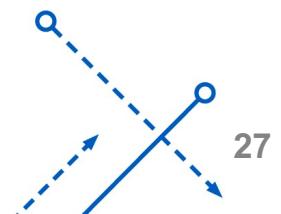
$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

# Box Filter

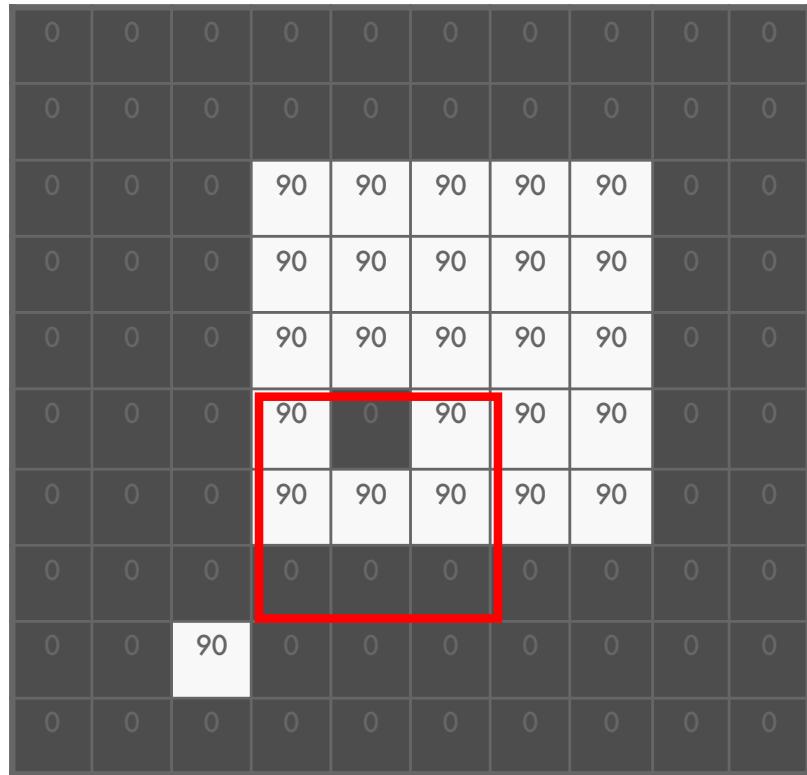
 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

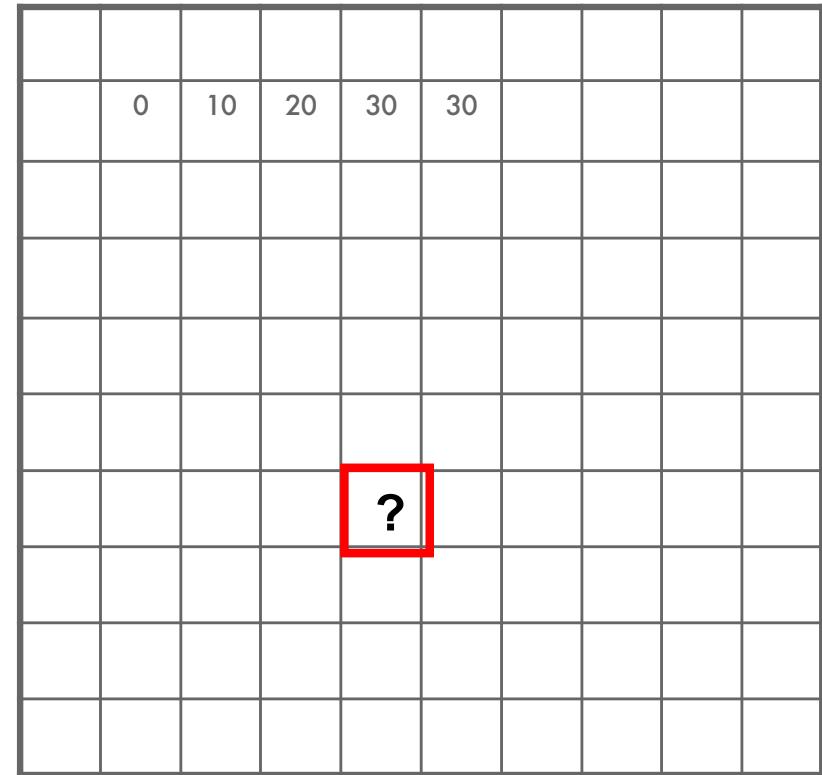
$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$



# Box Filter

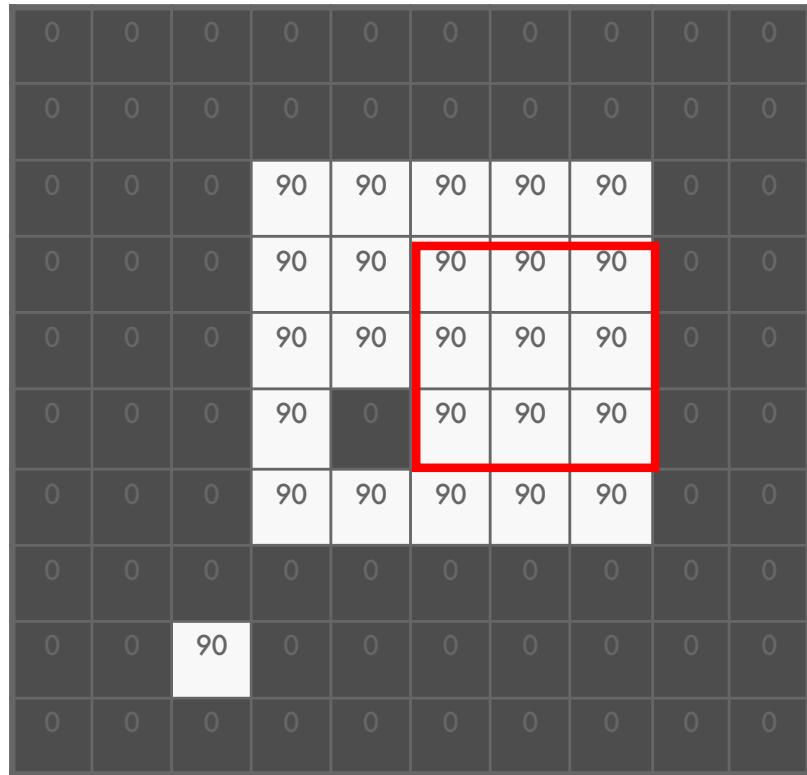
 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

1	1	1
1	1	1
1	1	1

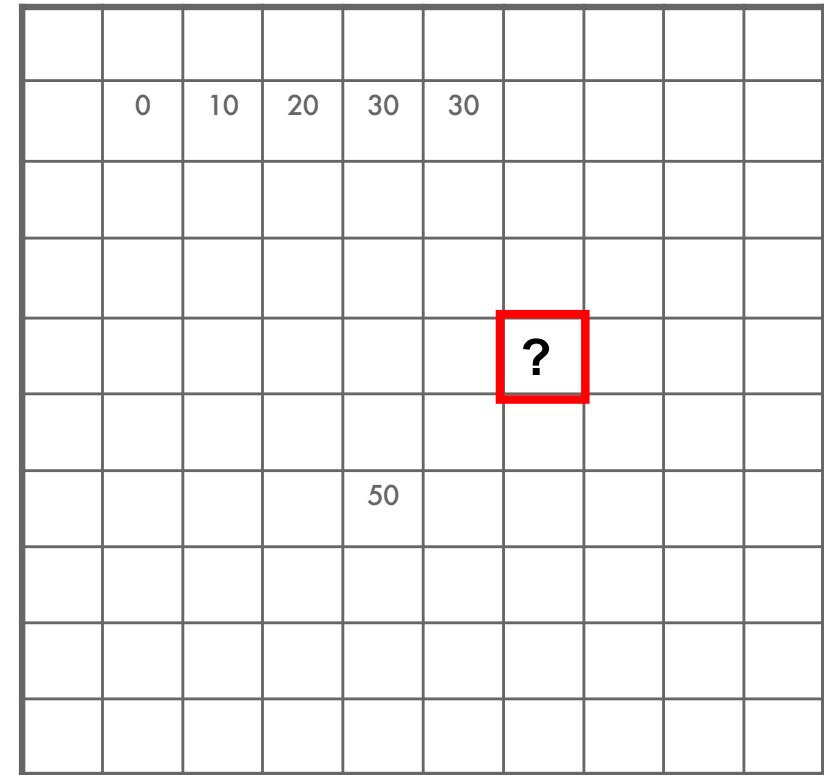


$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

# Box Filter

 $f[\cdot, \cdot]$  $g[\cdot, \cdot]$  $h[\cdot, \cdot]$ 

1	1	1
1	1	1
1	1	1



$$g[m, n] = \sum_{k,l} h[k, l] \ f[m + k, n + l]$$

# Box Filter

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

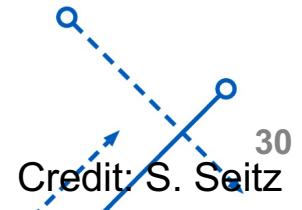
$g[\cdot, \cdot]$

$h[\cdot, \cdot]$

1	1	1
1	1	1
1	1	1

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

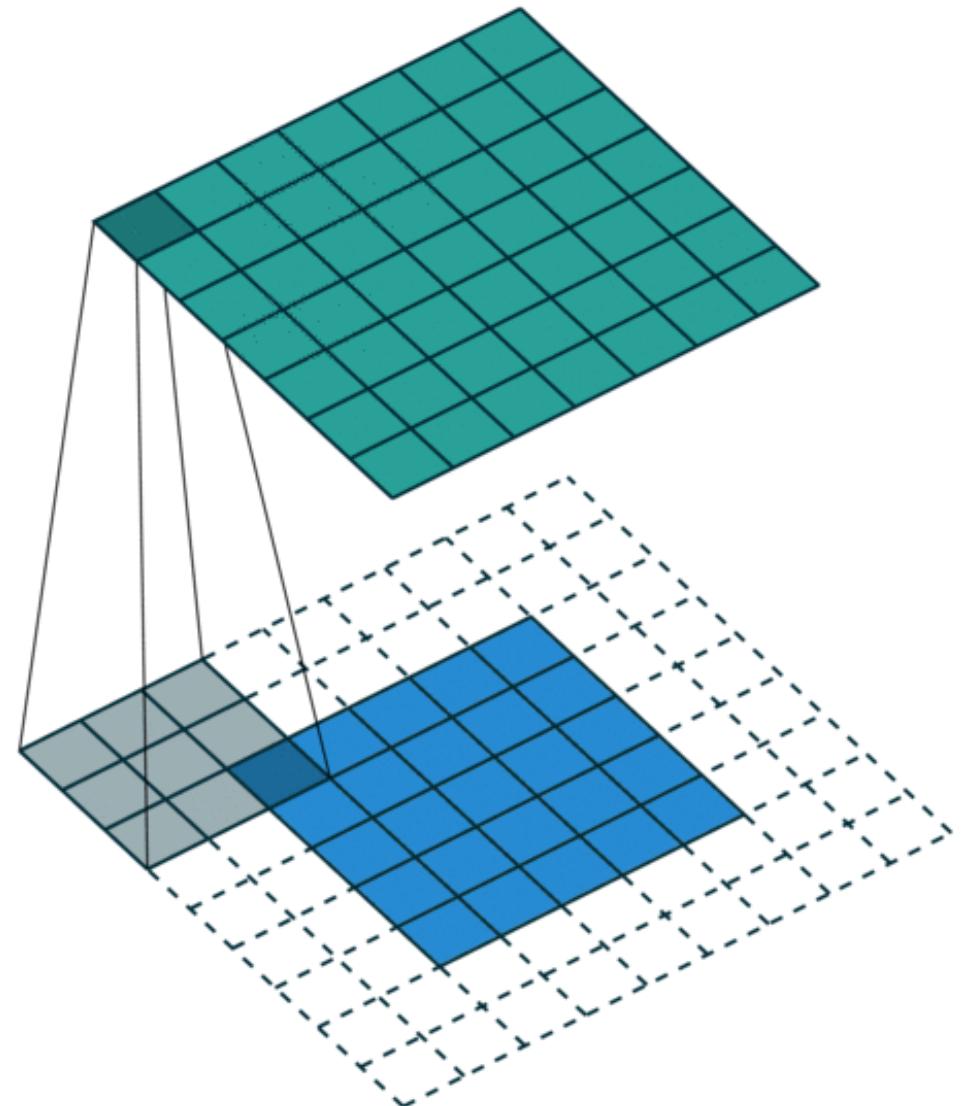
$$g[m, n] = \sum_{k,l} h[k, l] f[m + k, n + l]$$



# Correlation

---

- Correlation is the process of moving a filter mask over an image
- At each point in the image, one computes the sum of products at each location
- Filter is often referred to as the **Kernel or Mask**.
- Correlation is a function of the displacement



# Box Filter



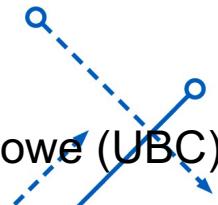
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$h[\cdot, \cdot]$

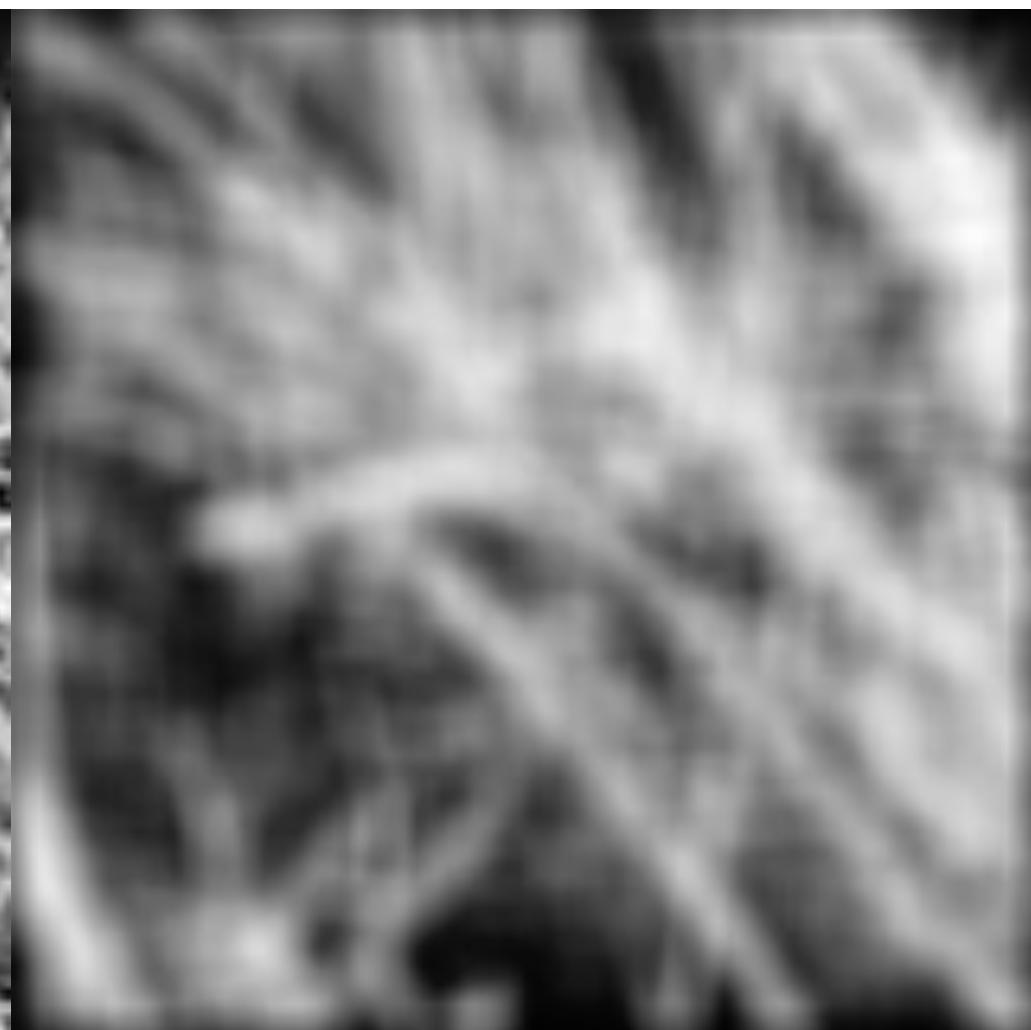
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

1	1	1
1	1	1
1	1	1



# Smoothing with box filter

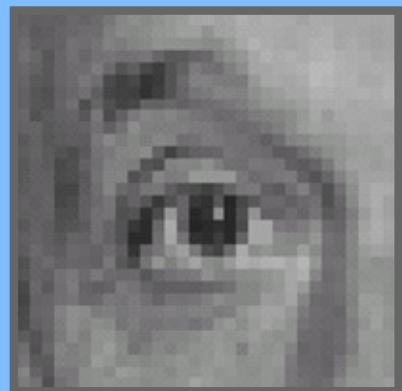
---



# Predict the filtered outputs

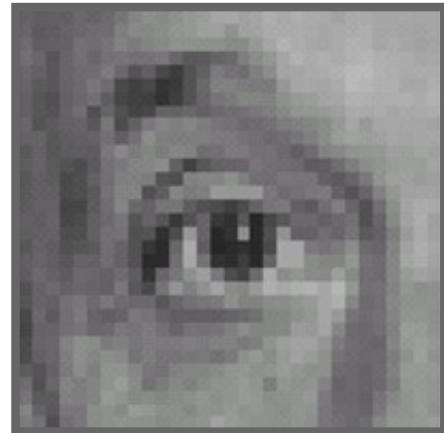

$$\text{Input Image} * \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} = ?$$


$$\text{Input Image} * \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} = ?$$


$$\text{Input Image} * \begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} = ?$$

# Practice with linear filters

---



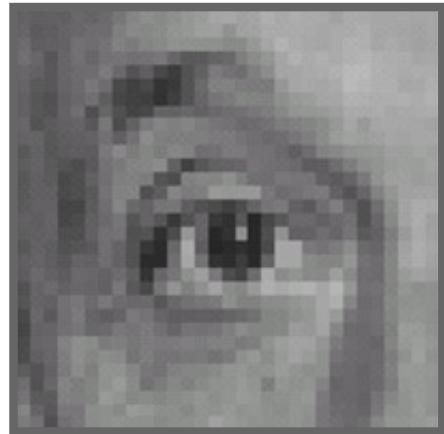
0	0	0
0	1	0
0	0	0

?

Original

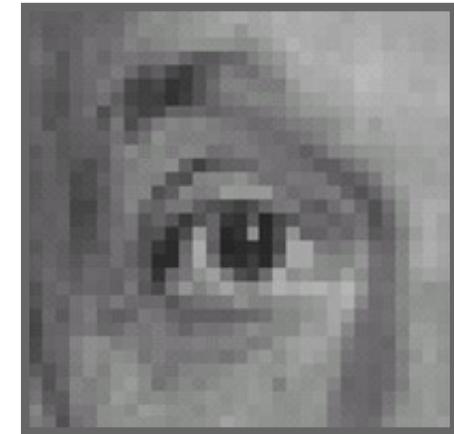
# Practice with linear filters

---



Original

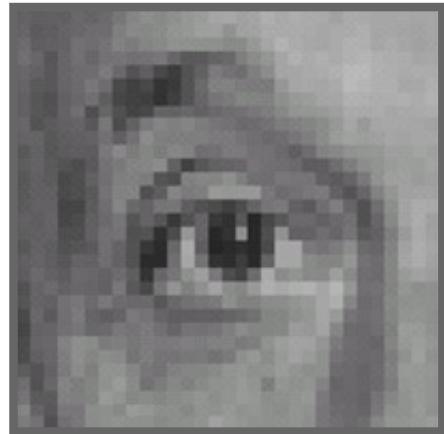
0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Practice with linear filters

---



0	0	0
0	0	1
0	0	0

?

Original

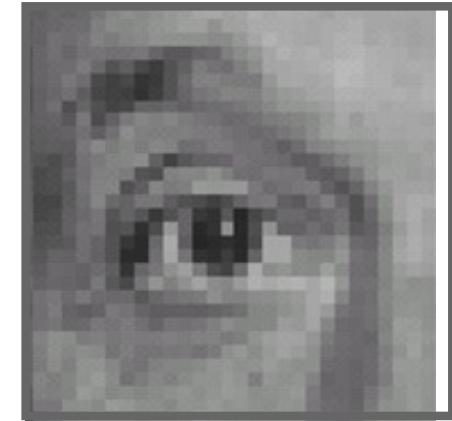
# Practice with linear filters

---



Original

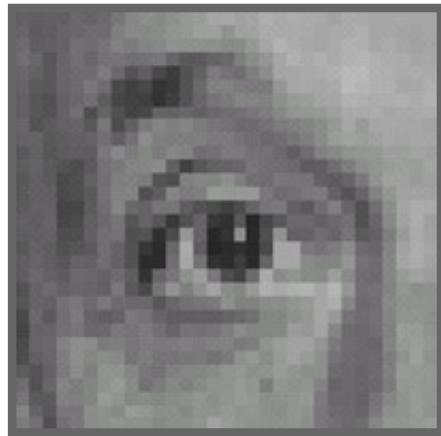
0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

# Practice with linear filters

---



Original

0	0	0
0	2	0
0	0	0

-

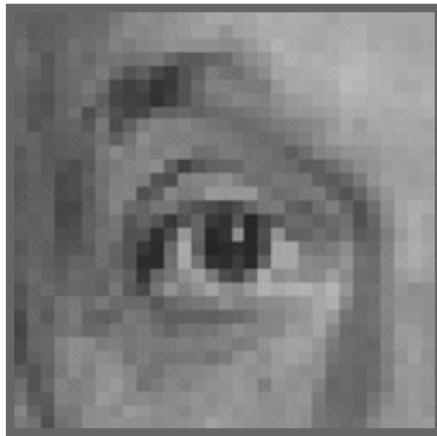
$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

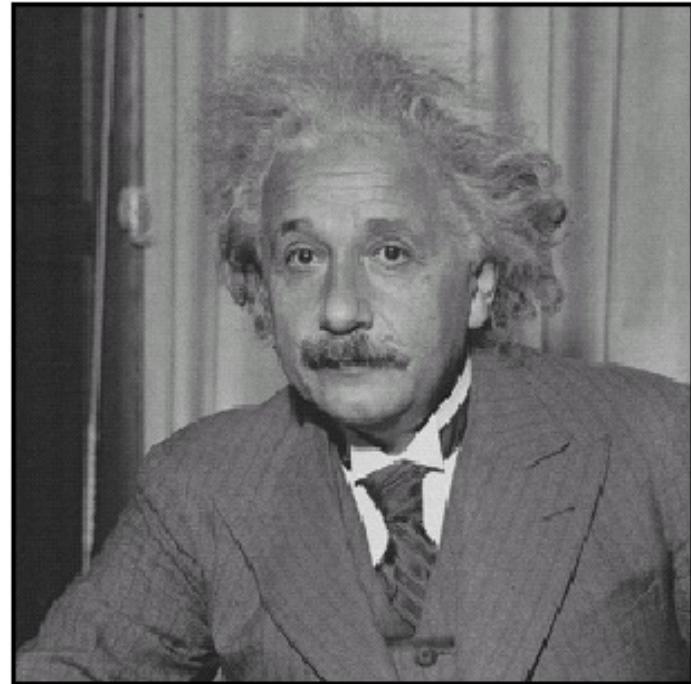


**Sharpening filter**  
Accentuates  
differences with local  
average

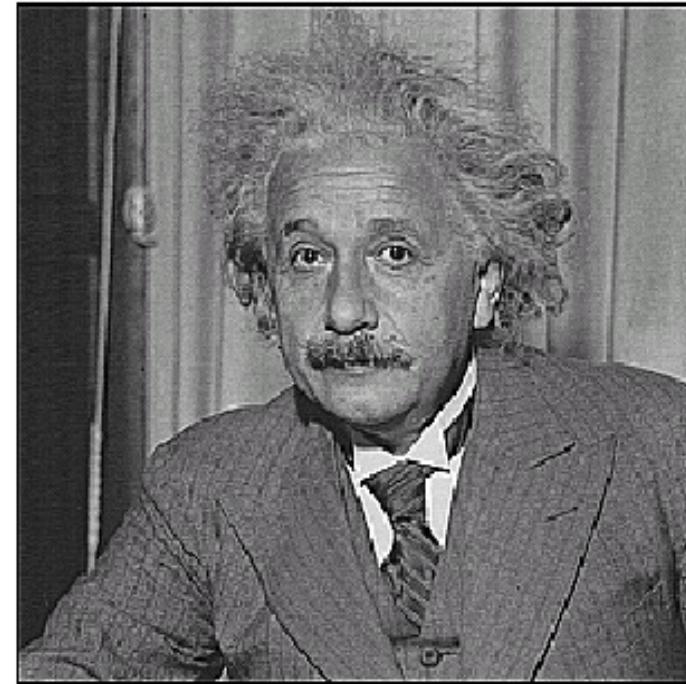
Source: D. Lowe  
40

# Sharpening

---



**before**



**after**

# Correlation (another name for filtering)

$$G[m, n] = \sum_{k,l} H[k, l] F[m + k, n + l]$$

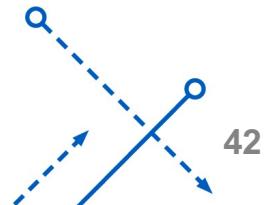
H=Filter      F=Signal

- 2D Convolution

- PyTorch implements convolution as correlation
- import torch.nn.functional as F
- `G = F.conv2d(I, f)`

I=image      f=filter

$$G[m, n] = \sum_{k,l} H[k, l] F[m - k, n - l]$$



# Correlation vs. Convolution

- 2D Correlation

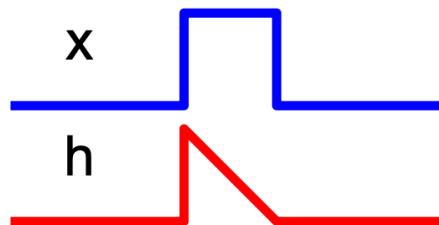
$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

- 2D Convolution

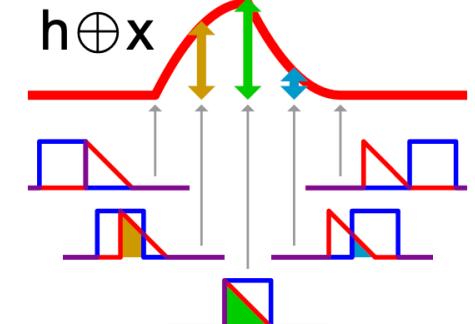
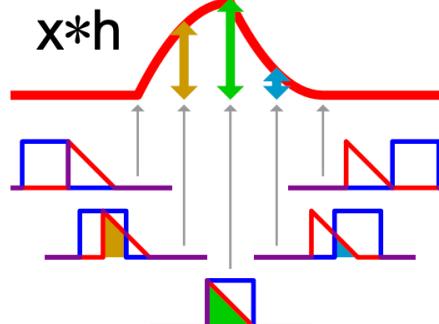
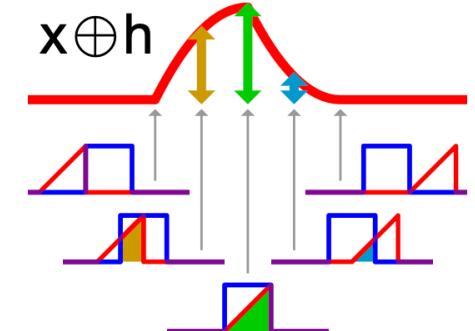
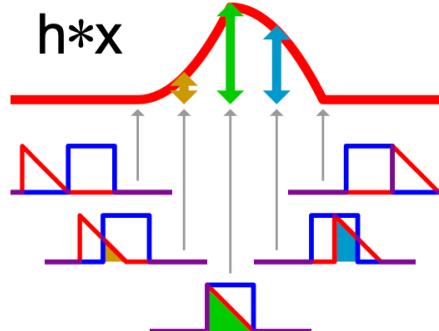
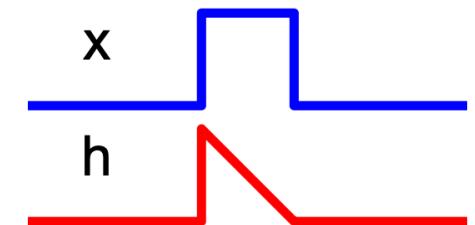
$$h[m, n] = \sum_{k, l} f[k, l] I[m - k, n - l]$$

- 1D Case

Cross-Correlation



Convolution



# Correlation filtering

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Say the averaging window size is  $(2k+1) \times (2k+1)$ :

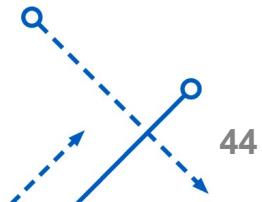
$$G[i, j] = \frac{1}{(2k+1)^2} \sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]$$

*Attribute uniform weight to each pixel*      *Loop over all pixels in neighborhood around image pixel  $F[i,j]$*

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

*Non-uniform weights*



# Correlation filtering

---

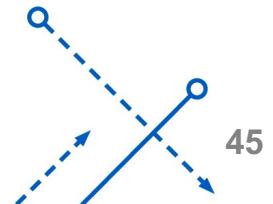
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

More formally, it is called cross-correlation, i.e.,

$$G = H \otimes F$$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter is also called “kernel”, “mask”, or a “window”.



# More Formal/Strict Terminology

---

- Cross-correlation:
  - Often referred as **correlation** in computer vision.

$$G[m, n] = \sum_{k,l} H[k, l] \ F[m + k, n + l]$$

- Correlation (in statistics):

$$\text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{\mathbb{E}[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}.$$

- Reading more:
  - Chen Wang, et al. "[Kernel Cross-correlator](#)." In Proceedings of the AAAI Conference on Artificial Intelligence, 2018.
  - Convolution vs Cross-correlation
    - Section 2.2 in “Chen Wang, [Kernel learning for visual perception](#)”. PhD thesis.

# Properties

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

- Commutative:

$$x \otimes y = y \otimes x$$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality

- Associative:

$$x \otimes (y \otimes z) = (x \otimes y) \otimes z$$

- Often apply several filters one after another:  $x \otimes y_1 \otimes y_2 \otimes y_3$
- This is equivalent to applying one filter:  $x \otimes (y_1 \otimes y_2 \otimes y_3)$

- Distributes over addition:

$$x \otimes (y + z) = (x \otimes y) + (x \otimes z)$$

- Scalars factor out:

$$ax \otimes y = x \otimes ay = a(x \otimes y)$$

- Identity (unit impulse), e.g.,  $e = [0, 0, 1, 0, 0]$ ,

$$x \otimes e = x$$



Source: S. Lazebnik 47

# Key properties of Linear Filters

---

- Assume  $f(x)$  is image filtering:  $f(x) = x \otimes w$ .

- Linearity (superposition property):

$$f(a \cdot x + b \cdot y) = af(x) + bf(y)$$

- Linear filter is **equivariant** (not **invariant**) to translation.

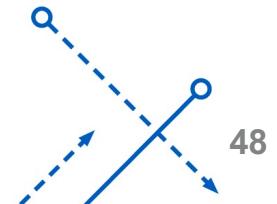
- Assume  $T(x)$  is image shift (translation):

- **Equivariance:**

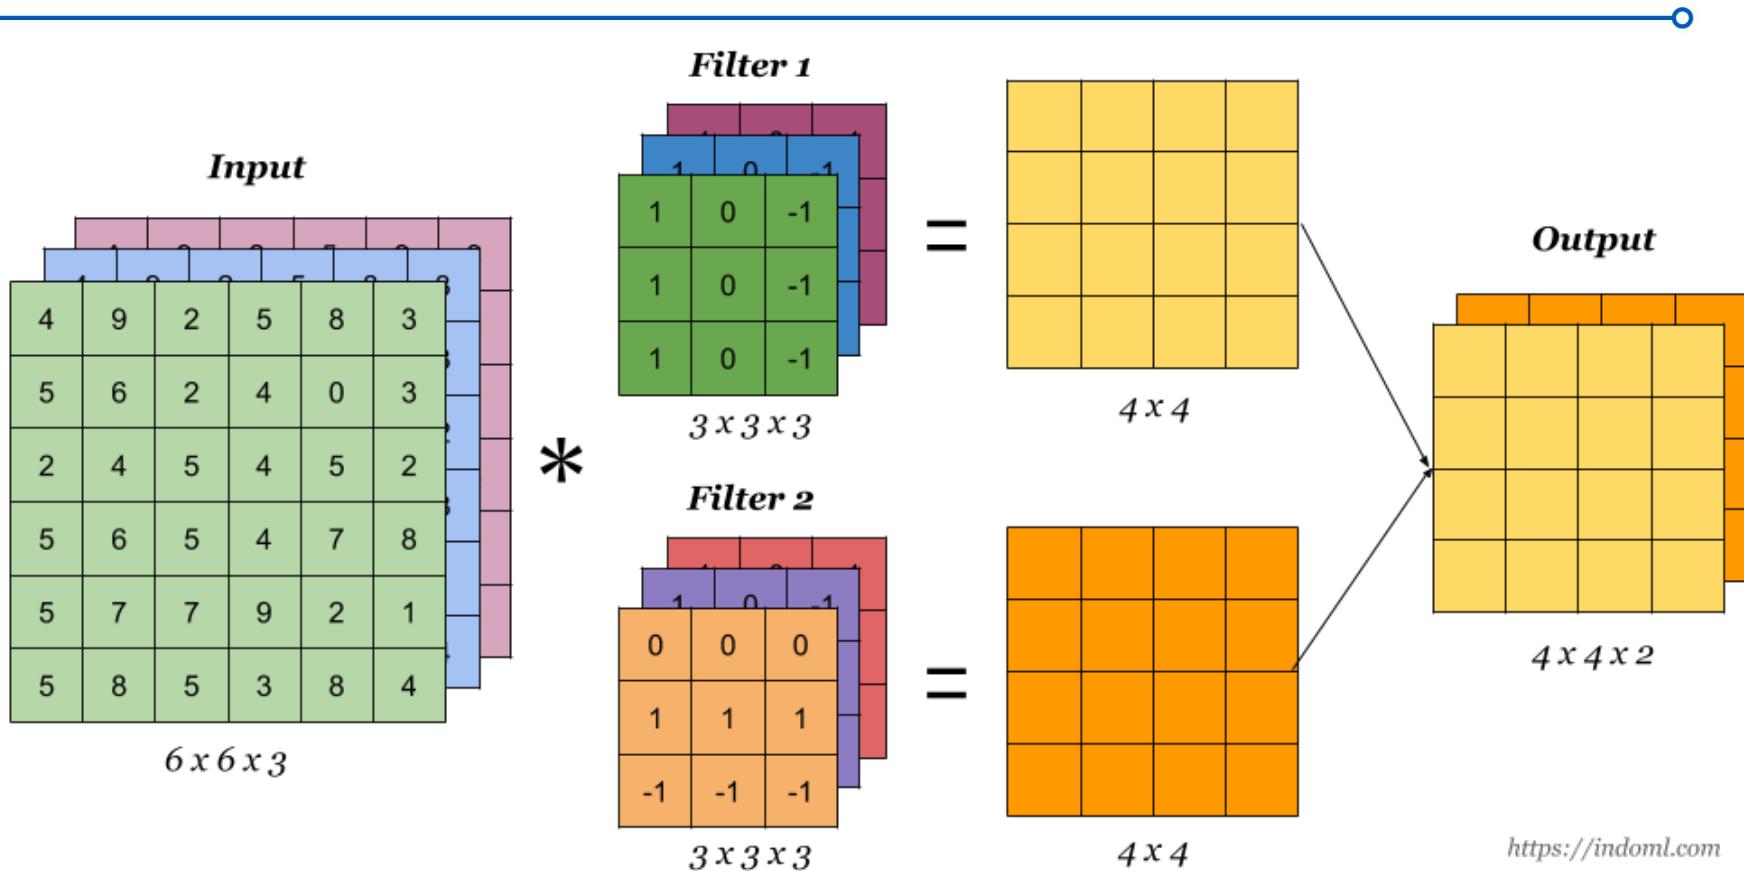
$$f(T(x)) = T(f(x))$$

- **Invariance:**

$$f(T(x)) = f(x)$$



# Multi-channel correlation

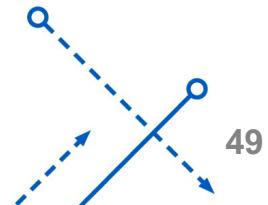


Examples:

```
>>> # With square kernels and equal stride
>>> filters = torch.randn(8, 4, 3, 3)
>>> inputs = torch.randn(1, 4, 5, 5)
>>> F.conv2d(inputs, filters, padding=1)
```

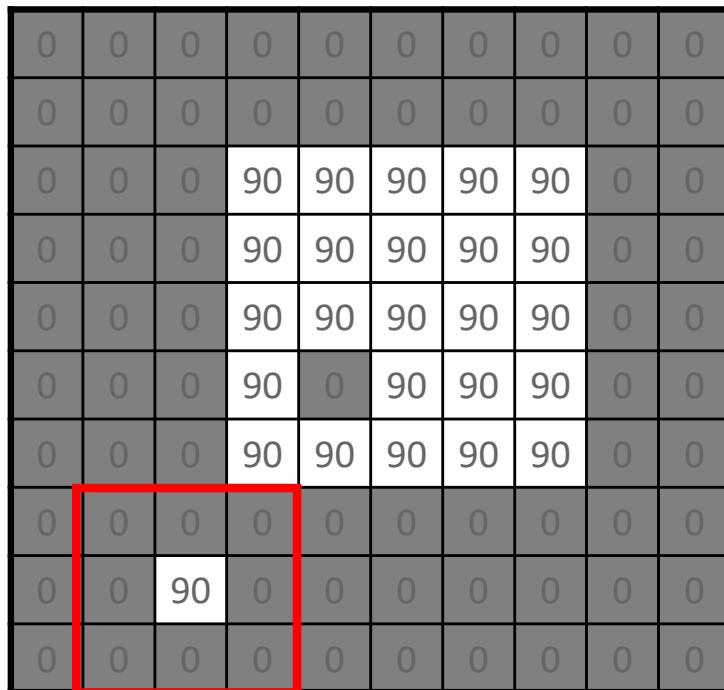
What is shape of the output  
in the left example?

(1, 8, 5, 5)



# Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?



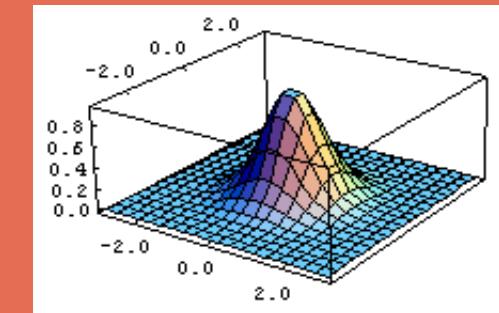
$F[x, y]$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

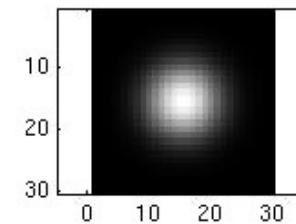
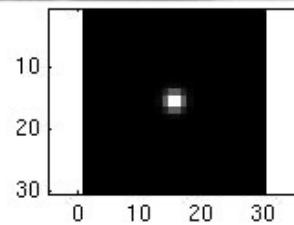
This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

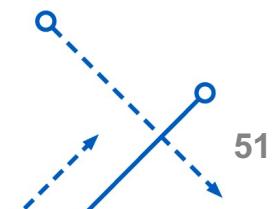
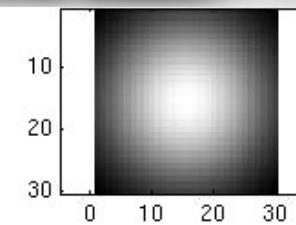


# Smoothing with a Gaussian

- Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.

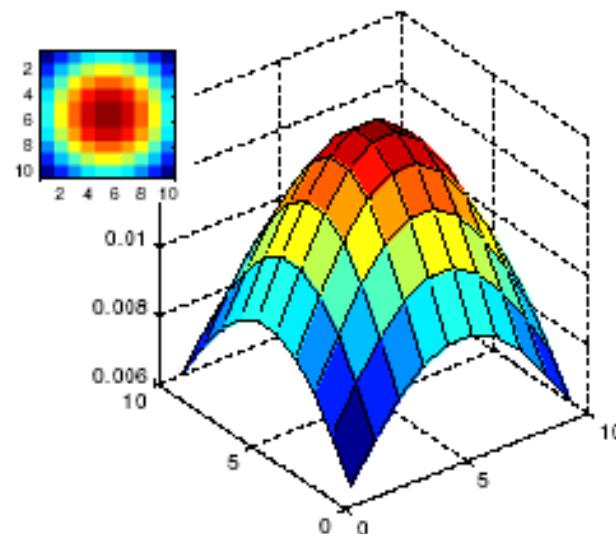


⋮ ⋮ ⋮

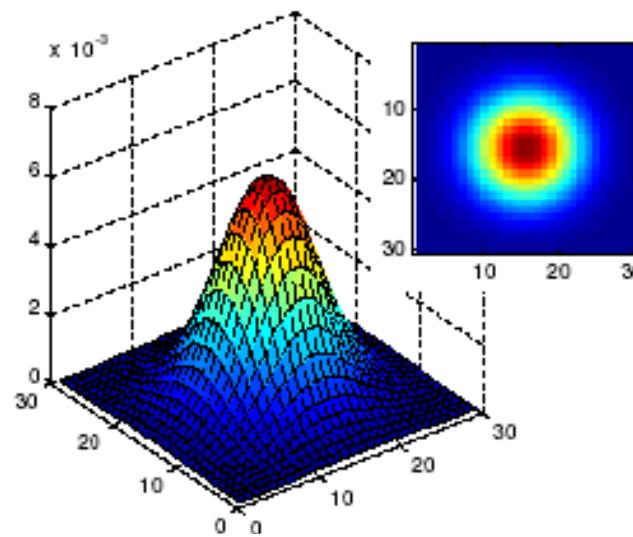


# Gaussian filters

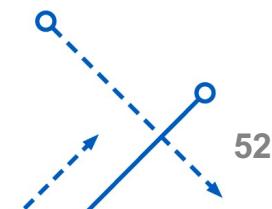
- What parameters matter here?
- **Size of kernel / mask**
  - Gaussian function has infinite support, but discrete filters use finite kernels



$\sigma = 5$  with  $10 \times 10$  kernel

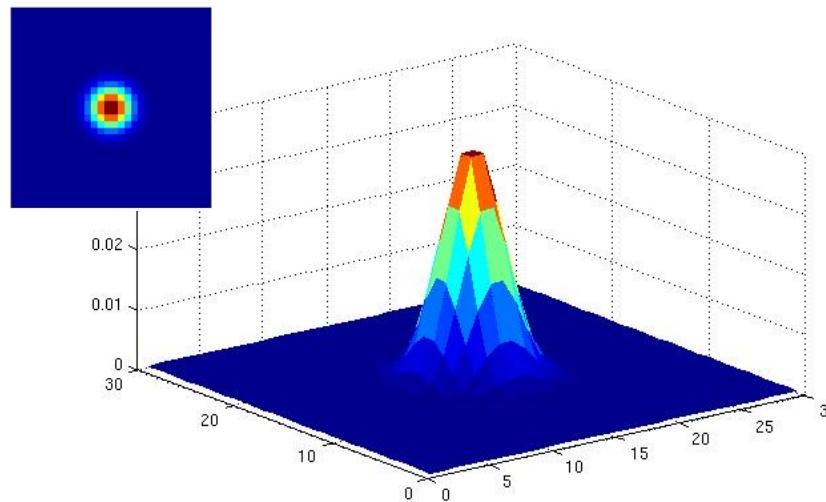


$\sigma = 5$  with  $30 \times 30$  kernel

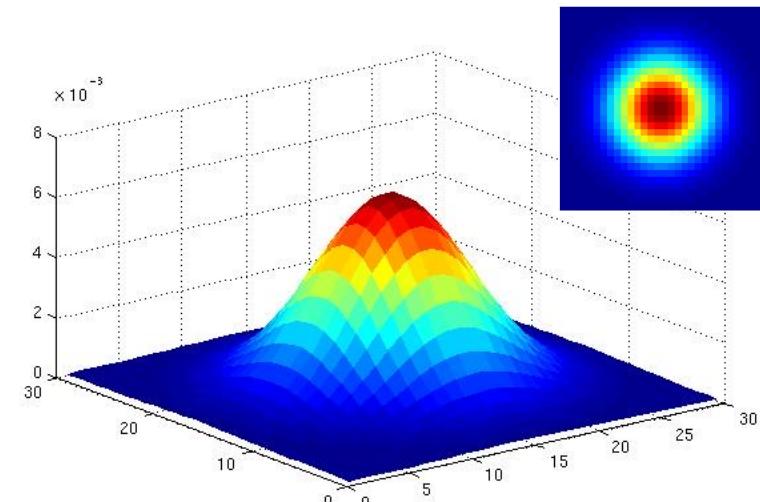


# Gaussian filters

- **Variance:** determines extent of smoothing



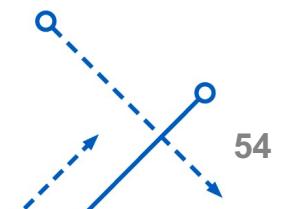
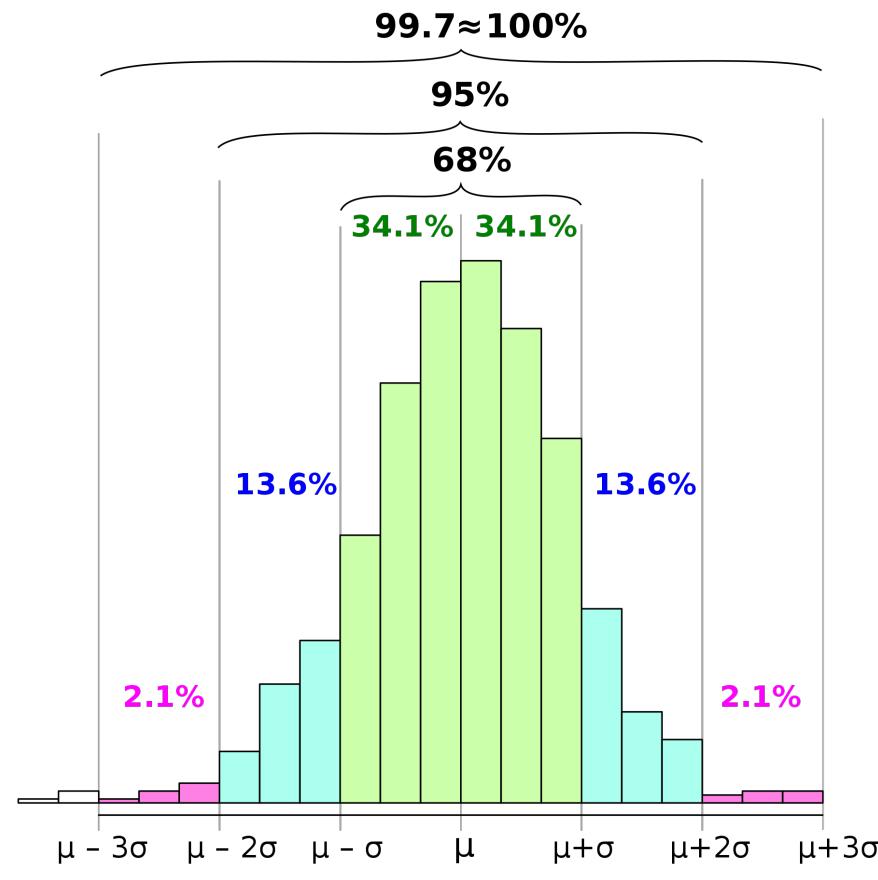
$\sigma = 2$  with  
30 x 30  
kernel



$\sigma = 5$  with  
30 x 30  
kernel

# How big should the filter be?

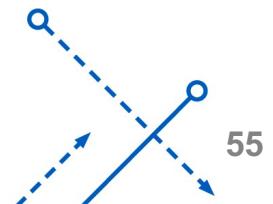
- Values at edges should be near zero
- Rule of thumb for Gaussian filter:
  - set filter half-width to about  $3\sigma$



# Gaussian filters

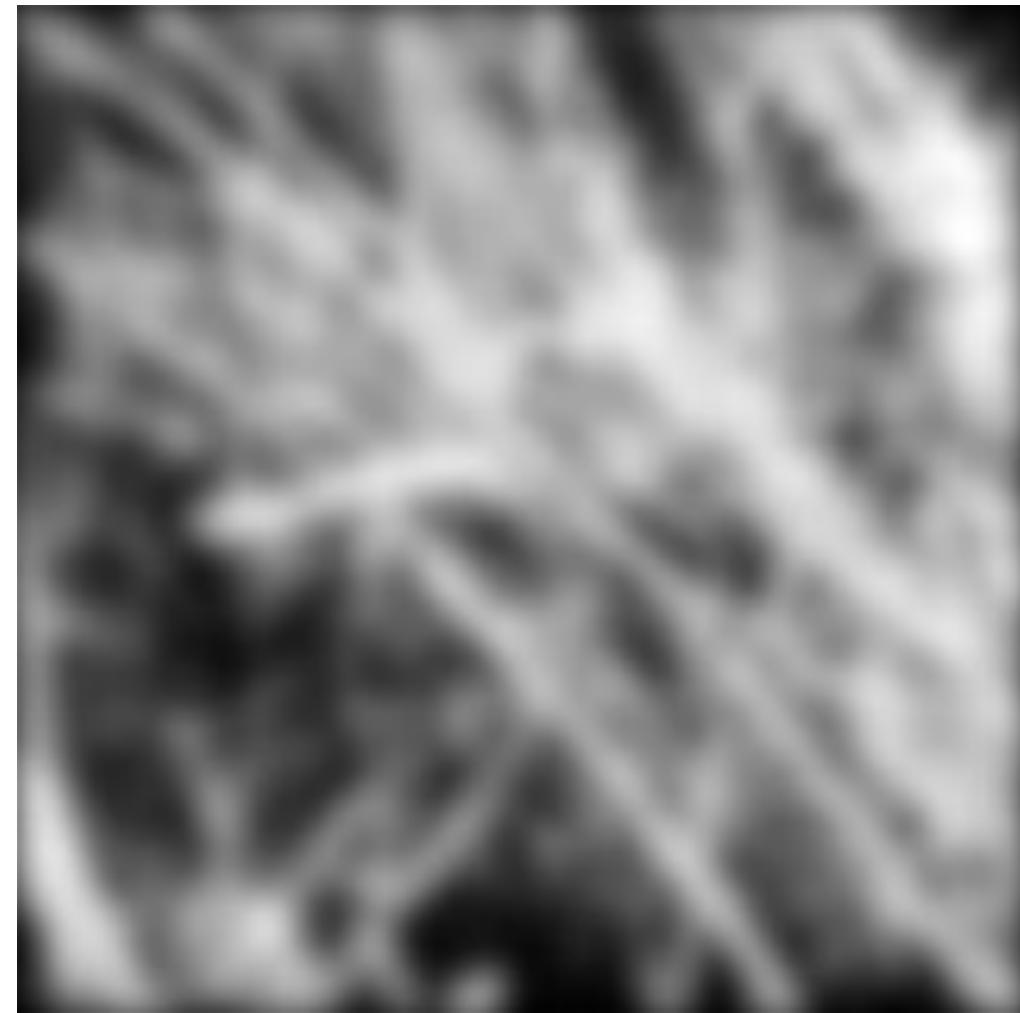
---

- Remove “high-frequency” components from image
  - Low-pass filter: Images become smoother.
- Recap of frequency in a signal (image)
  - High-frequency components
    - Fine details and edges
  - Low-frequency components
    - Large-scale structures and smooth regions



# Smoothing with Gaussian filter

---



# Smoothing with box filter

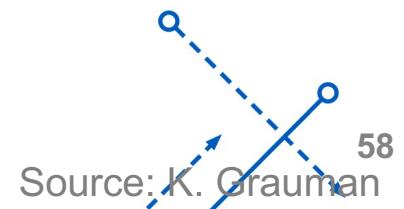
---



# Gaussian filters

---

- Convolution with self is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- Separable kernel
  - Factors into product of two 1D Gaussians



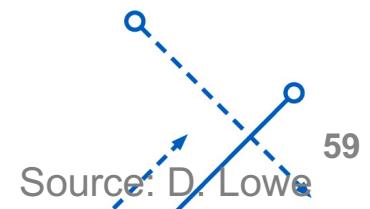
# Separability of the Gaussian filter

---

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian



# Separability

- In some cases, filter is separable, (factor into two steps):

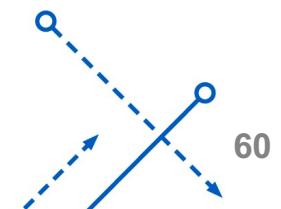
$$f \otimes (g \otimes h) = (f \otimes g) \otimes h$$

		h
g	1 2 1	2 3 3 3 5 5 4 4 6
		11 18 18

		f
f	1 2 1	11 18 18
		65

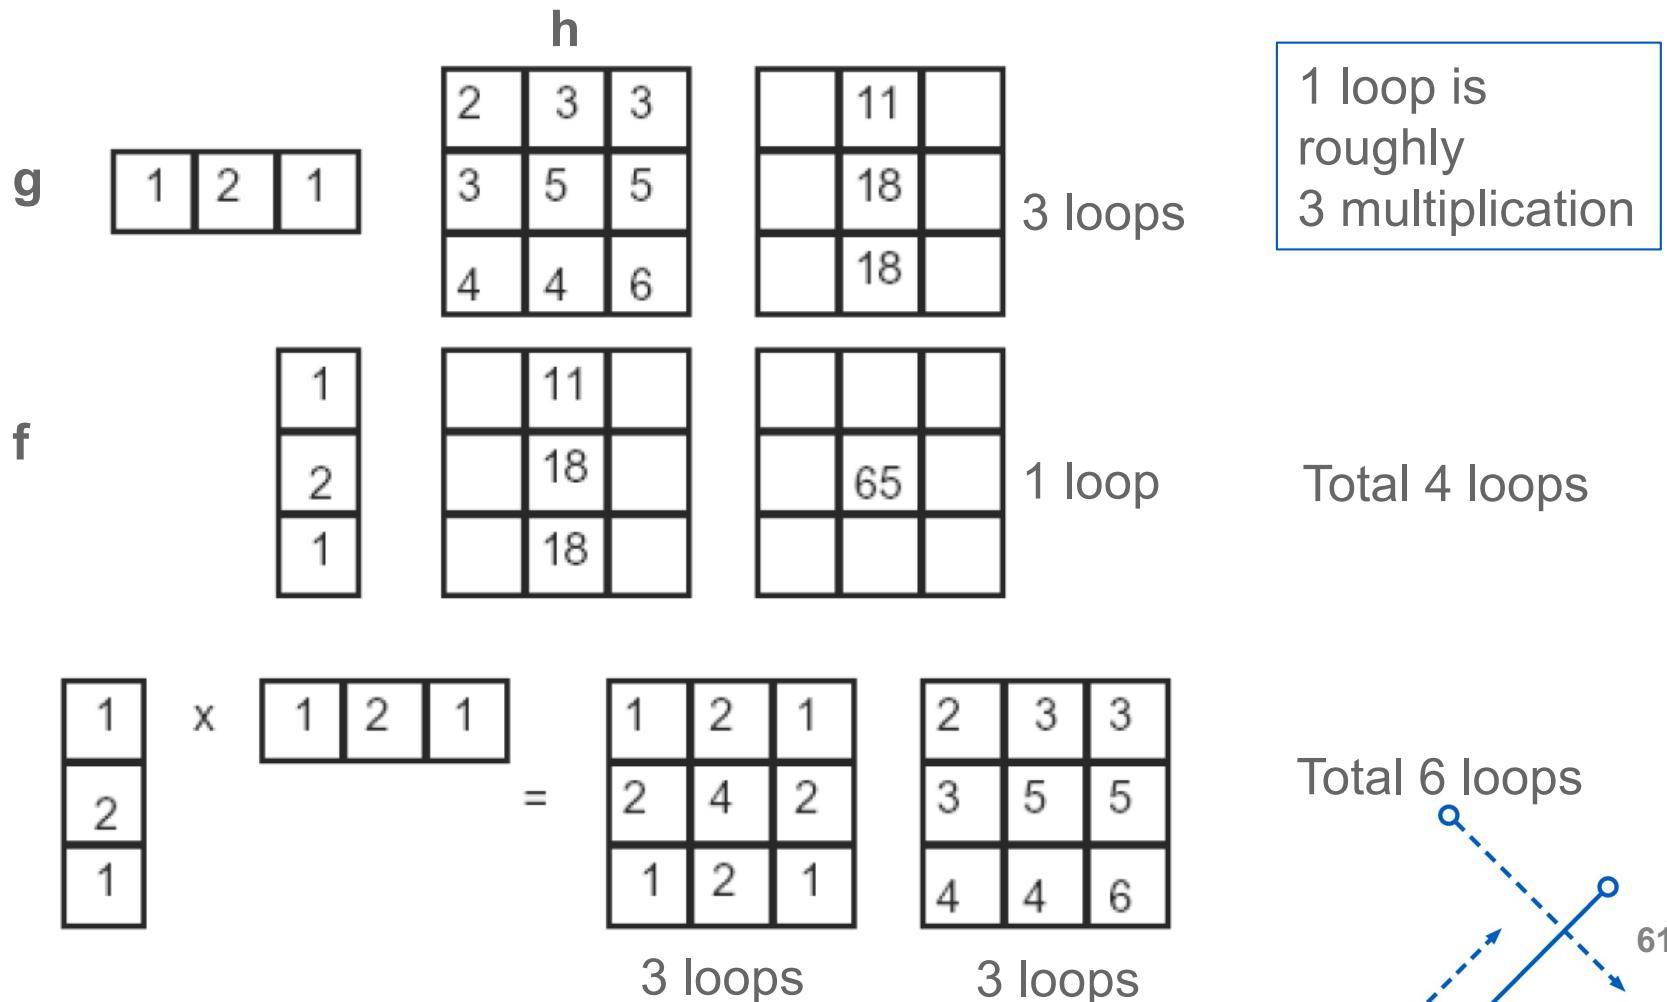
$$\begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} \quad \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} \quad \begin{matrix} 11 & & \\ & 18 & \\ & & 18 \end{matrix}$$

What is the computational complexity advantage for a separable filter of size  $k \times k$ , in terms of **number of operations per output pixel**?



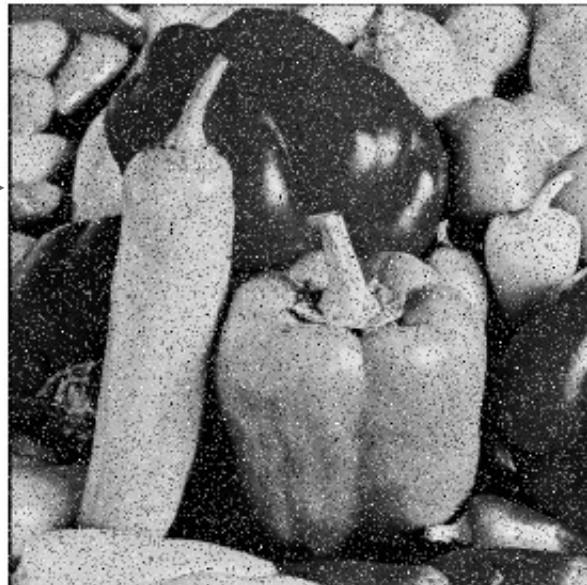
# Separability

- **Advantage:** much reduced computational cost.
- **Disadvantage:** requires an extra ram **memory** to store the intermediate image, problematic in certain applications.

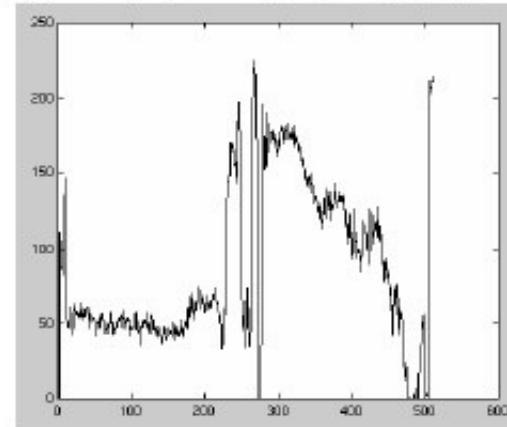
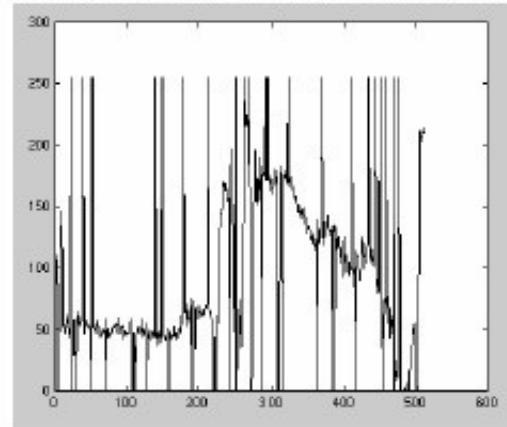


# Median filter

Salt and  
pepper  
noise



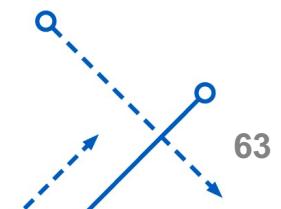
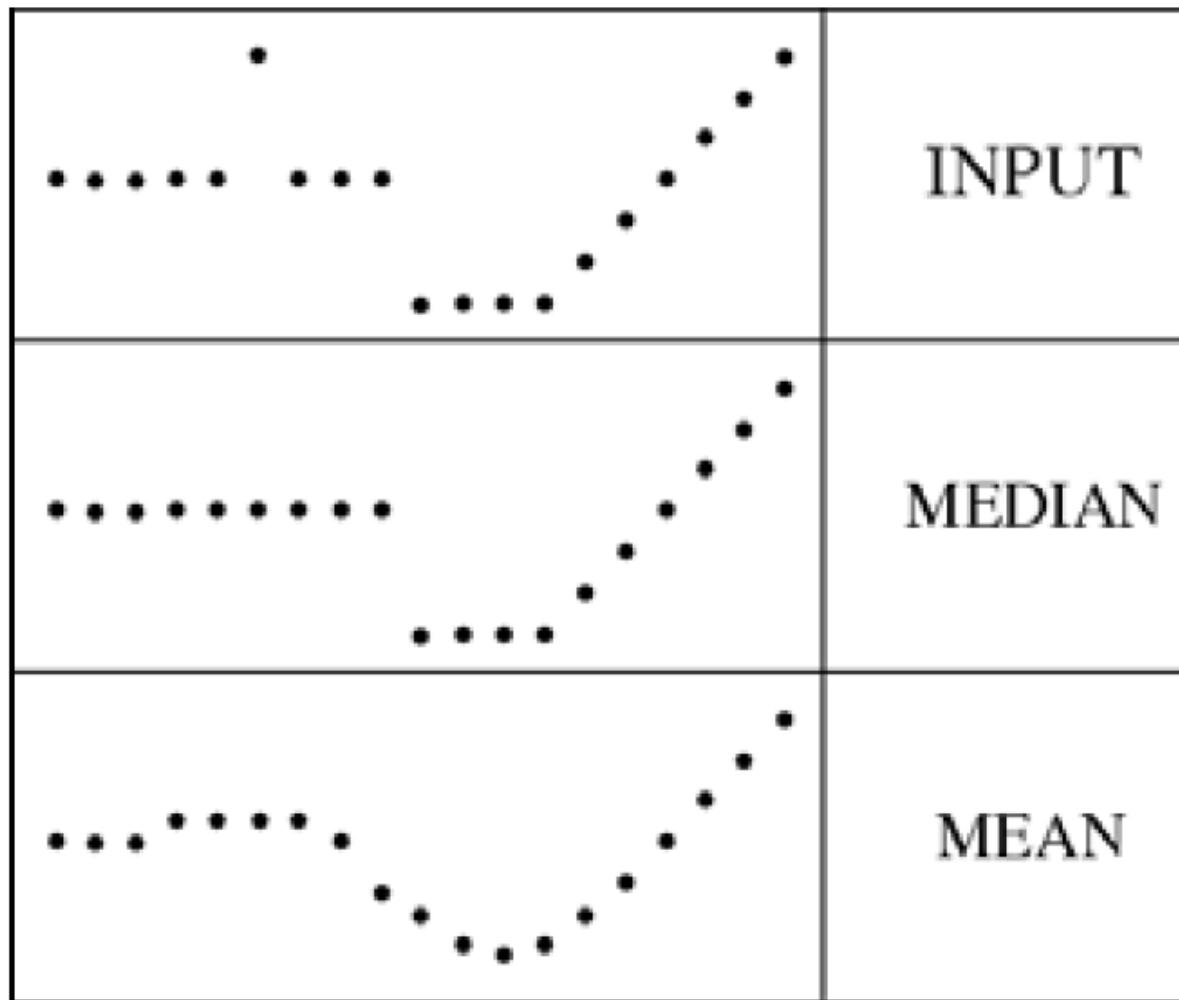
Median  
filtered



Plots of a row of the image

# Median filter

- Median filter is edge preserving
- It doesn't introduce new intensities, which is often expected.





# IMAGE PROCESSING

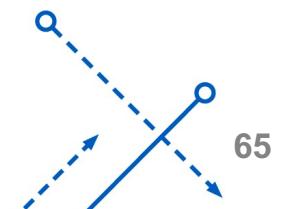
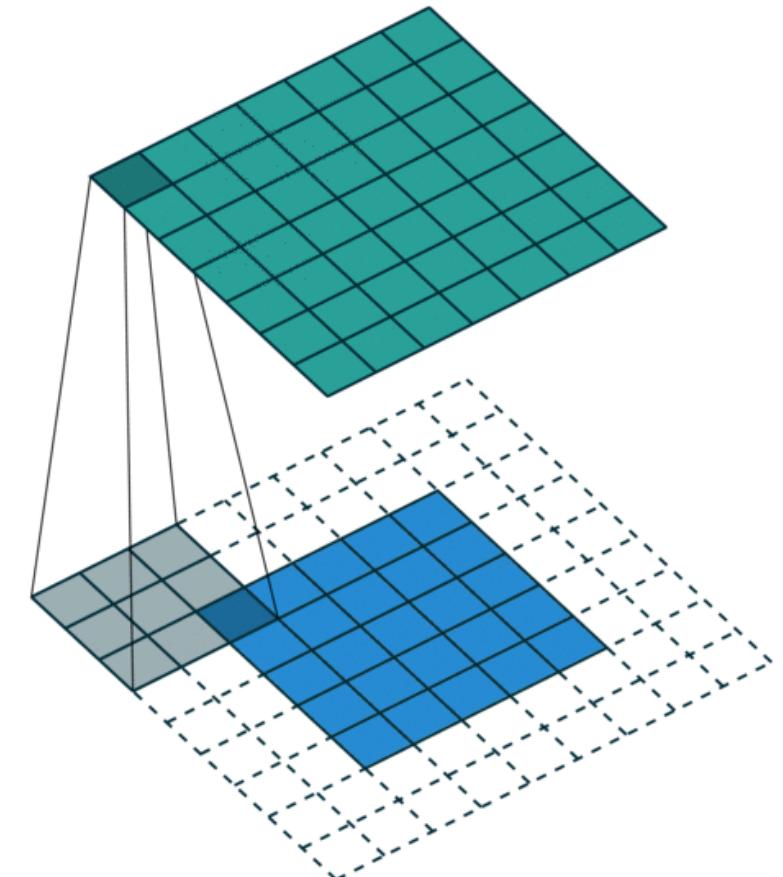
## Template Matching



# Similarity/Distance of Signals

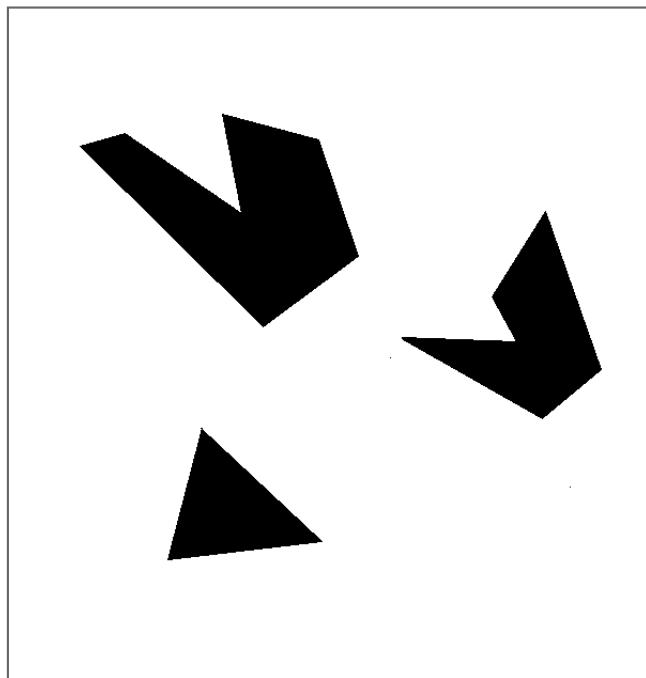
---

- L1-norm / Manhattan distance
  - $|x - w|_1$
- L2-norm / Euclidean distance
  - $\|x - w\|_2$
- Inner Product
  - $x \cdot w$
- Cosine Similarity
  - $$\frac{x \cdot w}{\|x\|_2 \|w\|_2}$$
- Filtering gives us a kind of similarity measurement, i.e., inner product.

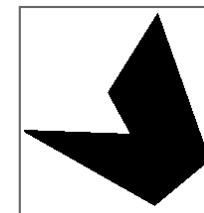


# Template matching

- Each element of the output is a similarity measure of a specific pattern, i.e., a filter or a template.
- Each similarity measure is also called a ``response''.
- This process is called template matching.

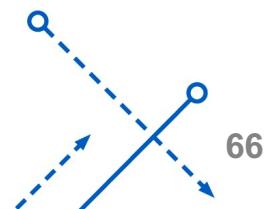


A toy example

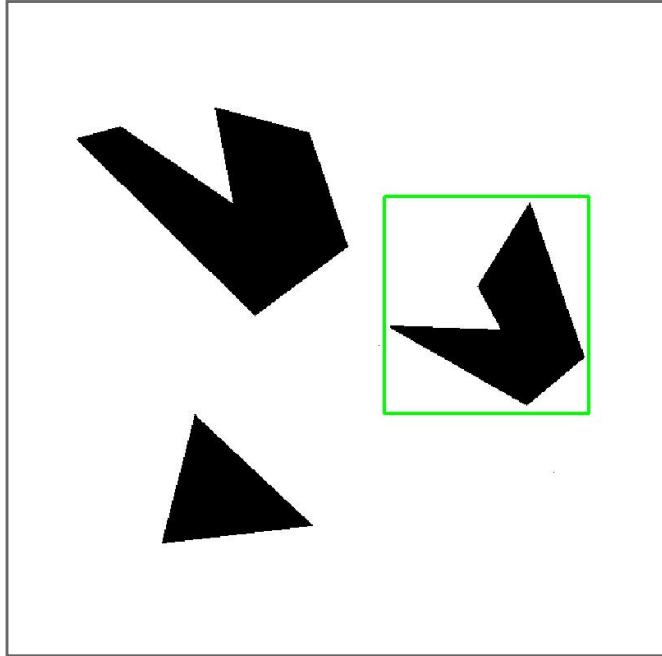


Template (mask)

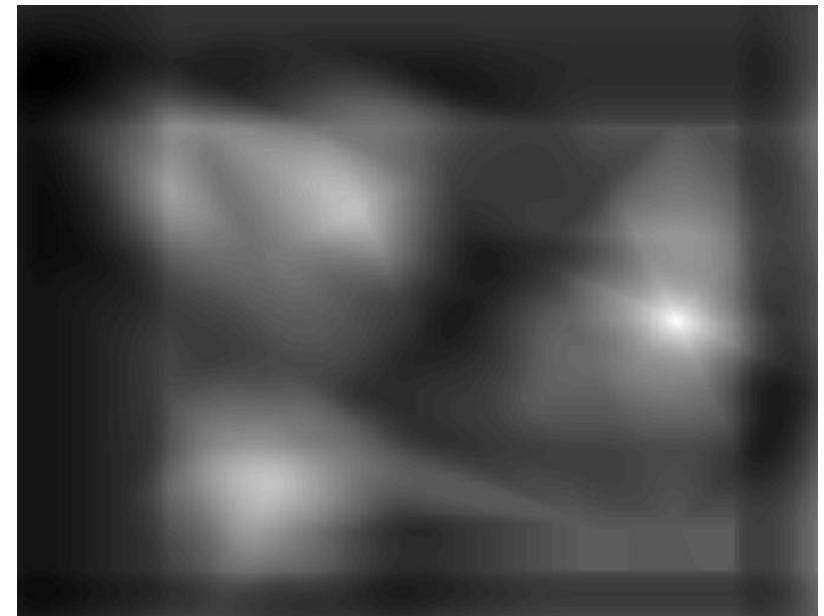
Is there only one match?  
What if the pattern is not exact?



# Correlation of Template and Image

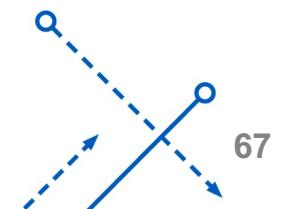


Detected template



Correlation map

Is there only one match?  
What if the pattern is not exact?



# Where's Waldo?

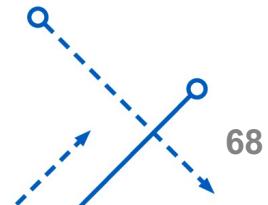
---



Scene

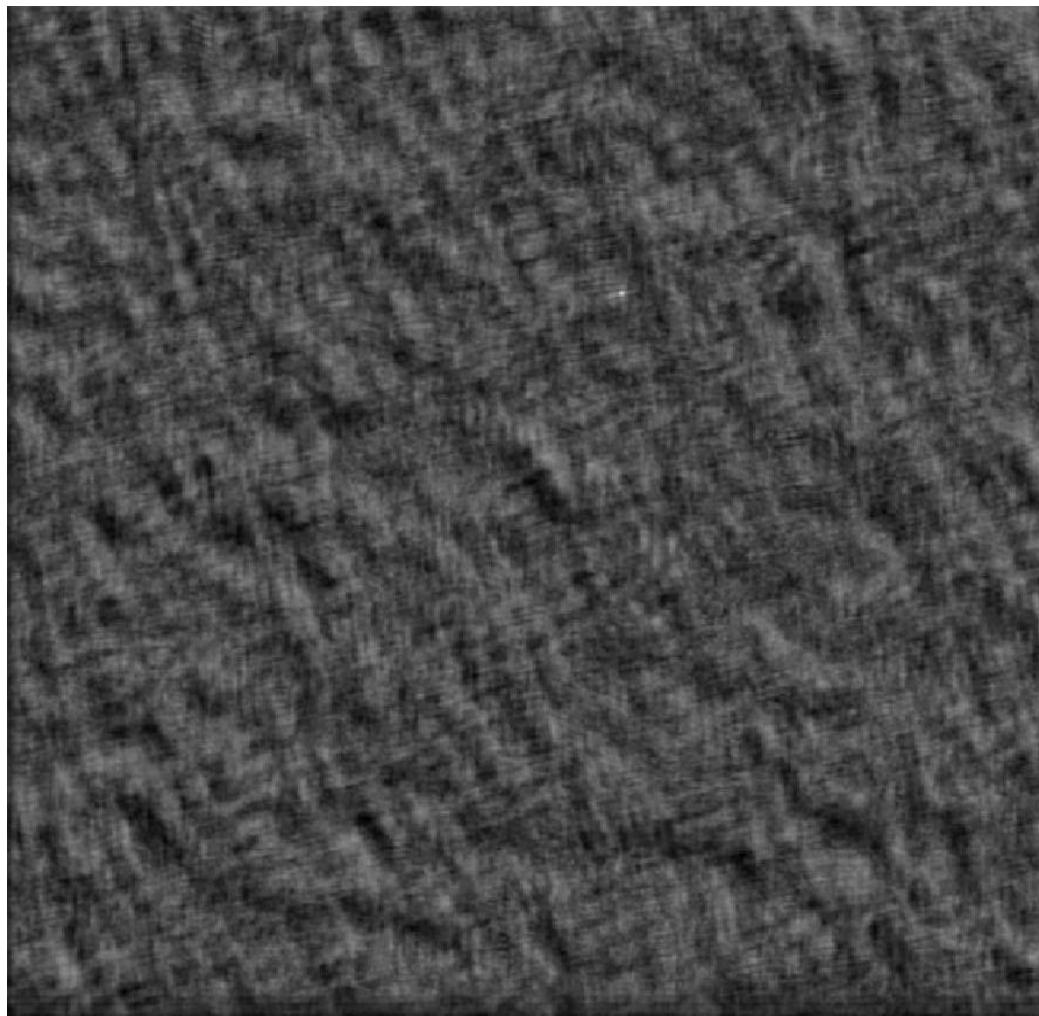


Template



# Where's Waldo?

---



Scene

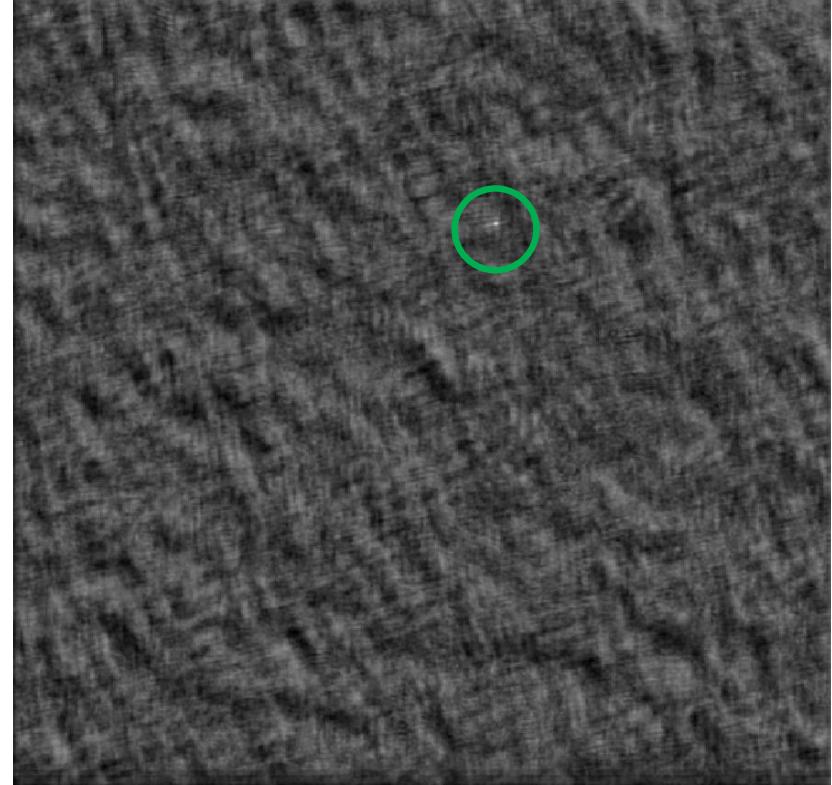


Template

# Where's Waldo?



Detected template



Correlation map

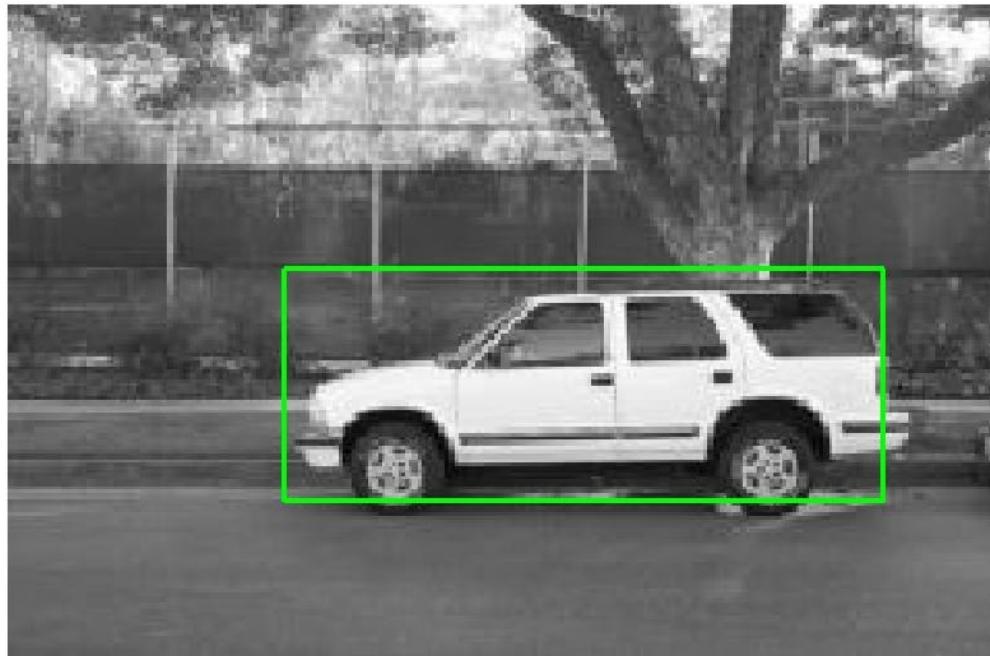
70

- Use normalized cross-correlation score to find a given pattern (template) in the image (Szeliski Eq. 8.11 in textbook).
- Normalization needed to control for relative brightness.



# Template matching

- Match can be meaningful, if **scale, orientation, and general appearance** is right.



Detected template



Template



# Template matching

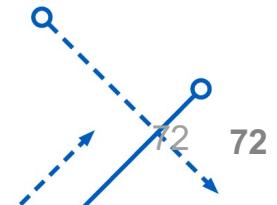
- What if the template is not identical to some sub image in the scene?



Scene



Template



# Template matching

---

- We need more flexible, powerful and forgiving representations.
  - Bolme, D. S., Beveridge, J. R., Draper, B. A., & Lui, Y. M. [Visual object tracking using adaptive correlation filters](#). CVPR, 2010.
    - Computational complexity:  $\mathcal{O}(N \log N)$
    - Equivariant to translation, robust to small appearance variance.
  - Wang, C., Zhang, L., Xie, L., & Yuan, J. (2018, April). [Kernel cross-correlator](#). AAAI, 2018.
    - Nonlinear cross-correlation with the kernel trick.
    - Equivariant to any transforms:
      - Translation, Scale, Rotation, Affine, etc.
    - Same computational complexity with linear filter:  $\mathcal{O}(N \log N)$





# IMAGE PROCESSING

## Edge Detection



# Filters for features

- Previously, filtering is a way to
  - Remove or reduce **noise**.
  - **Template matching**
- Filters also allows us to abstract higher-level “**features**”.
  - Map raw pixels to intermediate representations used for **subsequent processing**.
  - Reduce amount of data, discard redundancy, preserve useful information.



# Edge detection

- **Goal:** map image from 2D array of pixels to a set of **curves** or **line segments**, or **contours**.
- **Why?**

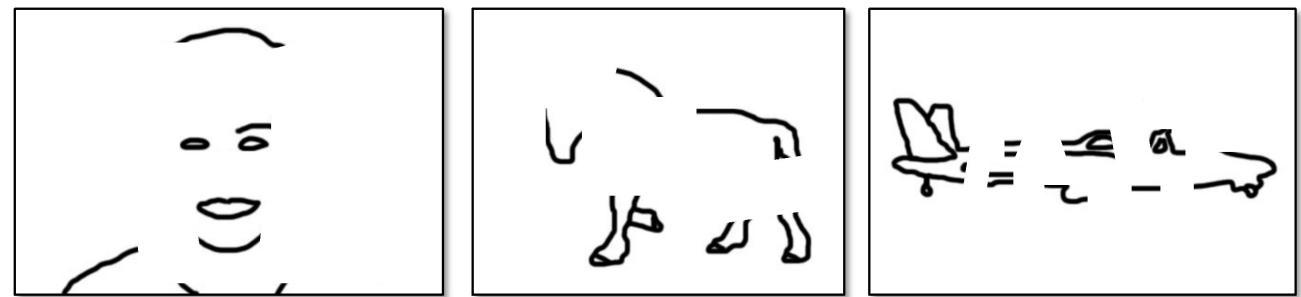


Figure from J. Shotton et al., PAMI 2007

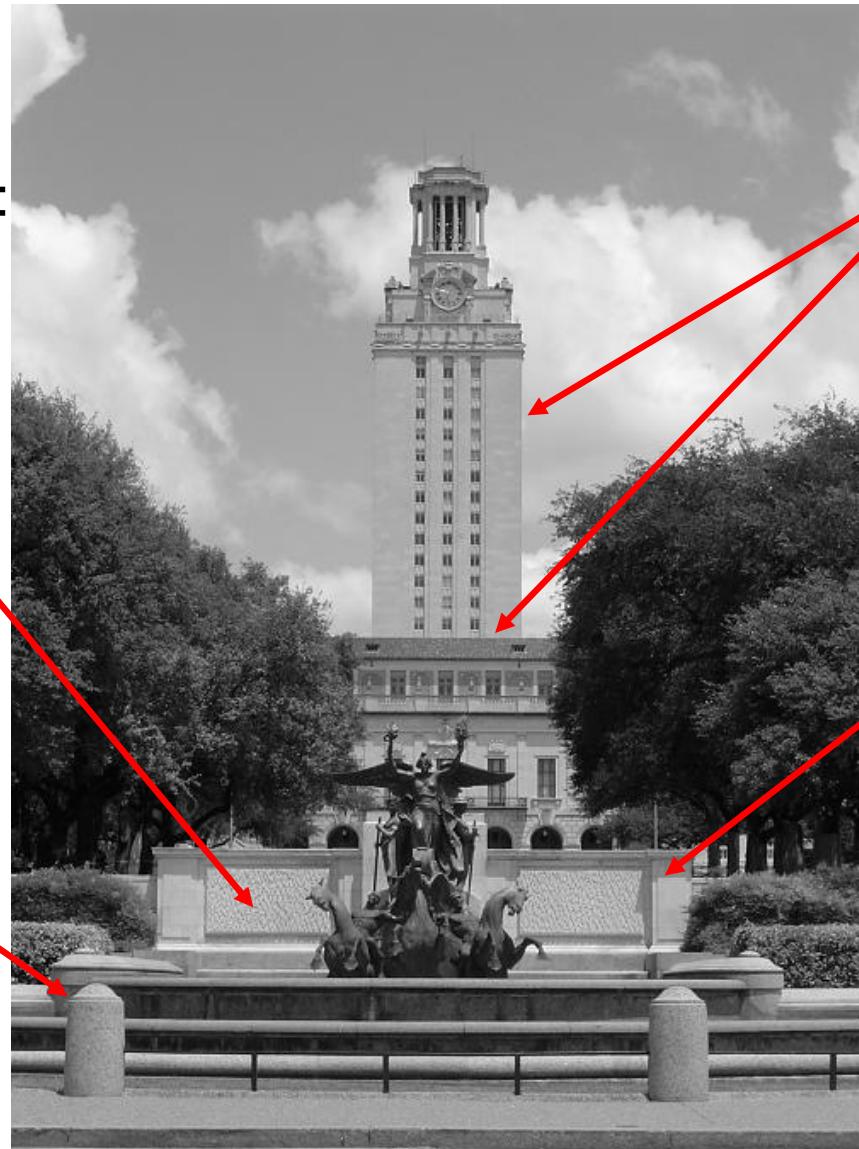
- **Even simple sketch of the objects are quite meaningful.**
- **Main idea:** look for strong gradients, post-process.



# What can cause an edge?

Reflectance change:  
appearance  
information, texture

Change in surface  
orientation: shape



Depth discontinuity:  
object boundary

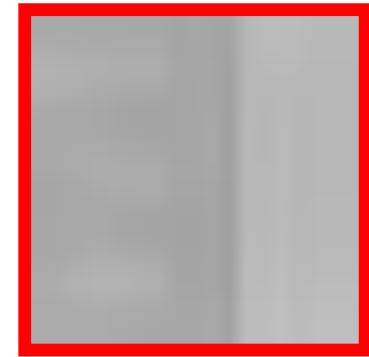
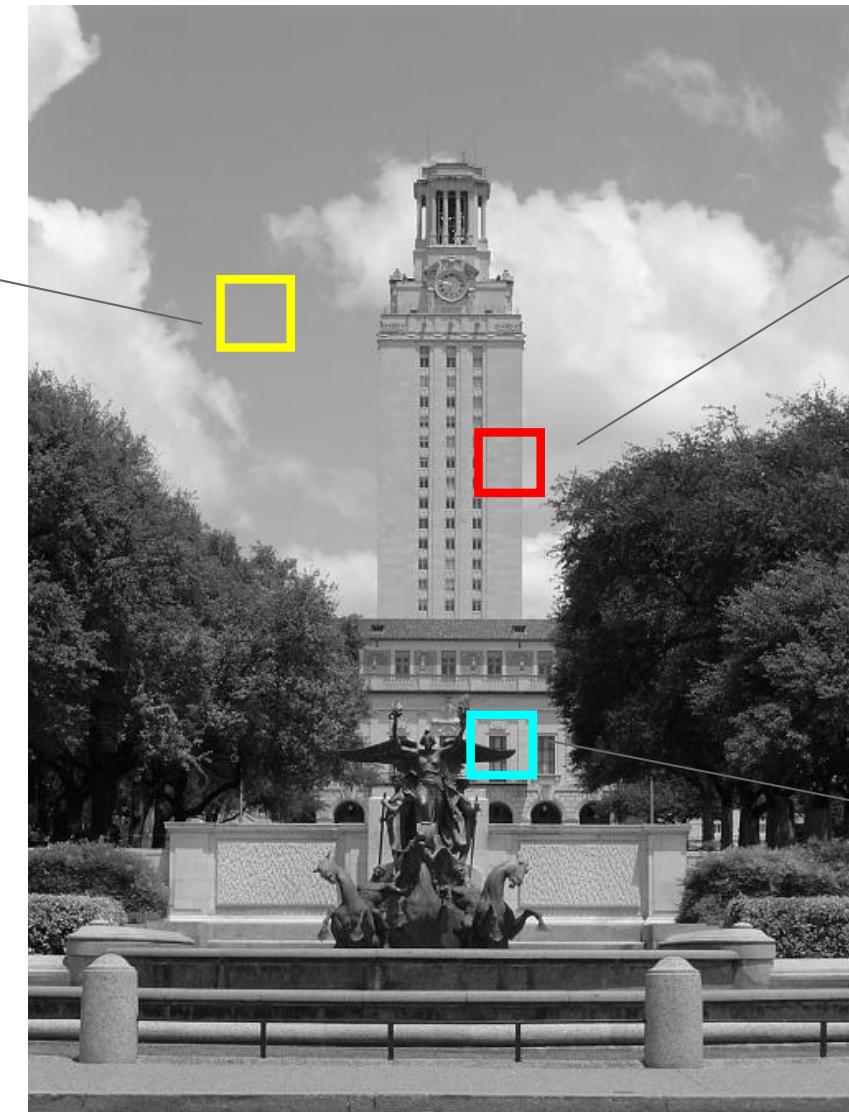
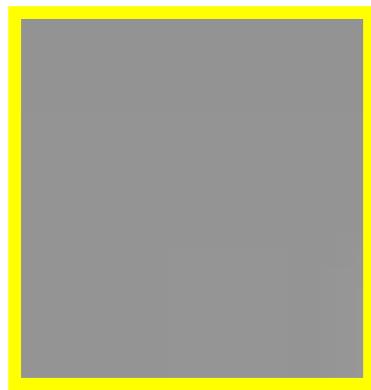
Cast shadows

77

77

# Contrast and invariance

---

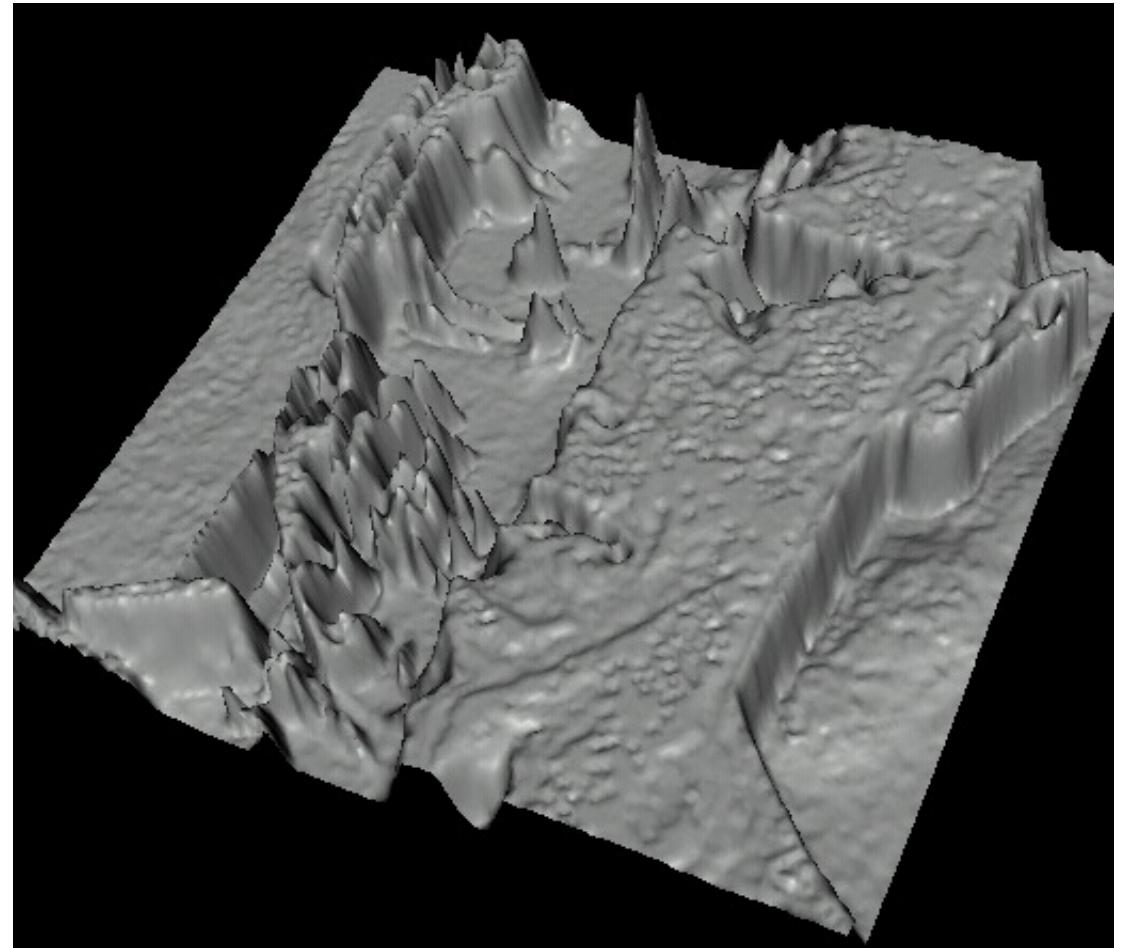


78  
78

A dashed blue line with arrows at both ends, connecting the two '78' labels, likely indicating a comparison or transformation between the two cyan-highlighted patches.

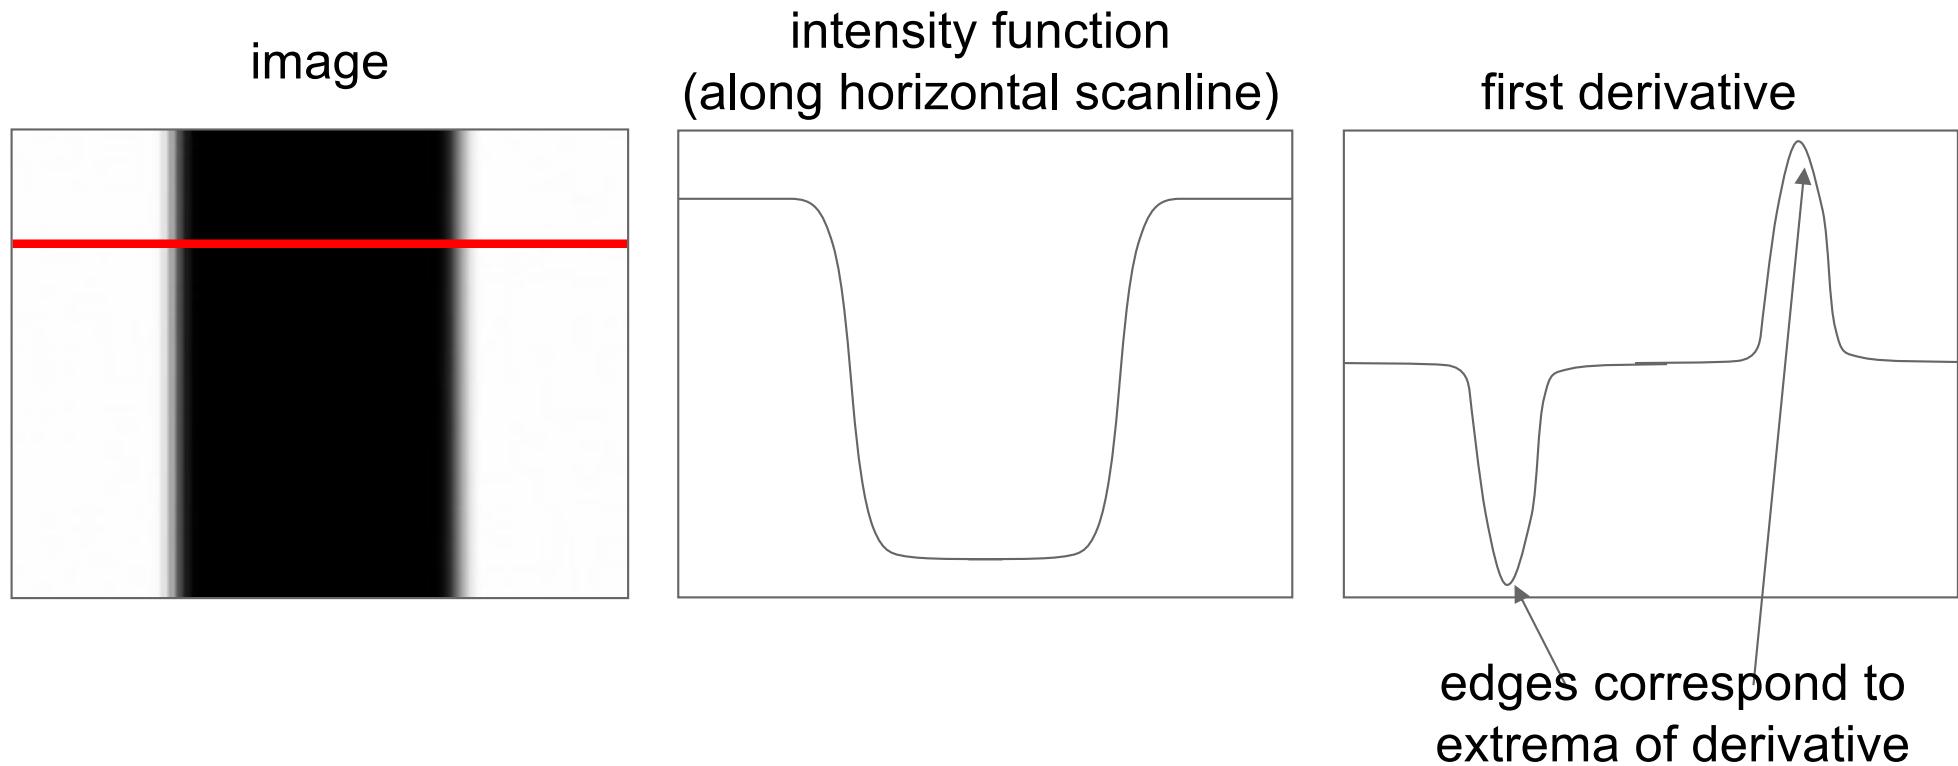
# Edges look like steep cliffs

---



# Derivatives and edges

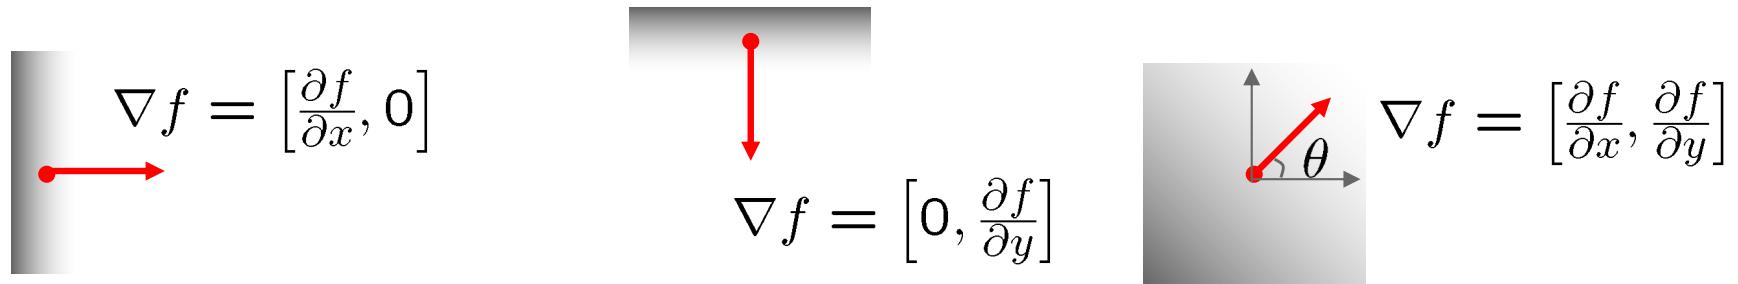
An edge is a place of rapid change in the image intensity function.



# Image gradient

The gradient points in the direction of most rapid change in intensity

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



The gradient direction (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The edge *strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Differentiation and convolution

---

For 2D function,  $f(x, y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement above as correlation, what would be the filter?



# Partial derivatives of an image

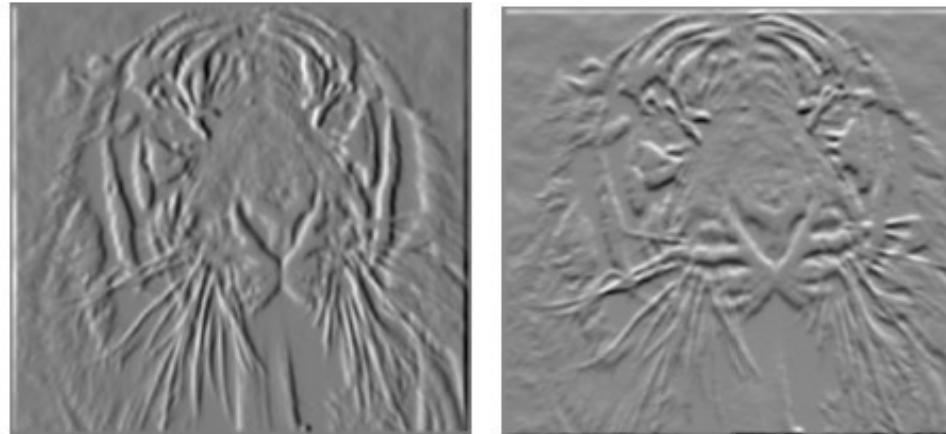
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



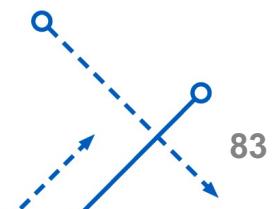
$$\frac{\partial f(x, y)}{\partial y}$$

-1	?	1
1	or	-1



Which shows changes with respect to x?

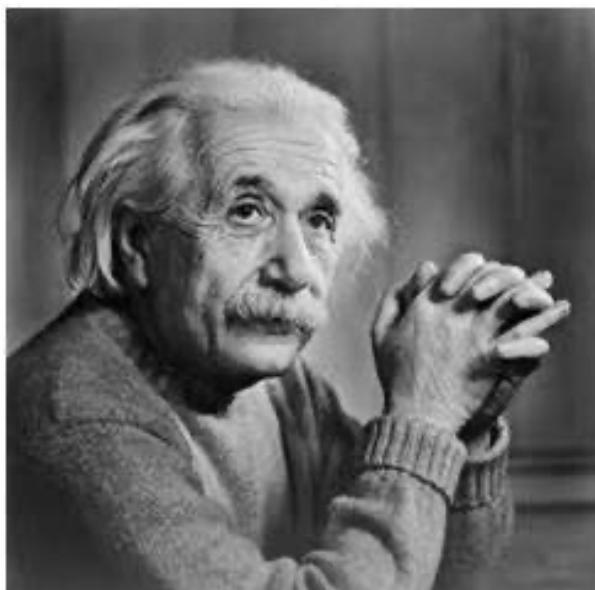
(showing flipped filters)



# Prewitt operator

$$G_x = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$
$$G_y = \begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Fig. 1. The horizontal and vertical Prewitt edge detection masks.



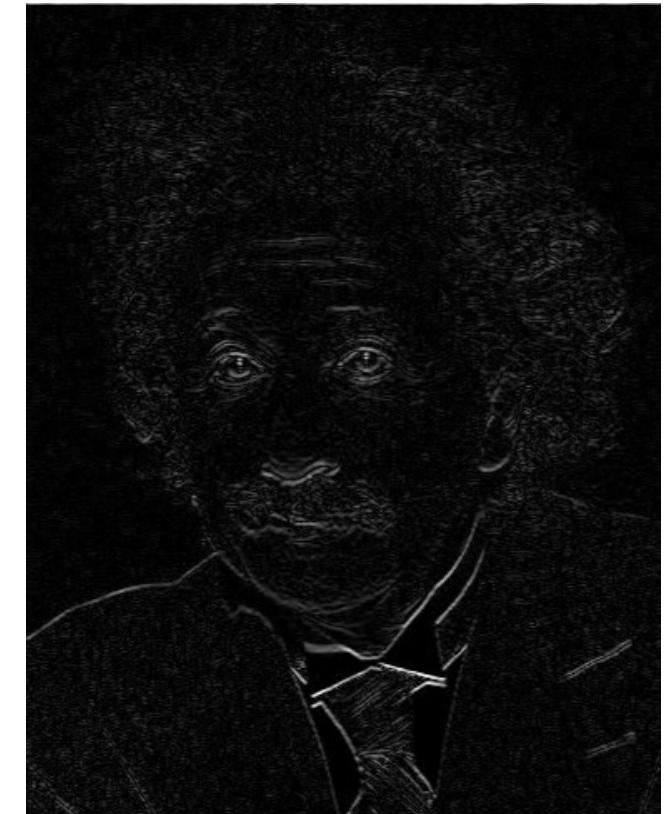
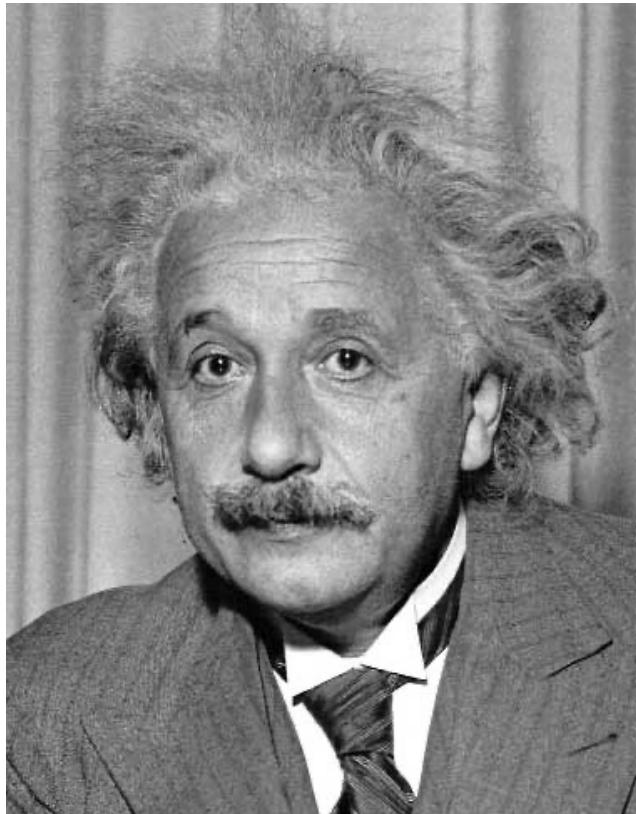
# Sobel Operator

$G_x =$

1	0	-1
2	0	-2
1	0	-1

$G_y =$

1	2	1
0	0	0
-1	-2	-1



85

# Roberts Operator

+1	0
0	-1

Gx

0	+1
-1	0

Gy

$$|G| = \sqrt{G_x^2 + G_y^2}$$

$$|G| = |G_x| + |G_y|$$

$$|G| = |P_1 - P_4| + |P_2 - P_3|$$

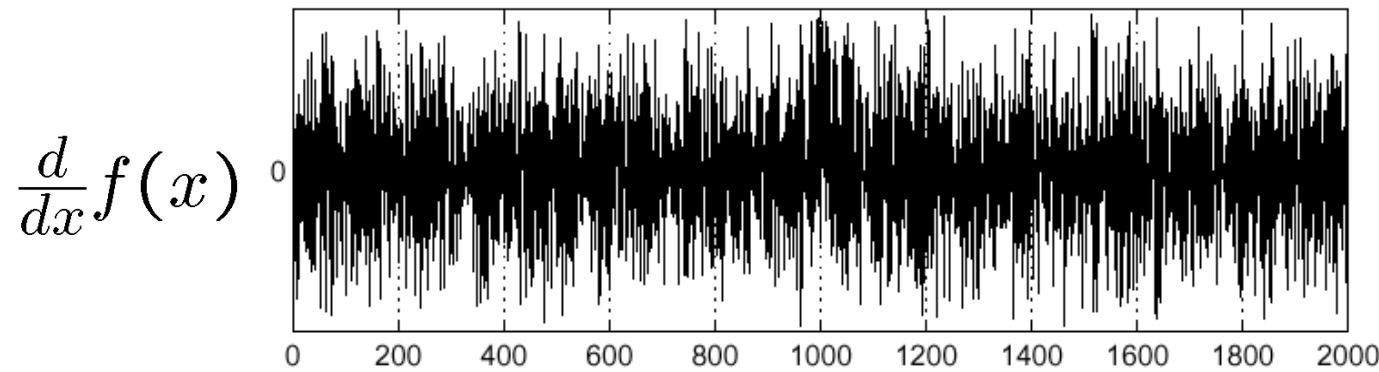
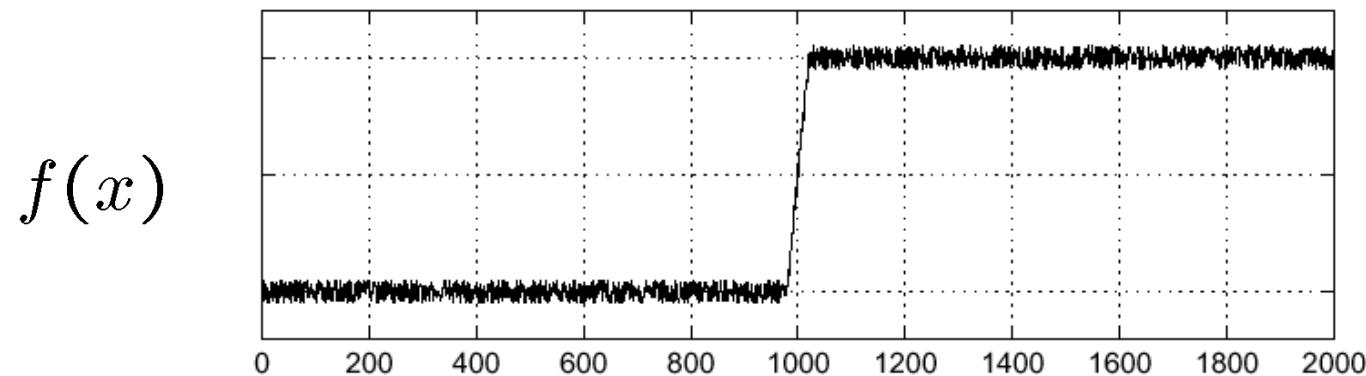
P <sub>1</sub>	P <sub>2</sub>
P <sub>3</sub>	P <sub>4</sub>



# Effects of noise

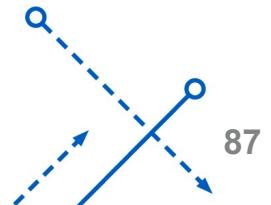
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



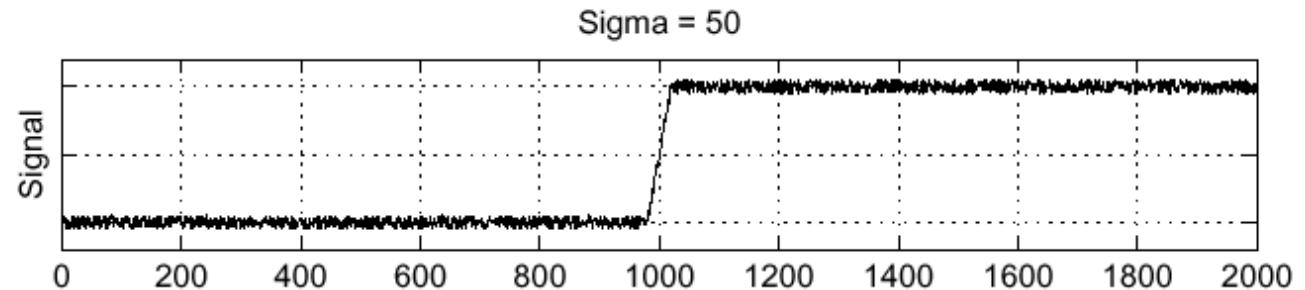
87

Where is the edge?

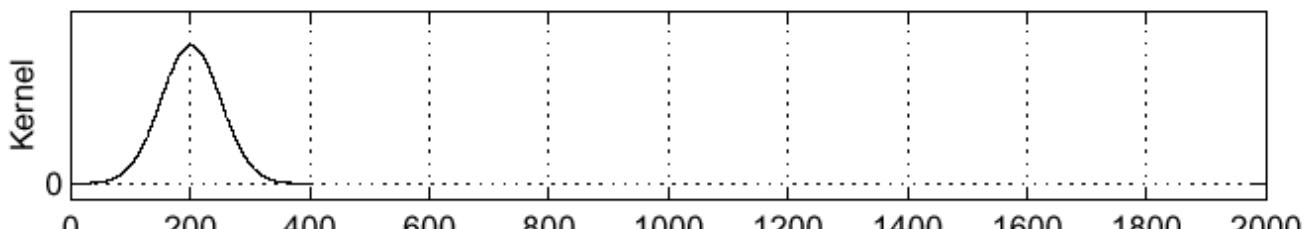


# Solution: smooth first

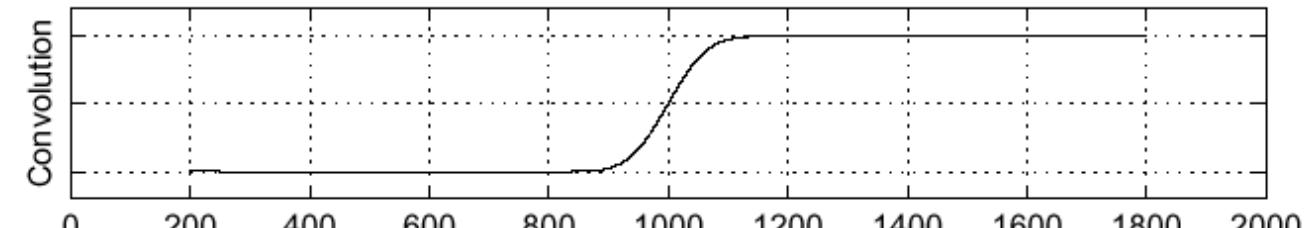
$f$



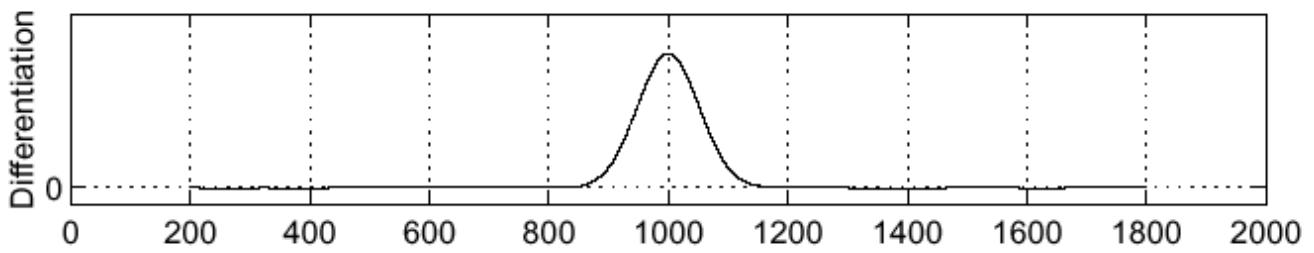
$h$



$h \star f$



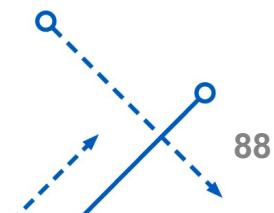
$\frac{\partial}{\partial x}(h \star f)$



Where is the edge?

Look for peaks in

$\frac{\partial}{\partial x}(h \star f)$

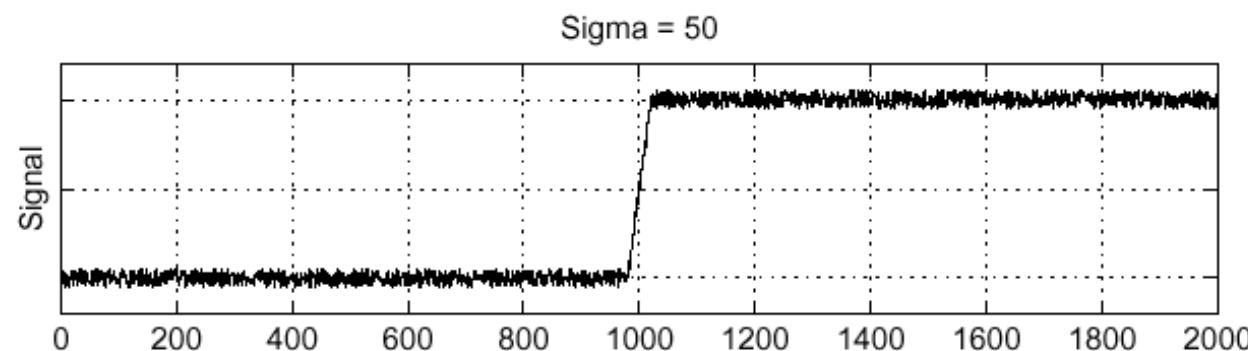


# Derivative theorem of convolution

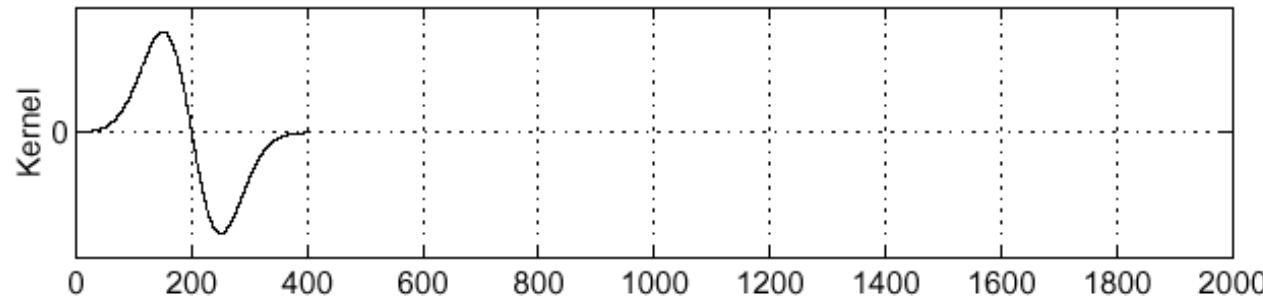
Differentiation property of convolution.

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

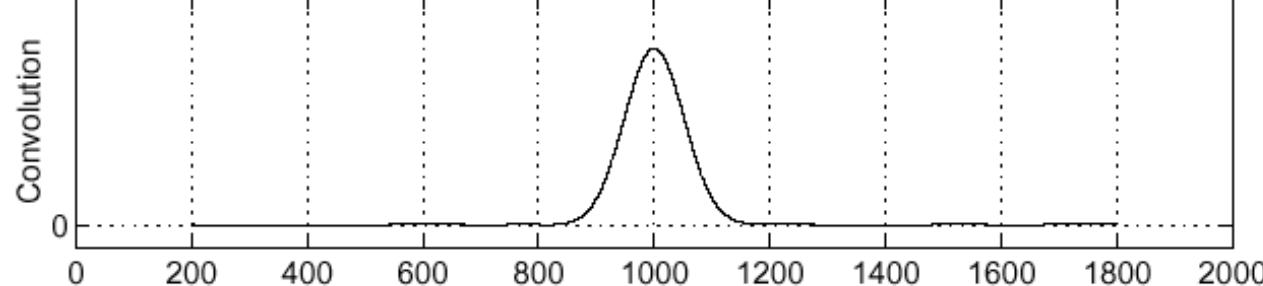
$f$



$\frac{\partial}{\partial x}h$



$(\frac{\partial}{\partial x}h) \star f$

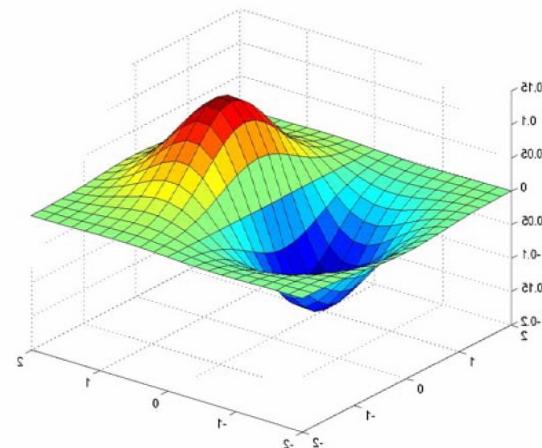


# Derivative of Gaussian filter

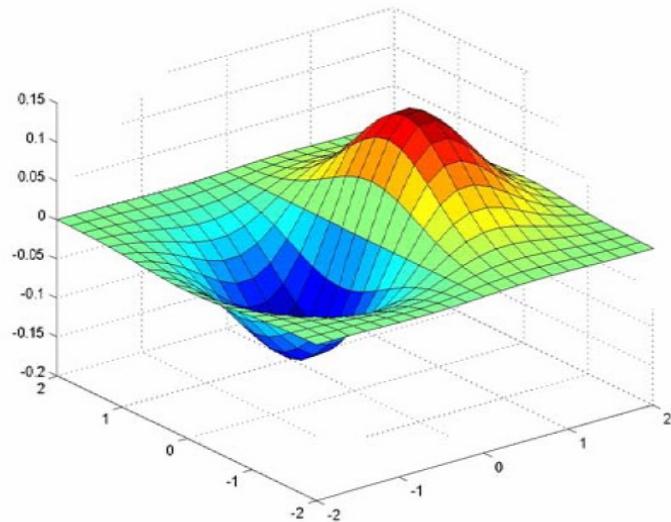
$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

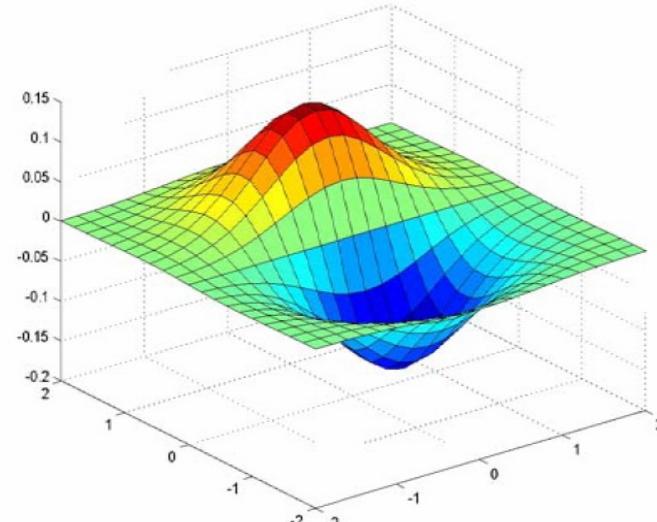
$$\otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$



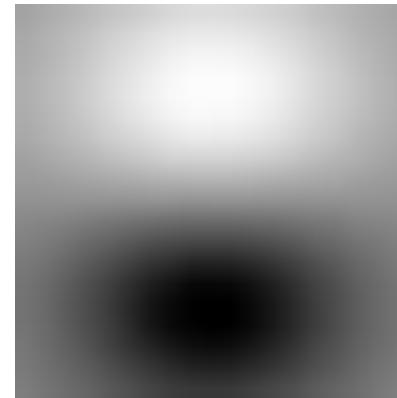
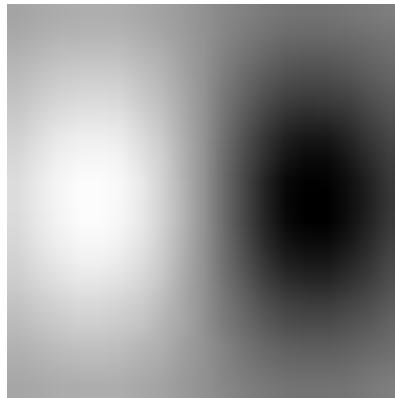
# Derivative of Gaussian filters



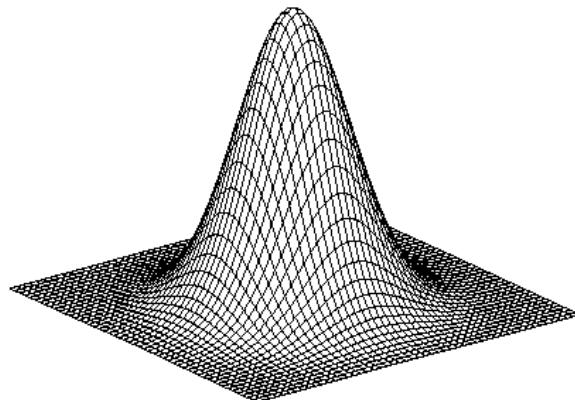
x-direction



y-direction

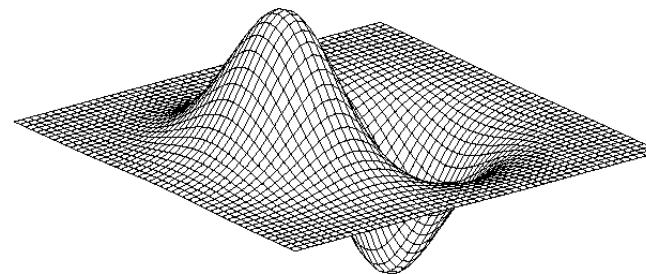


# 2D edge detection filters



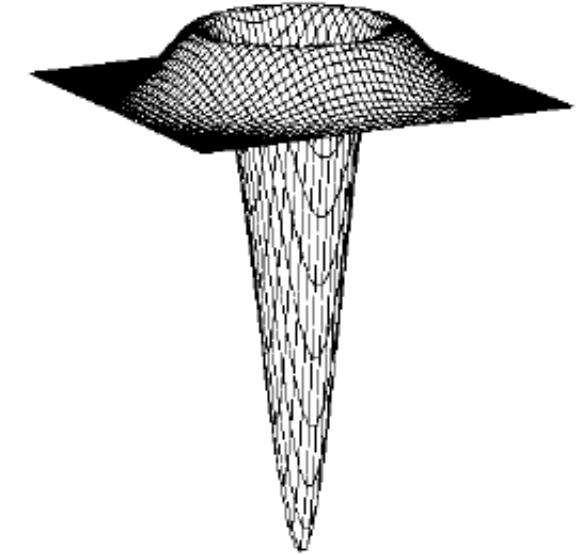
Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$



Laplacian of Gaussian

$$\nabla^2 h_\sigma(u, v)$$

- $\nabla^2$  is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$



# Laplacian of Gaussian

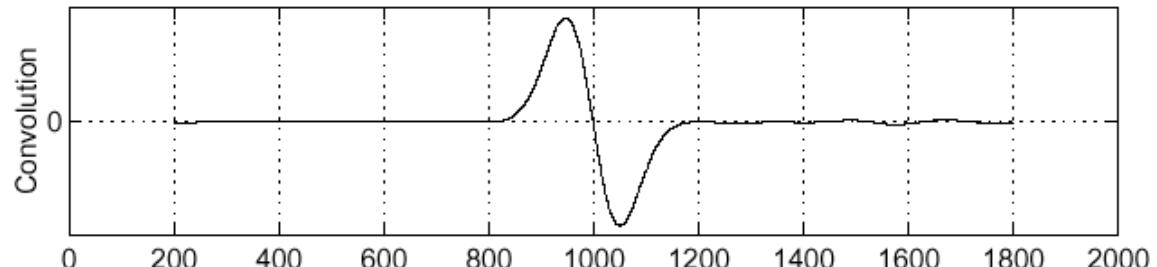
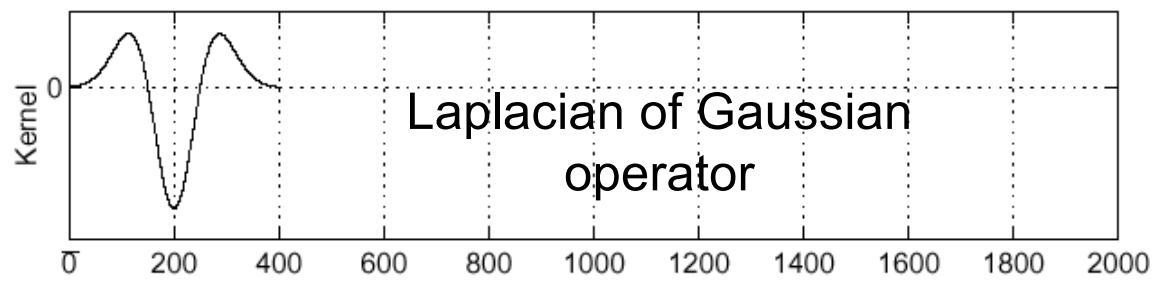
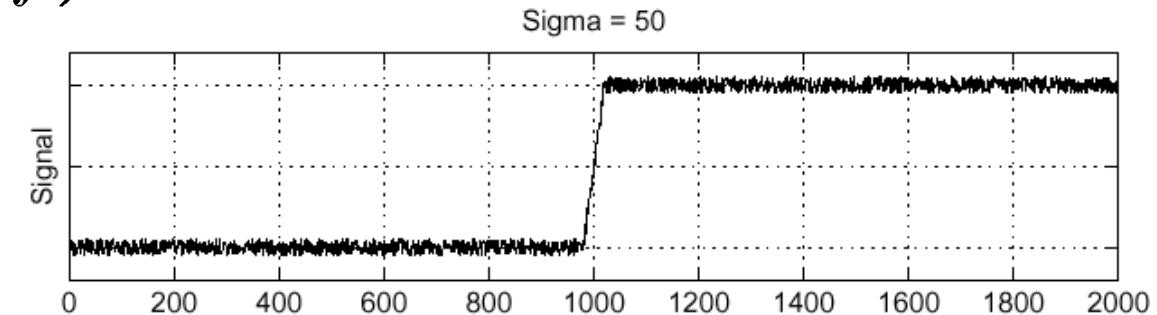
Consider  $\frac{\partial^2}{\partial x^2}(h \star f)$

$f$

$\frac{\partial^2}{\partial x^2}h$

$(\frac{\partial^2}{\partial x^2}h) \star f$

Where is the edge?



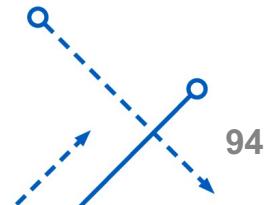
Zero-crossings of bottom graph



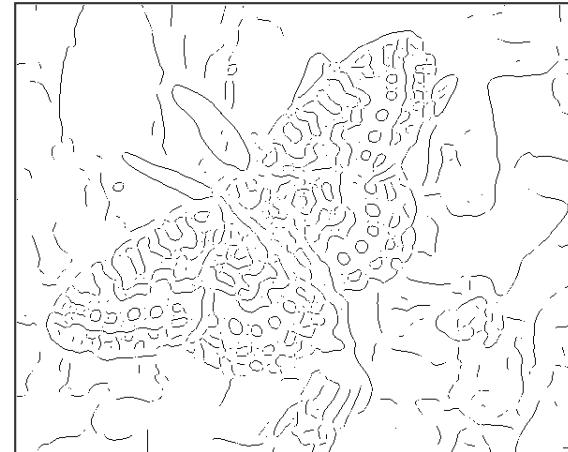
# Mask/Filter/Kernel Properties

---

- Smoothing
  - Values positive
  - **Sum to 1:** constant regions → same as input (no change)
  - Amount of smoothing proportional to mask size
  - Remove “high-frequency” components; “low-pass” filter
- Template Matching
  - Dot product as correlation (inner product).
  - Highest response for regions “look the most like the filter”
- Derivatives
  - Opposite signs get high response in high contrast regions
  - **Sum to 0:** constant regions → no response
  - High contrast → high absolute values



# Gradients -> edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: Filter for contrast
3. Edge localization

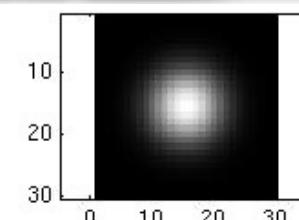
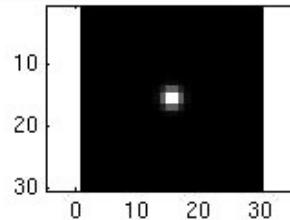
Determine which local maxima from filter output are actually edges vs. noise

- Threshold, Thin

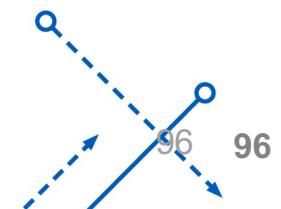
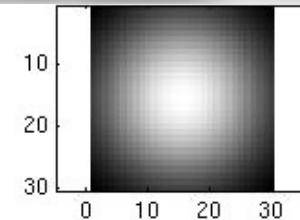


# Smoothing with a Gaussian

Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



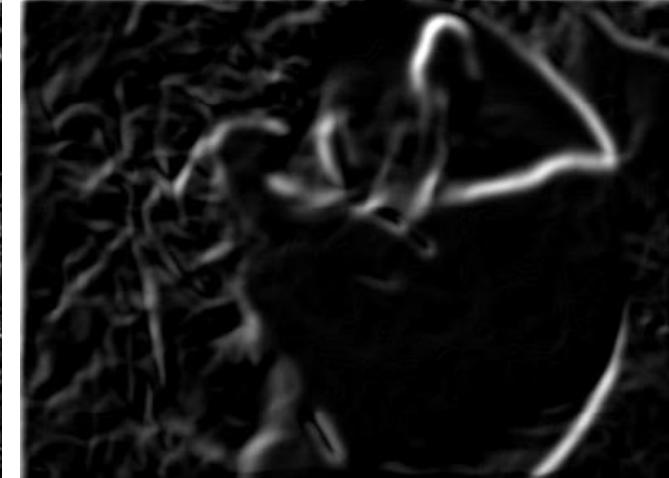
...



# Effect of $\sigma$ on derivatives



$\sigma = 1$  pixel



$\sigma = 3$  pixels

The apparent structures differ depending on Gaussian's scale/width parameter.

Larger values: larger scale edges detected  
Smaller values: finer features detected



# So, what scale to choose?

It depends what we're looking for.



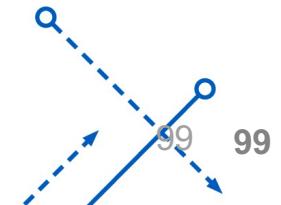
Too small of a scale...can't see the forest for the trees.

Too big of a scale...can't tell the maple grain from the cherry.

# Thresholding

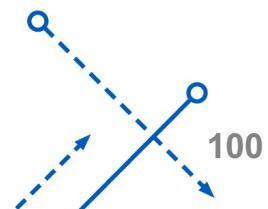
---

- Choose a threshold value  $t$ .
- Set any pixels less than  $t$  to zero (off)
- Set any pixels greater than or equal to  $t$  to one (on)



# Gradient magnitude image

---



# Thresholding gradient

lower threshold



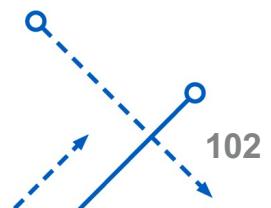
higher threshold



# Canny edge detector

---

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Local non-maximum suppression:**
  - Thin multi-pixel wide “ridges” down to single pixel width
- Linking and thresholding (**hysteresis**):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them.

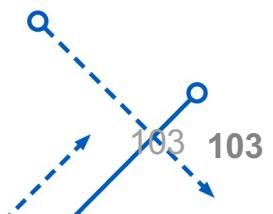


Source: D. Lowe, L. Fei-Fei

# The Canny edge detector

---

original image (Lena)

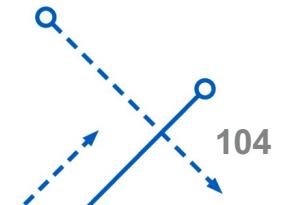


# The Canny edge detector

---



norm of the gradient

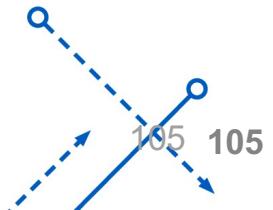


# The Canny edge detector

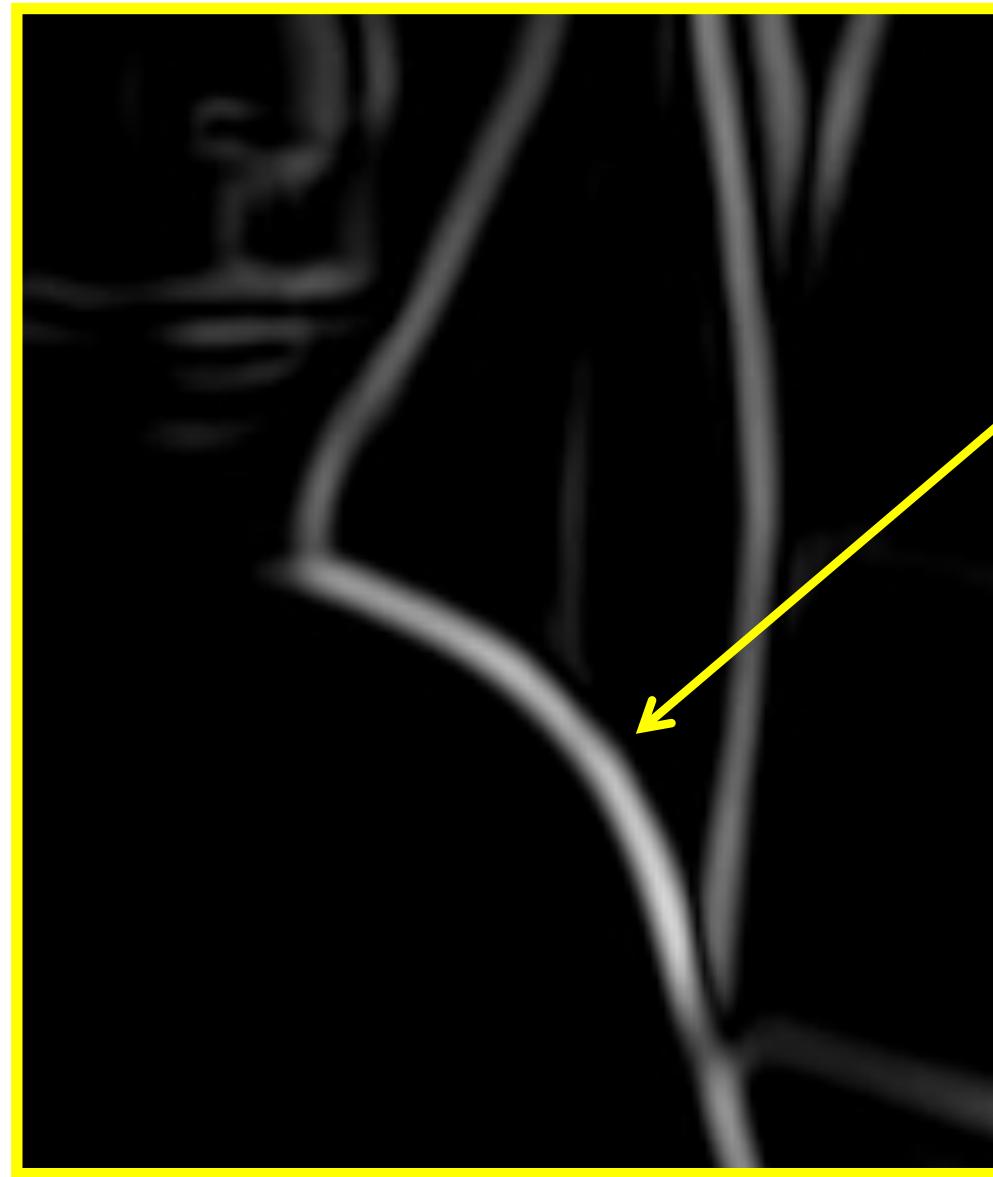
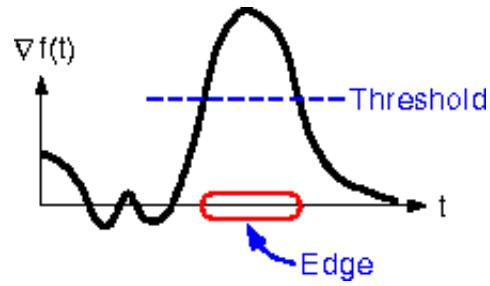
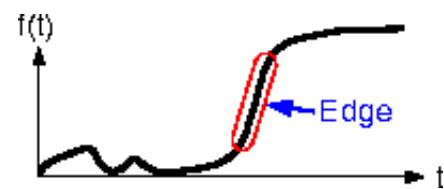
---



thresholding

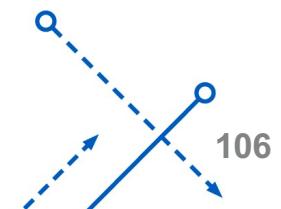


# The Canny edge detector



How to turn  
these thick  
regions of the  
gradient into  
curves?

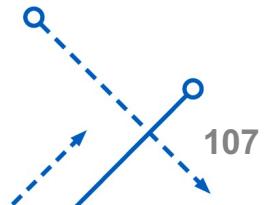
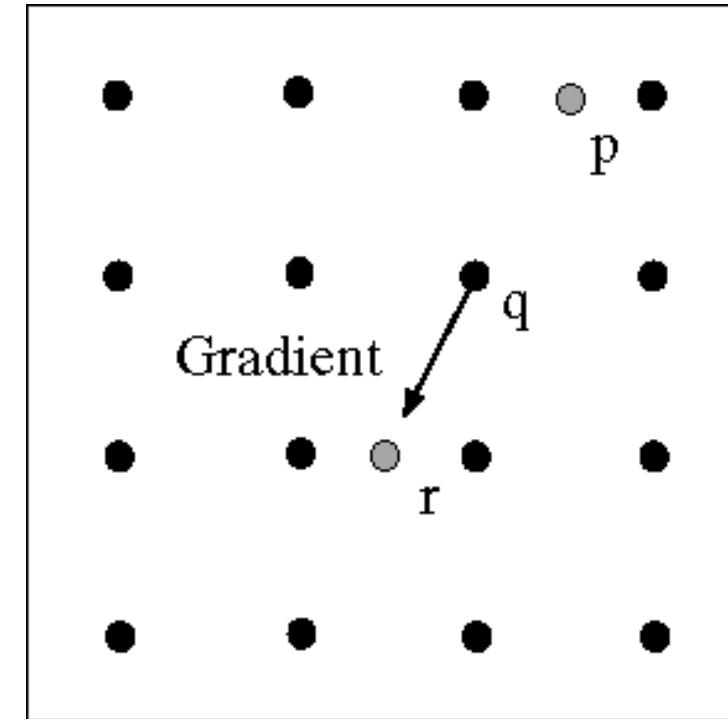
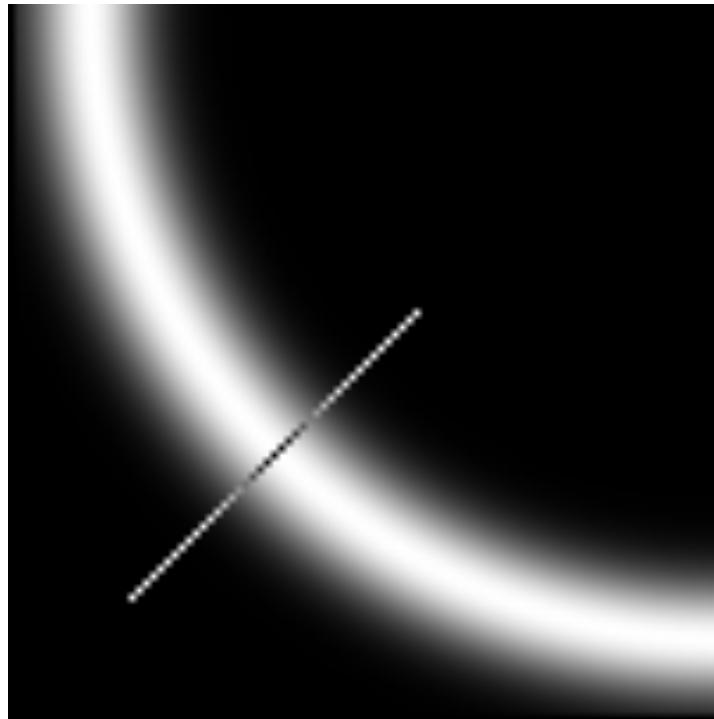
thresholding



# Non-maximum suppression

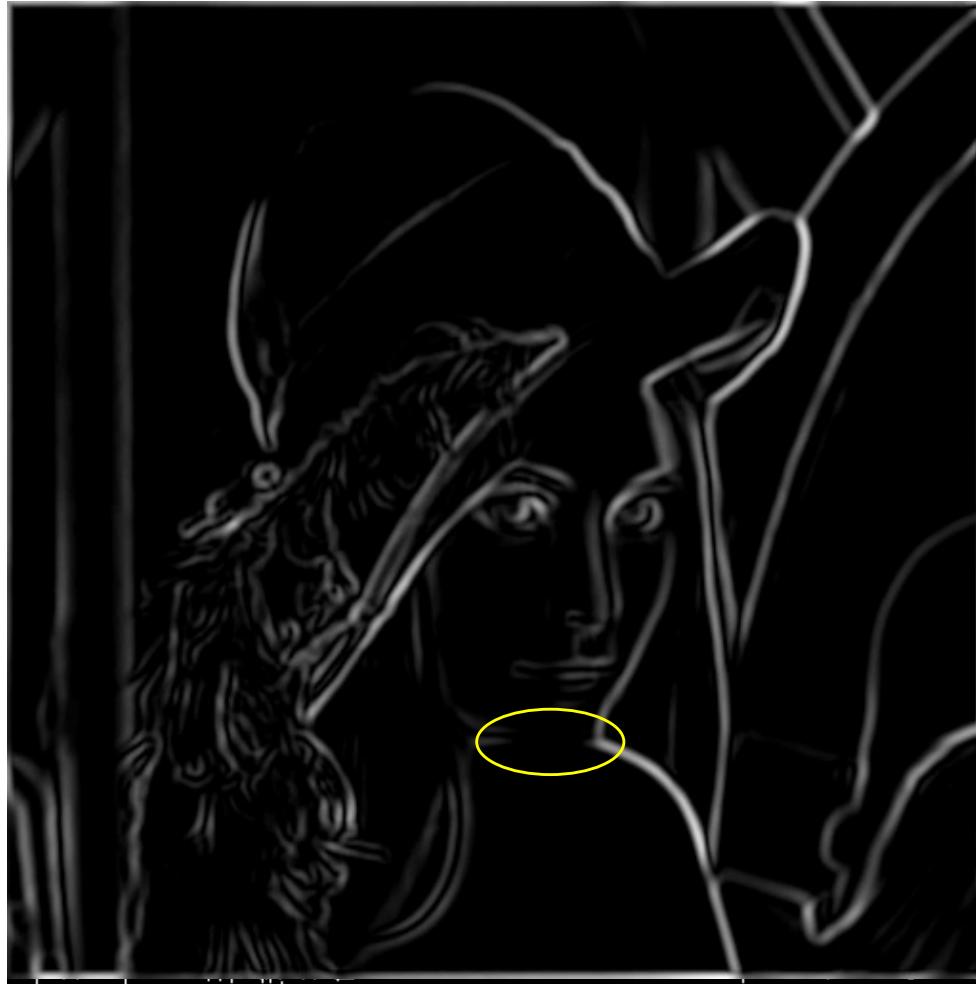
Check if pixel is local maximum along gradient direction, select single max across width of the edge

- requires checking interpolated pixels p and r



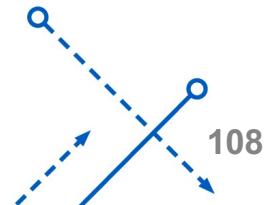
# The Canny edge detector

---



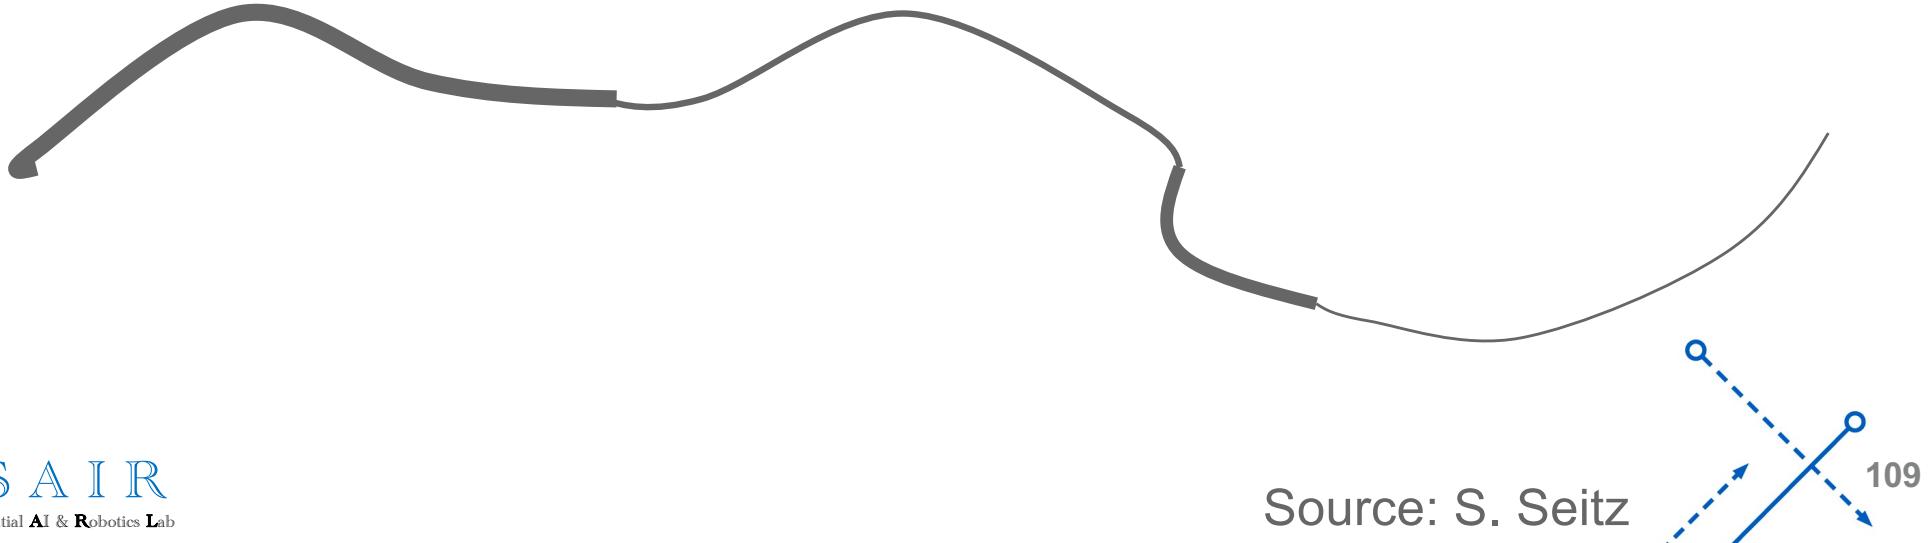
thinning  
(non-maximum suppression)

Problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding



# Hysteresis thresholding

- Check that maximum value of gradient value is sufficiently large
  - drop-outs? use **hysteresis**
    - use a high threshold to start edge curves and a low threshold to continue them.



Source: S. Seitz

# Hysteresis thresholding



original image



high threshold  
(strong edges)



low threshold  
(weak edges)

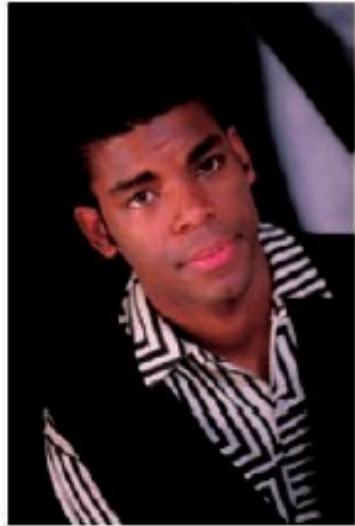


hysteresis threshold

Source: L. Fei-Fei



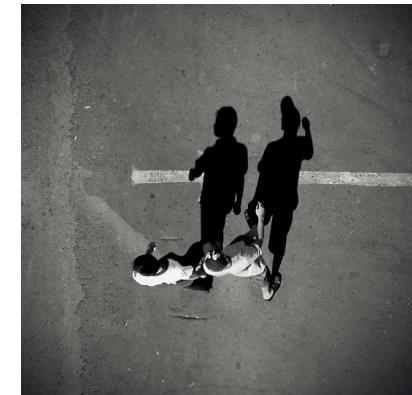
# Object boundaries vs. edges



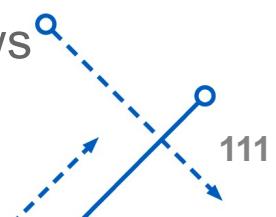
Background



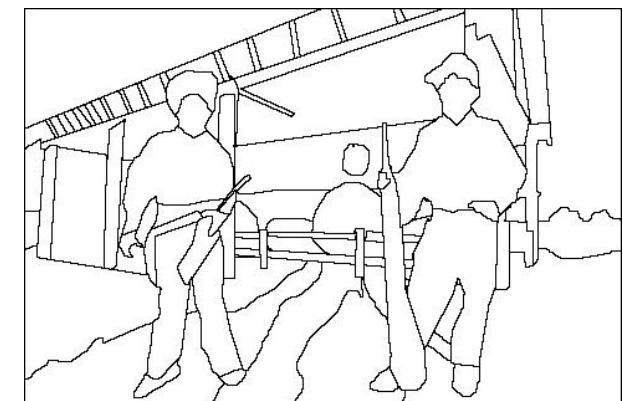
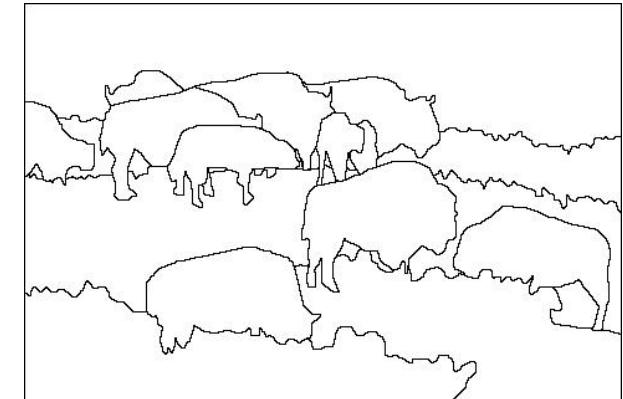
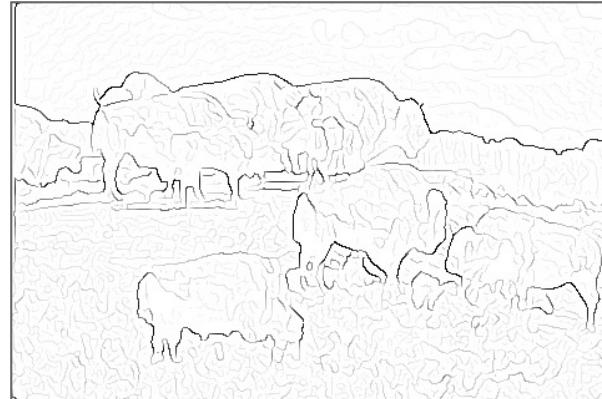
Texture



Shadows



# Edge detection is just the beginning...



image

gradient magnitude

human segmentation

# Important Concepts

---

- Filtering
  - Linear filters, Correlation, Equivariance, Invariance
  - Gaussian Filter, Smoothing, Median filter
- Template Matching
  - Cross-correlation
- Edge Detection
  - Image differentiation and gradient
  - Derivative theorem of convolution
  - 2D edge detection filters, Sobel operator
  - Canny edge detector, Hysteresis thresholding

