



SAIR

Spatial AI & Robotics Lab

CSE 473/573-A

W5: HOUGH TRANSFORM & ALIGNMENT

Chen Wang

Spatial AI & Robotics Lab

Department of Computer Science and Engineering

UB University at Buffalo The State University of New York



Content

- Hough Transform
 - Line Parameterization,
 - Line Detection,
- Alignment
 - Homography
 - Interpolation
 - Stitching, Panorama

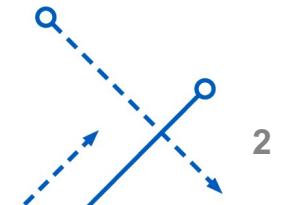




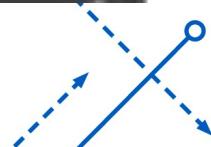
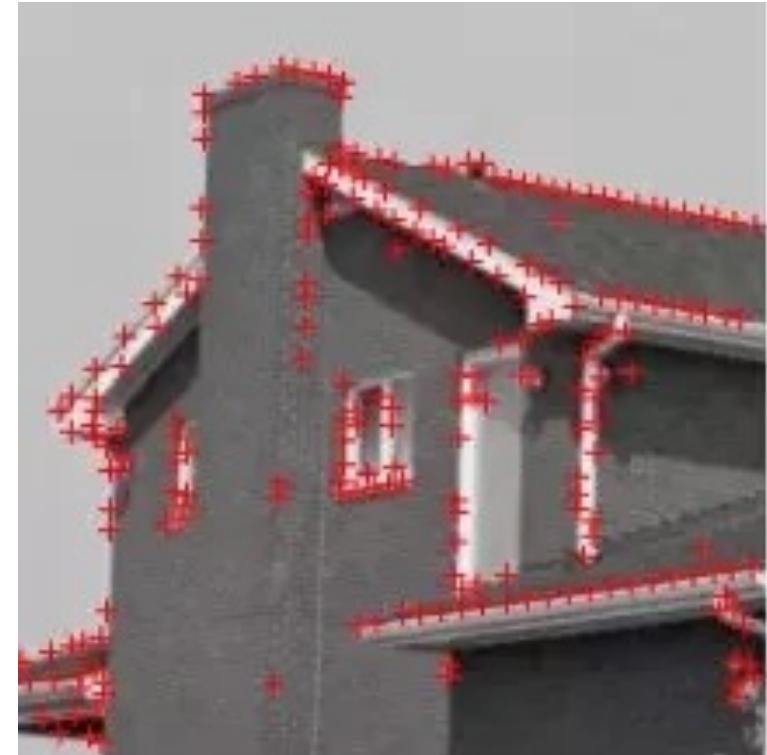
IMAGE PROCESSING

Hough Transform



Hough Transform

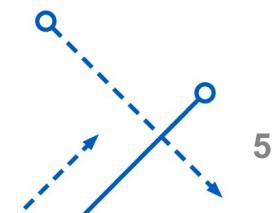
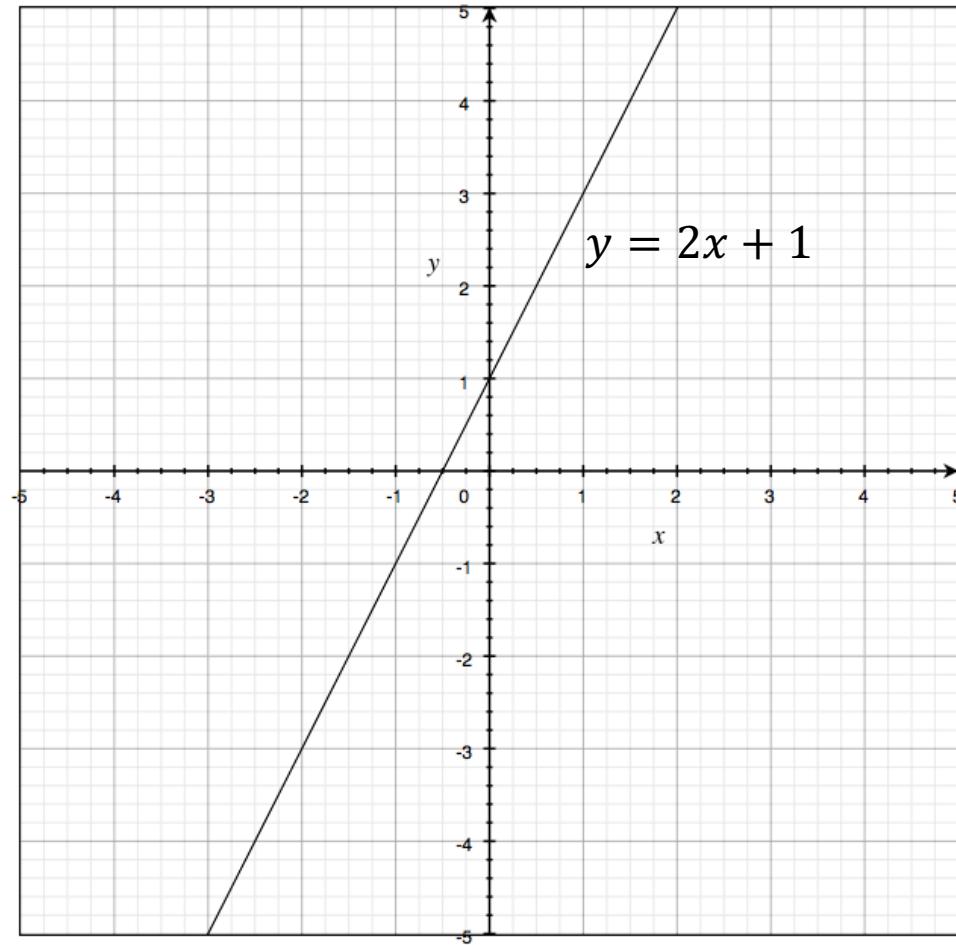
- Hough Transform can detect basic shapes
 - Detect points/edges → Find shapes.
 - Lines, Circles, etc.
- Line parameterizations
 - Slope intercept form
 - Double intercept form
 - Normal Form



Slope intercept form

$$y = mx + b$$

slope y-intercept



Double intercept form

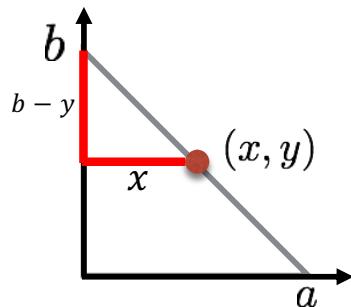
$$\frac{x}{a} + \frac{y}{b} = 1$$

x-intercept

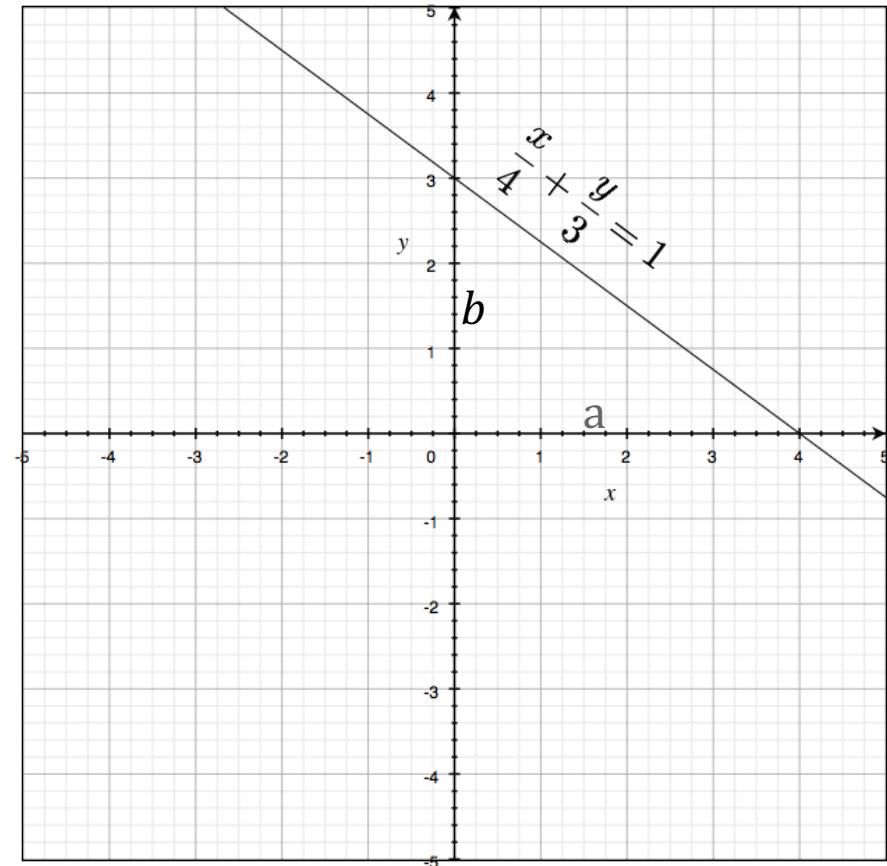
y-intercept

Derivation:

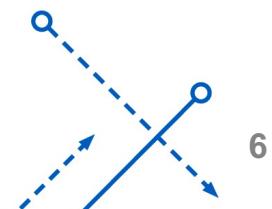
(Similar Triangles)



$$\frac{x}{a} = \frac{b - y}{b}$$



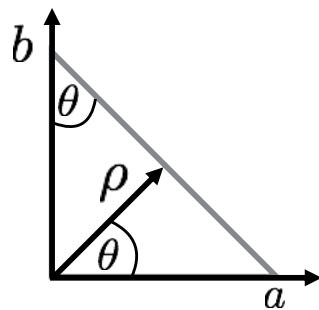
What are a and b ?



Normal Form

$$x \cos \theta + y \sin \theta = \rho$$

Derivation:

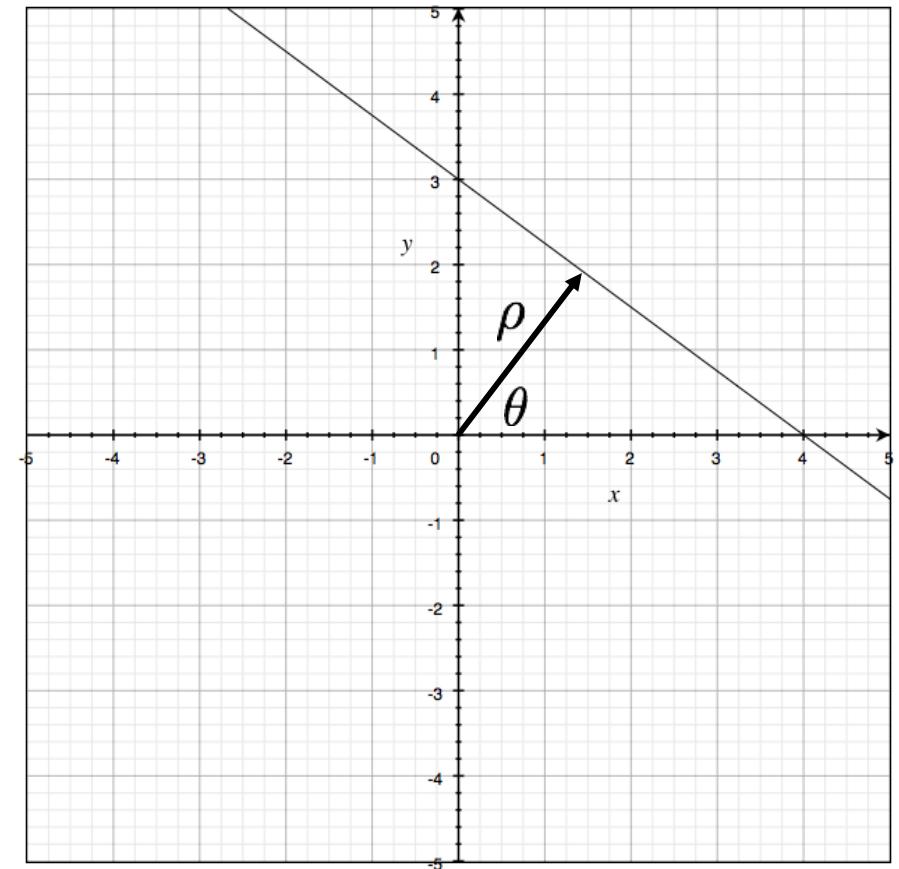


$$a = \frac{\rho}{\cos \theta}$$

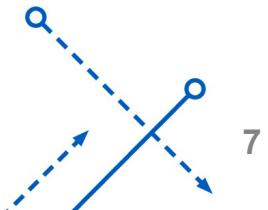
$$b = \frac{\rho}{\sin \theta}$$

plug into:

$$\frac{x}{a} + \frac{y}{b} = 1$$



What are ρ and θ ?



Hough Transform

- Slope intercept form
- Normal Form

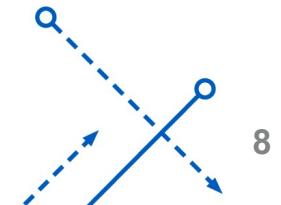
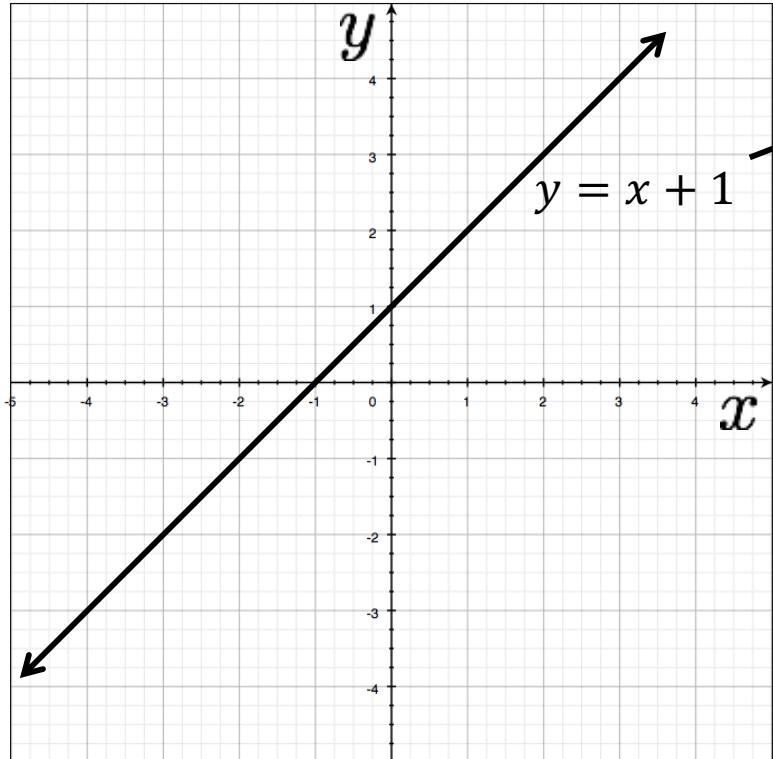


Image and parameter space

variables
 $y = mx + b$
parameters

variables
 $b = -xm + y$
parameters



$m = 1$
 $b = 1$
a line becomes a point

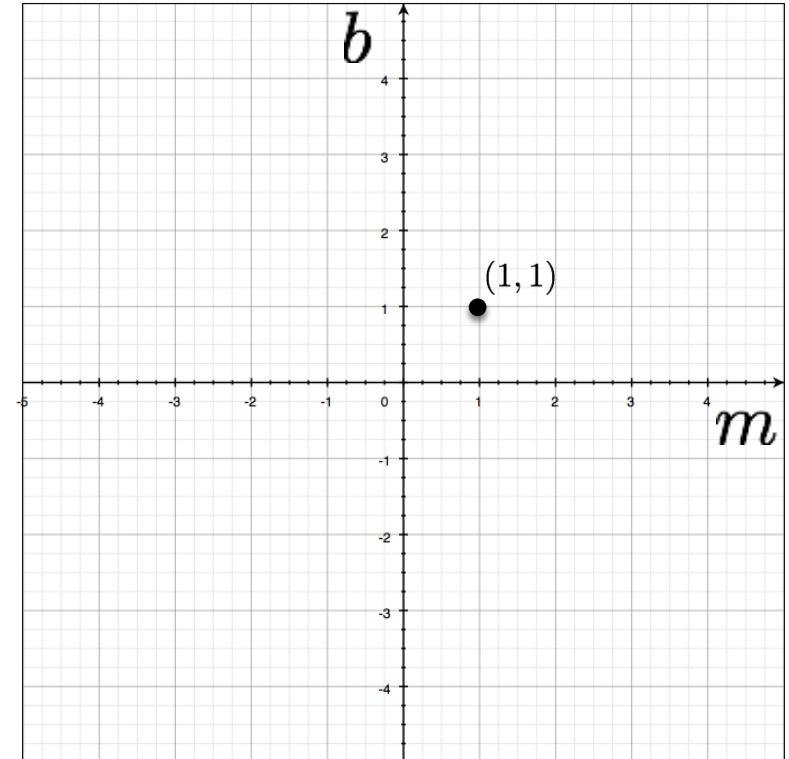


Image space

Parameter space

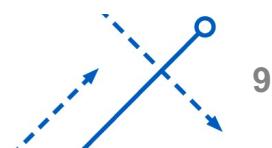


Image and parameter space

$$y = mx + b$$

variables
parameters

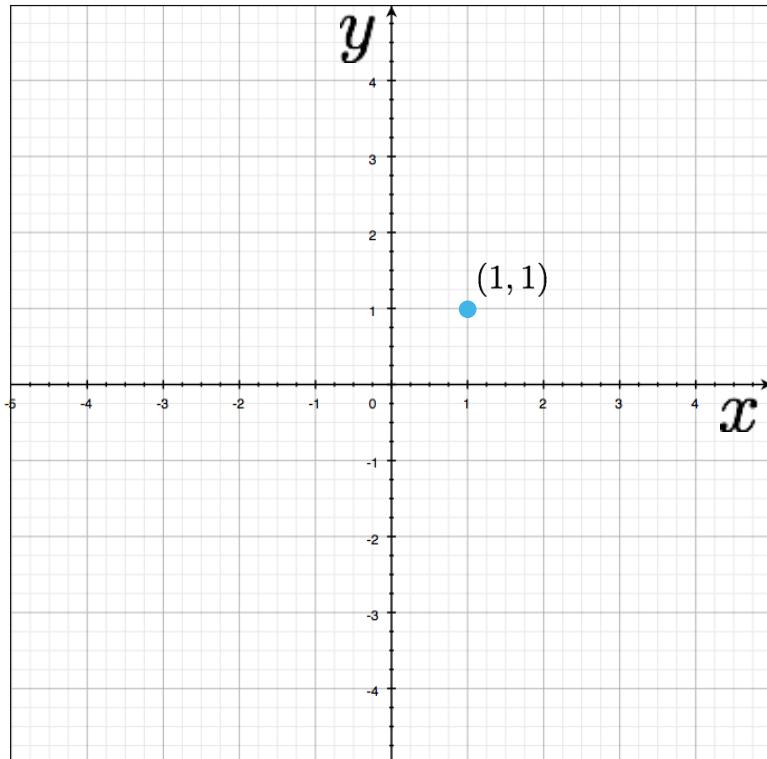


Image space

What would a point in image space become in parameter space?

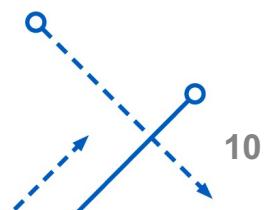
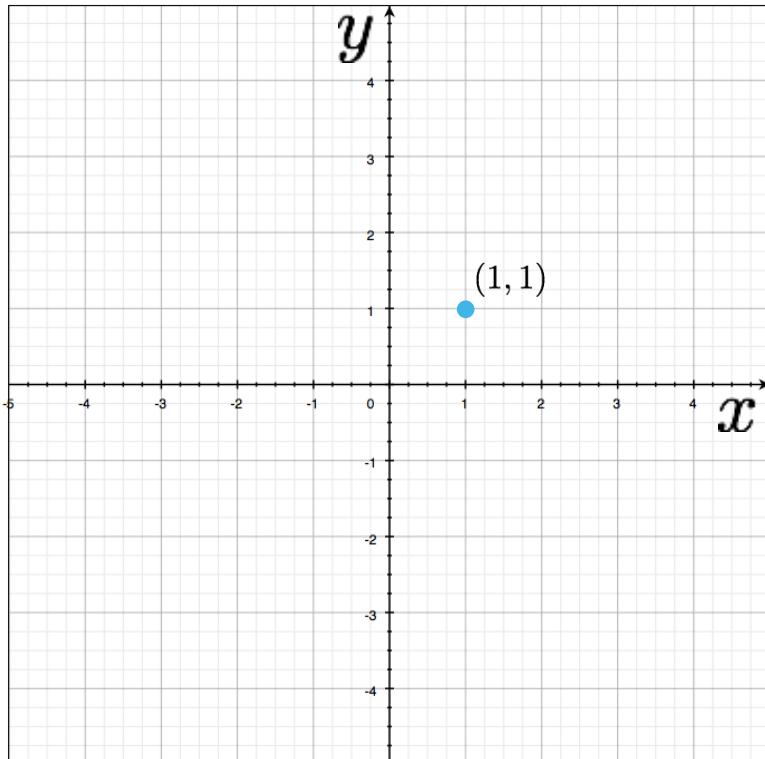


Image and parameter space

variables
 $y = mx + b$
parameters



a point becomes a line

variables
 $b = -xm + y$
parameters

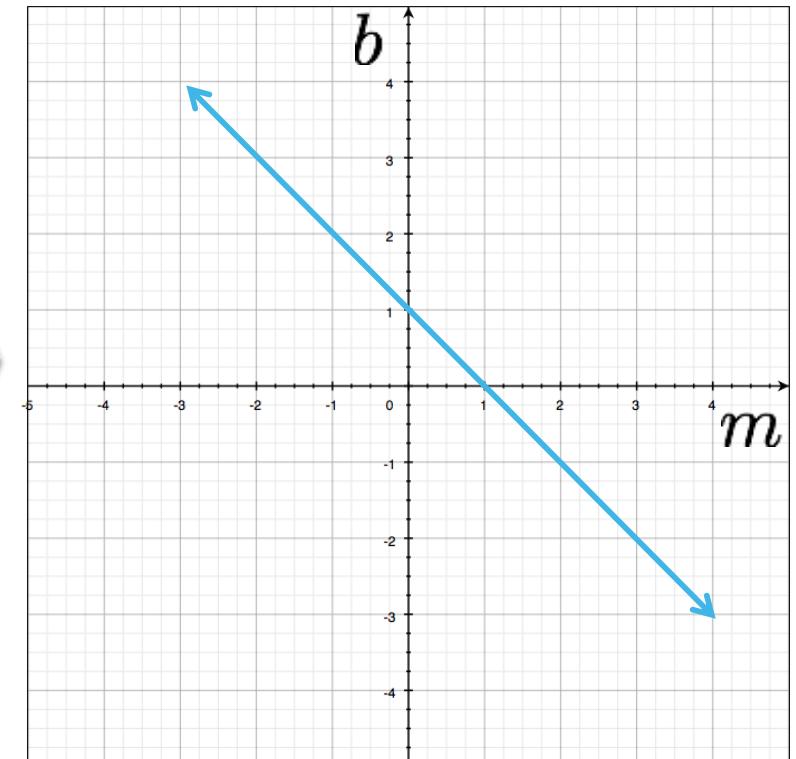
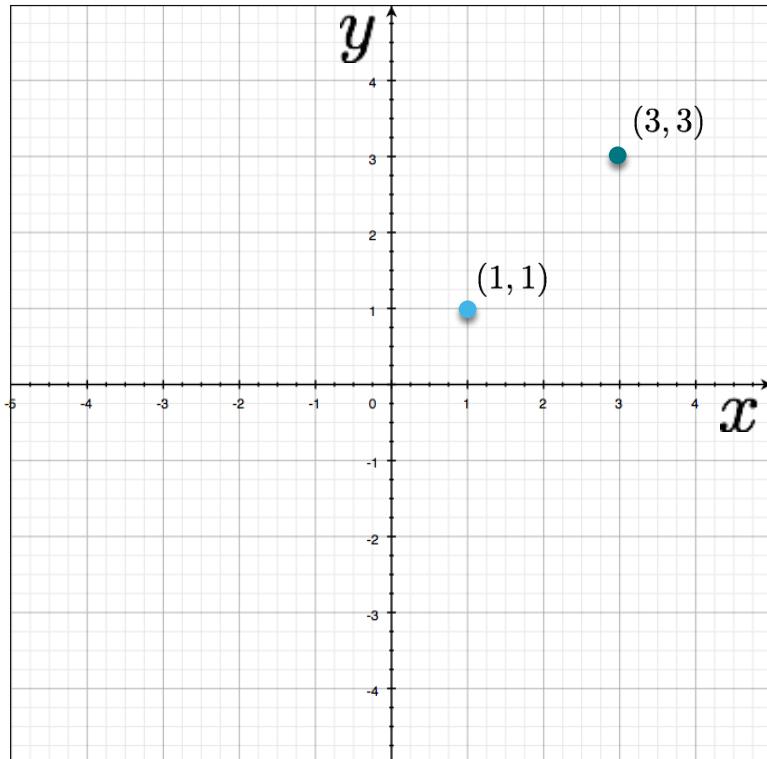


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



two points
become
?

variables
 $b = -xm + y$
parameters

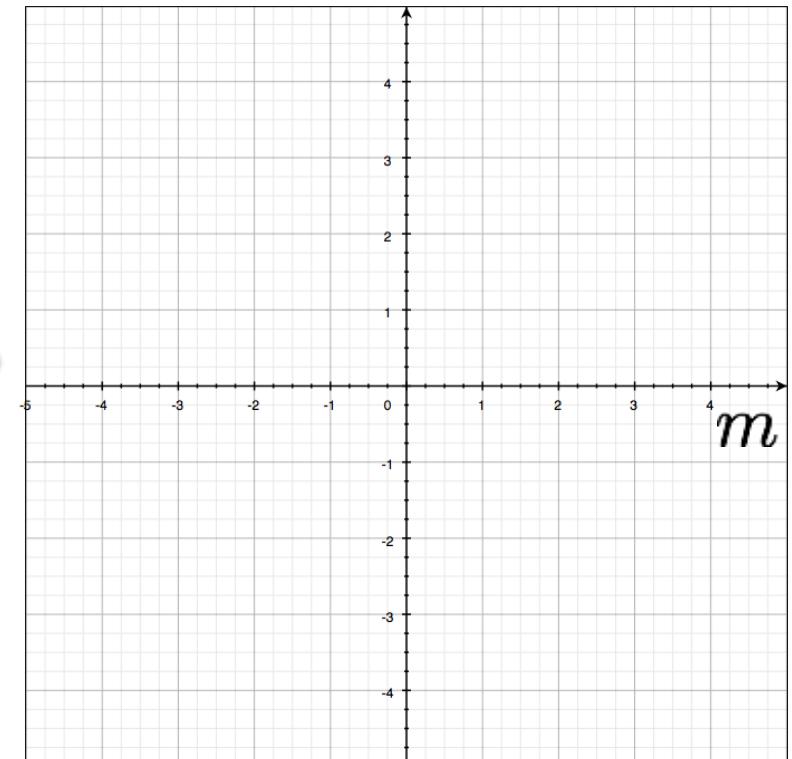
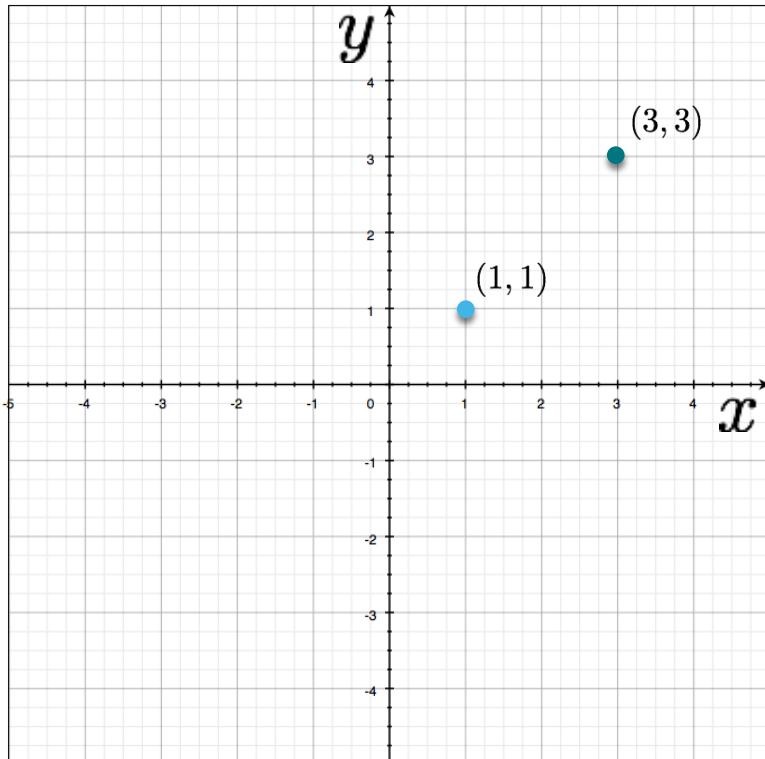


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



two points
become
?

variables
 $b = -xm + y$
parameters

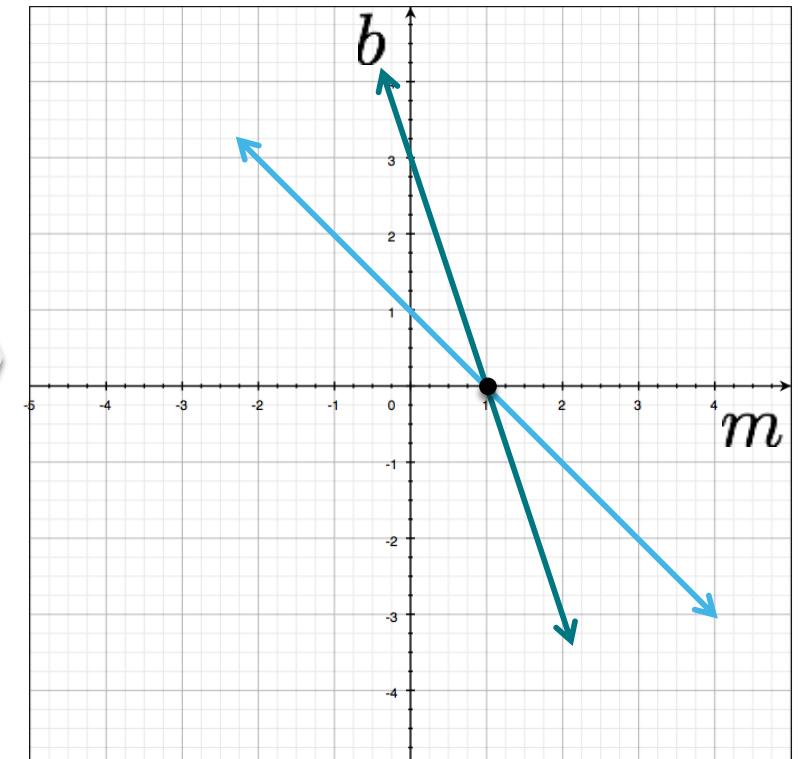
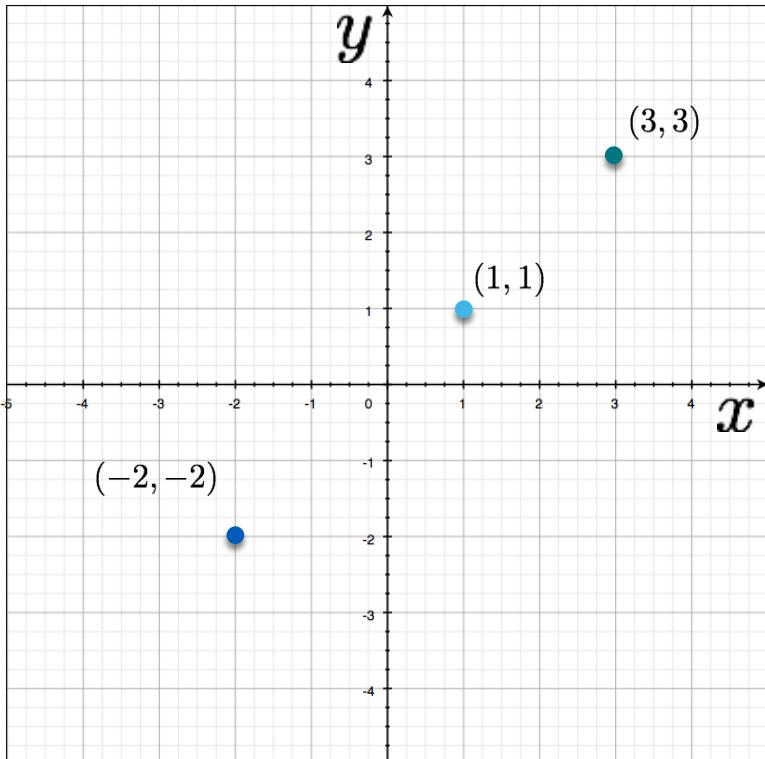


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



three points become?
?

variables
 $b = -xm + y$
parameters

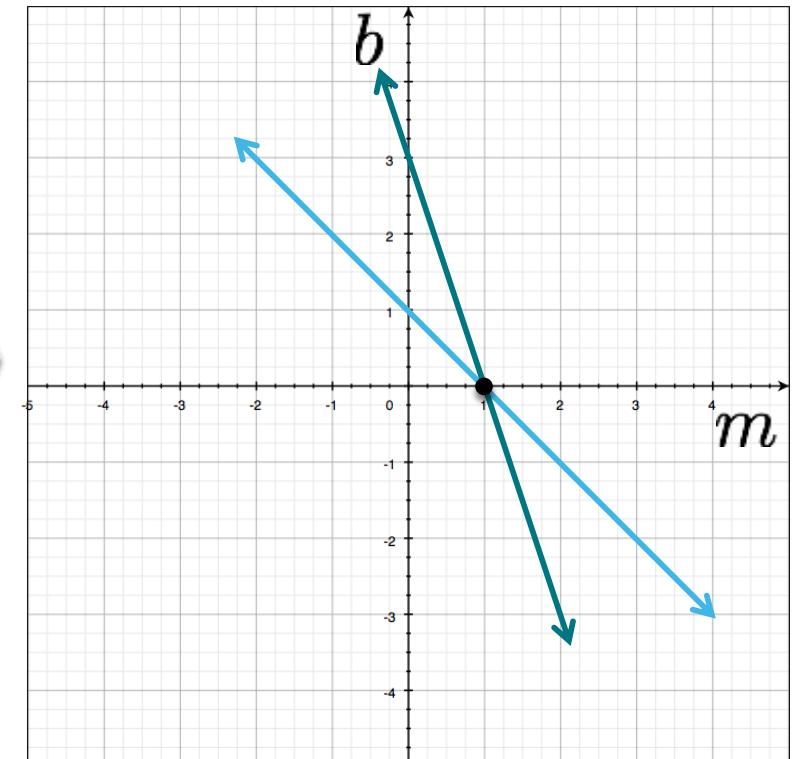


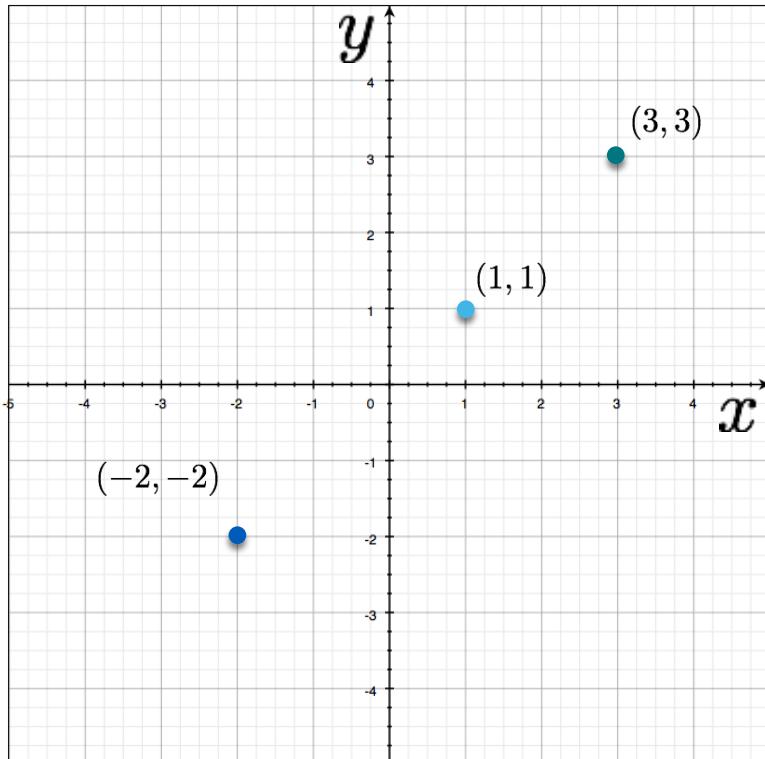
Image space

Parameter space

Image and parameter space

$$y = mx + b$$

variables
parameters



three points become ?

$$b = -xm + y$$

variables
parameters

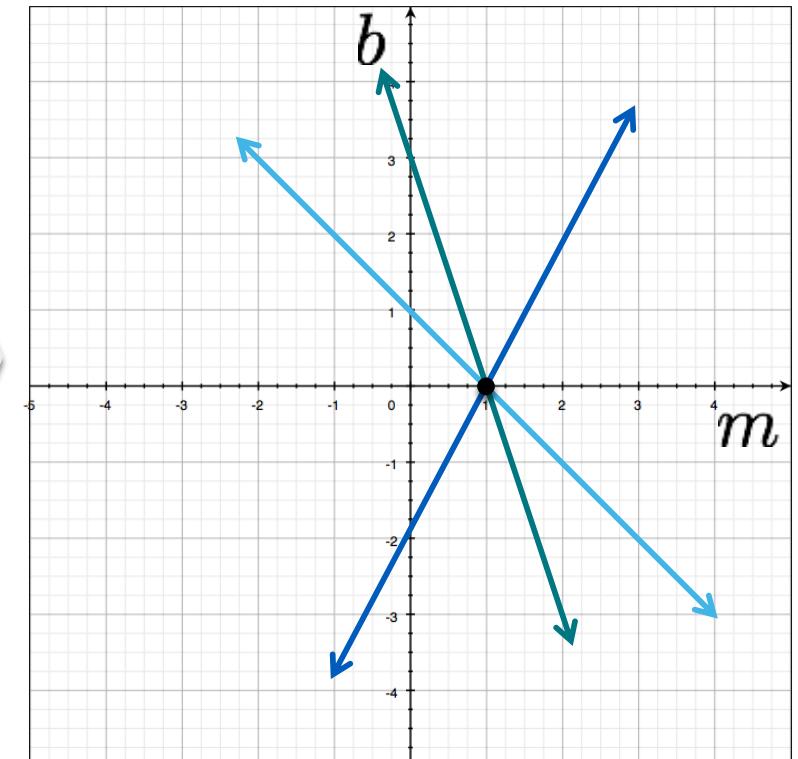
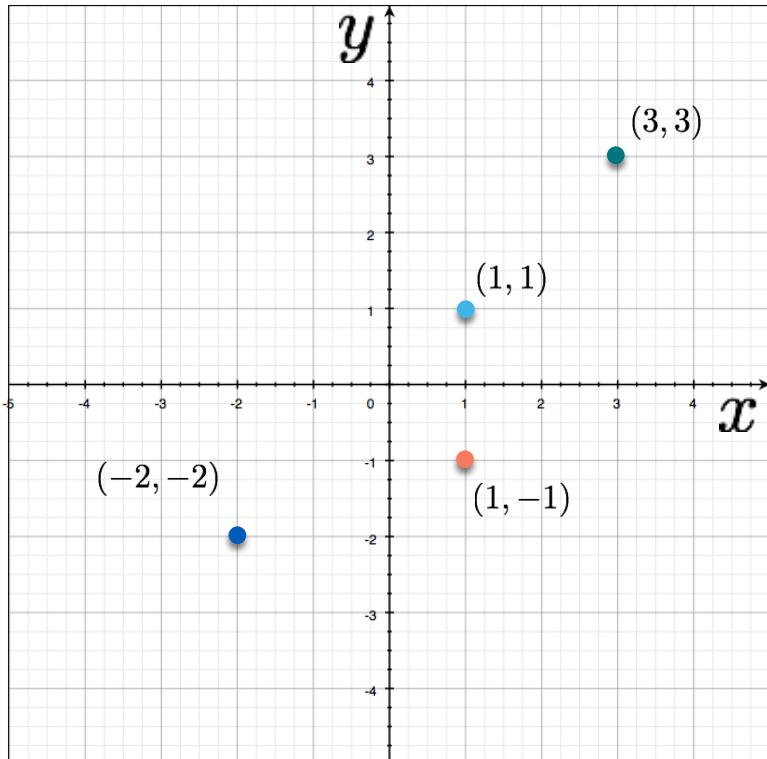


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



four points
become?
?

variables
 $b = -xm + y$
parameters

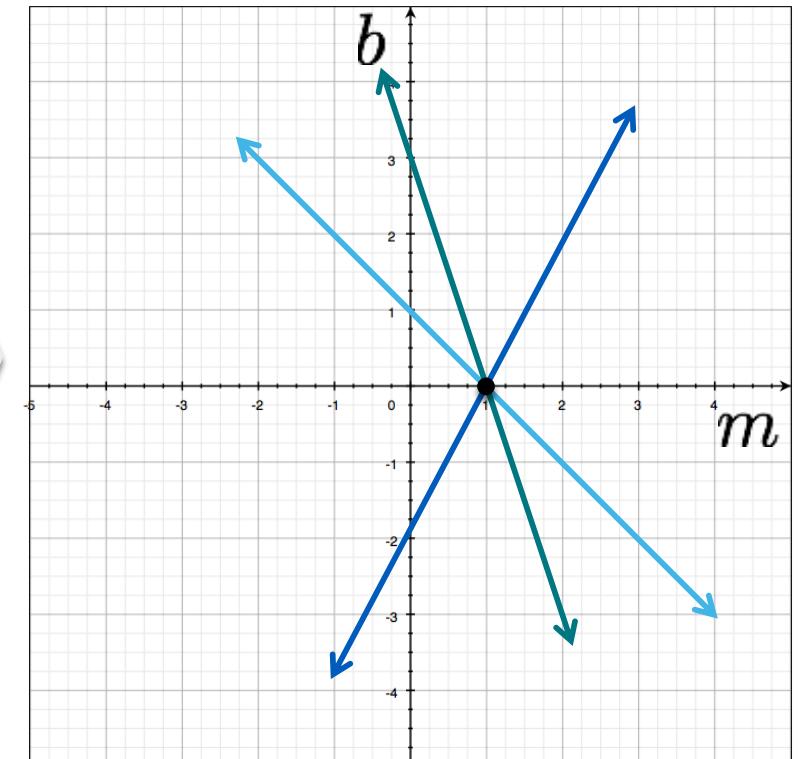
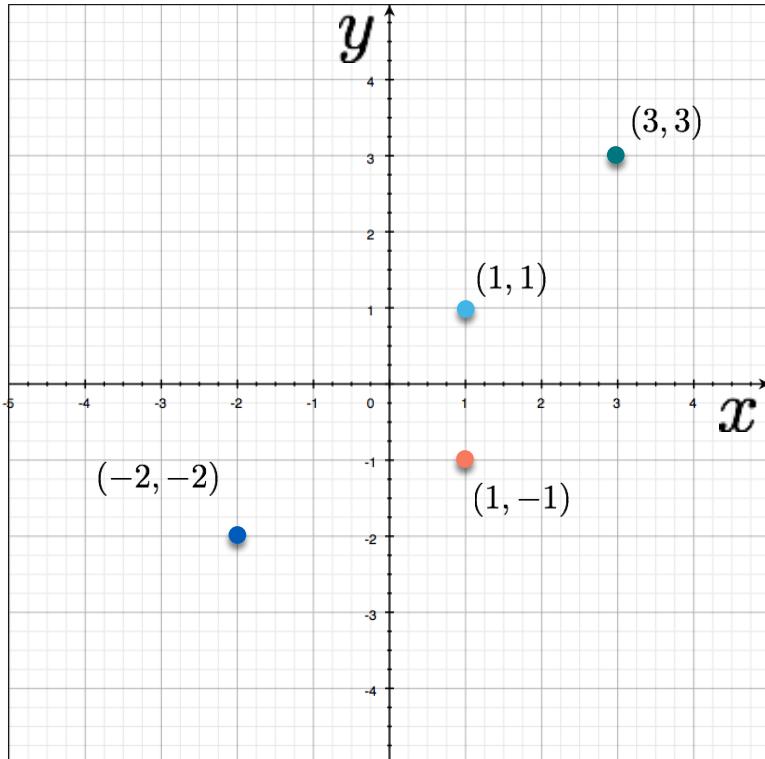


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



four points
become
?

variables
 $b = -xm + y$
parameters

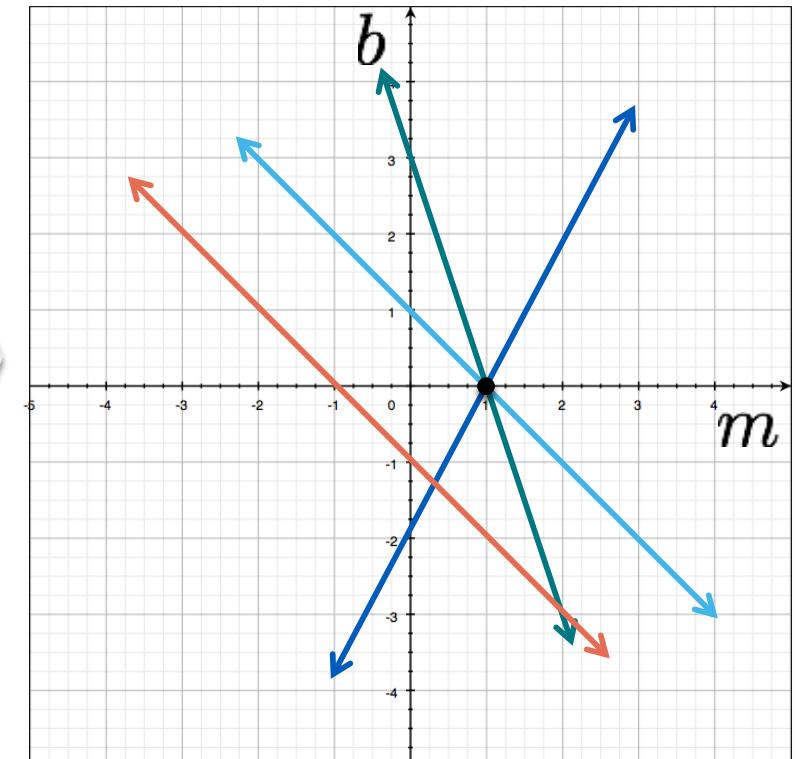


Image space

Parameter space

Is this method robust to noise?

Hough Voting

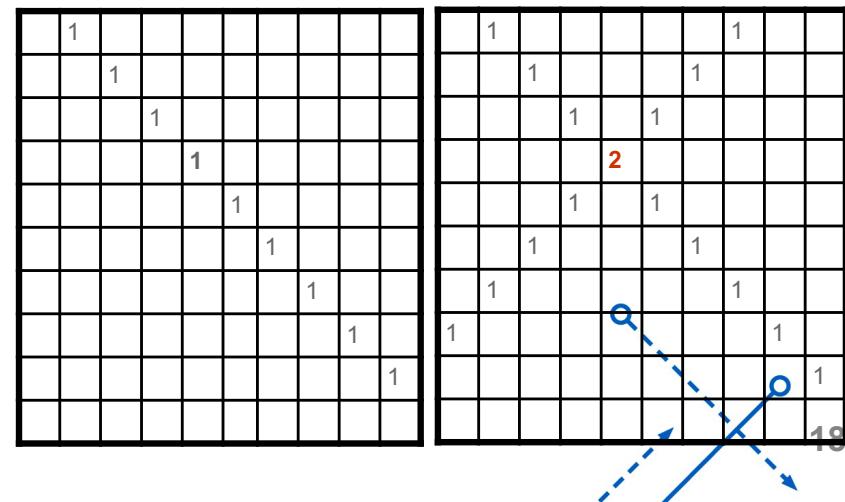
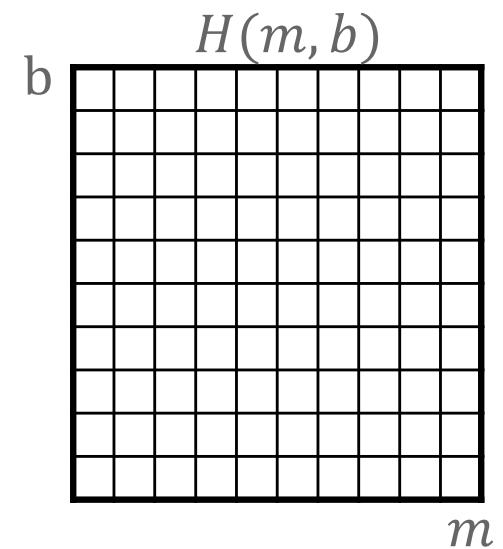
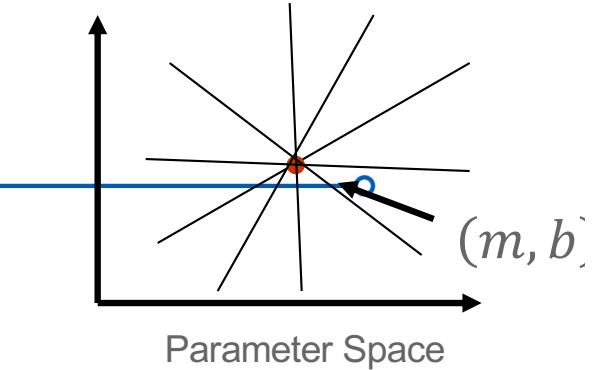
Line Detection Algorithm:

1. Quantize Parameter Space (m, b) .
2. Create Hough Space Array $H(m, b) = 0$.
3. For each image point (x_i, y_i) :
For all points (m, b) on $b = -x_i m + y_i$:
$$H(m, b) = H(m, b) + 1$$
4. Find local maxima in $H(m_m, b_m)$.
5. The detected line: $y = m_m x + b_m$.

Is it able to detect multiple lines?

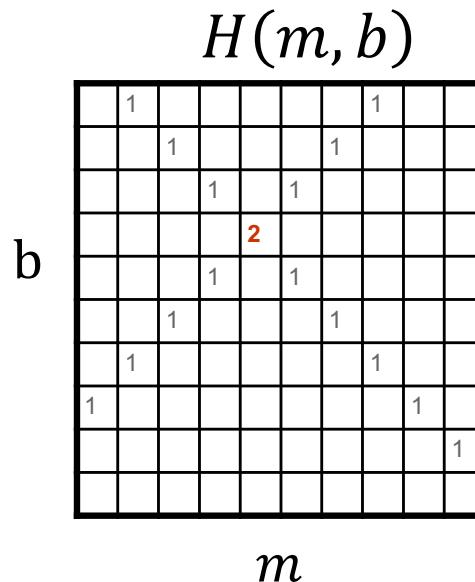


Is this solution good enough?



Problems with slope intercept form

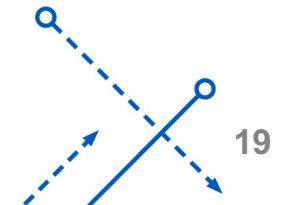
How big does the Hough array have to be?



The space of m is huge! The space of b is huge!

$$-\infty \leq m \leq \infty$$

$$-\infty \leq b \leq \infty$$



Hough Transform with Normal Form

Use normal form:

$$x \cos \theta + y \sin \theta = \rho$$

The Hough space become $H(\rho, \theta)$

Hough Space

$$0 \leq \theta \leq \pi$$

$$0 \leq \rho \leq \rho_{max}$$

(Finite Hough Array Size)

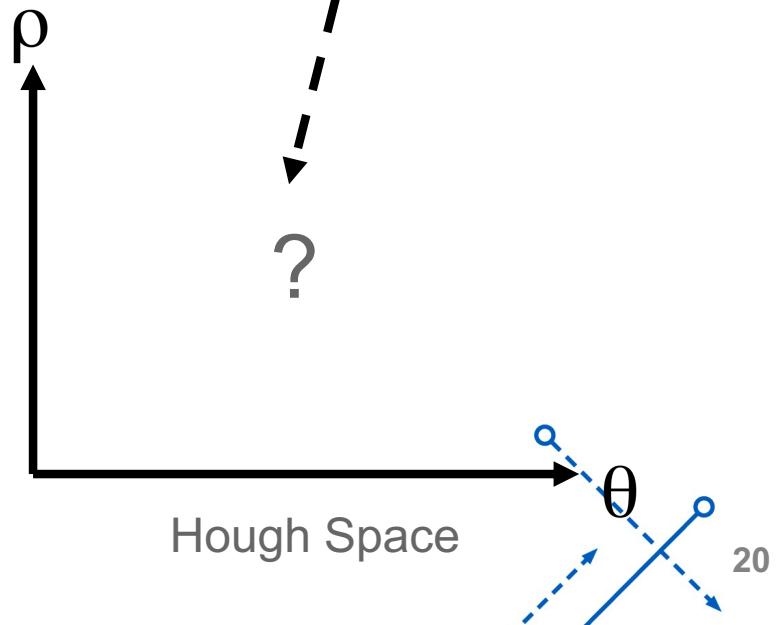
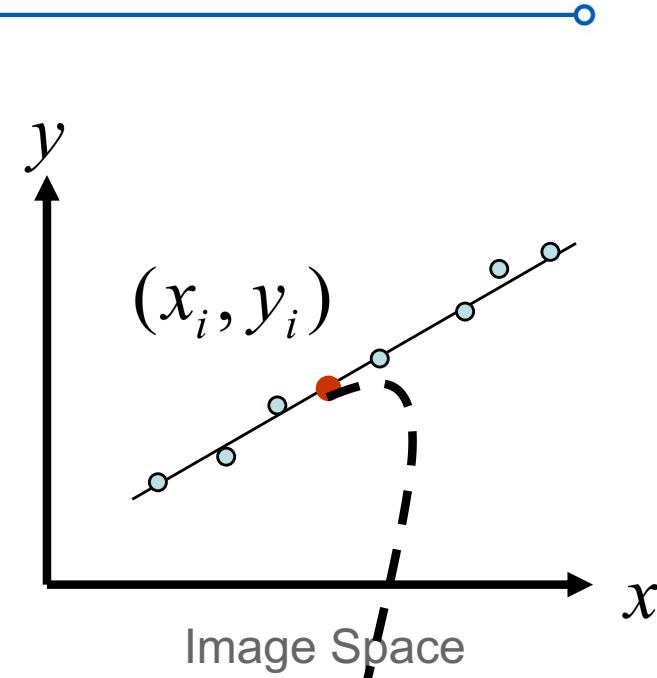
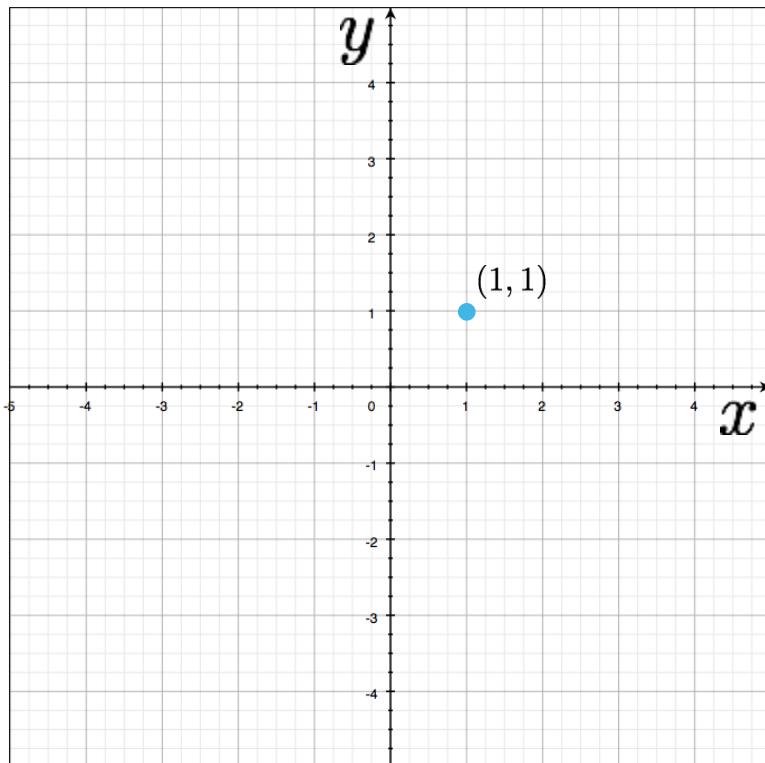


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

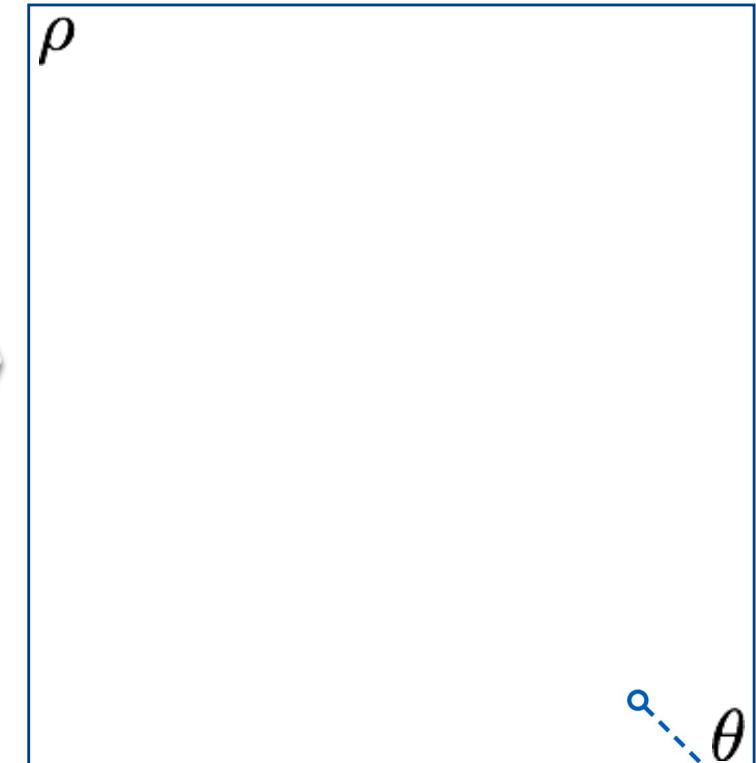
variables
parameters



a point becomes?

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

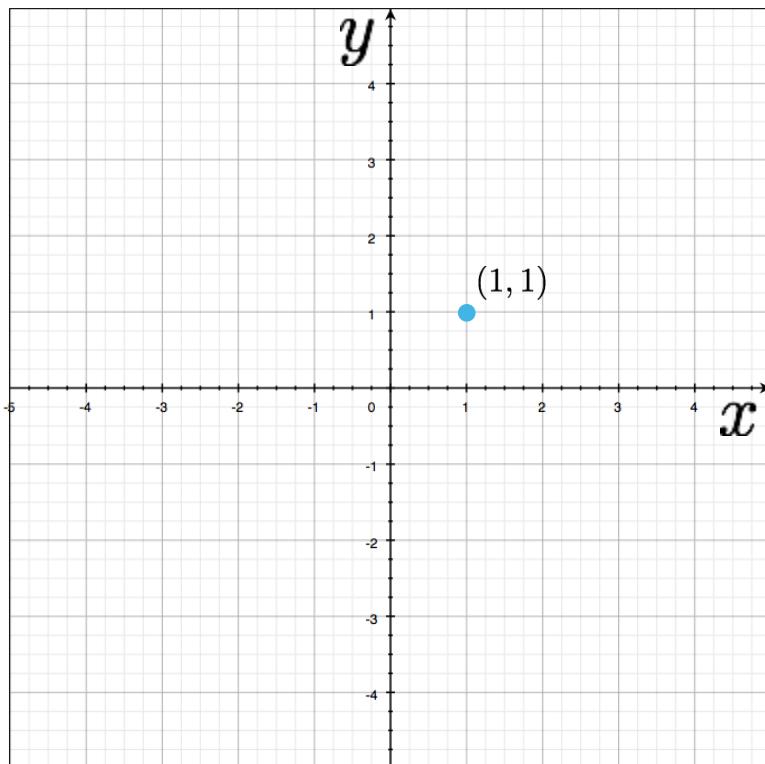


Parameter space

Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



a point becomes a wave

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

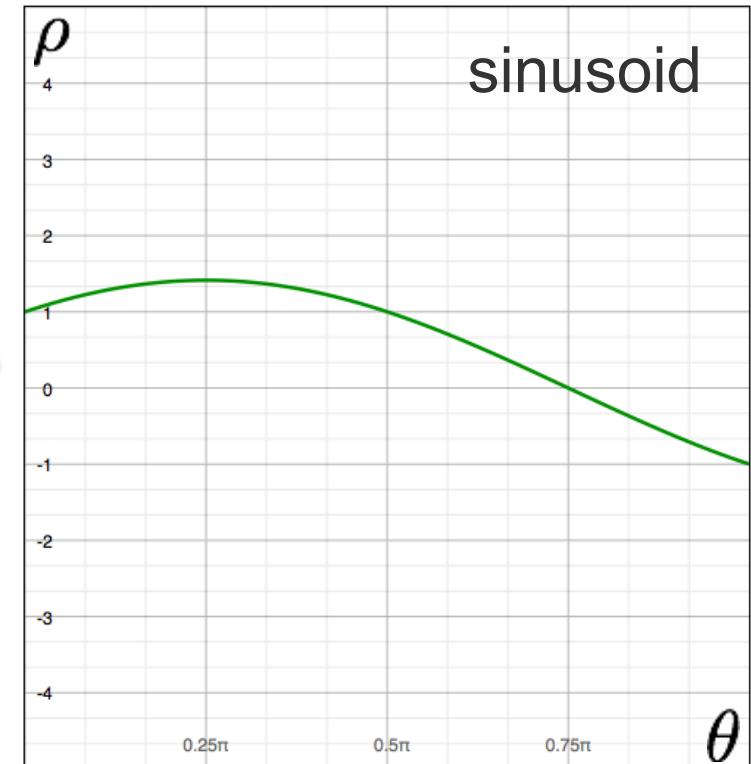
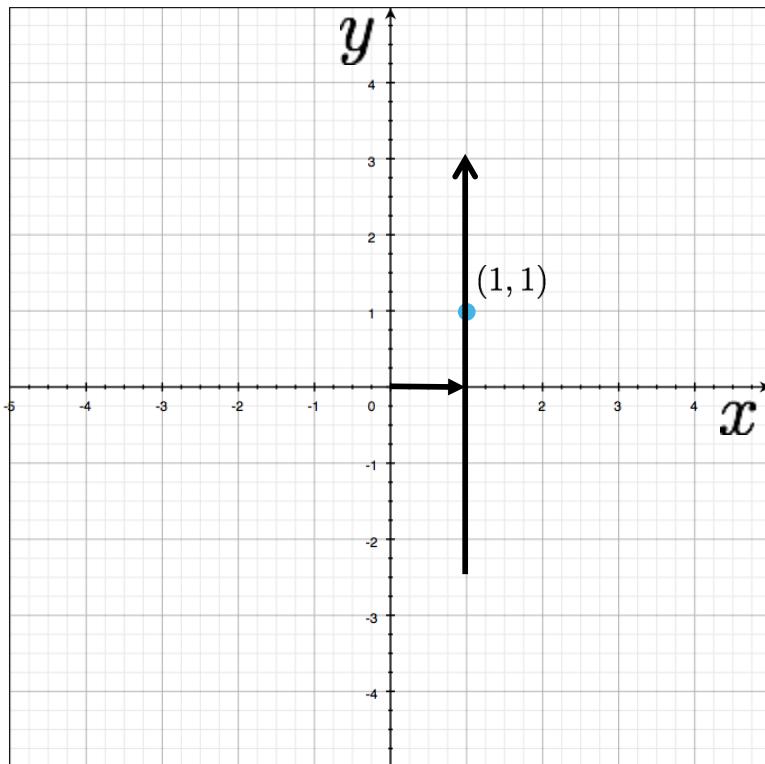


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



a line becomes?

$$x \cos \theta + y \sin \theta = \rho$$

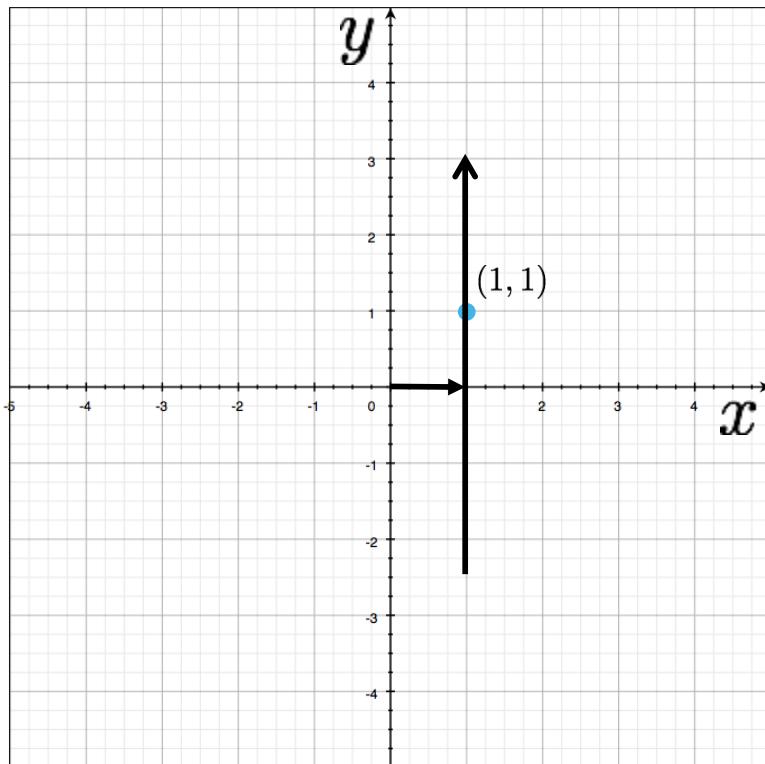
parameters
variables



Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



a line becomes a point

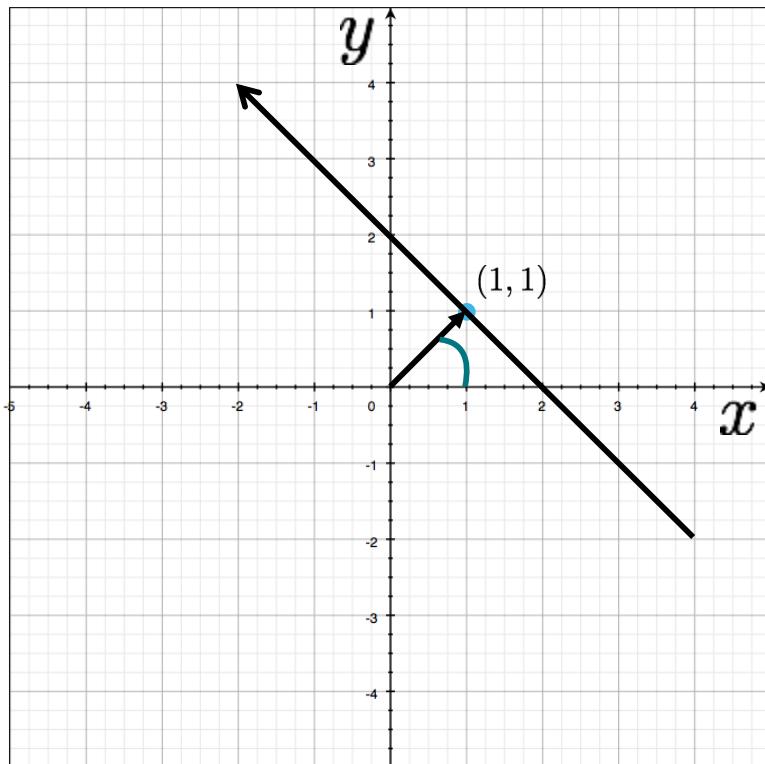
$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables



Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



a line becomes?

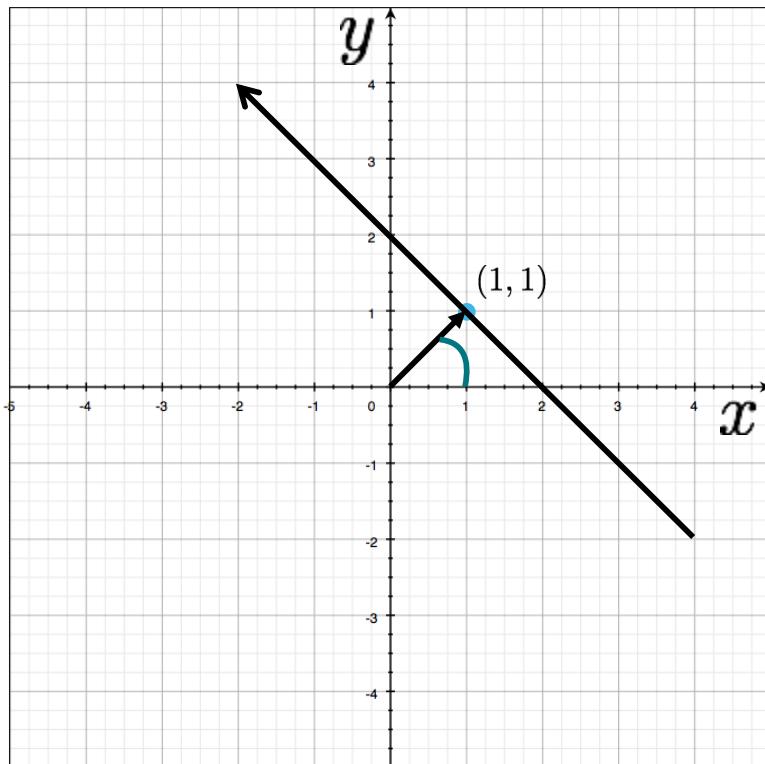
parameters
 $x \cos \theta + y \sin \theta = \rho$
variables



Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



a line becomes a point

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

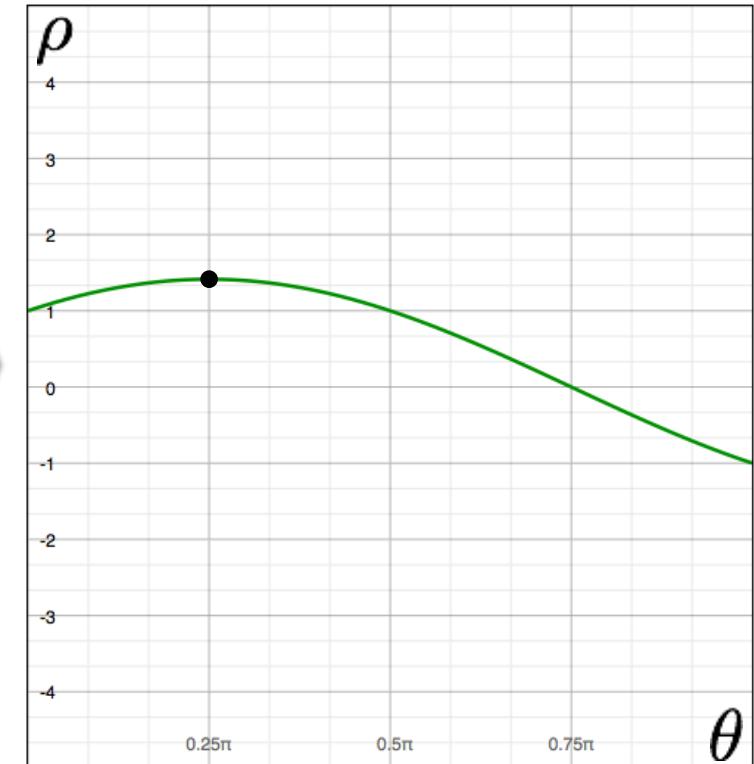
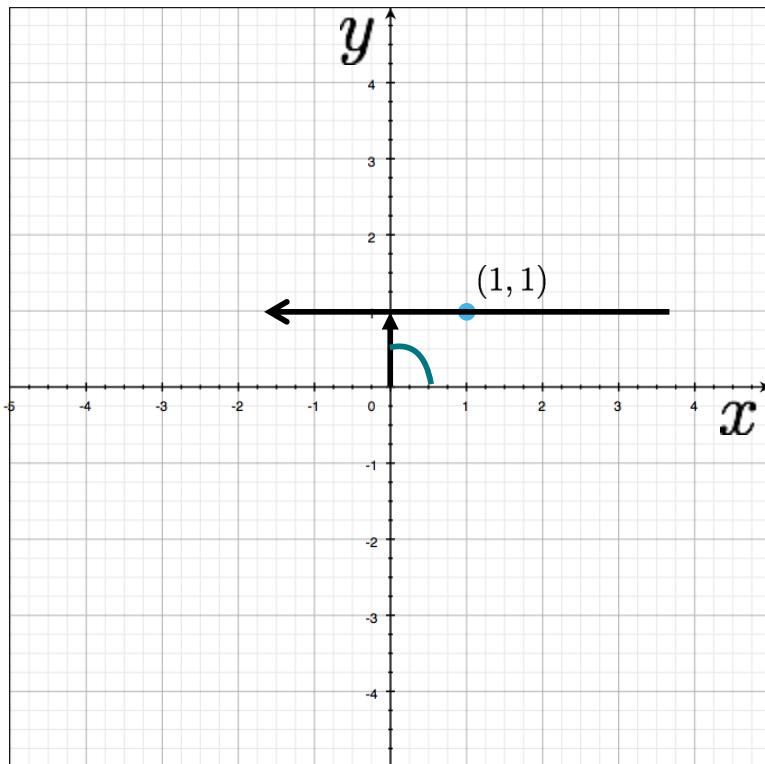


Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



a line becomes a point

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables

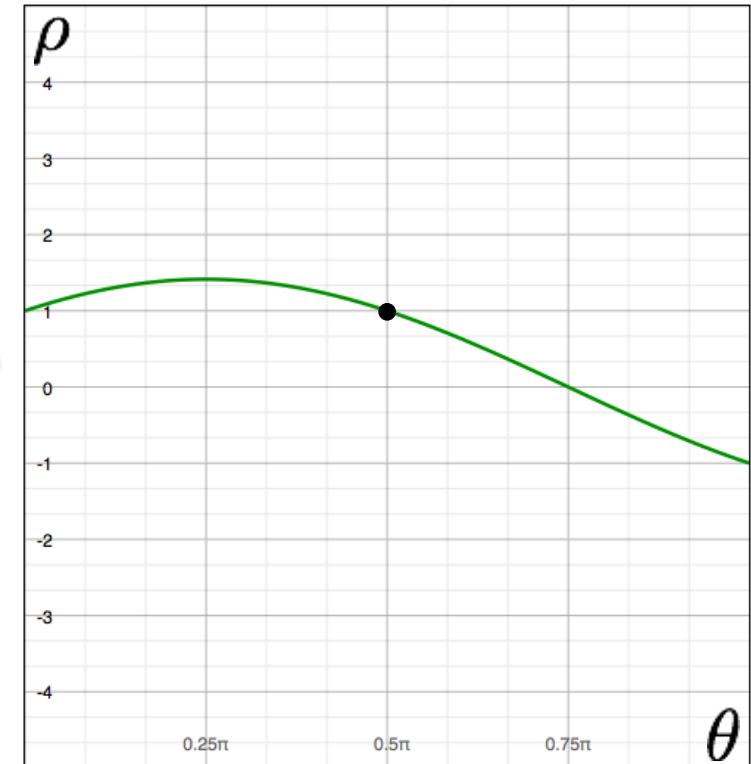
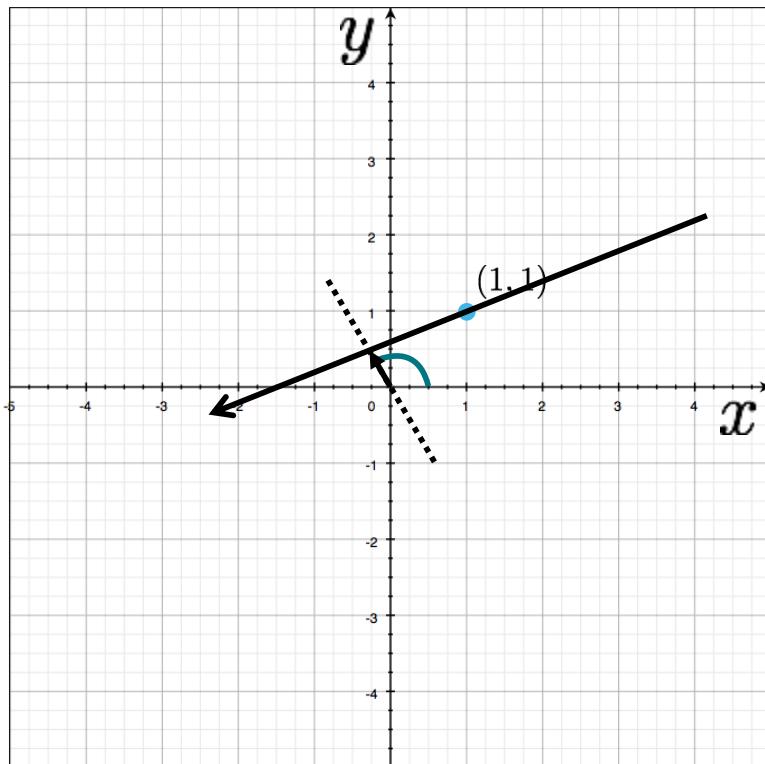
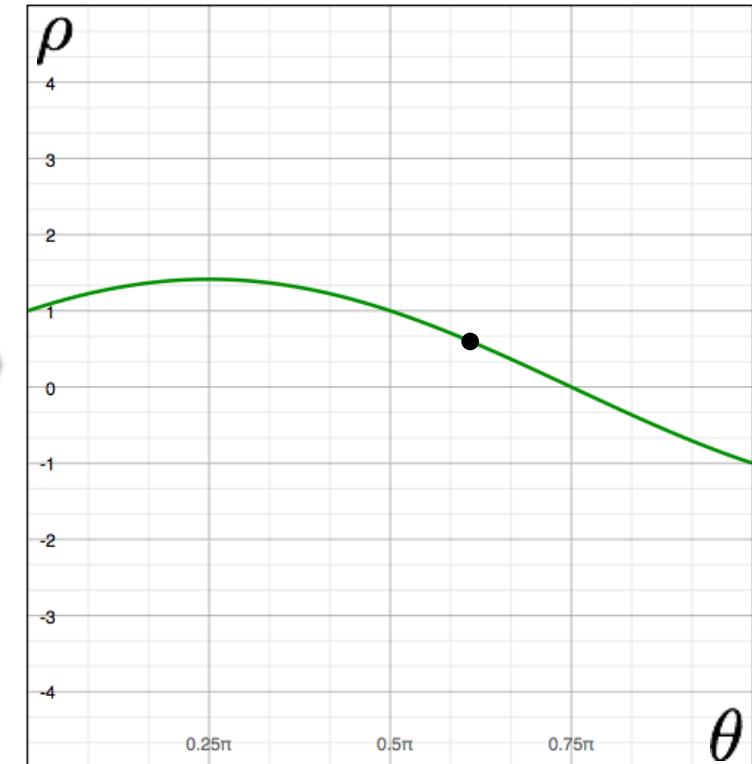


Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



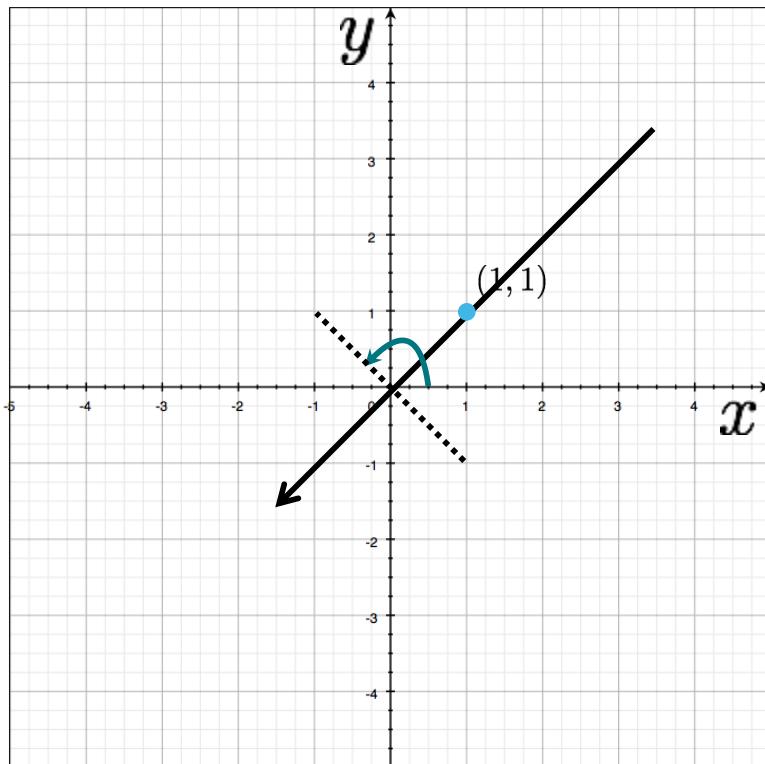
parameters
 $x \cos \theta + y \sin \theta = \rho$
variables



a line becomes a point

Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



a line becomes a point

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables

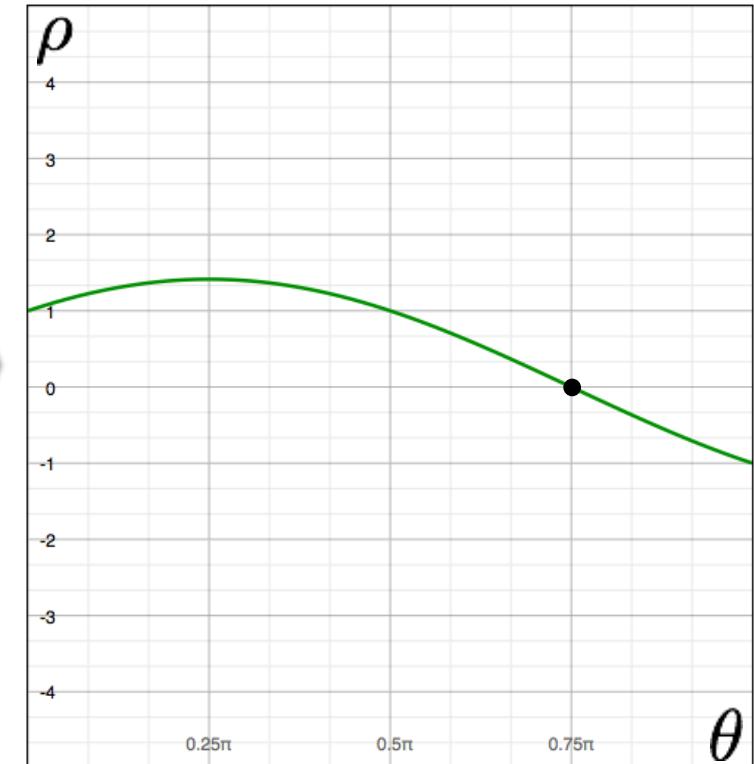
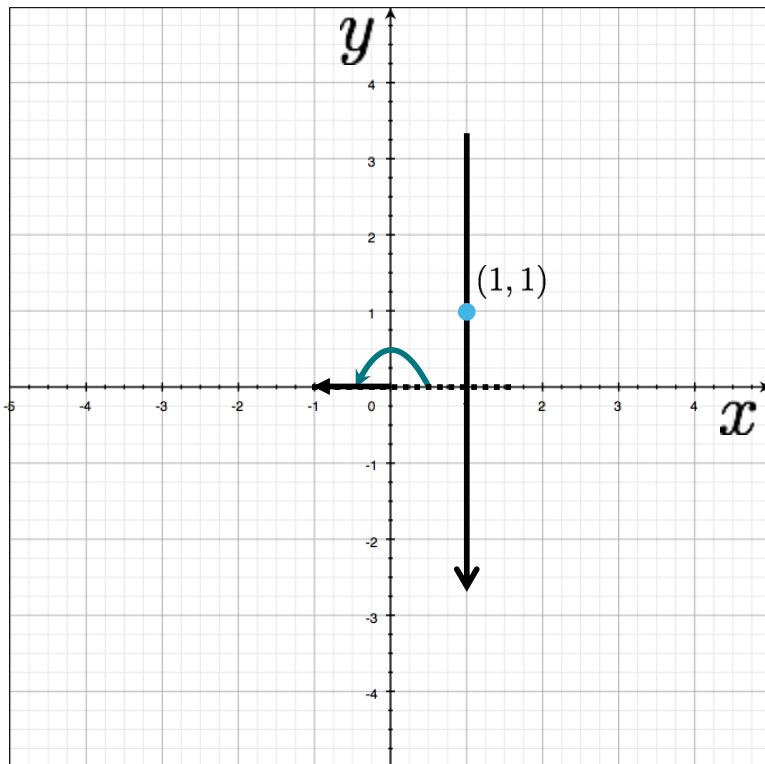


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



a line becomes a point

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

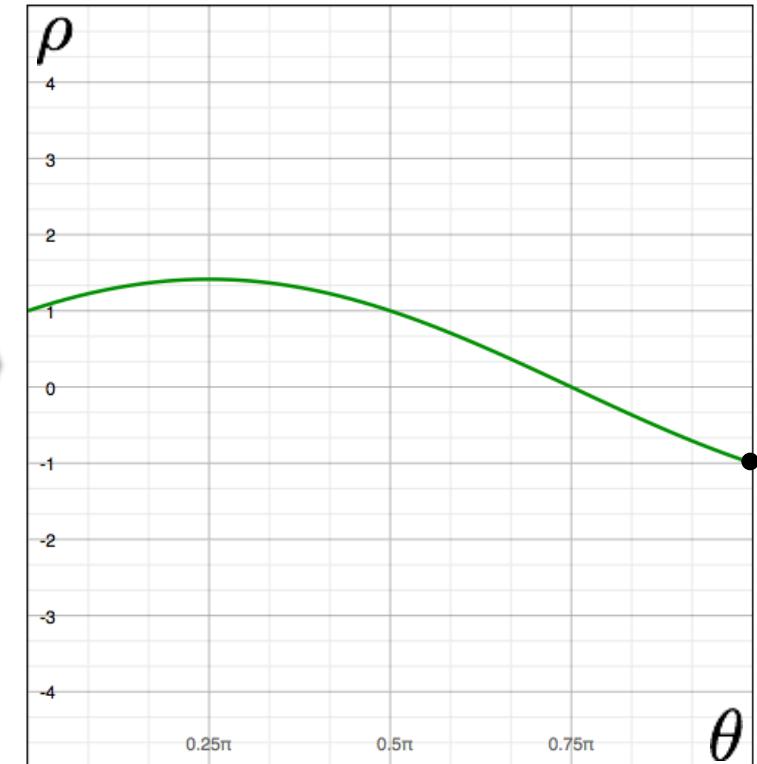


Image space

Parameter space

Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables

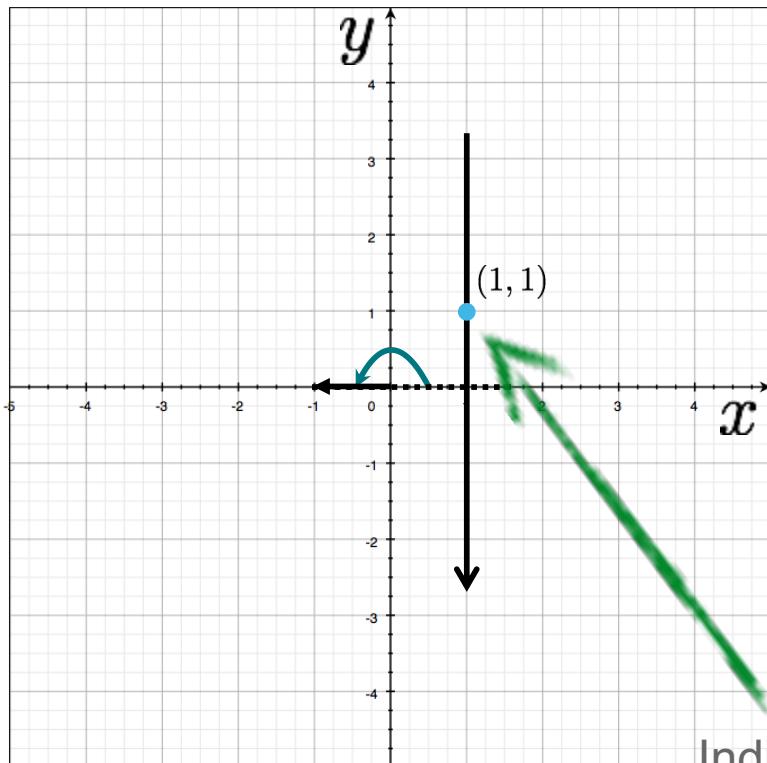
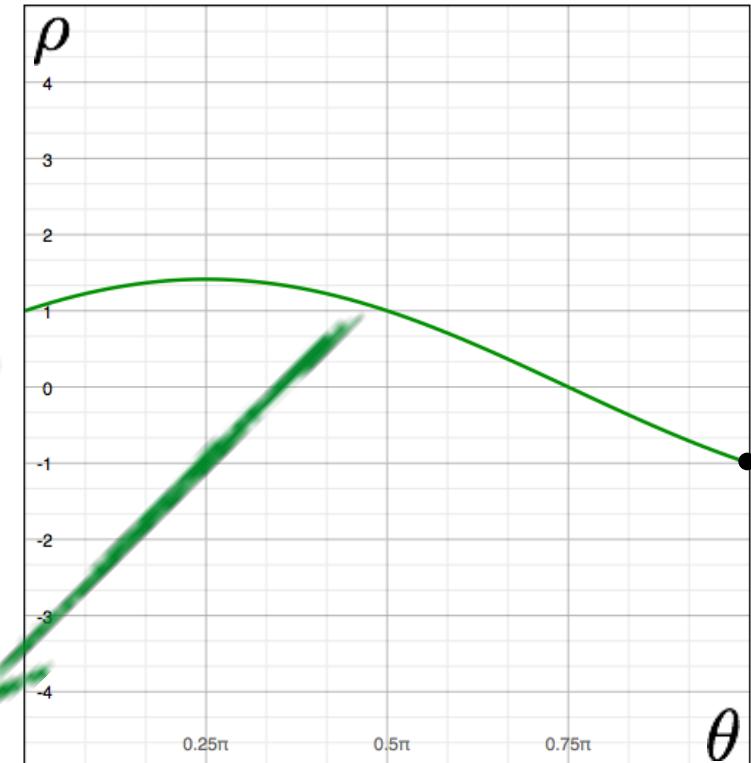


Image space

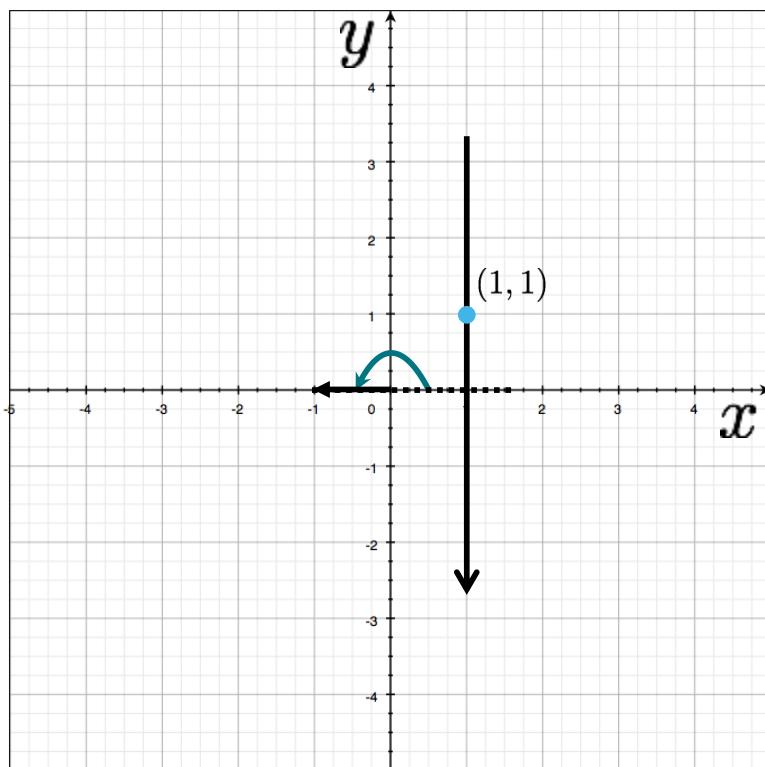
Indicate all lines
pass through (1,1)



Parameter space

Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



a line becomes a point

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables

Wait ... why is ρ negative?

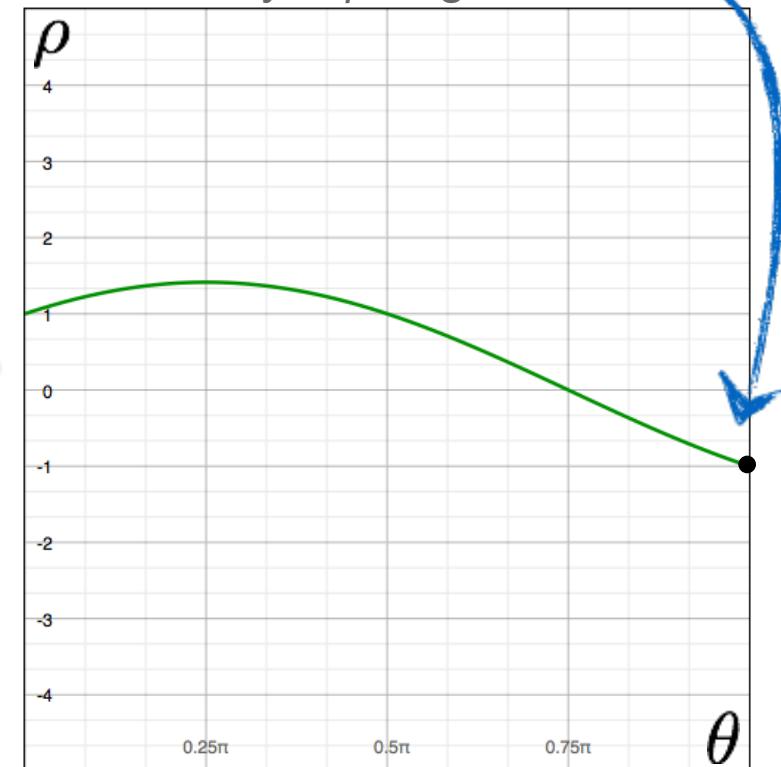
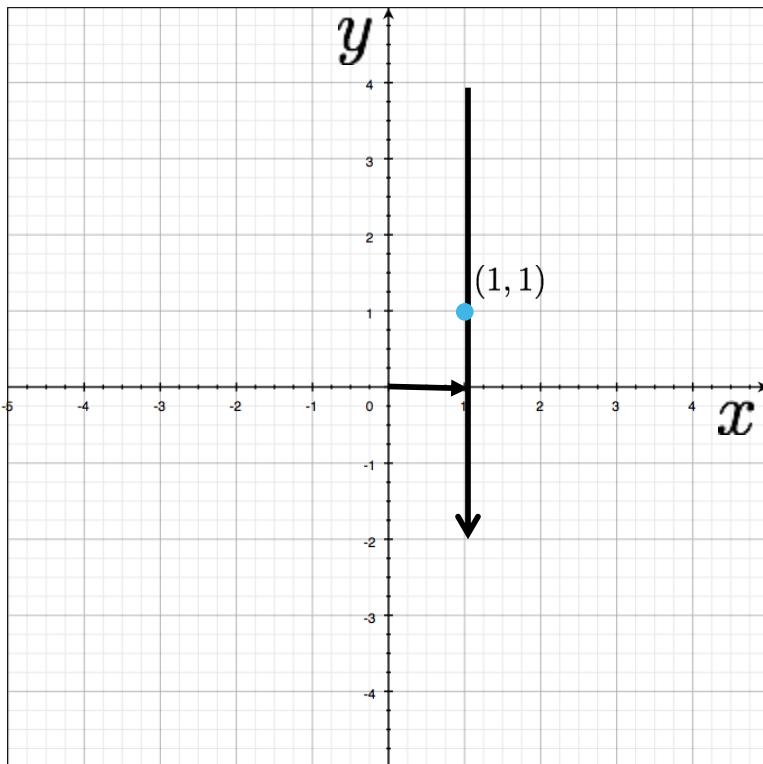


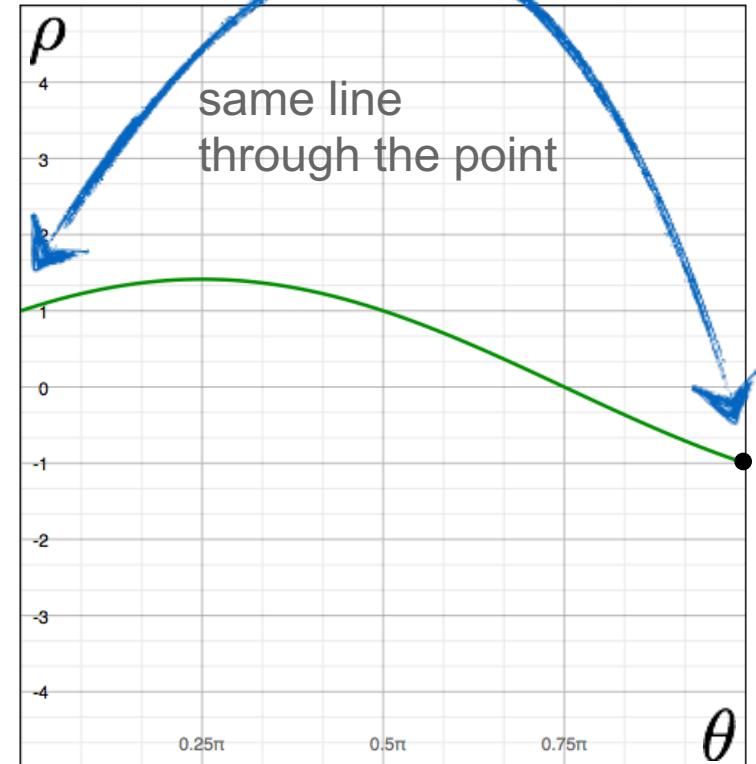
Image and parameter space

variables
 $x \cos \theta + y \sin \theta = \rho$
parameters



a line becomes a point

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables



There are two ways to write the same line

Positive ρ version:

$$x \cos \theta + y \sin \theta = \rho$$

Negative ρ version:

$$x \cos(\theta + \pi) + y \sin(\theta + \pi) = -\rho$$

Recall:

$$\sin(\theta) = -\sin(\theta + \pi)$$

$$\cos(\theta) = -\cos(\theta + \pi)$$

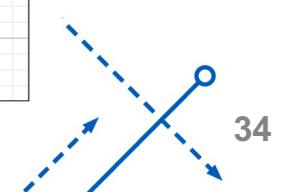
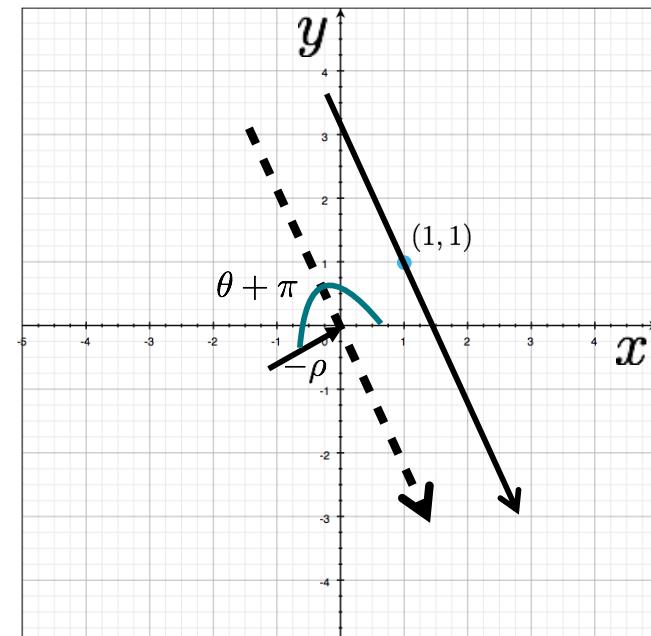
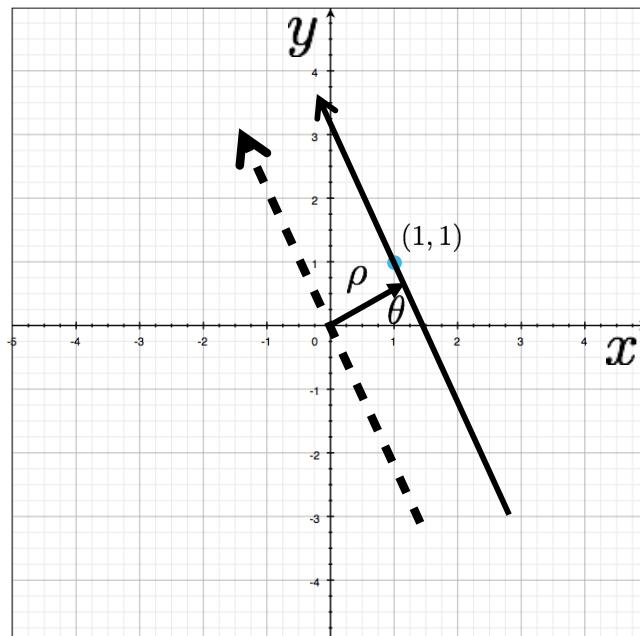
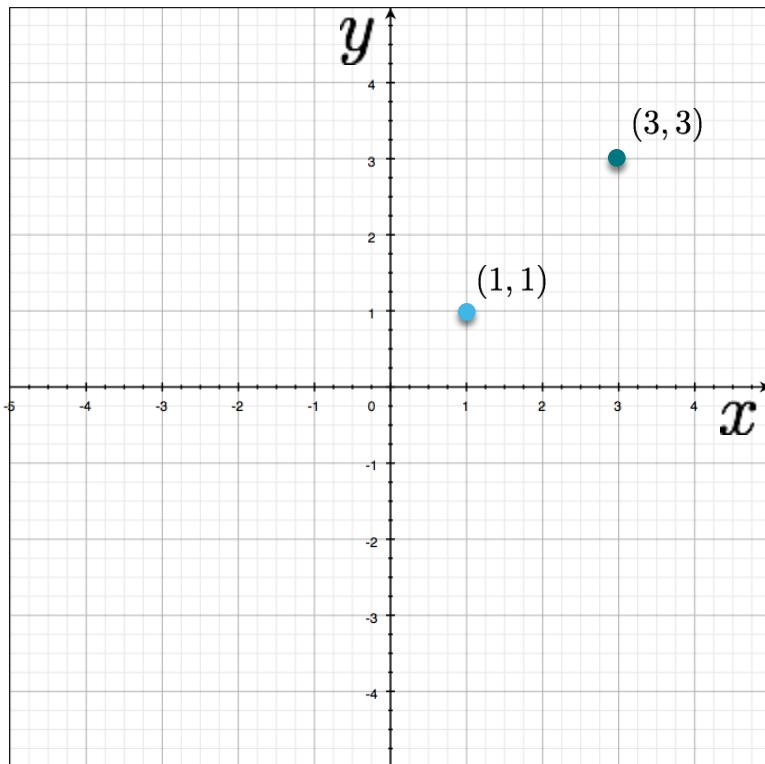


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



two points
become
?

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

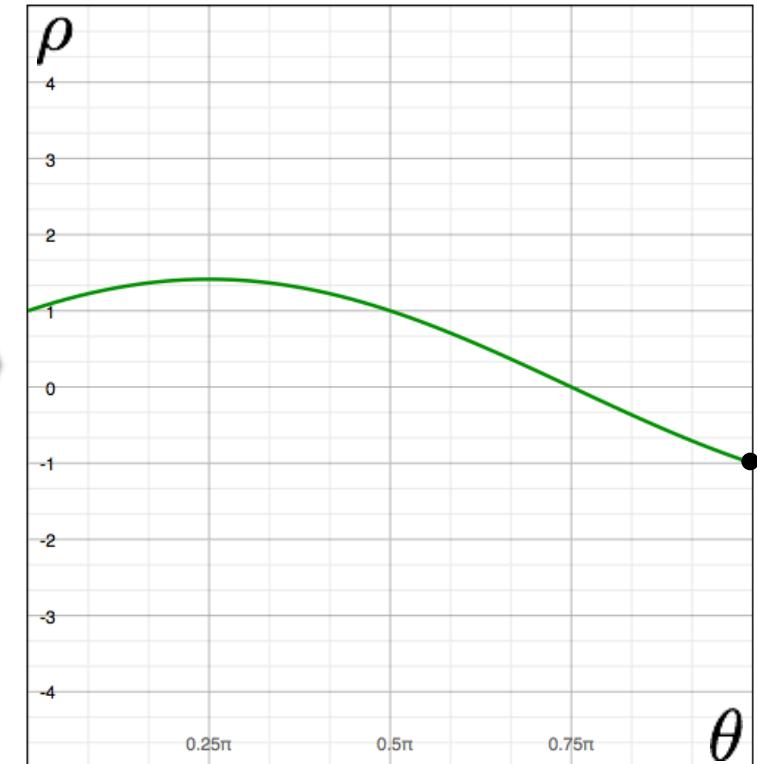
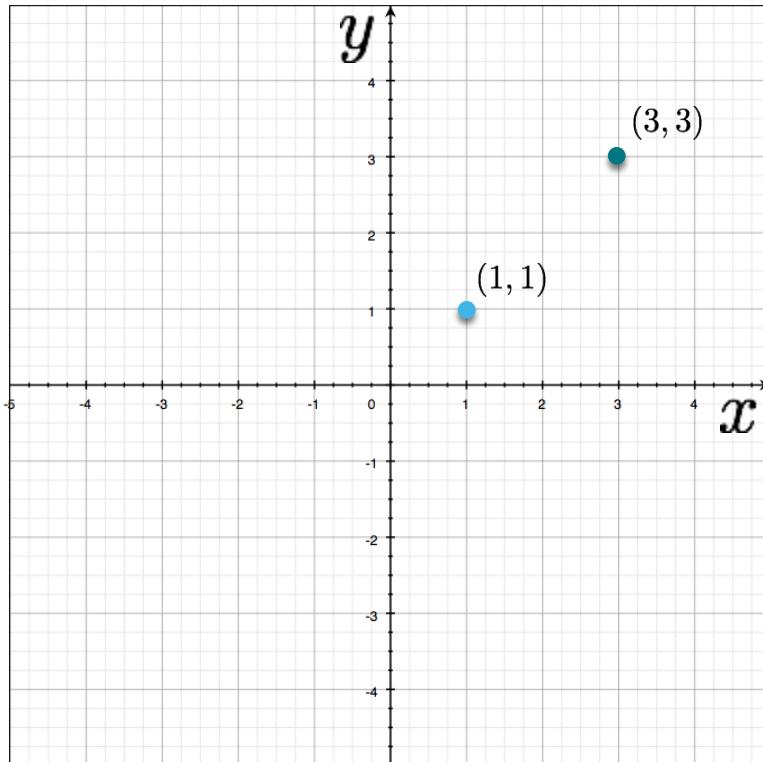


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

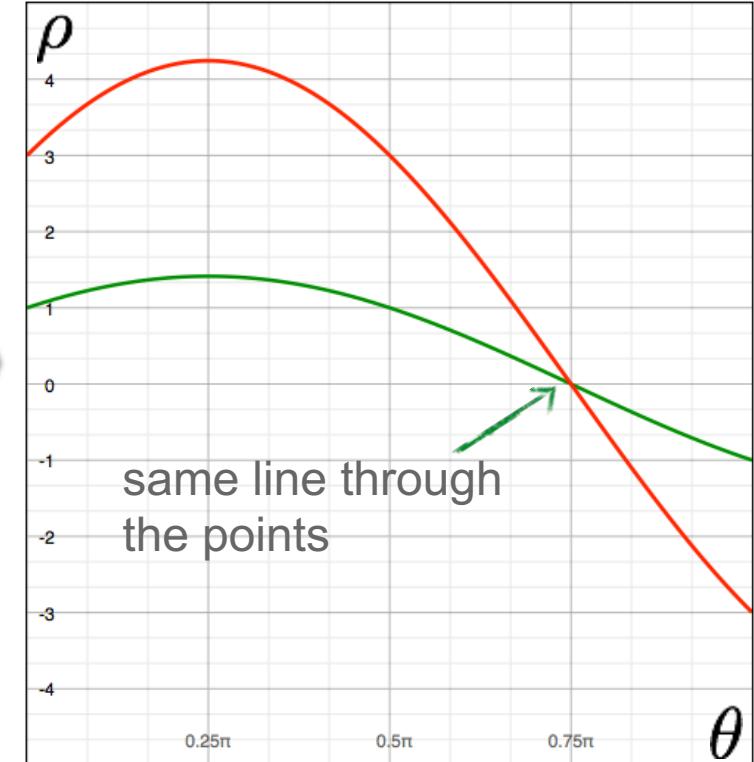
variables
parameters



two points
become
?

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

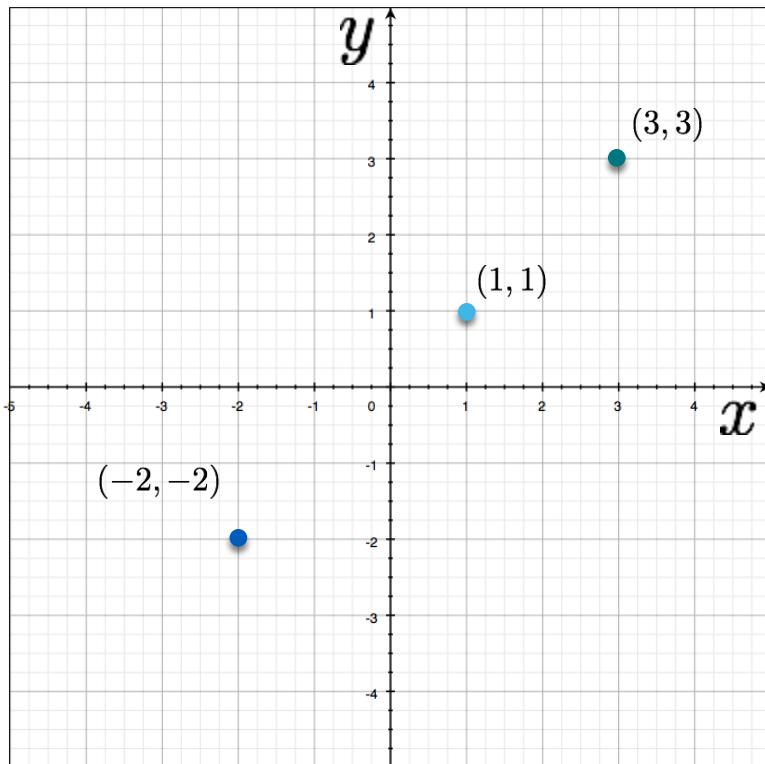


Parameter space

Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

variables
parameters



three points
become
?

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

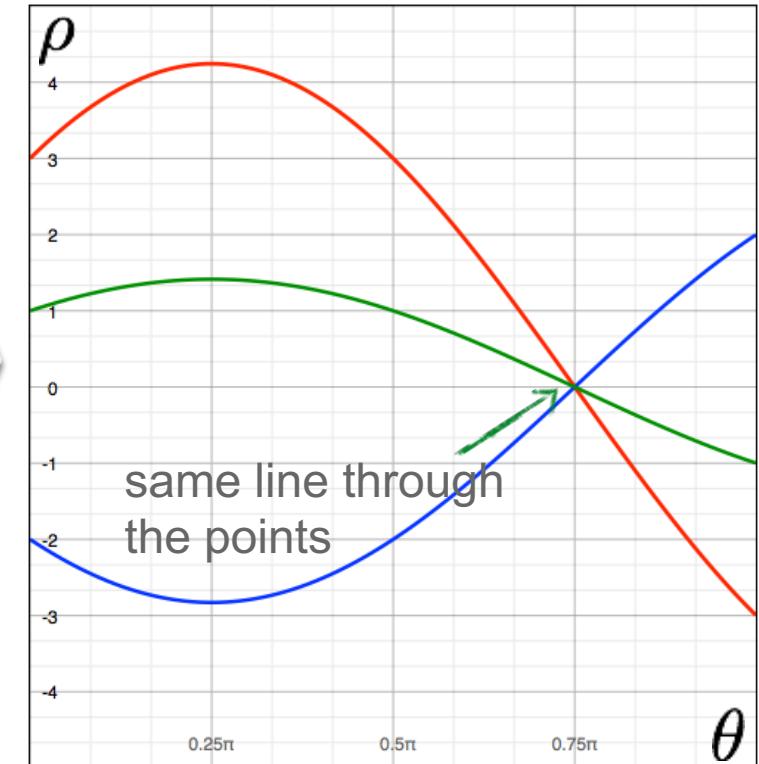
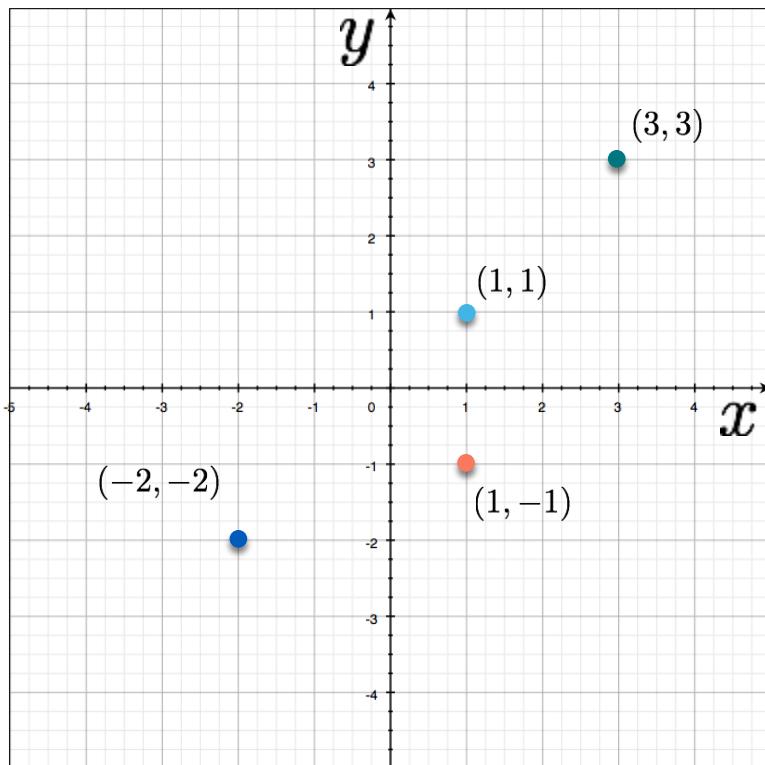


Image and parameter space

$$x \cos \theta + y \sin \theta = \rho$$

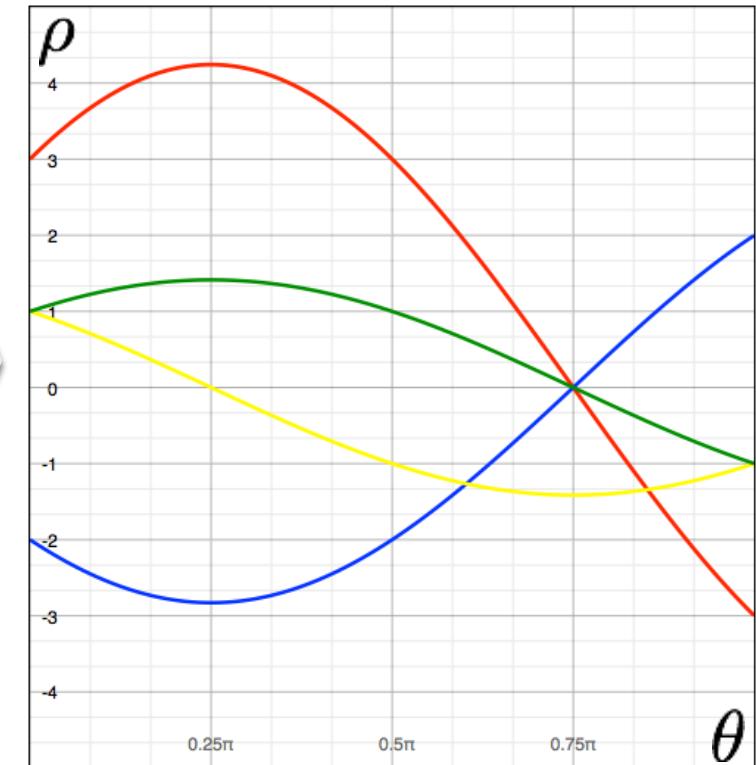
variables
parameters



four points
become
?

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables

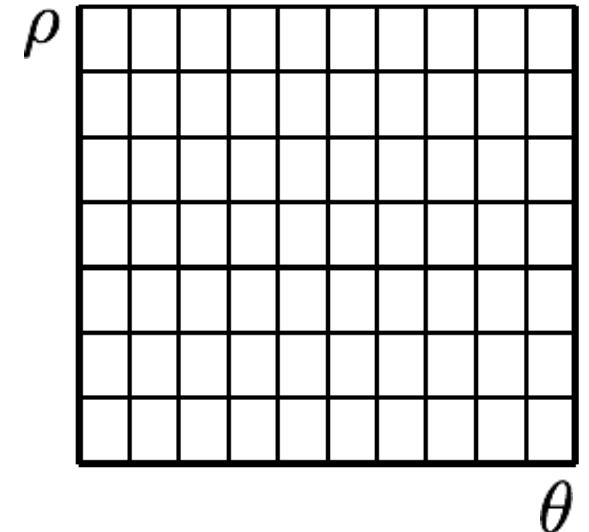


Line Detection by Hough Voting

Algorithm:

1. Quantize Parameter Space (θ, ρ) .
2. Create Hough Space Array $H(\theta, \rho) = 0$.
3. For each image point (x_i, y_i) :
For all points (θ, ρ) on $\rho = x_i \cos \theta + y_i \sin \theta$:
$$H(\theta, \rho) = H(\theta, \rho) + 1$$
4. Find local maxima $H(\theta_m, \rho_m)$.
5. The detected line: $x \cos \theta_m + y \sin \theta_m = \rho_m$

H: accumulator array (votes)



Line Detection by Hough Voting

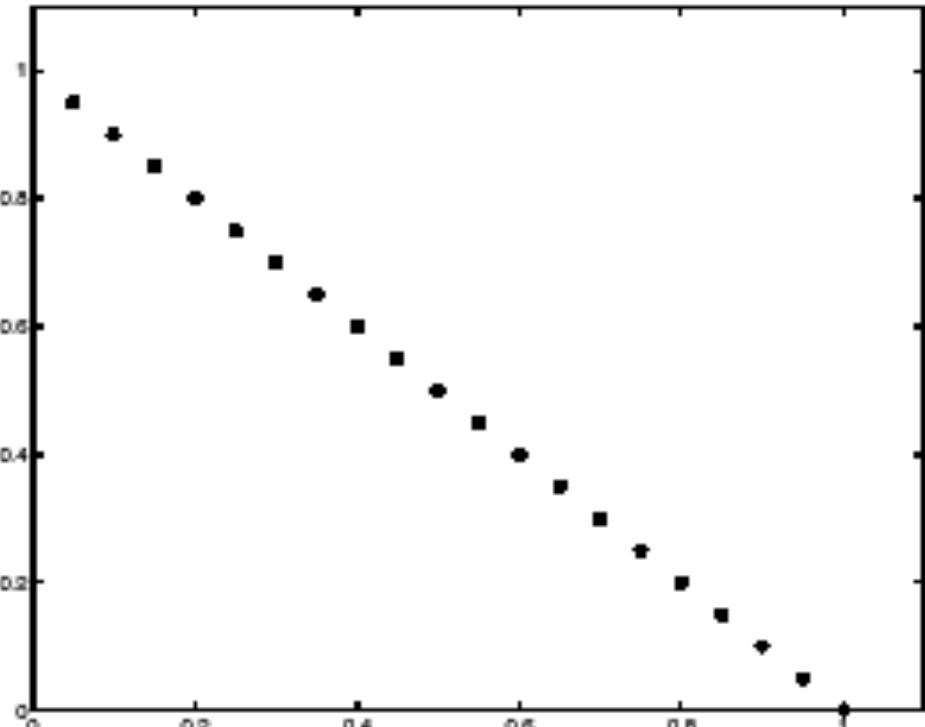
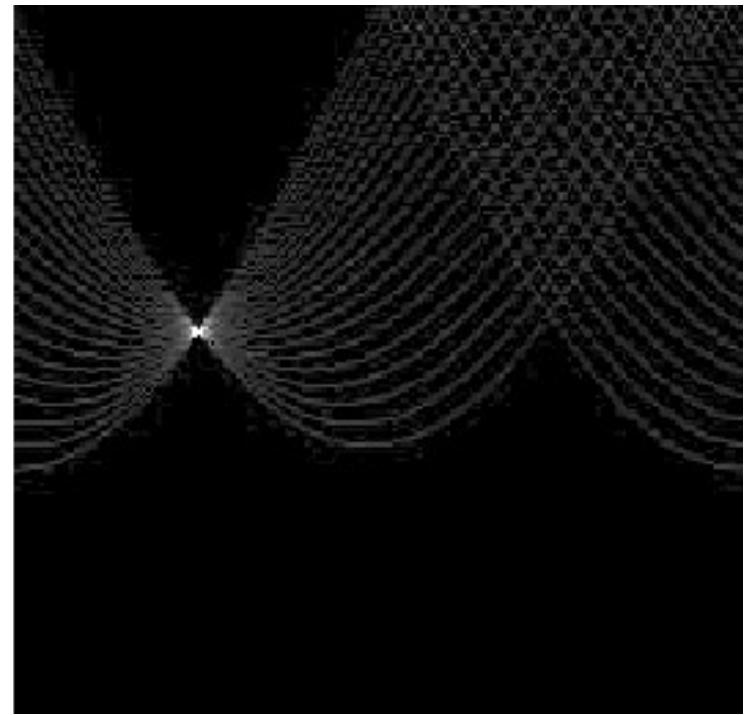
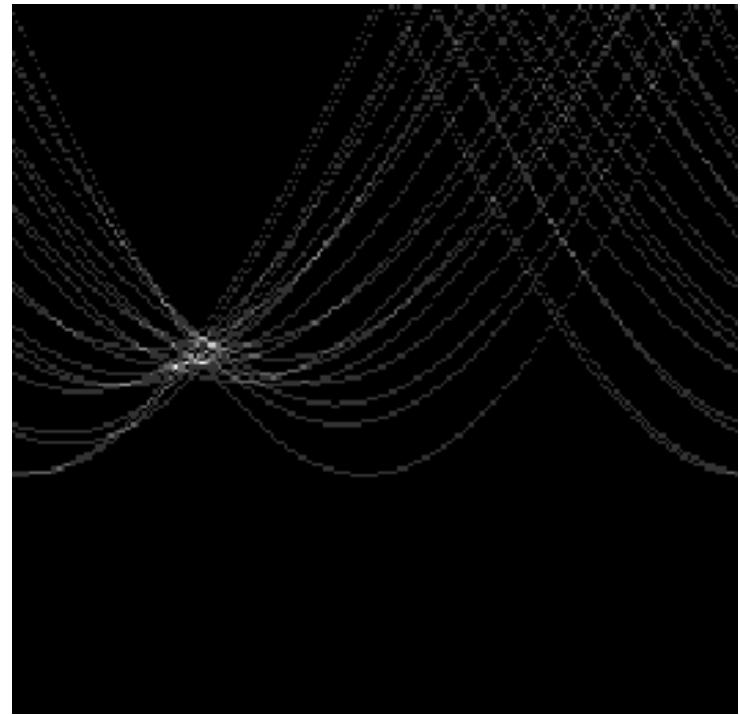
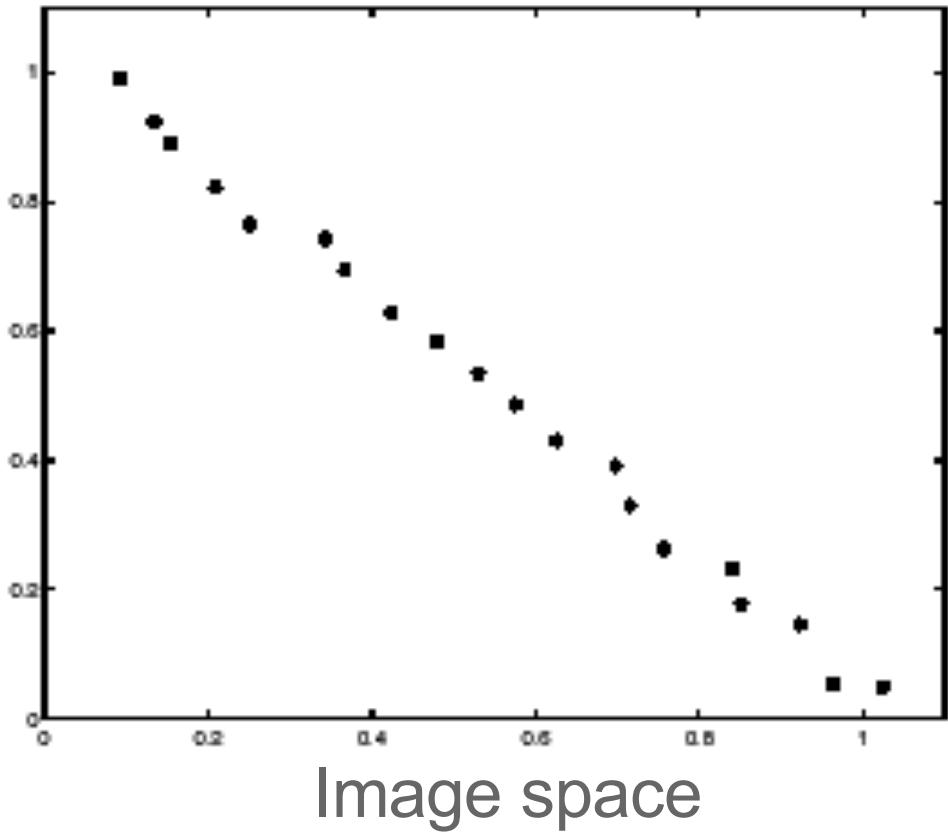


Image space



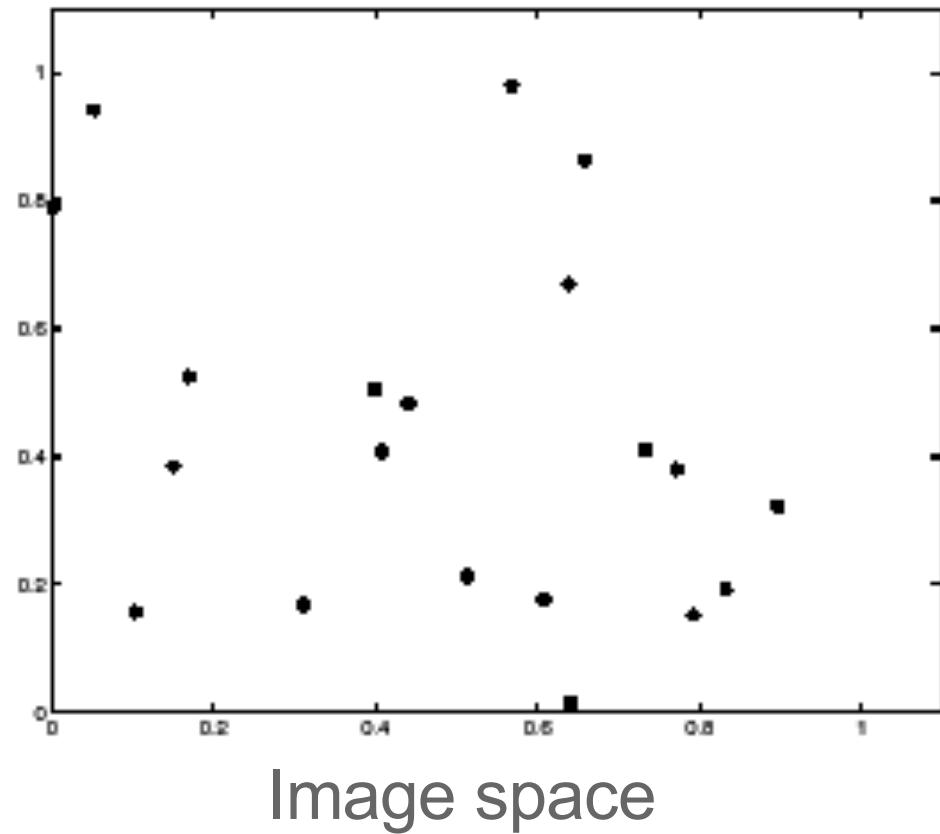
Votes

If images are noisy...

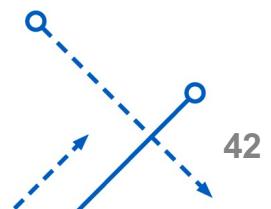
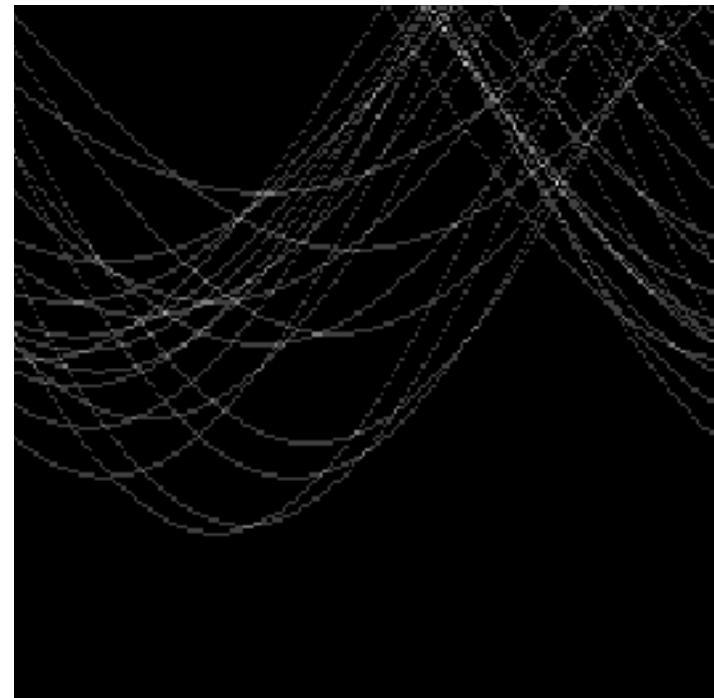


Votes

Too much noise



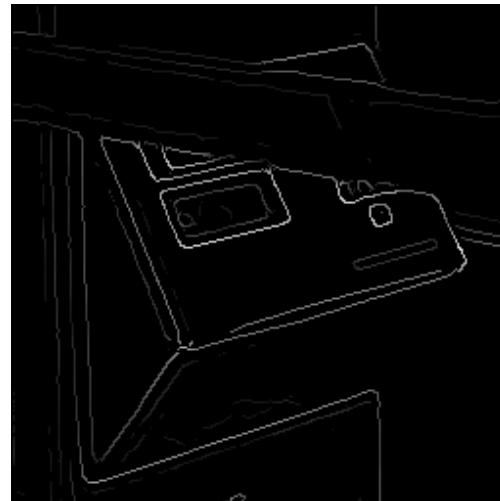
Votes



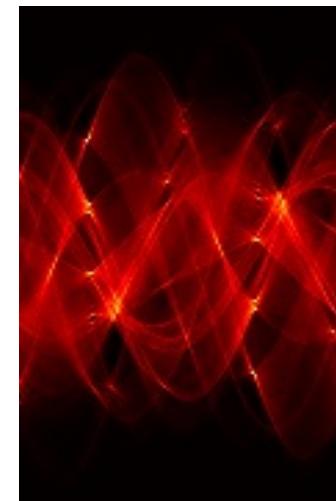
Real-world example



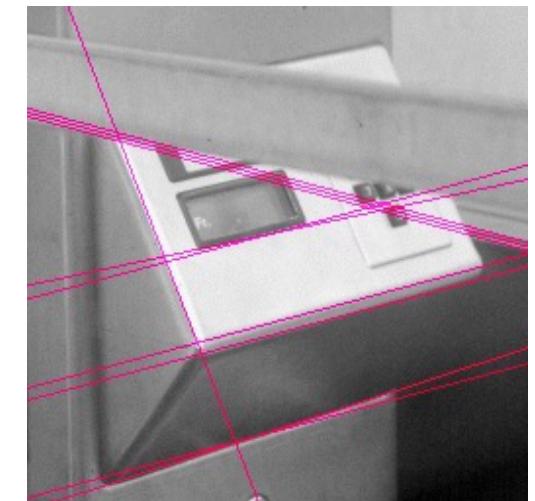
Original



Edges

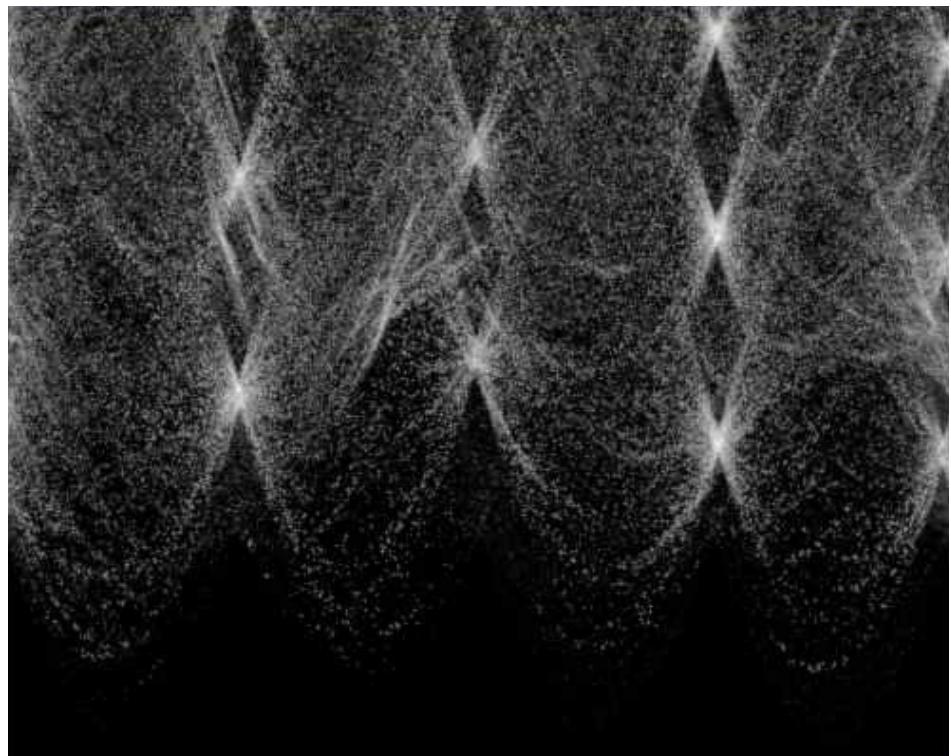


Parameter
Space



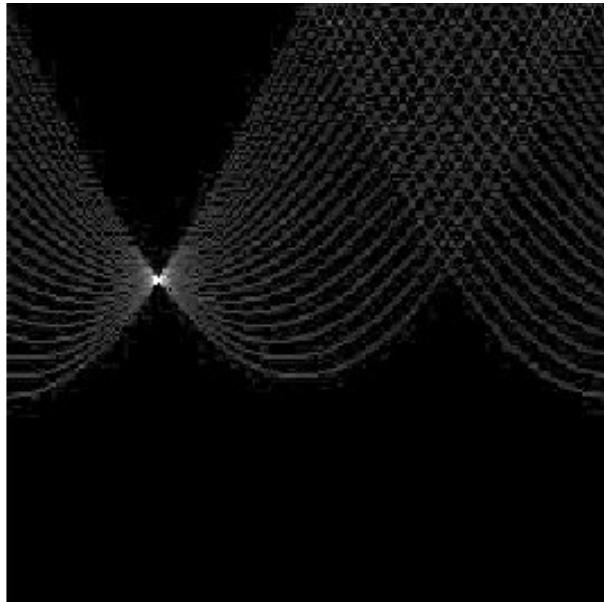
Hough Lines

More complex image

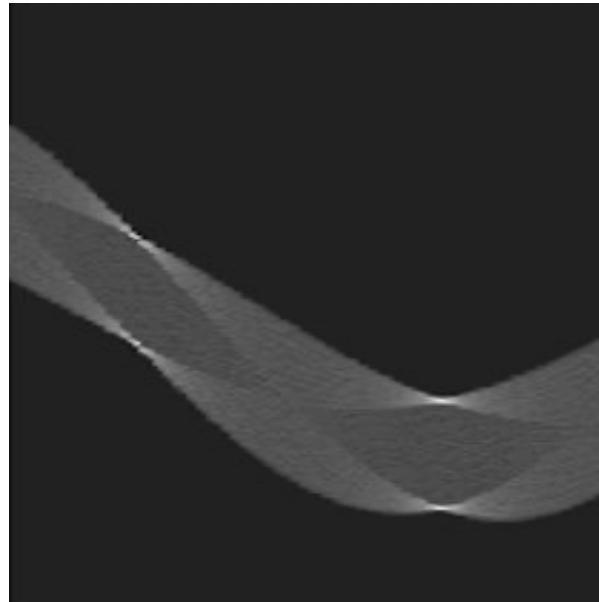


Basic Shapes

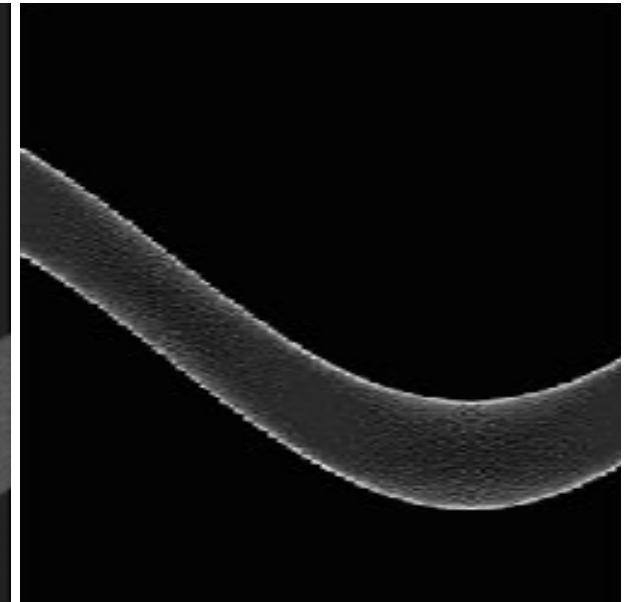
Parameter space



Line

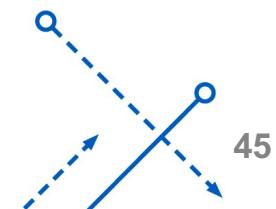


rectangle
(parallelogram)



Circle

Can you guess the shape in image space?



Hough Circles

Let's assume known radius

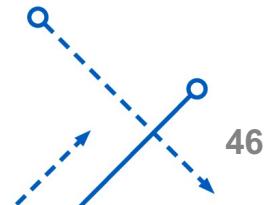
$$(x - a)^2 + (y - b)^2 = r^2$$

parameters
variables
Fixed

$$(x - a)^2 + (y - b)^2 = r^2$$

parameters
variables
Fixed

What is the dimension of the parameter space?



Hough Circles

parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

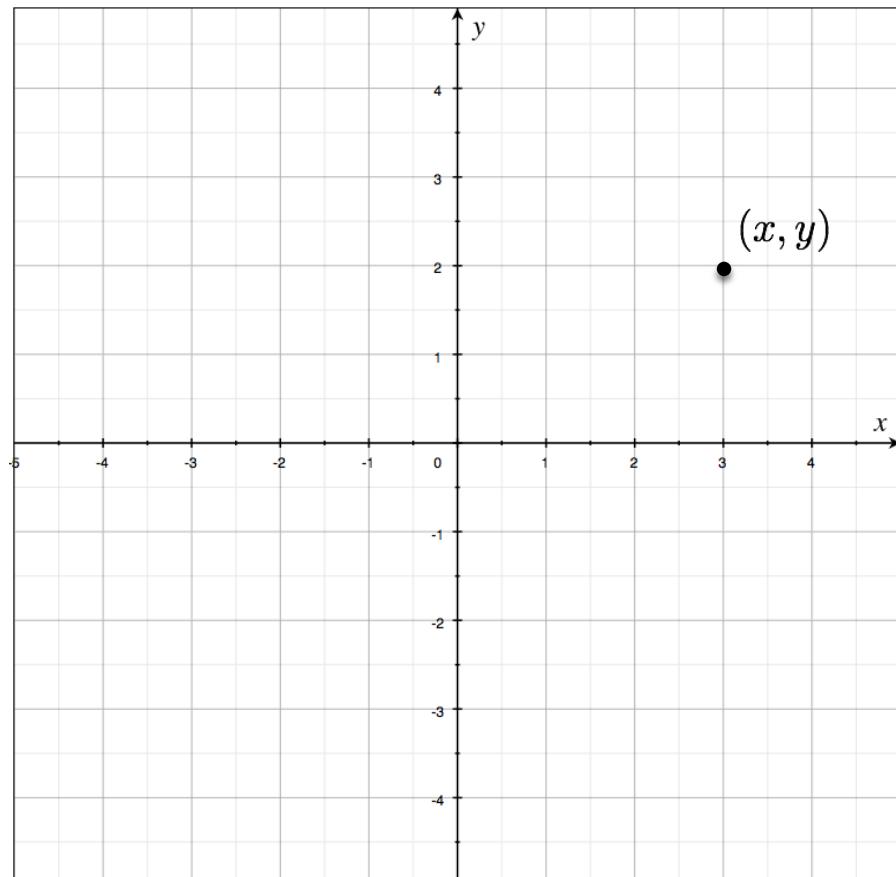
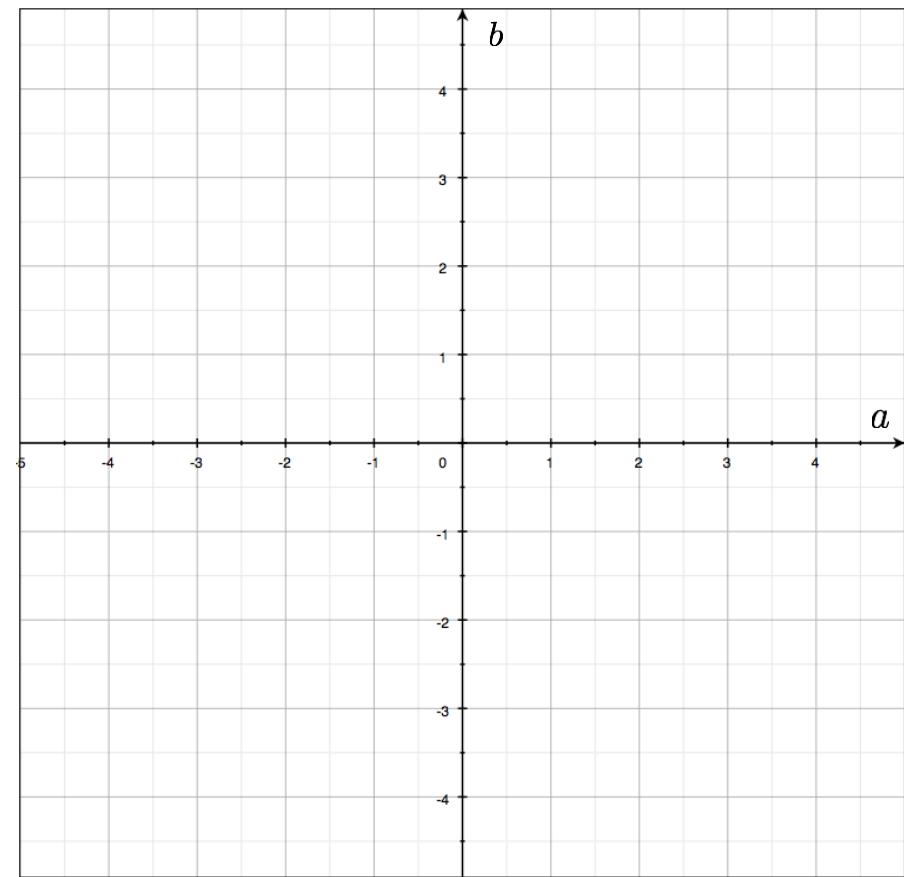


Image space

parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

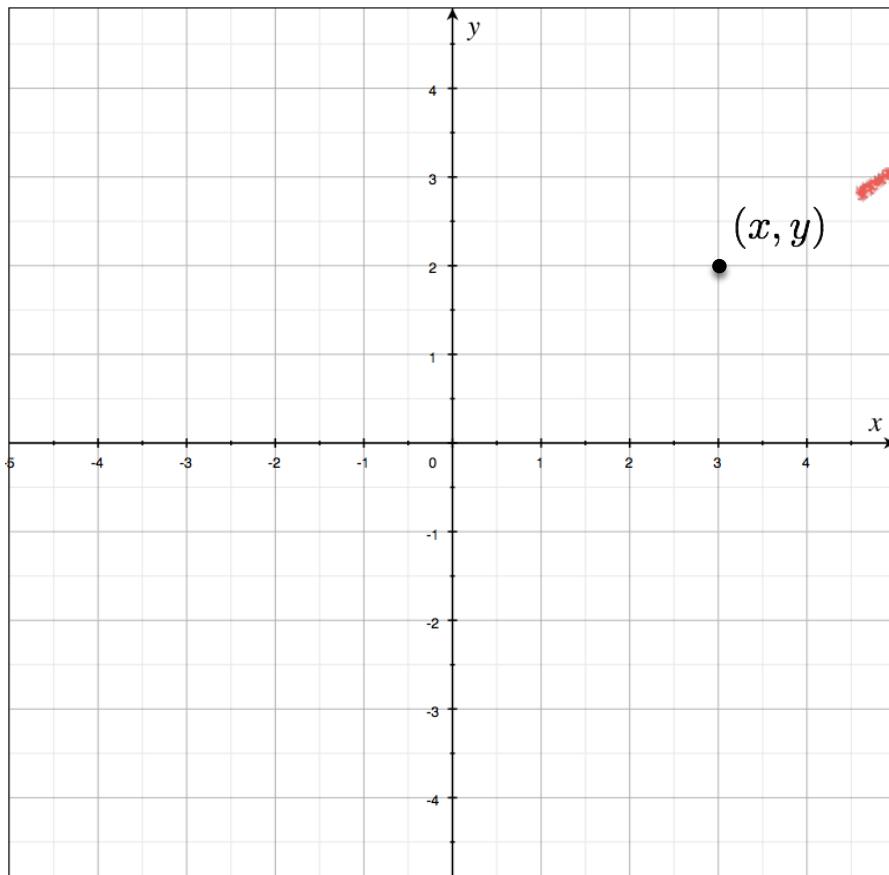


Parameter space

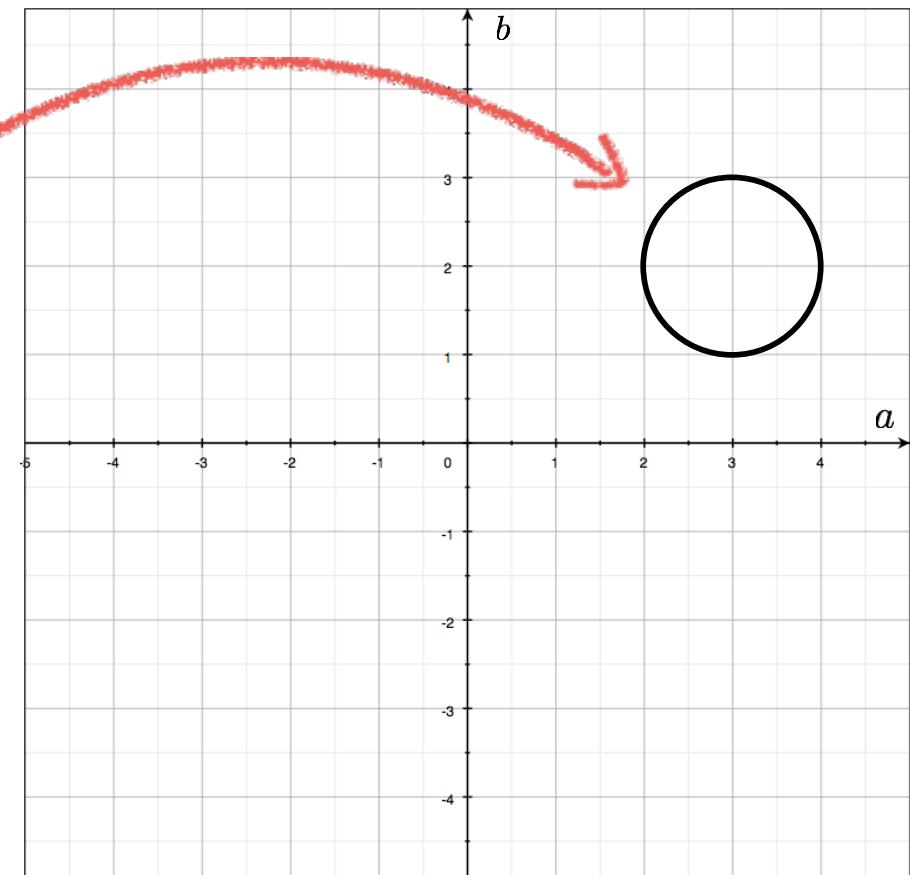
What does a point in image space correspond to in parameter space?

Hough Circles

parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

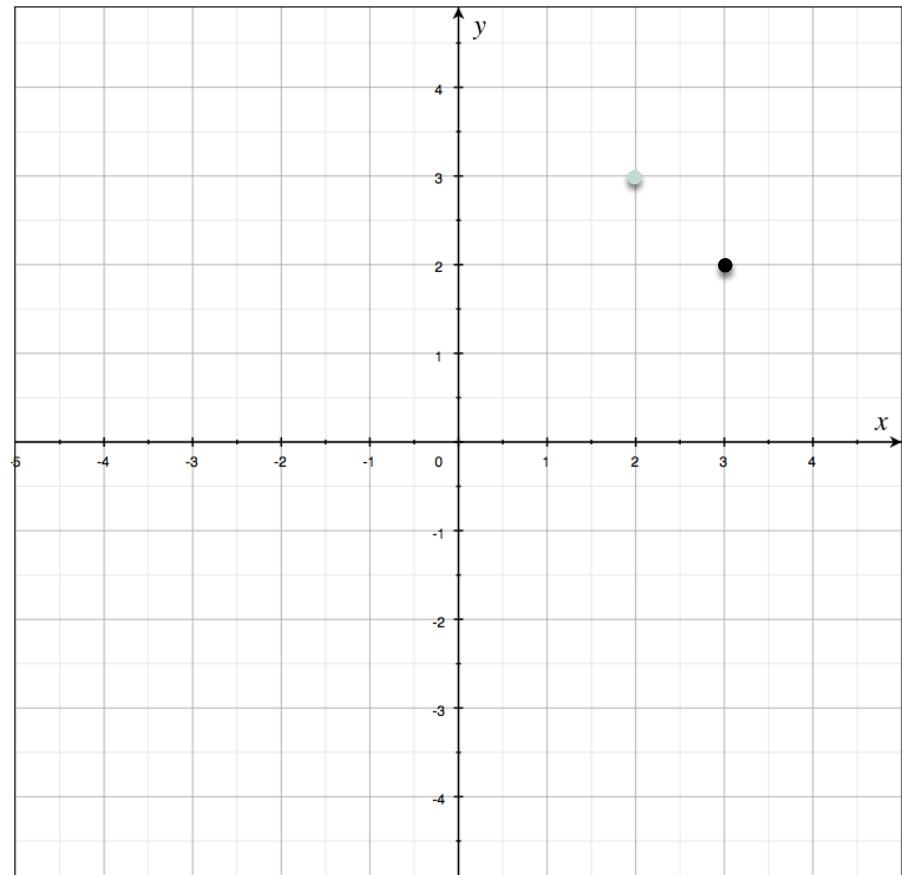


parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

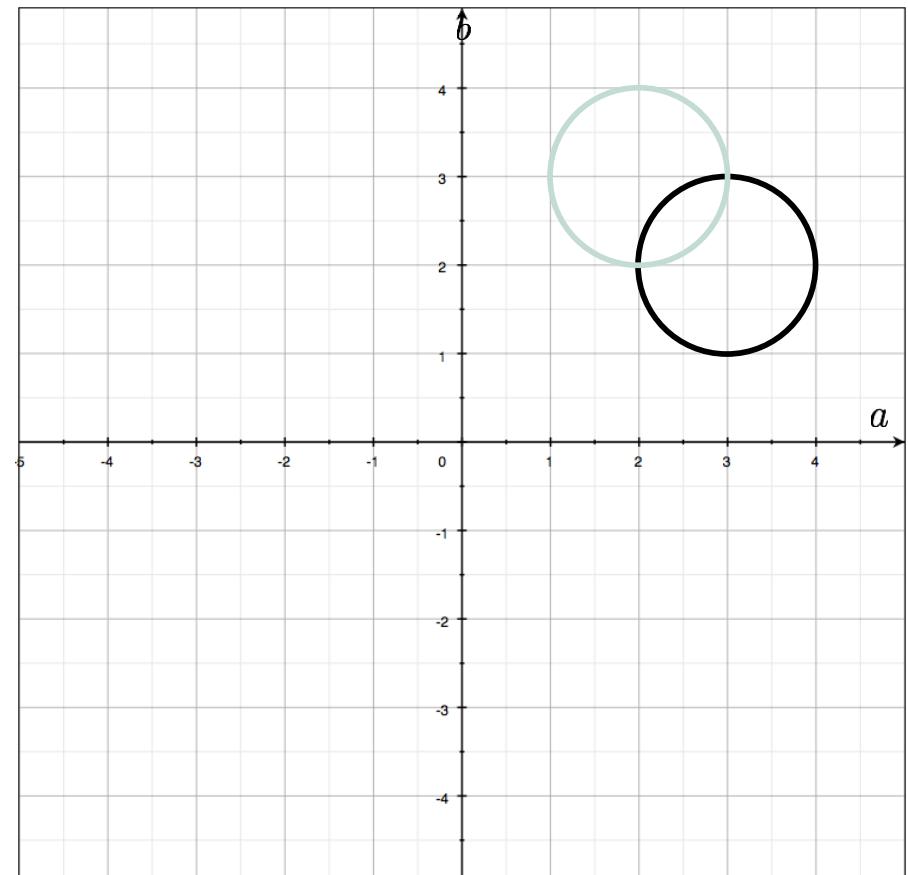


Hough Circles

parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

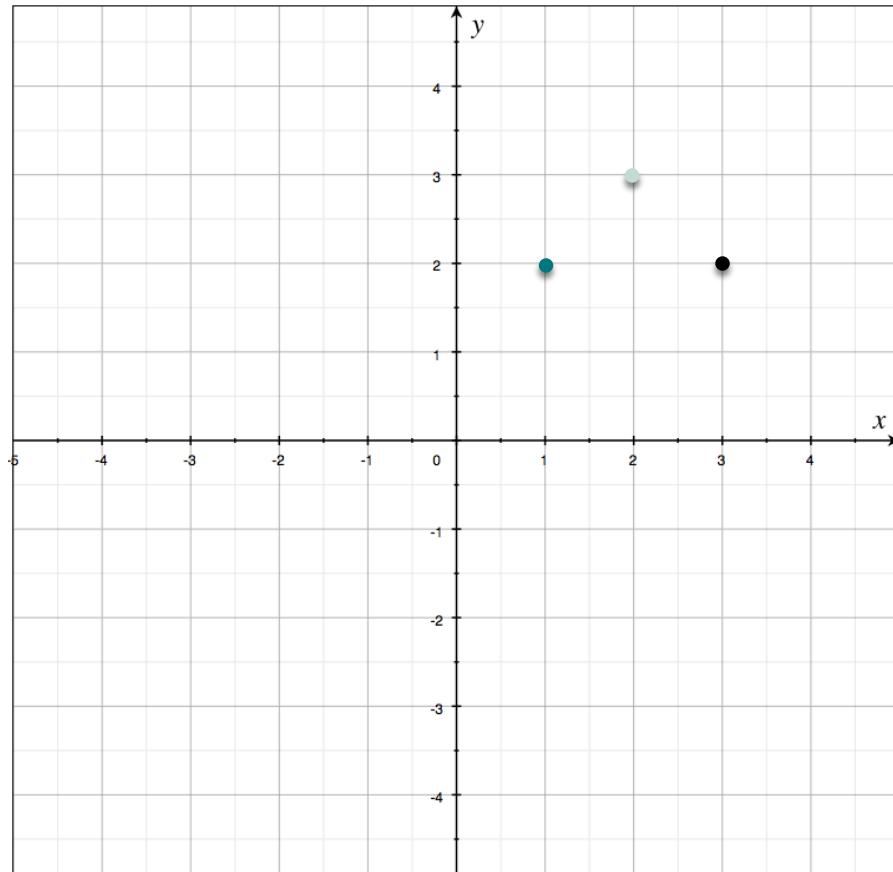


parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

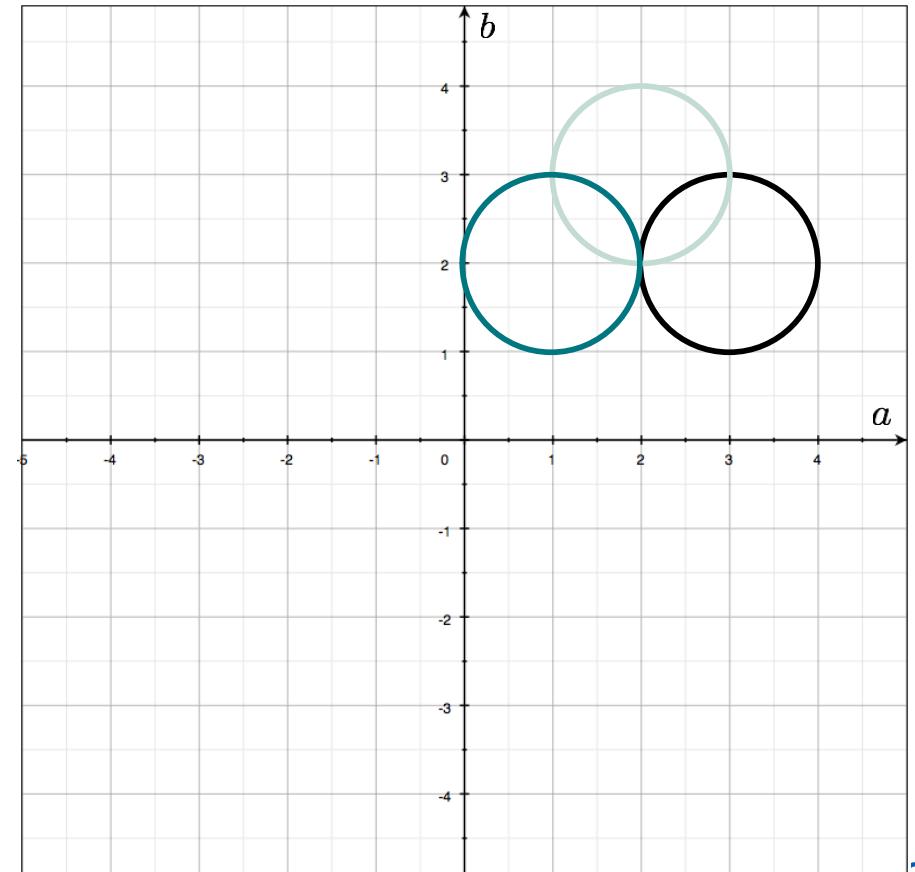


Hough Circles

parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

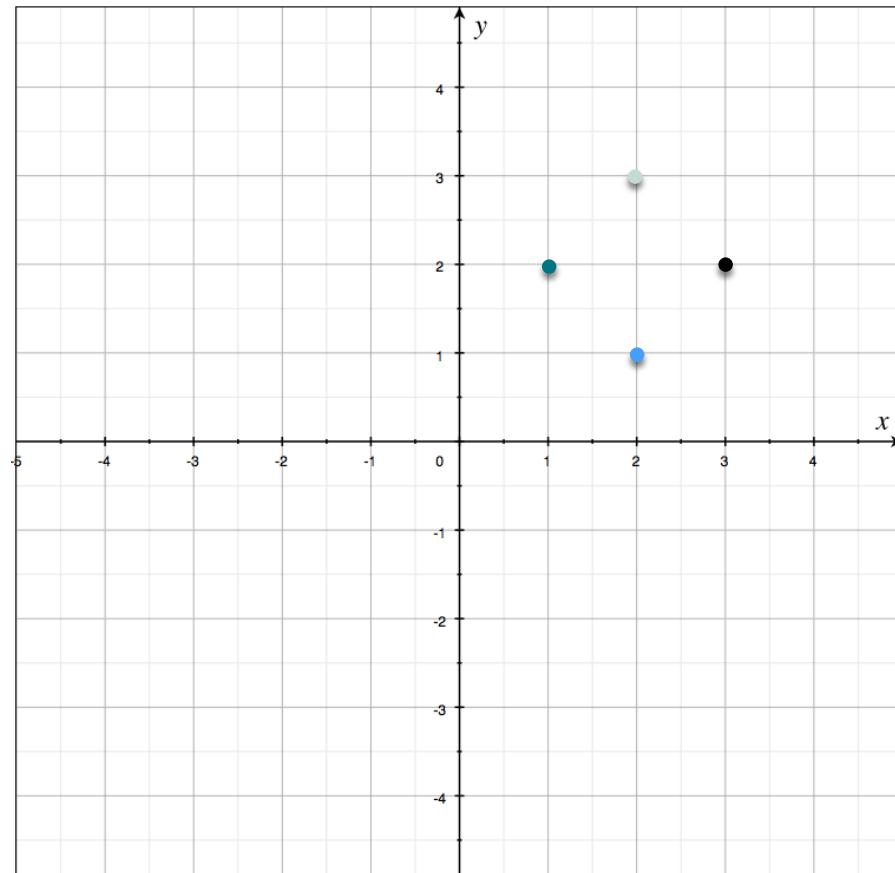


parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

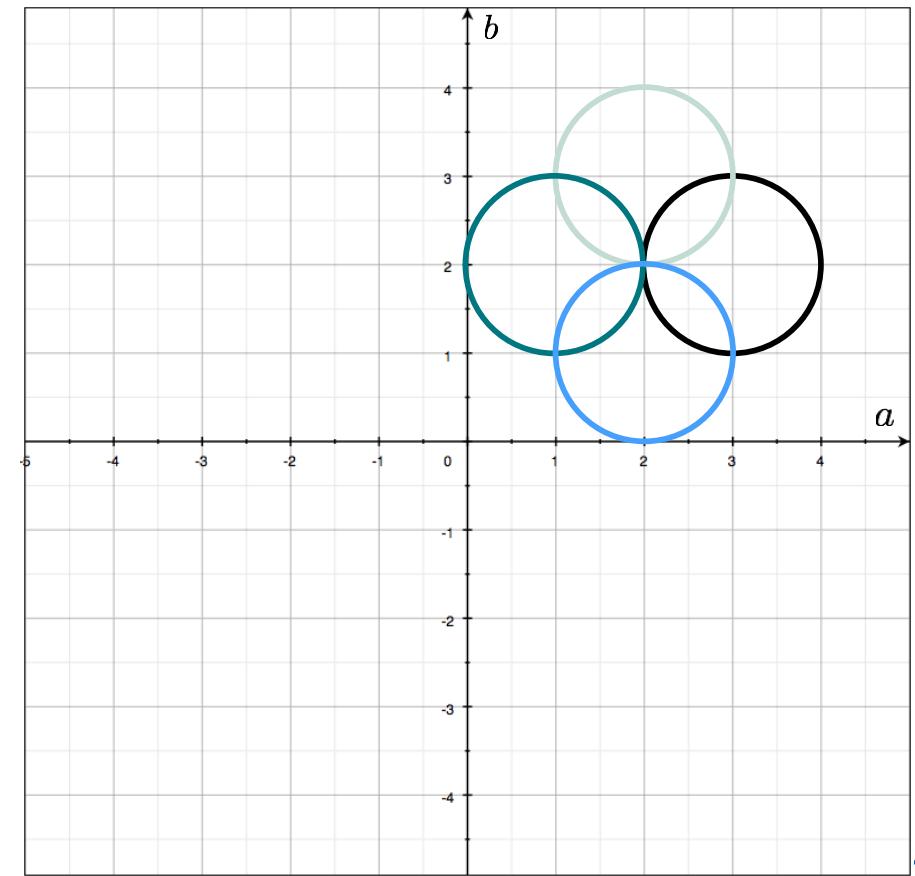


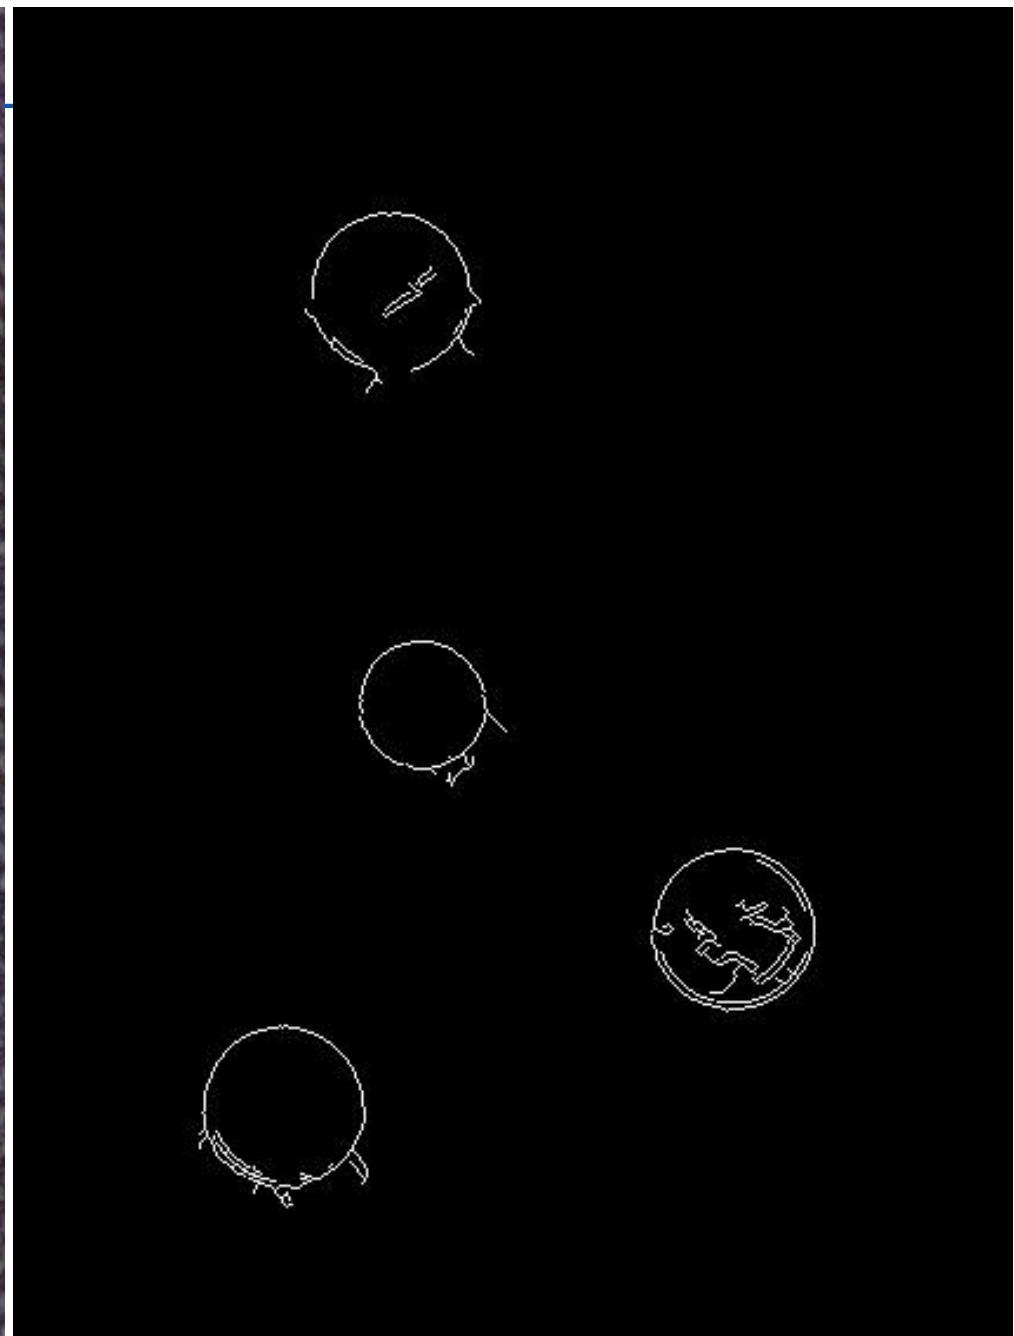
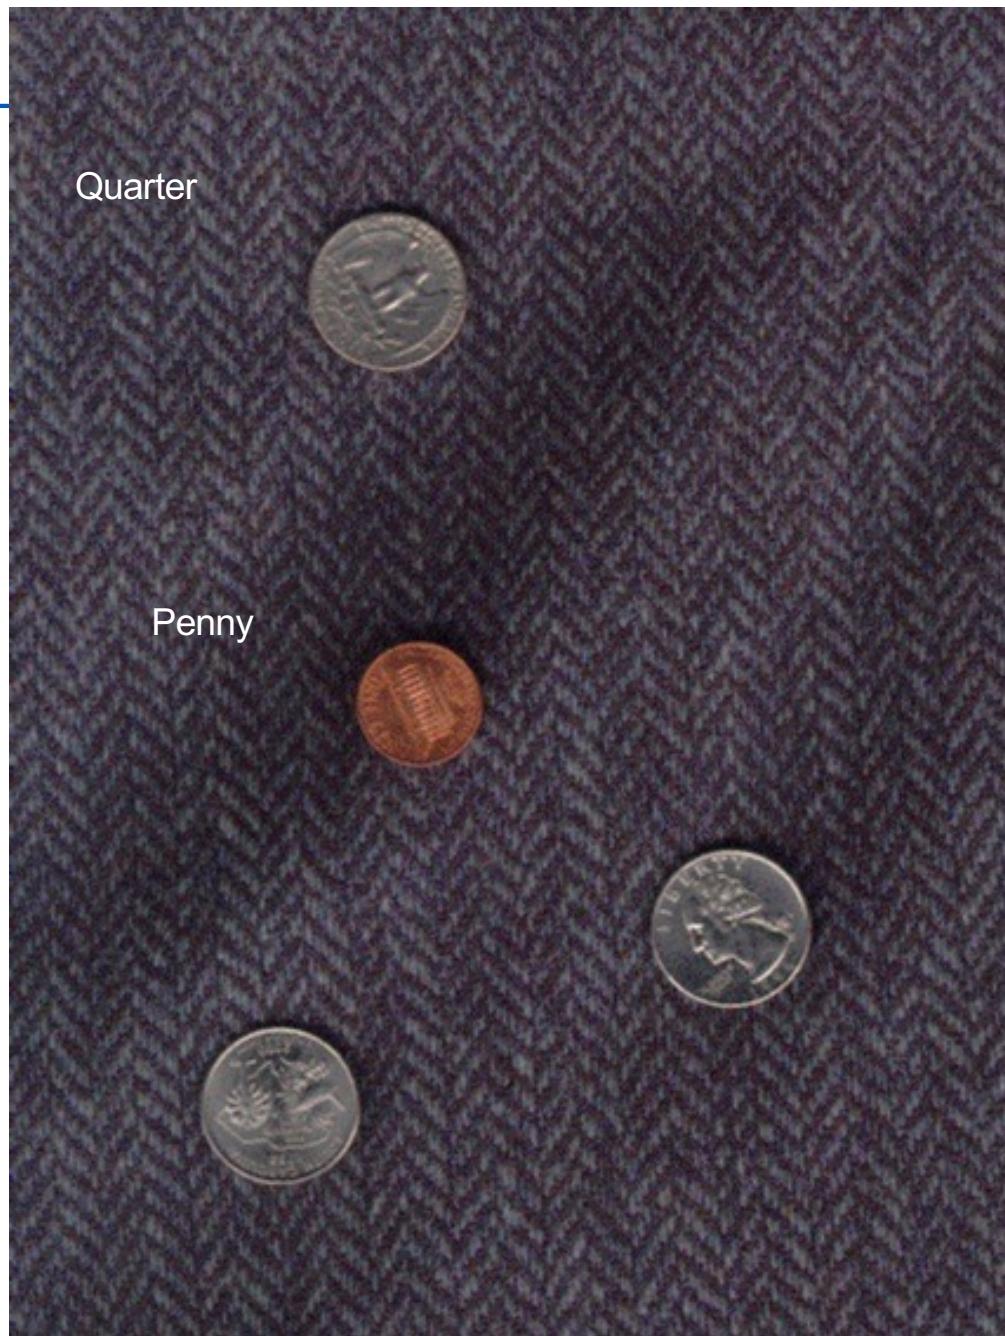
Hough Circles

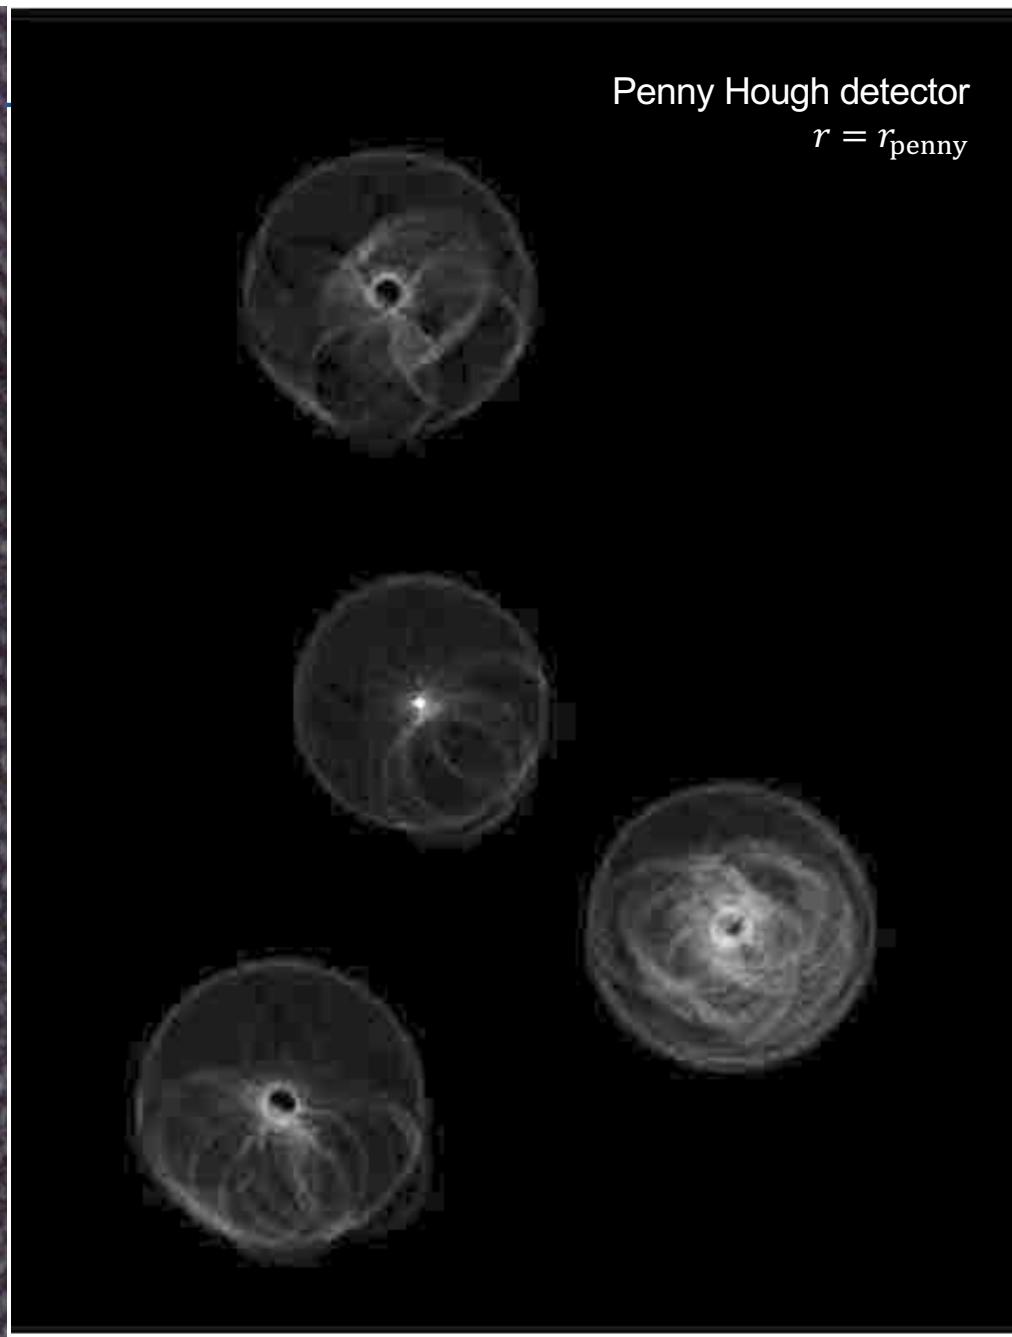
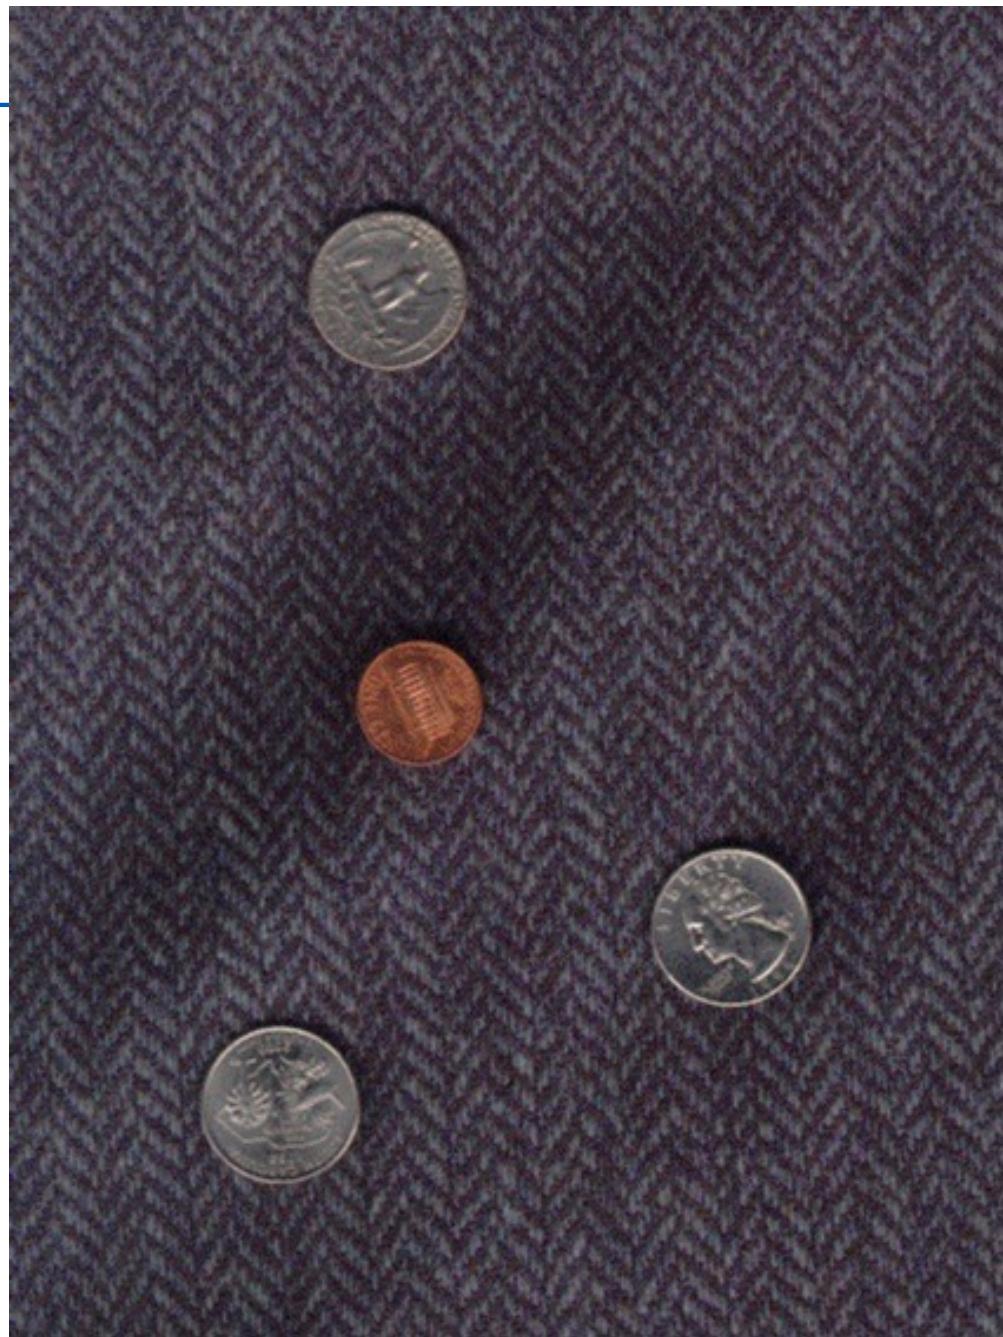
parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables

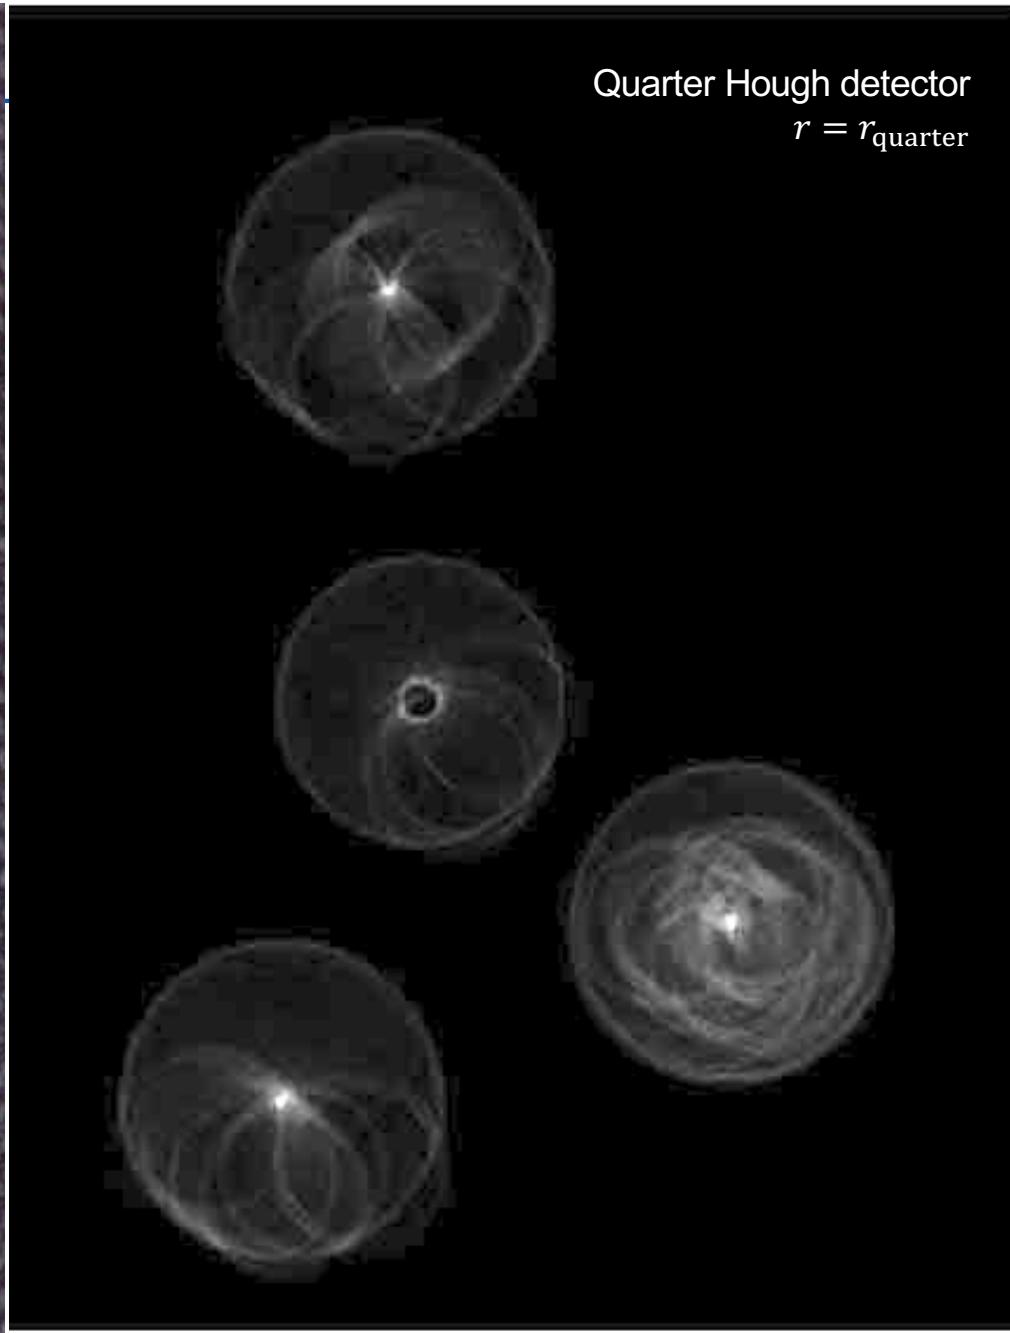


parameters
 $(x - a)^2 + (y - b)^2 = r^2$
variables



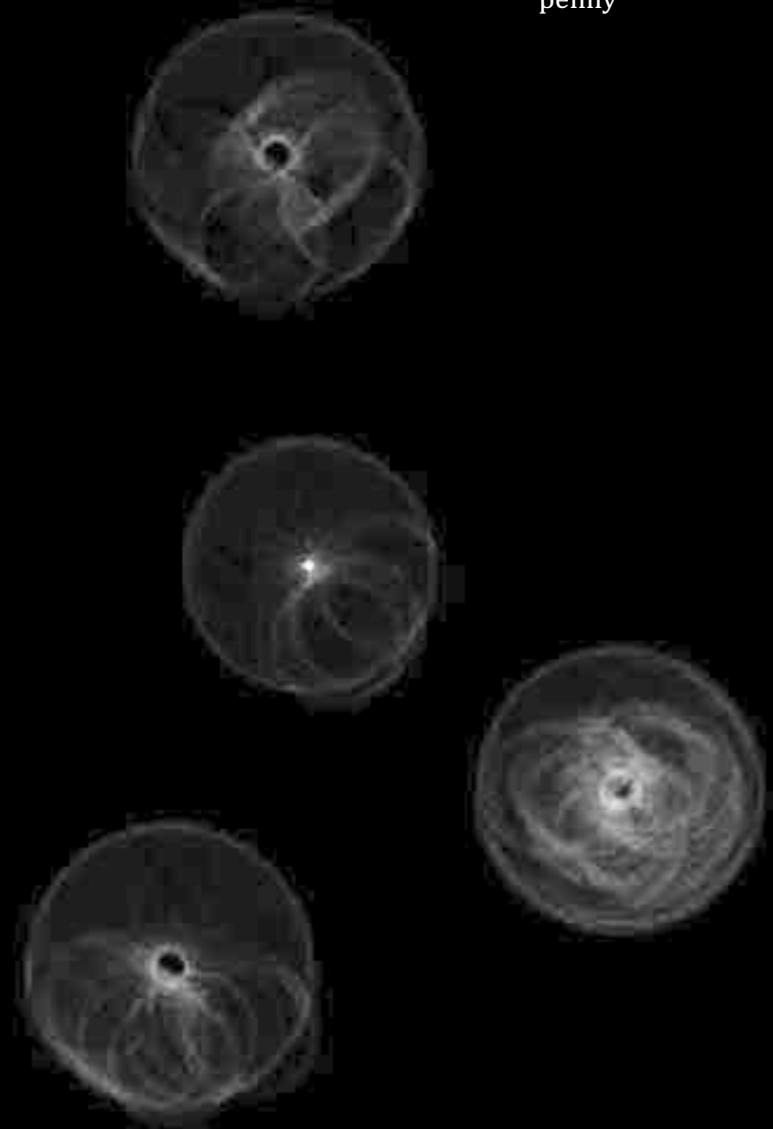






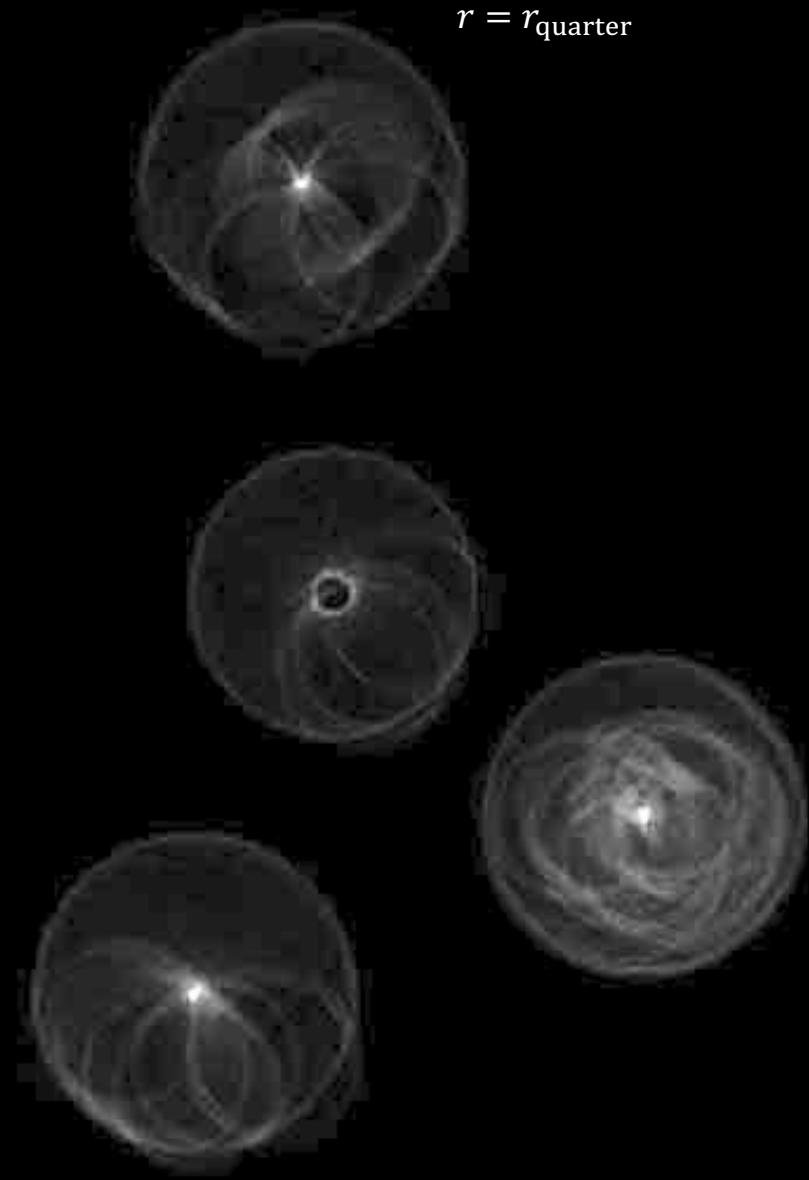
Penny Hough detector

$$r = r_{\text{penny}}$$



Quarter Hough detector

$$r = r_{\text{quarter}}$$



Penny Hough detector

$$r = r_{\text{penny}}$$



Quarter Hough detector

$$r = r_{\text{quarter}}$$



What if radius is unknown?

$$(x - a)^2 + (y - b)^2 = r^2$$

parameters
variables

$$(x - a)^2 + (y - b)^2 = r^2$$

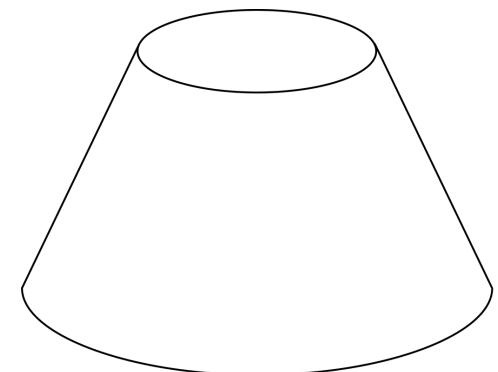
parameters
variables

If radius is unknown:

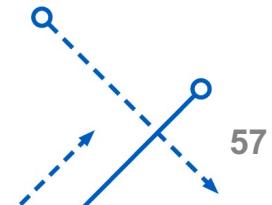
3D Hough Space!

Use Hough array $H(a, b, r)$.

Surface shape in Hough space is complicated.



Frustum of cone



Other Shapes?

Vertical Ellipse:

$$\frac{(x - x_0)^2}{a^2} + \frac{(y - y_0)^2}{b^2} = 1$$

↑ ↑

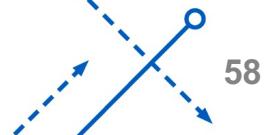
parameters

$H(x_0, y_0, a, b)$

Ellipse:

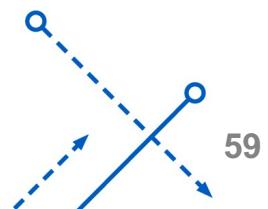
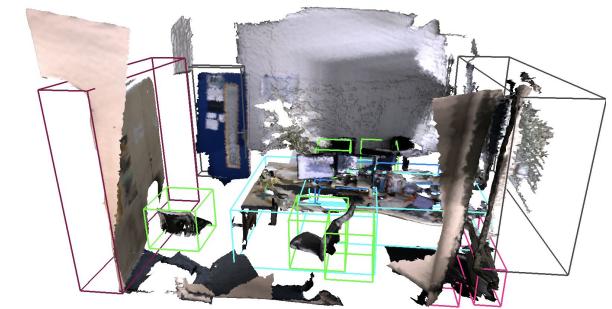
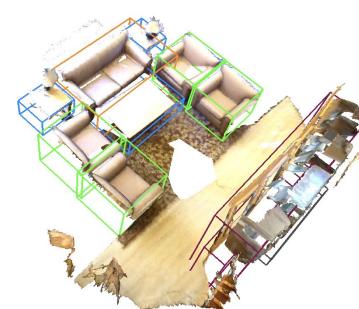
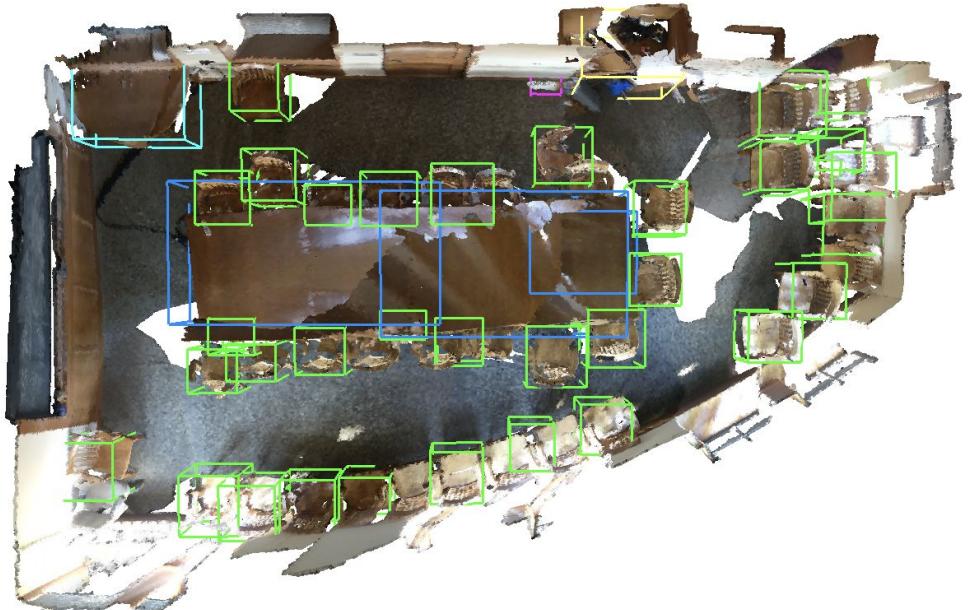
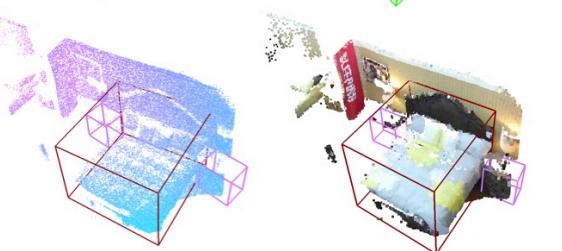
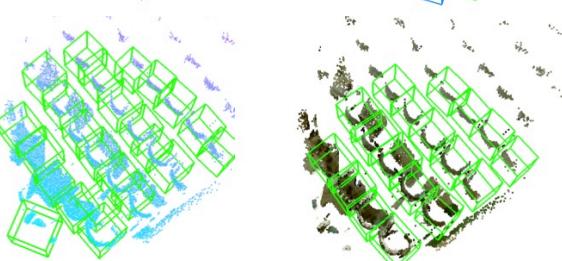
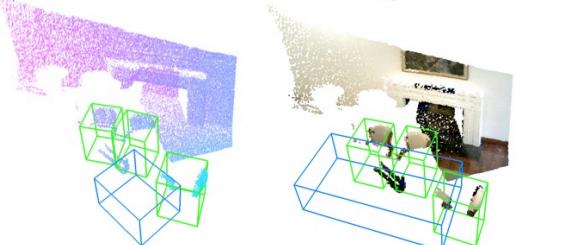
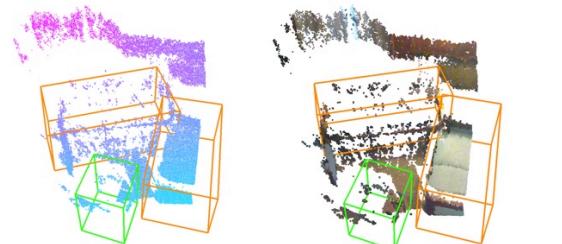
$$\frac{[(x - x_0) \cos \theta + (y - y_0) \sin \theta]^2}{a^2} + \frac{[-(x - x_0) \sin \theta + (y - y_0) \cos \theta]^2}{b^2} = 1$$

$H(x_0, y_0, a, b, \theta)$



Applications of Hough Voting

Scenes Prediction Ground Truth



Conclusion

Is the following correct about Hough transform ...



- Detects multiple instances (lines/circles)?



- Robust to noise?



- Can be used for other shapes beyond lines/circles?



- Good computational complexity?



- Deals with occlusion well?

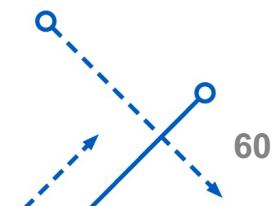




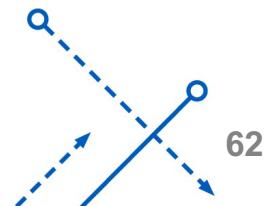
IMAGE PROCESSING

Alignment and Fitting



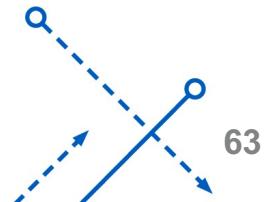
What are Fitting and Alignment?

- Fitting
 - find the parameters of a model that best fit the data
- Alignment
 - find the parameters of the transformation that best align matched points

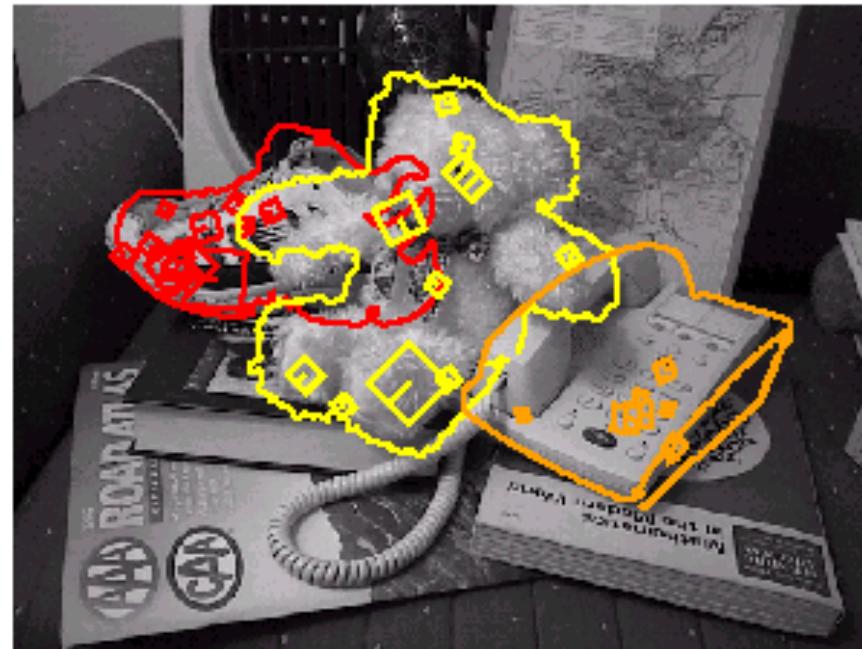


Fitting and Alignment: Methods

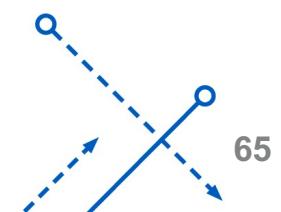
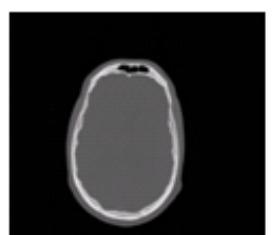
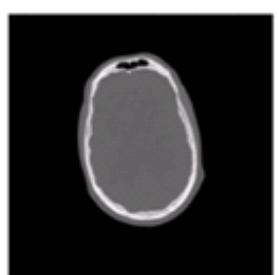
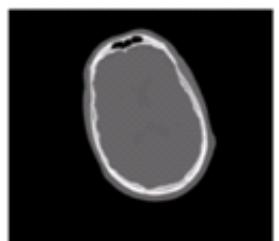
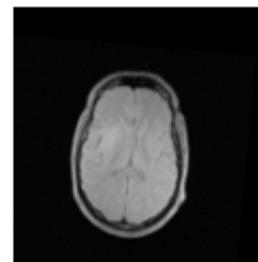
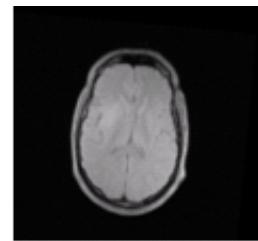
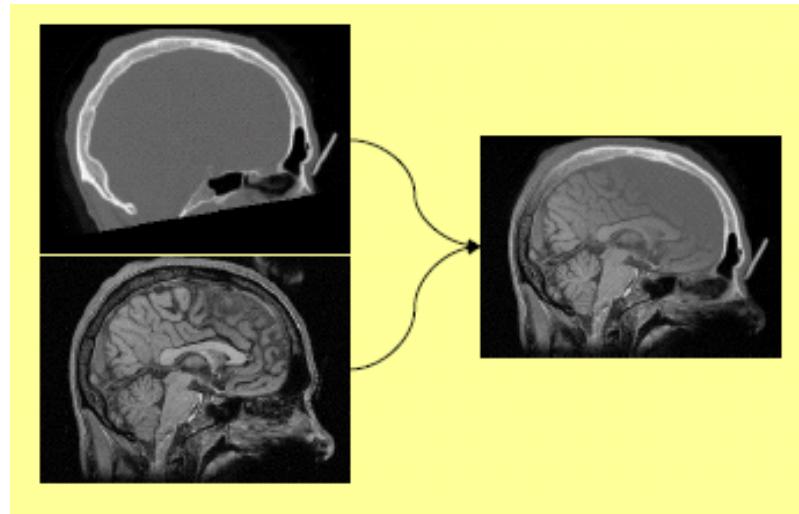
- Hypothesize and test
 - Generalized Hough transform
- General Alignment
 - Homographies
 - Rotational Panoramas
 - RANSAC
 - Global alignment
 - Warping
 - Blending
- Global optimization / Search for parameters
 - Least squares fit
 - Robust least squares
 - Other parameter search methods



Motivation: Recognition



Motivation: Medical image registration



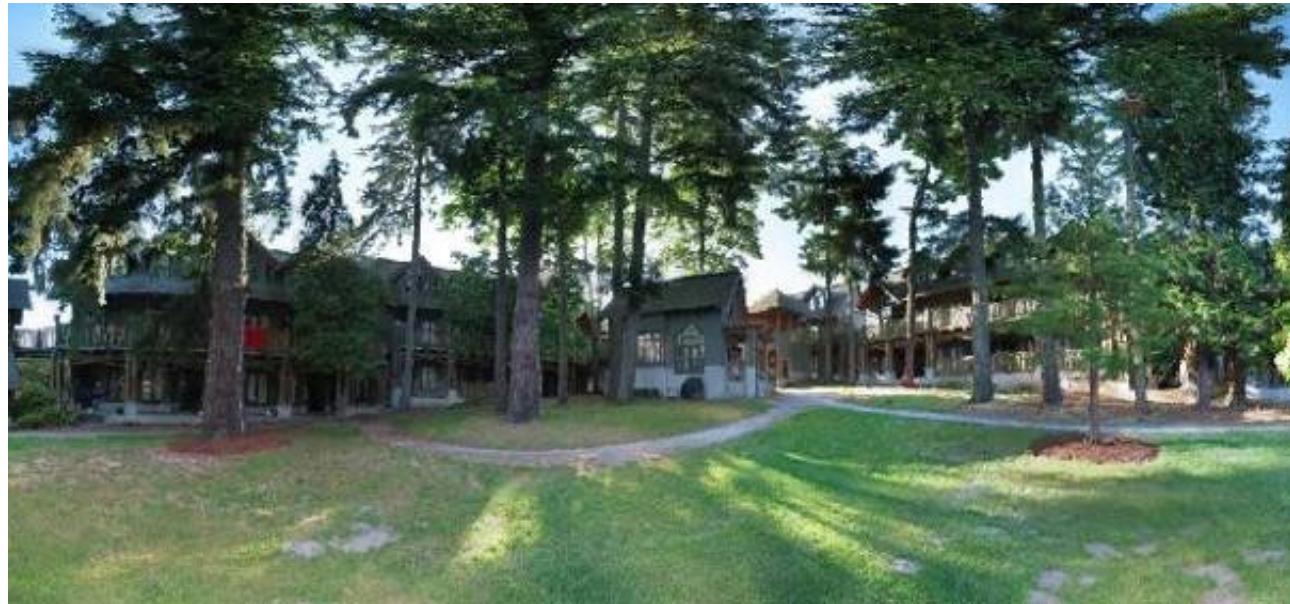
Motivation: Mosaics

- Getting the whole picture
 - Typical camera: $50^\circ \times 35^\circ$



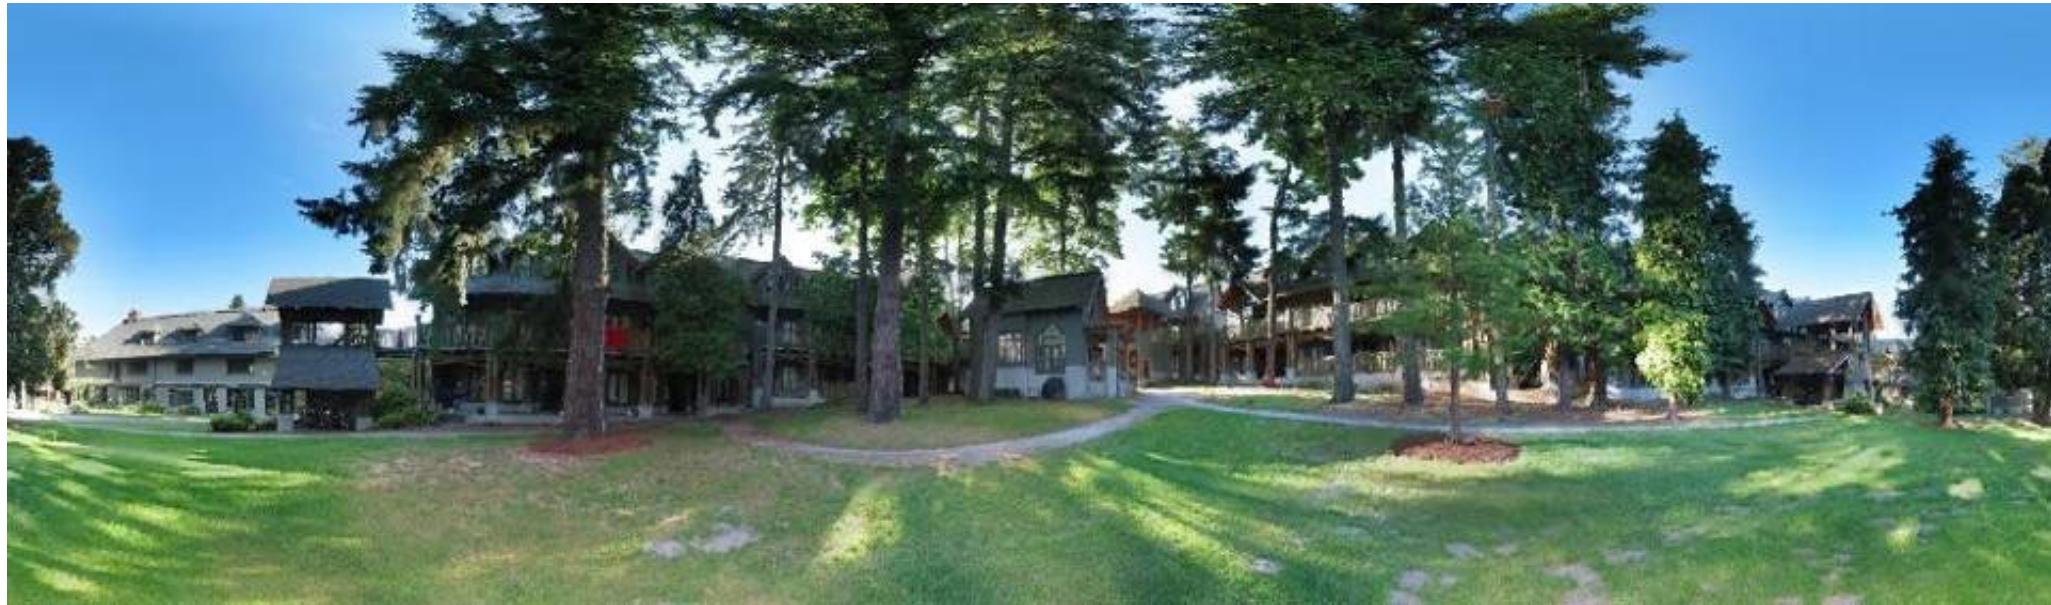
Motivation: Mosaics

- Getting the whole picture
 - Typical camera: $50^\circ \times 35^\circ$
 - Human Vision: $176^\circ \times 135^\circ$



Motivation: Mosaics

- Getting the whole picture
 - Typical camera: $50^\circ \times 35^\circ$
 - Human Vision: $176^\circ \times 135^\circ$



Alignment

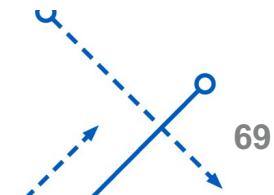
- Homography
- Rotational Panoramas
- RANSAC
- Global alignment
- Warping
- Blending



(a)



(b)



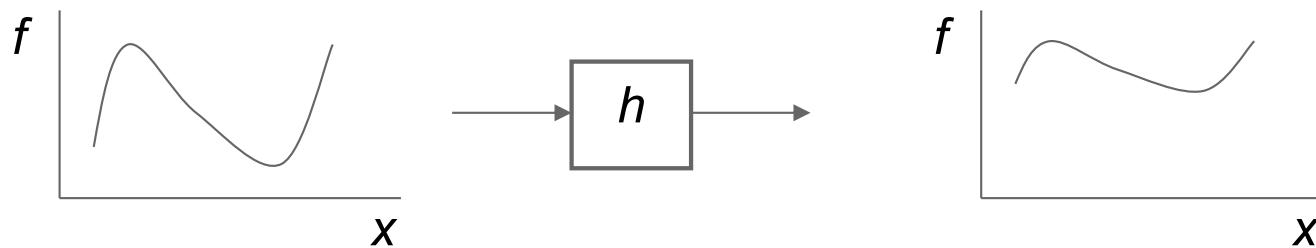
Motion models

- What happens when we take two images with a camera and try to align them?
- translation?
- rotation?
- scale?
- affine?
- perspective?



Image Warping (Recap)

- image filtering: change *range* of image
 - $g(x) = h(f(x))$



- image warping: change *domain* of image
 - $g(x) = f(h(x))$

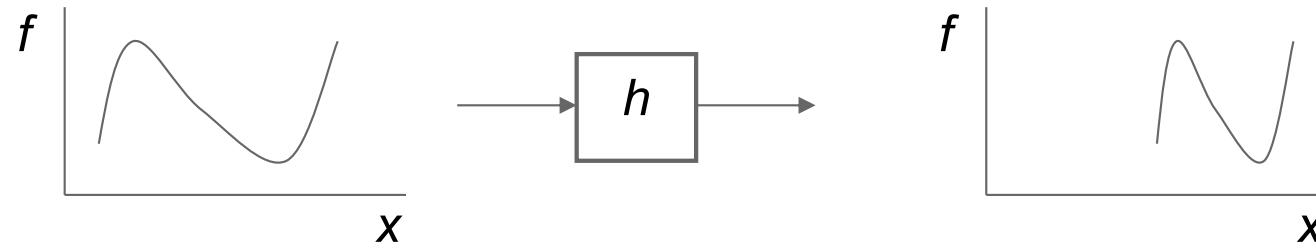
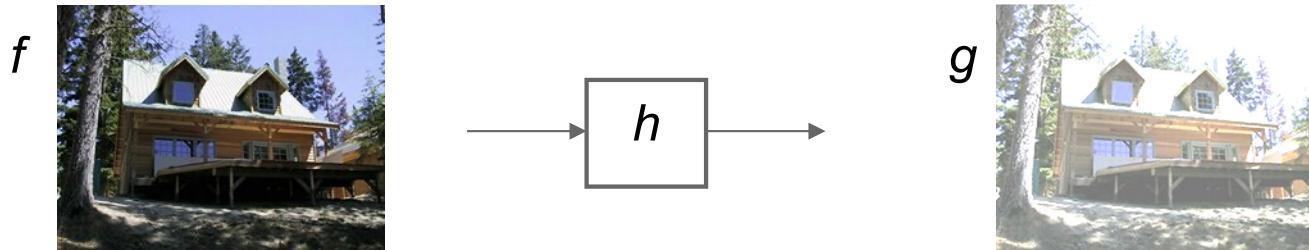
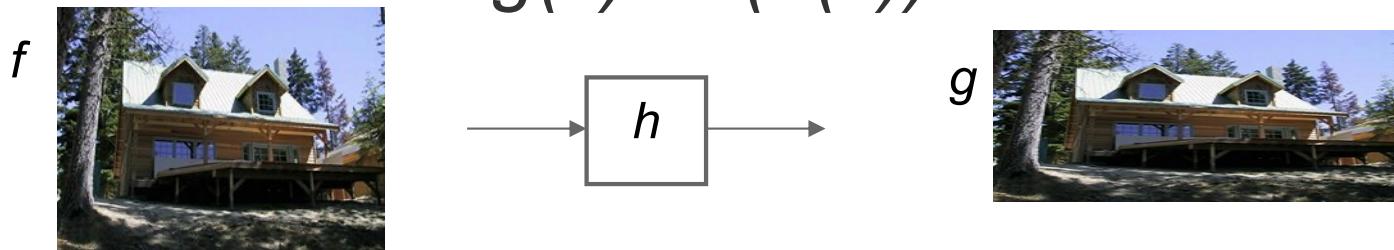


Image Warping

- image filtering: change *range* of image
 - $g(x) = h(f(x))$



- image warping: change *domain* of image



Parametric (global) warping

- Examples of parametric warps:



translation



rotation



aspect



affine



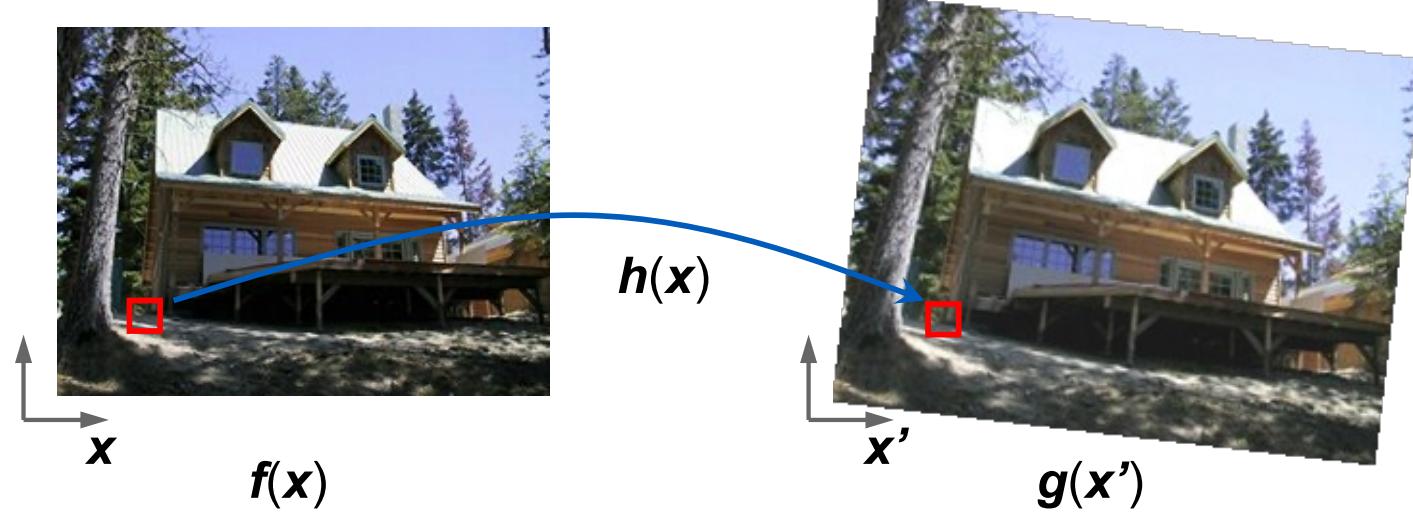
perspective



cylindrical

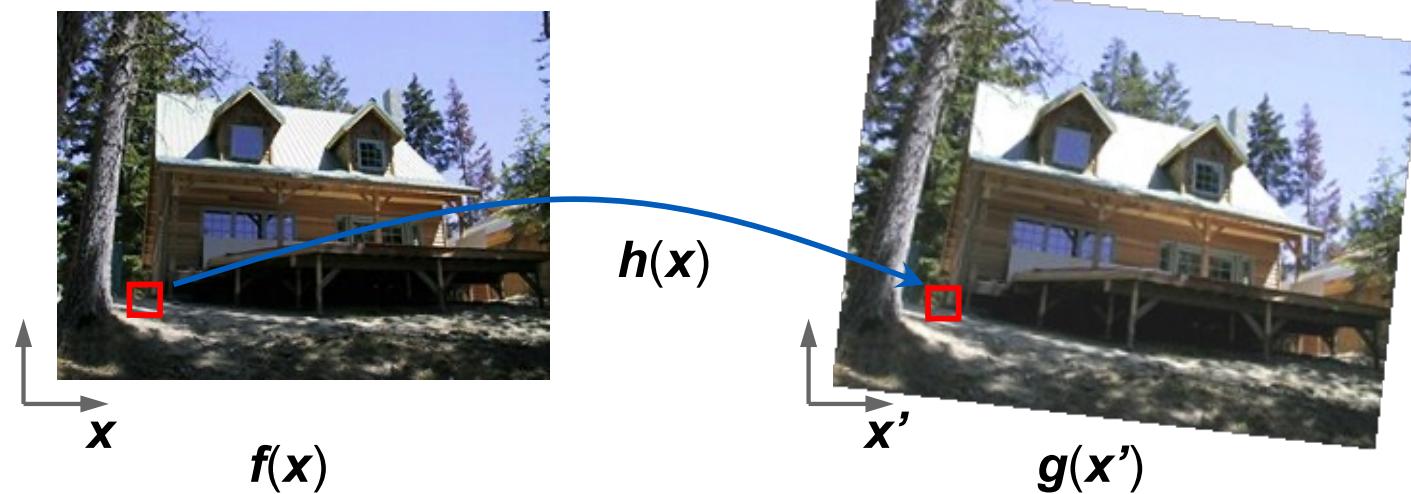
Image Warping

- Given a coordinate transform $\mathbf{x}' = h(\mathbf{x})$ and a source image $f(\mathbf{x})$, how do we compute a transformed image $g(\mathbf{x}') = f(h(\mathbf{x}))$?



Forward Warping

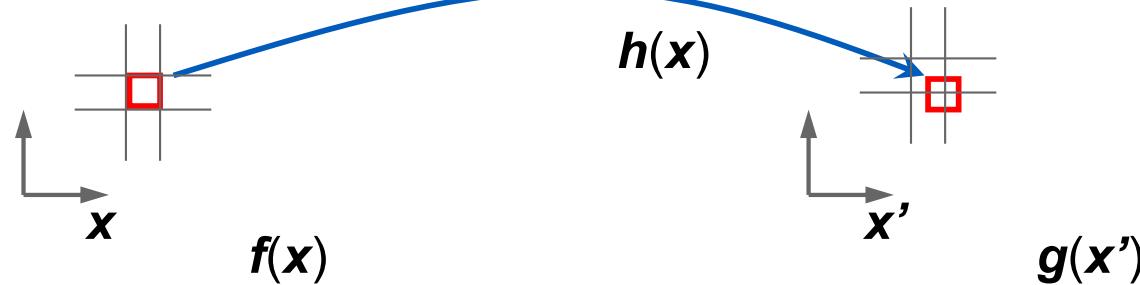
- Send each pixel $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$
- What if pixel lands “between” two pixels?



Forward Warping

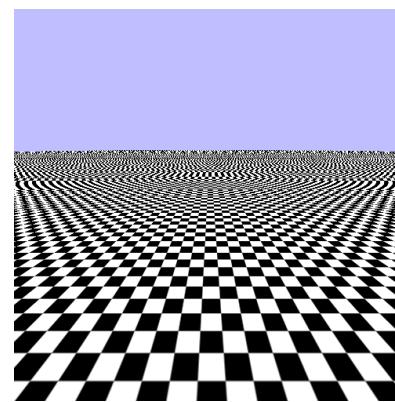
- Send each pixel $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$

- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels, normalize later



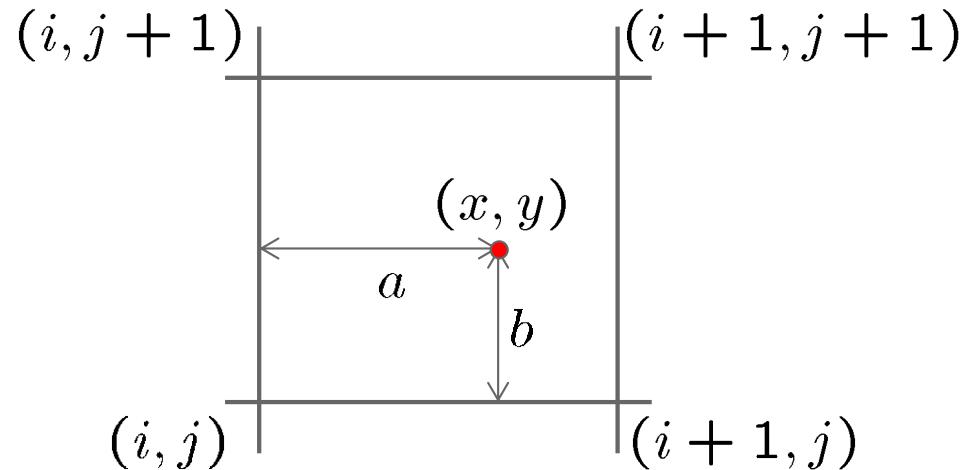
Interpolation

- Possible interpolation filters:
 - **Nearest Neighbor**
 - **Bilinear**
 - **Bicubic**
 - **Sinc / FIR (band-limited)**
- Needed to prevent “jaggies” and “texture crawl”



Bilinear interpolation

Sampling at $f(x, y)$:

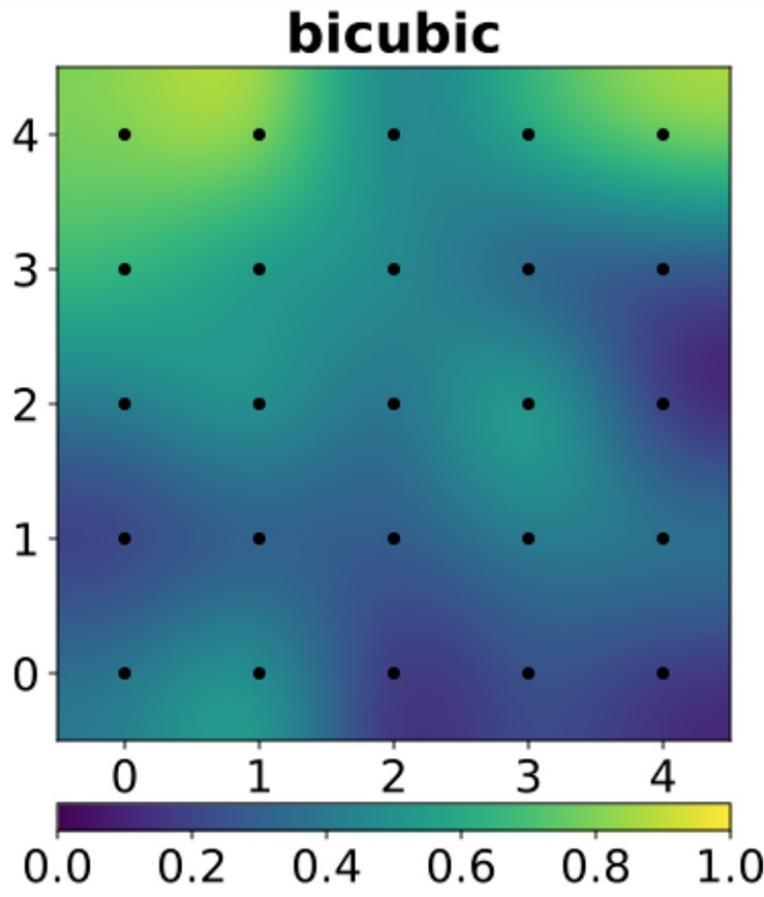


$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$



Bicubic interpolation

- Bilinear interpolation processes 2x2 (4 pixels) squares
- Bicubic interpolation processes 4x4 (16 pixels) squares.



$$p(x, y) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij} x^i y^j.$$

This requires determining the 16 coefficients.

Consider 4 corners of the unit square.

$$(0, 0) (1, 0) (0, 1) (1, 1)$$

1. $f(0, 0) = p(0, 0) = a_{00}$,
2. $f(1, 0) = p(1, 0) = a_{00} + a_{10} + a_{20} + a_{30}$,
3. $f(0, 1) = p(0, 1) = a_{00} + a_{01} + a_{02} + a_{03}$,
4. $f(1, 1) = p(1, 1) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij}$.

If the quality is of concern, bicubic would be the best choice.

2D coordinate transformations

- translation: $\mathbf{x}' = \mathbf{x} + \mathbf{t}$ $\mathbf{x} = (x, y)$

- rotation: $\mathbf{x}' = \mathbf{R} \mathbf{x} + \mathbf{t}$

- similarity: $\mathbf{x}' = s \mathbf{R} \mathbf{x} + \mathbf{t}$

- affine: $\mathbf{x}' = \mathbf{A} \mathbf{x} + \mathbf{t}$

- perspective: $\underline{\mathbf{x}}' \approx \mathbf{H} \underline{\mathbf{x}}$ $\underline{\mathbf{x}} = (x, y, 1)$

($\underline{\mathbf{x}}$ is a *homogeneous* coordinate)

- These all form a nested group (closed w/ inv.)



Basic 2D Transformations

- Basic 2D transformations as 3x3 matrices

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

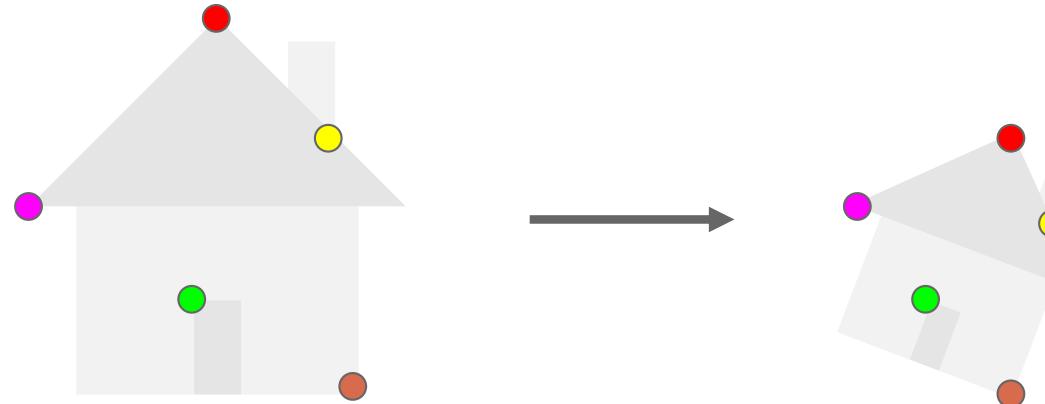
Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear (Skew)

Image alignment

- Two broad approaches:
 - Direct (pixel-based) alignment
 - Search for alignment where most pixels agree
 - Feature-based alignment
 - Search for alignment where *extracted features* agree
 - Can be verified using pixel-based alignment



Fitting an affine transformation

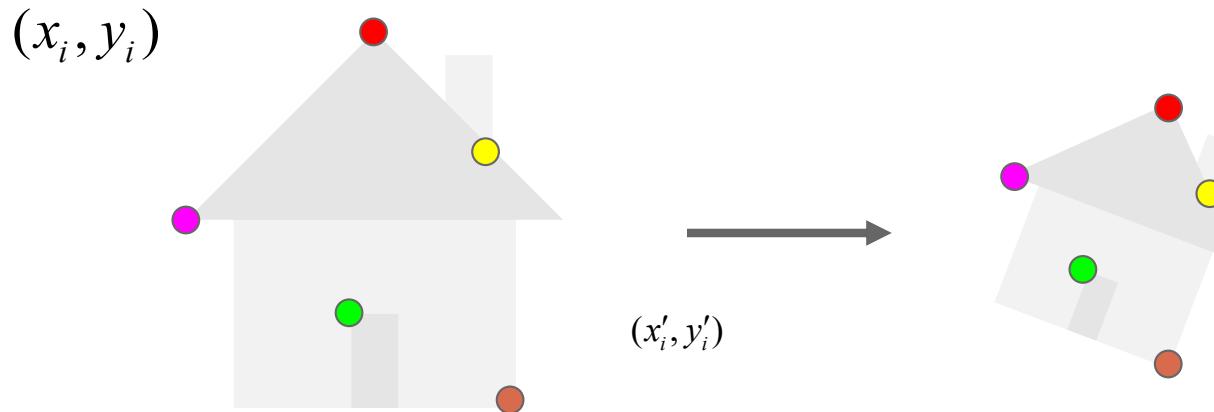


Affine model approximates perspective projection of planar objects.

Figures from David Lowe, ICCV 1999

Fitting an affine transformation

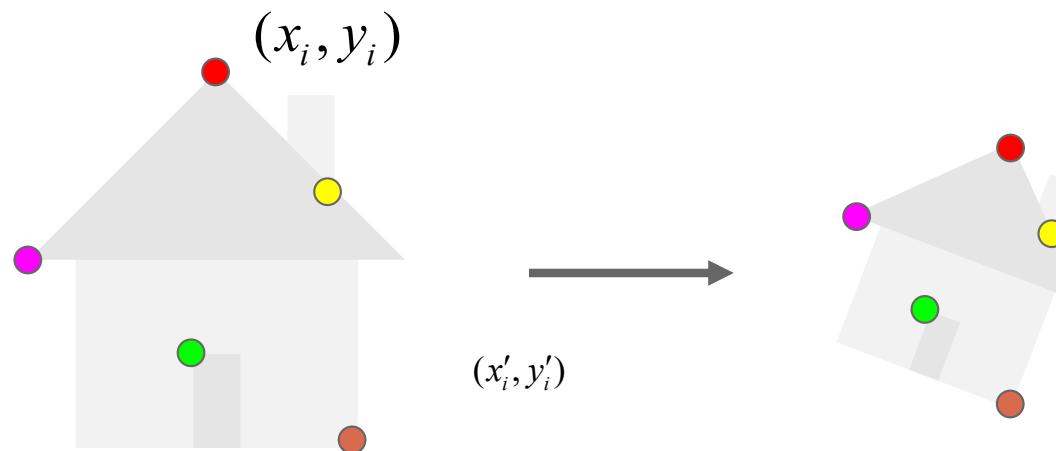
- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \quad \\ \quad \\ \quad \\ \quad \\ \quad \\ \quad \end{bmatrix}$$

A dashed blue line with arrows at the ends connects the vector t_1 to the bottom-left corner of the matrix and the vector t_2 to the bottom-right corner of the matrix.

Fitting an affine transformation

$$\begin{bmatrix} & & \dots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ & & \dots & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

- How many matches (correspondence pairs) do we need to solve for the transformation parameters?
- Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for (x_{new}, y_{new}) ?



Panoramas

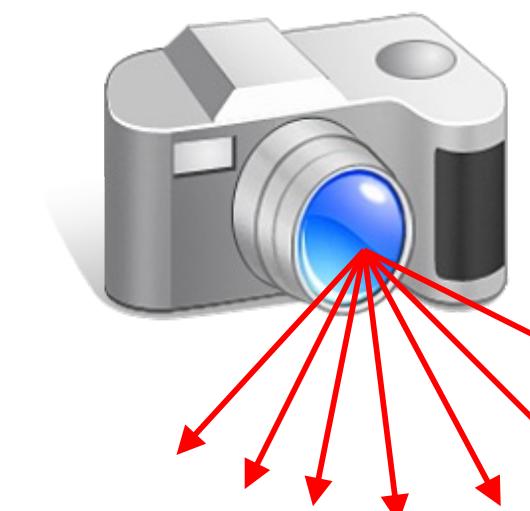
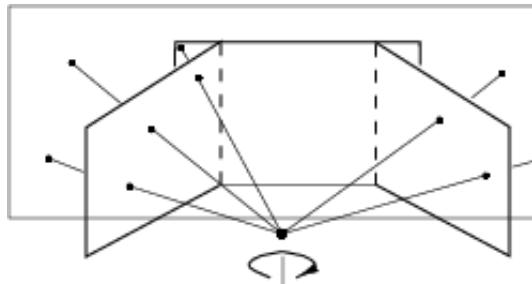
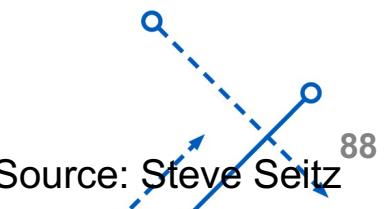


Image from S. Seitz

Obtain a wide angle view by combining multiple images.

How to stitch together a panorama?

- Basic Procedure
 - Take a sequence of images from the same position
 - Rotate the camera about its **optical center**.
 - Compute transformation between 2nd and 1st image
 - Transform the 2nd image to overlap with the 1st
 - Blend the two together to create a mosaic
 - (If there are more images, repeat)
- Why should this work at all?
 - What about the 3D geometry of the scene?
 - Why aren't we using it?



Correspondence

- Allows us to map image back to some real space

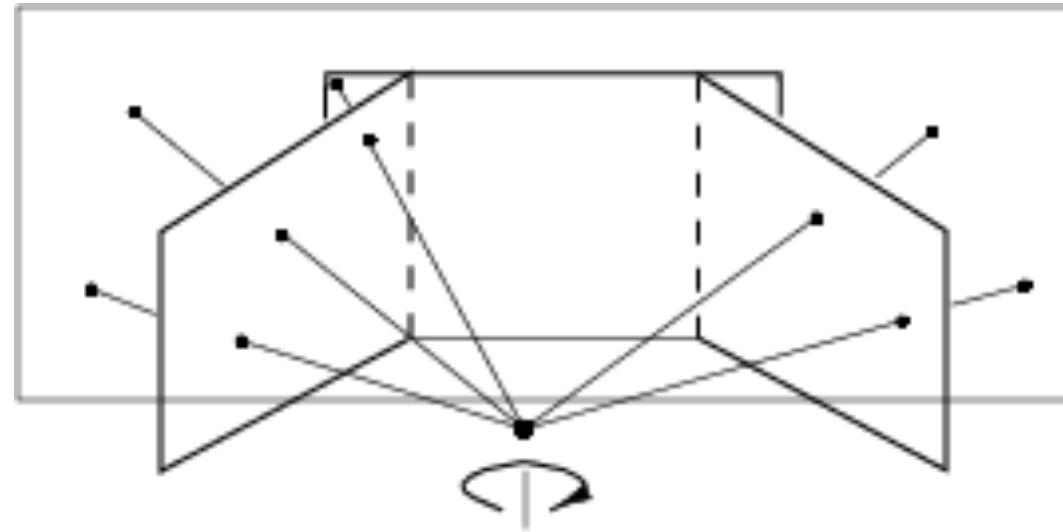
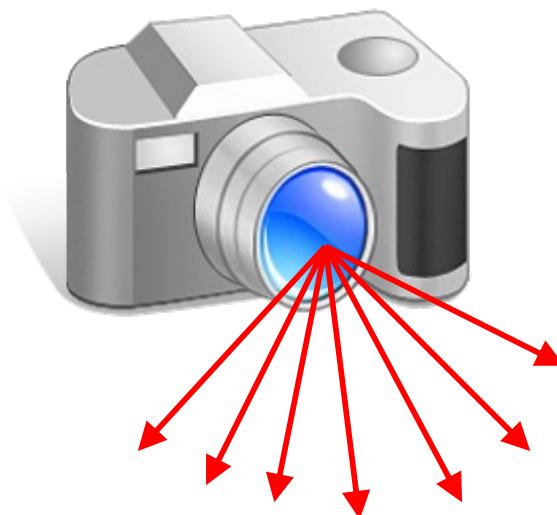
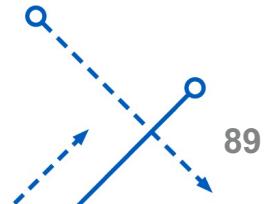
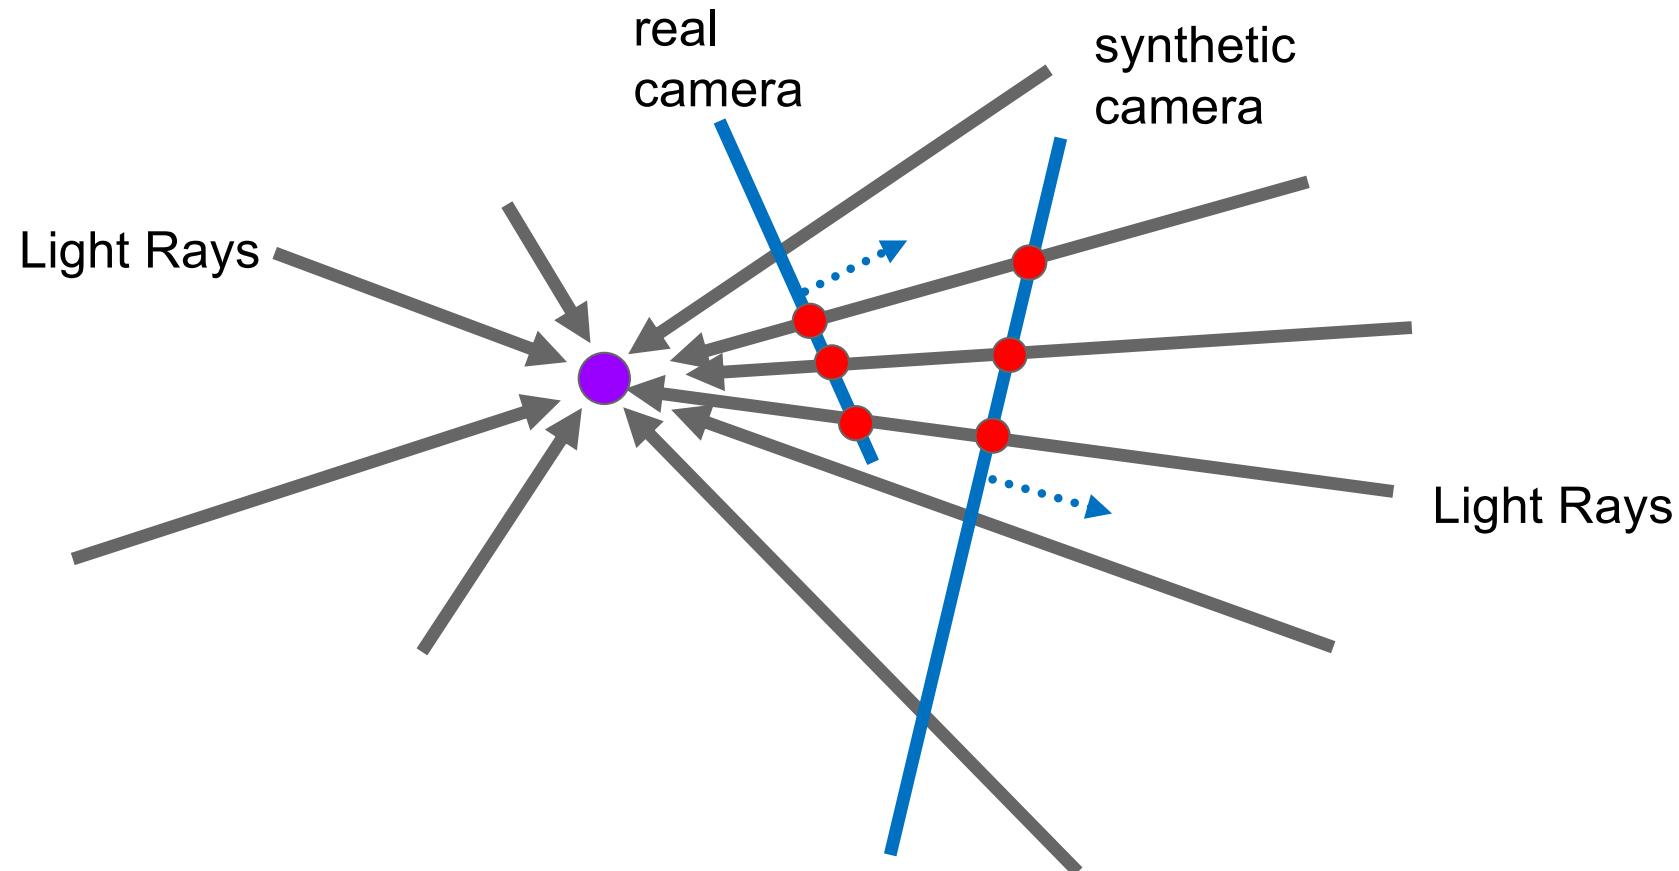


image from S. Seitz



Panoramas: generating synthetic views

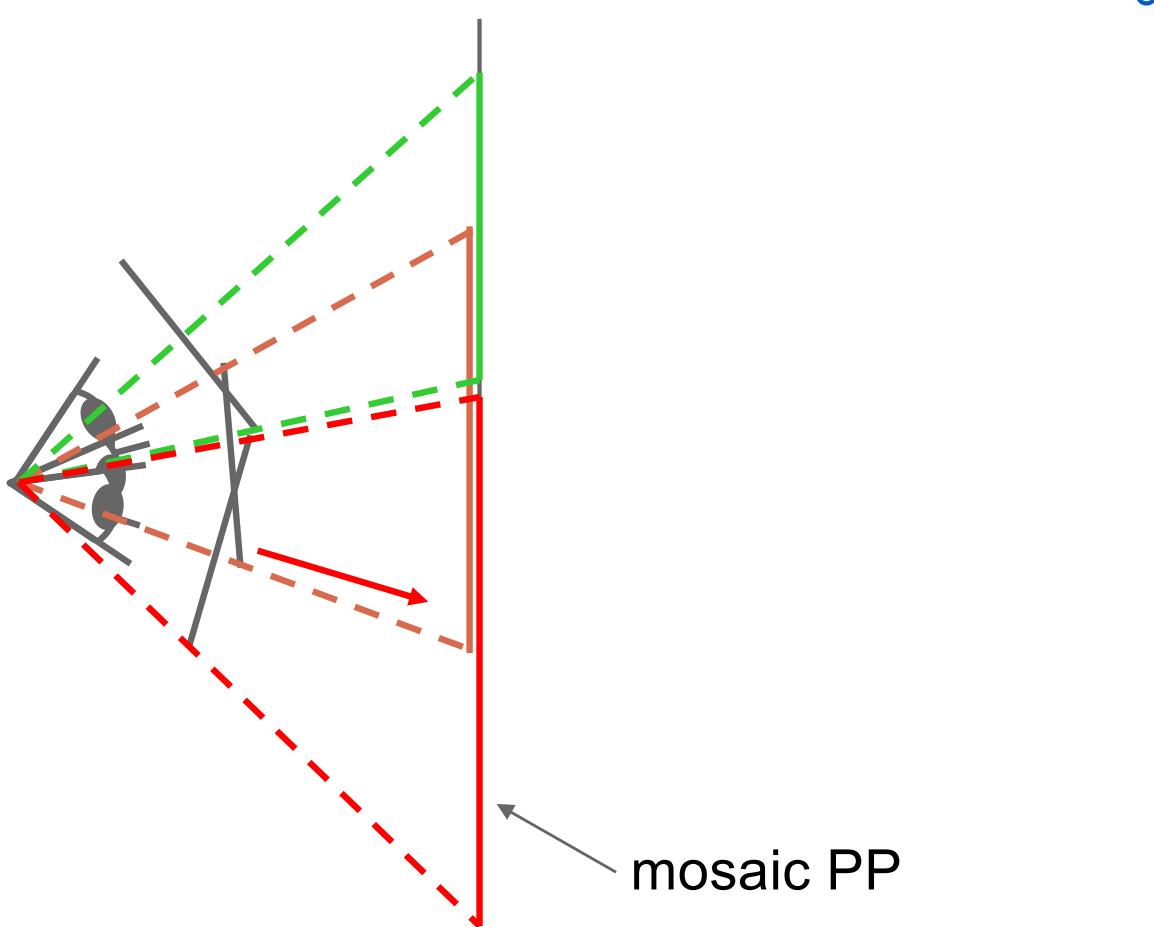


Can generate any synthetic camera view
as long as it has **the same center of projection!**



Source: Alyosha Efros

Image reprojection



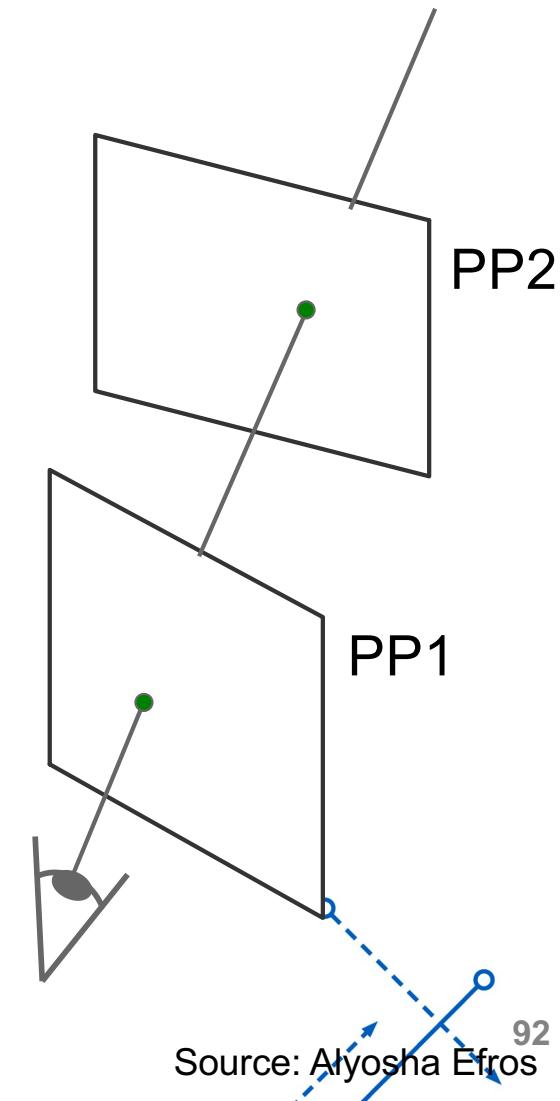
- The mosaic has a natural interpretation in 3D
 - The images are reprojected onto a common plane
 - The mosaic is formed on this plane
 - Mosaic is a *synthetic wide-angle camera*

Homography

- How to relate two images from the same camera center?
 - how to map a pixel from PP1 to PP2?
- Take as a 2D **image warp** using projective transform.
- A **projective transform** is a mapping between any two PPs with the same center of projection
 - rectangle should map to arbitrary quadrilateral
 - parallel lines aren't preserved.
 - but straight lines are preserved.
- Called **Homography**

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

H



Source: Alyosha Efros 92

Solving for homographies

$$(x, y)$$



$$\left(\frac{wx'}{w}, \frac{wy'}{w} \right)$$

$$= (x', y')$$



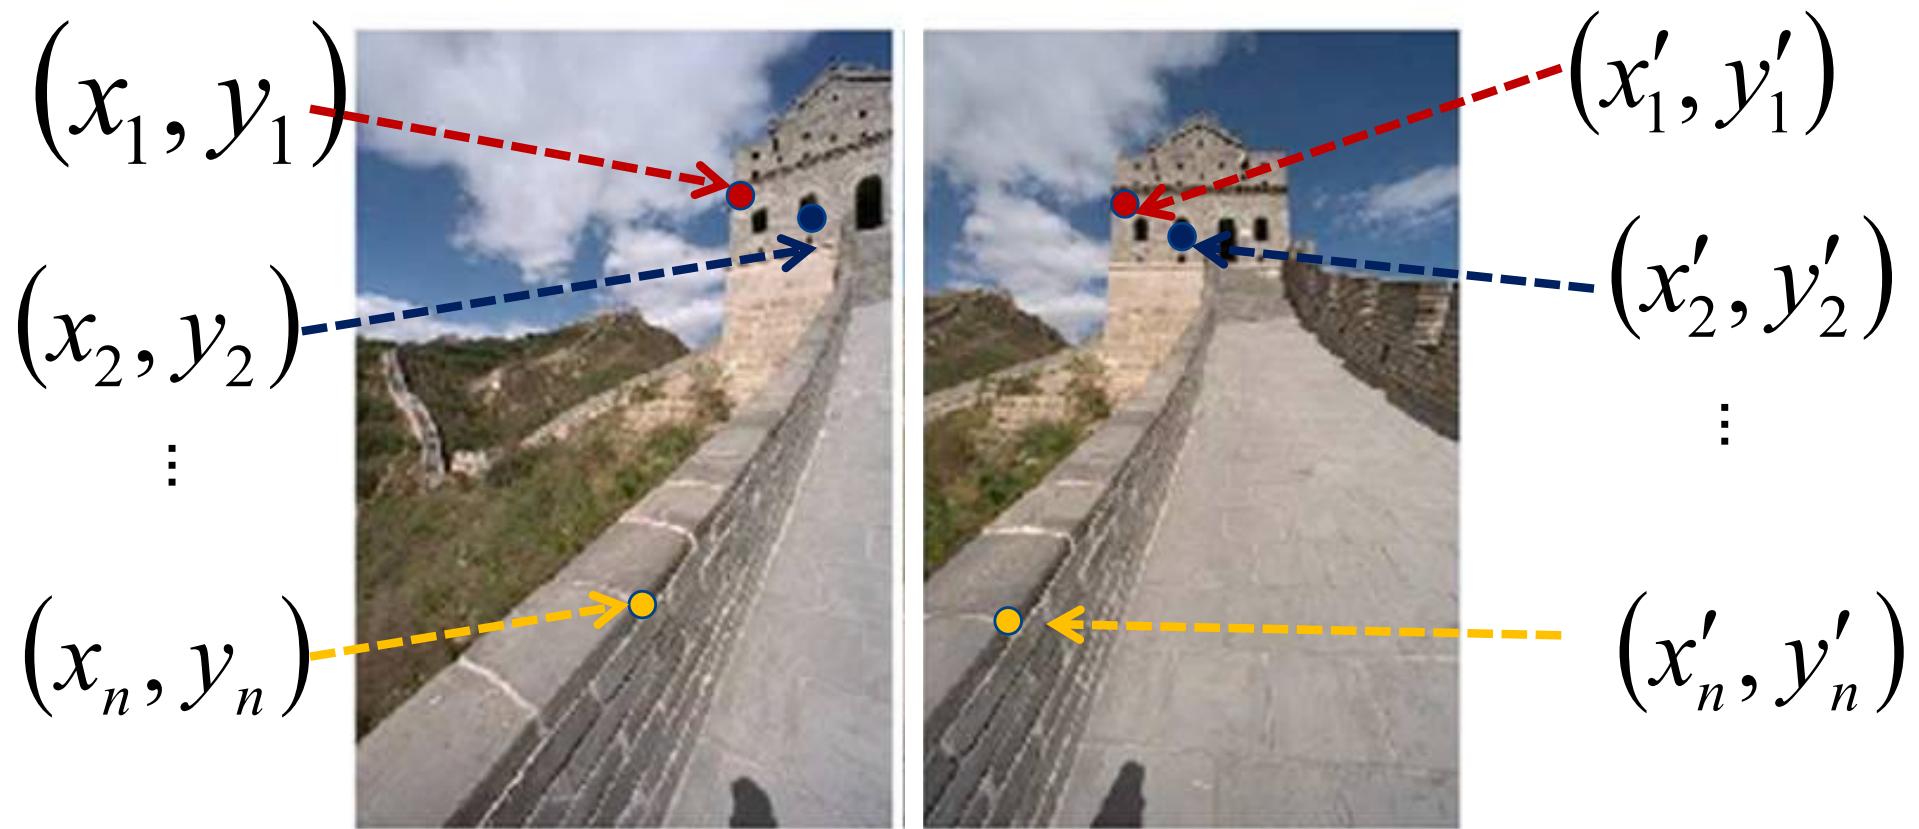
To apply a given homography \mathbf{H}

- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates

$$\begin{bmatrix} wx' \\ wy' \\ w \\ p' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

93

Solving for homographies



To **compute** the homography given pairs of corresponding points, we need to set up an equation where the parameters of H are the unknowns...

Solving for homographies

$$\mathbf{p}' = \mathbf{H}\mathbf{p}$$
$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Can set scale factor $i=1$ or $\|\mathbf{H}\| = 1$. So, there are 8 unknowns.
- Set up a system of linear equations:

$$\bullet \mathbf{A}\mathbf{h} = \mathbf{b}$$

where vector of unknowns $\mathbf{h} = [a, b, c, d, e, f, g, h]^T$

- Need at least 8 equations, but the more the better...
- Solve for \mathbf{H} . If over-constrained, solve using least-squares:

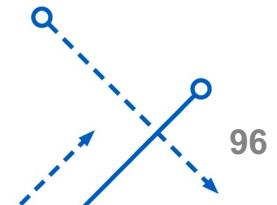
$$\min \|A\mathbf{h} - \mathbf{b}\|^2$$

$$\mathbf{h} = (A^T A)^{-1} A^T \mathbf{b}$$



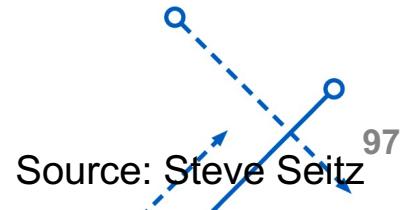
Proof of least squares

- $F(h) = ||Ah - b||^2 = (Ah - b)^T (Ah - b)$
- $F(h) = h^T A^T Ah - h^T A^T b - b^T Ah + b^T b$
- $\frac{\partial}{\partial h} F(h) = 2A^T Ah - A^T b - (b^T A)^T$
- Setting derivative to 0: $\frac{\partial}{\partial h} F(h) = 0$
- $A^T Ah = A^T b$
- $h = (A^T A)^{-1} A^T b$



How to stitch together a panorama?

- Take a sequence of images from the same position
 - Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- If there are more images, repeat



Source: Steve Seitz