



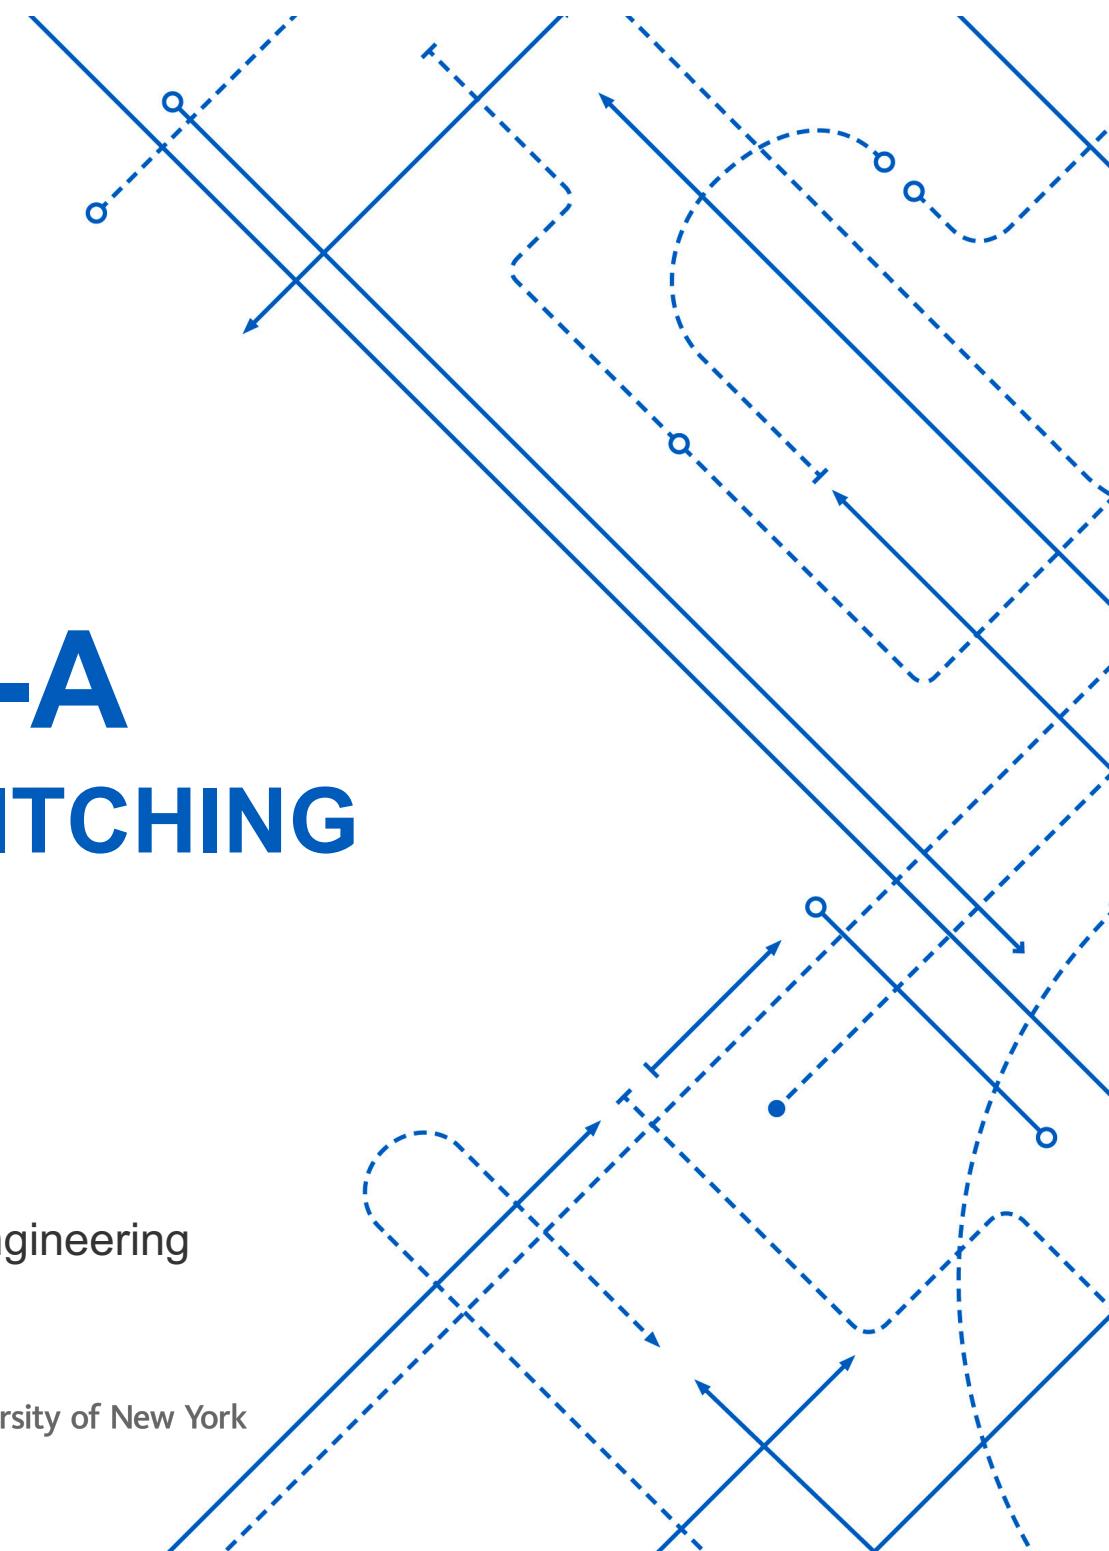
**SAIR**  
Spatial AI & Robotics Lab

# CSE 473/573-A

## W6: RANSAC & STITCHING

Chen Wang  
Spatial AI & Robotics Lab  
Department of Computer Science and Engineering

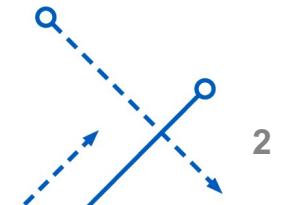
**UB** University at Buffalo The State University of New York



# Content

---

- RANSAC
  - Sampling
  - Consensus
- Stitching
  - Panorama, Blending, Deghosting
  - Feathering, Pyramid Blending
  - Laplacian Blending





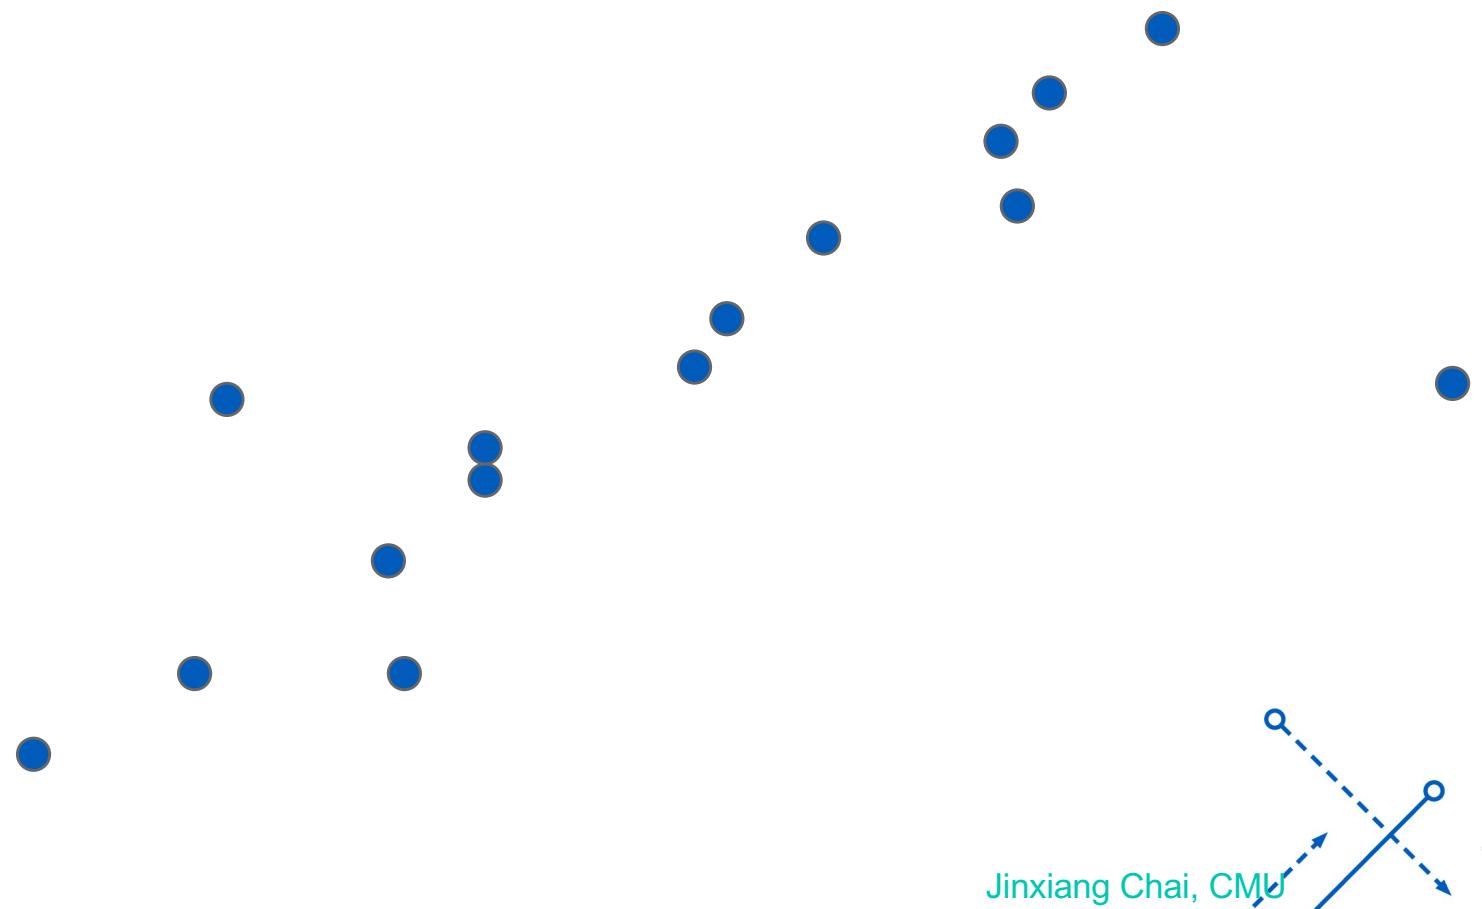
# IMAGE PROCESSING

## RANSAC



# Problem: Line Fitting

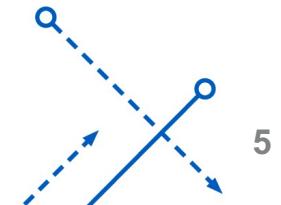
- Not all data may be representative
- There may be “outliers”
- If you use them, your result may not be accurate



# Solutions?

---

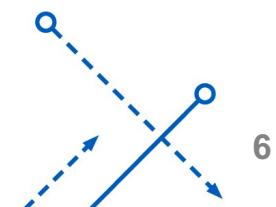
- Least Squares Fit?
  - Closed for solution...
  - Sensitive to outliers
- Hypothesize and Test
  - Try out as many lines as we want
  - Keep the best lines
  - But which are the best?



# RANSAC

---

- **RAN**dom Sample Consensus
  - An iterative method for estimating a mathematical model from a data set that contains outliers.
- Motivation: we want to avoid the impact of outliers, so let's look for “inliers”, and use those only.
- Idea: if an outlier is chosen to compute the current model, then the model won't have much support from rest of the points.

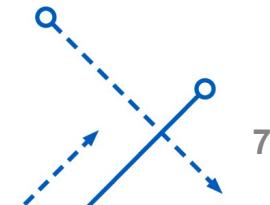


# RANSAC

---

## RANSAC loop:

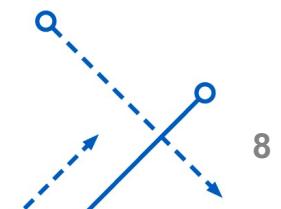
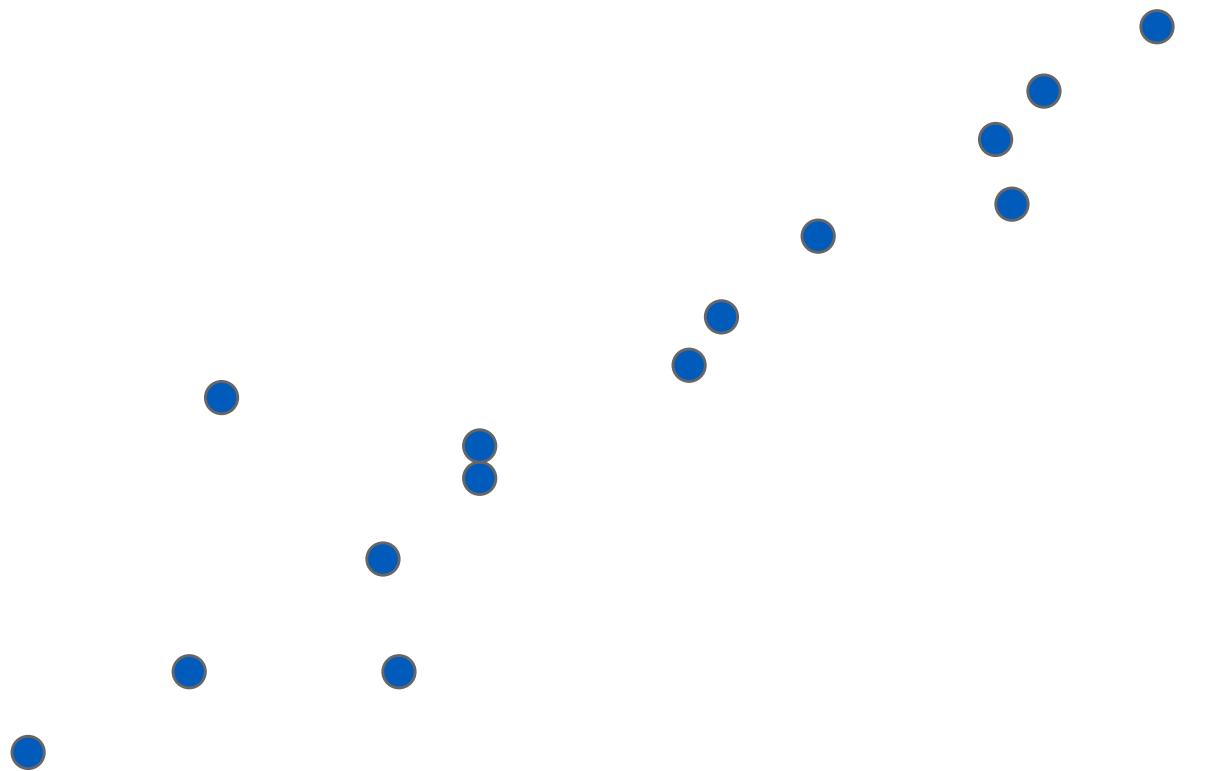
1. Randomly select a **seed group** of points on which to base transformation estimate (e.g., a group of matches)
  2. Compute transformation (**model**) from seed group
  3. Find **inliers** to this transformation
  4. If the number of inliers is sufficiently large, re-compute least-squares estimate on all inliers.
- 
- Keep the model with the largest number of inliers



# RANSAC Line Fitting Example

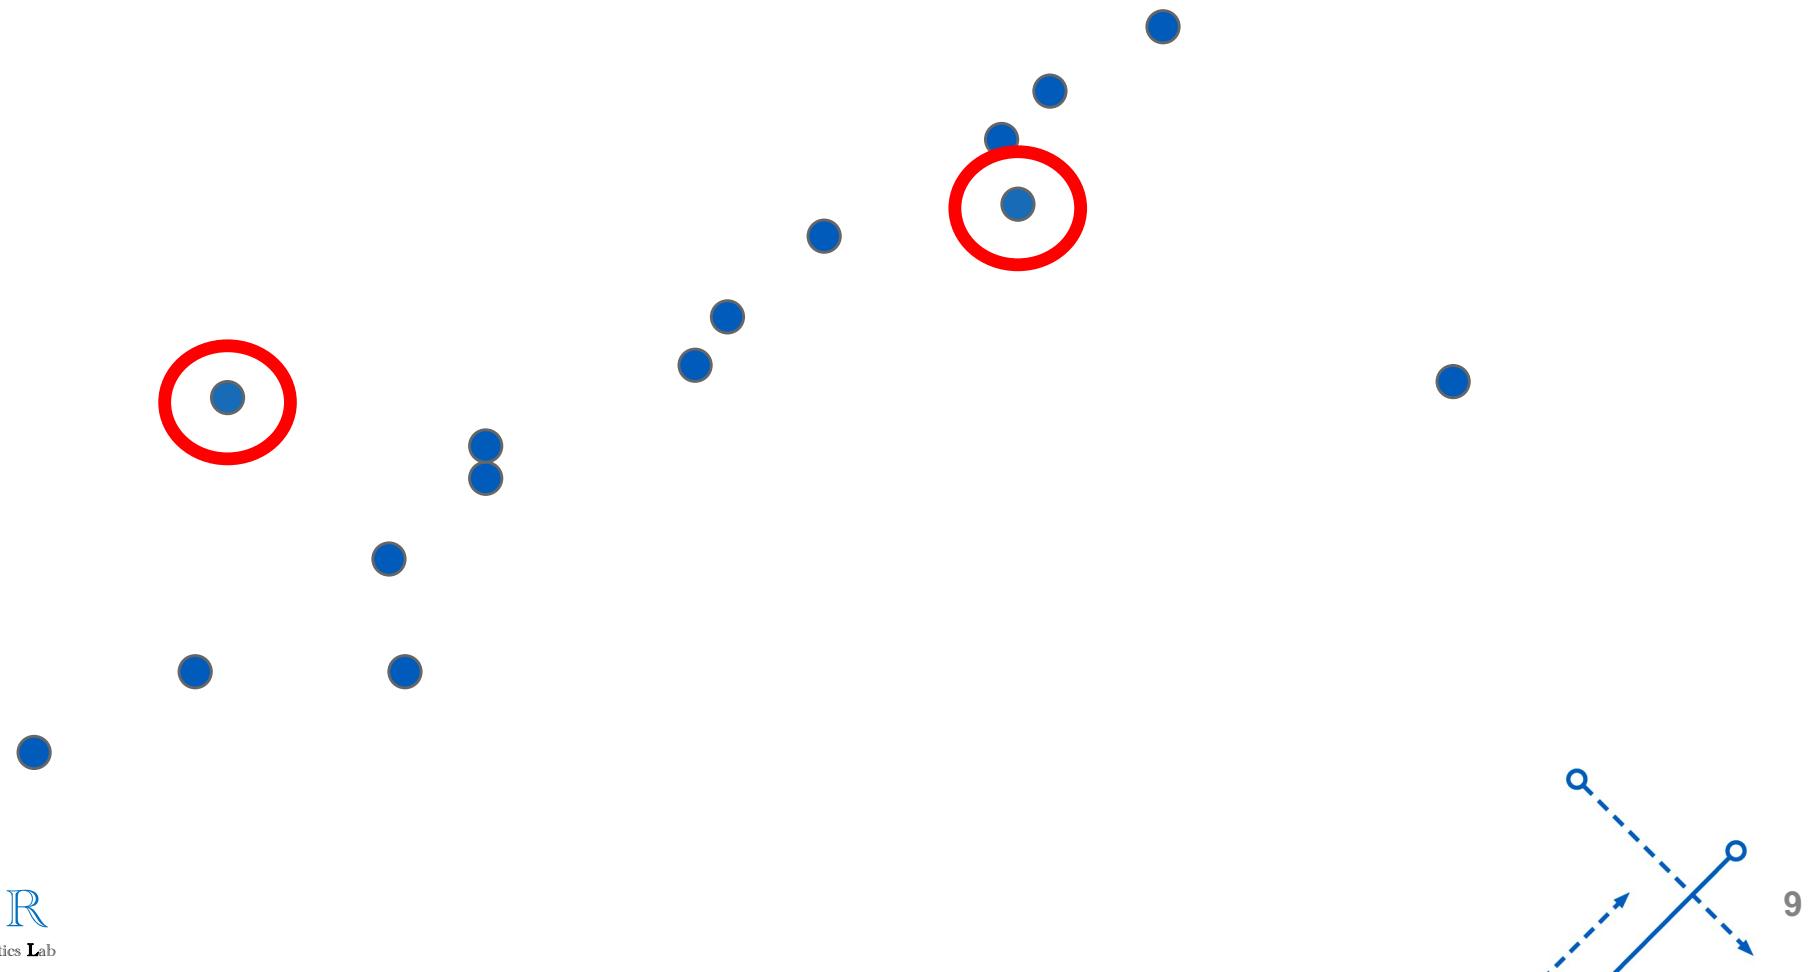
Task:

Estimate best line



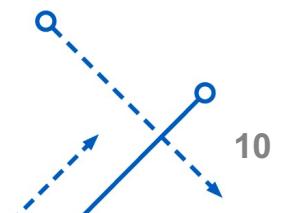
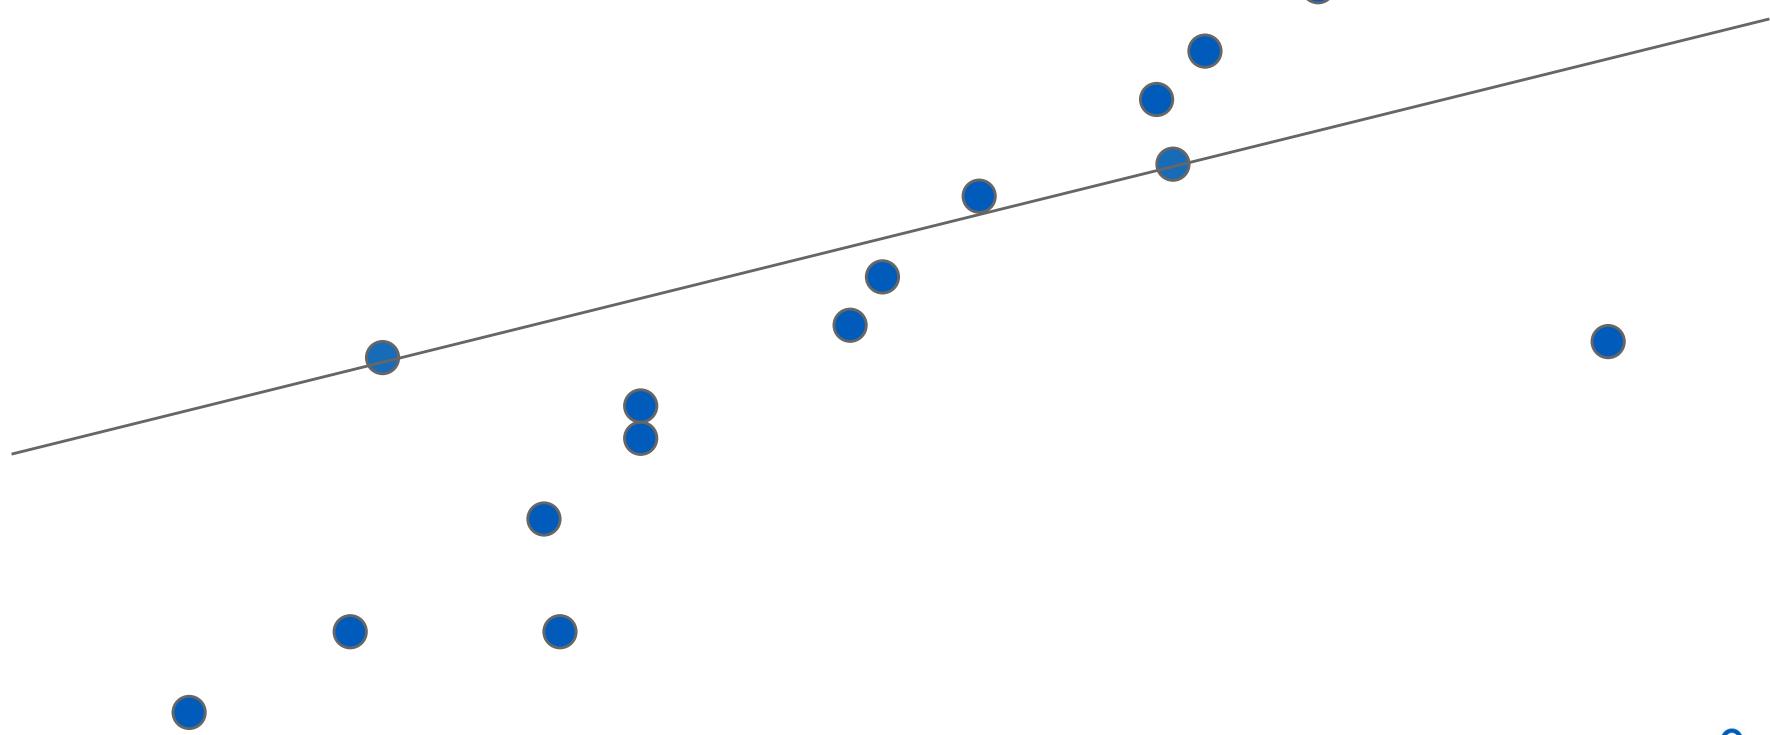
# RANSAC Line Fitting Example

Sample two points



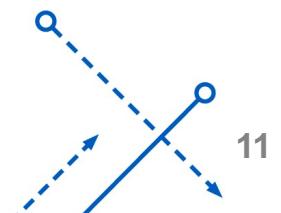
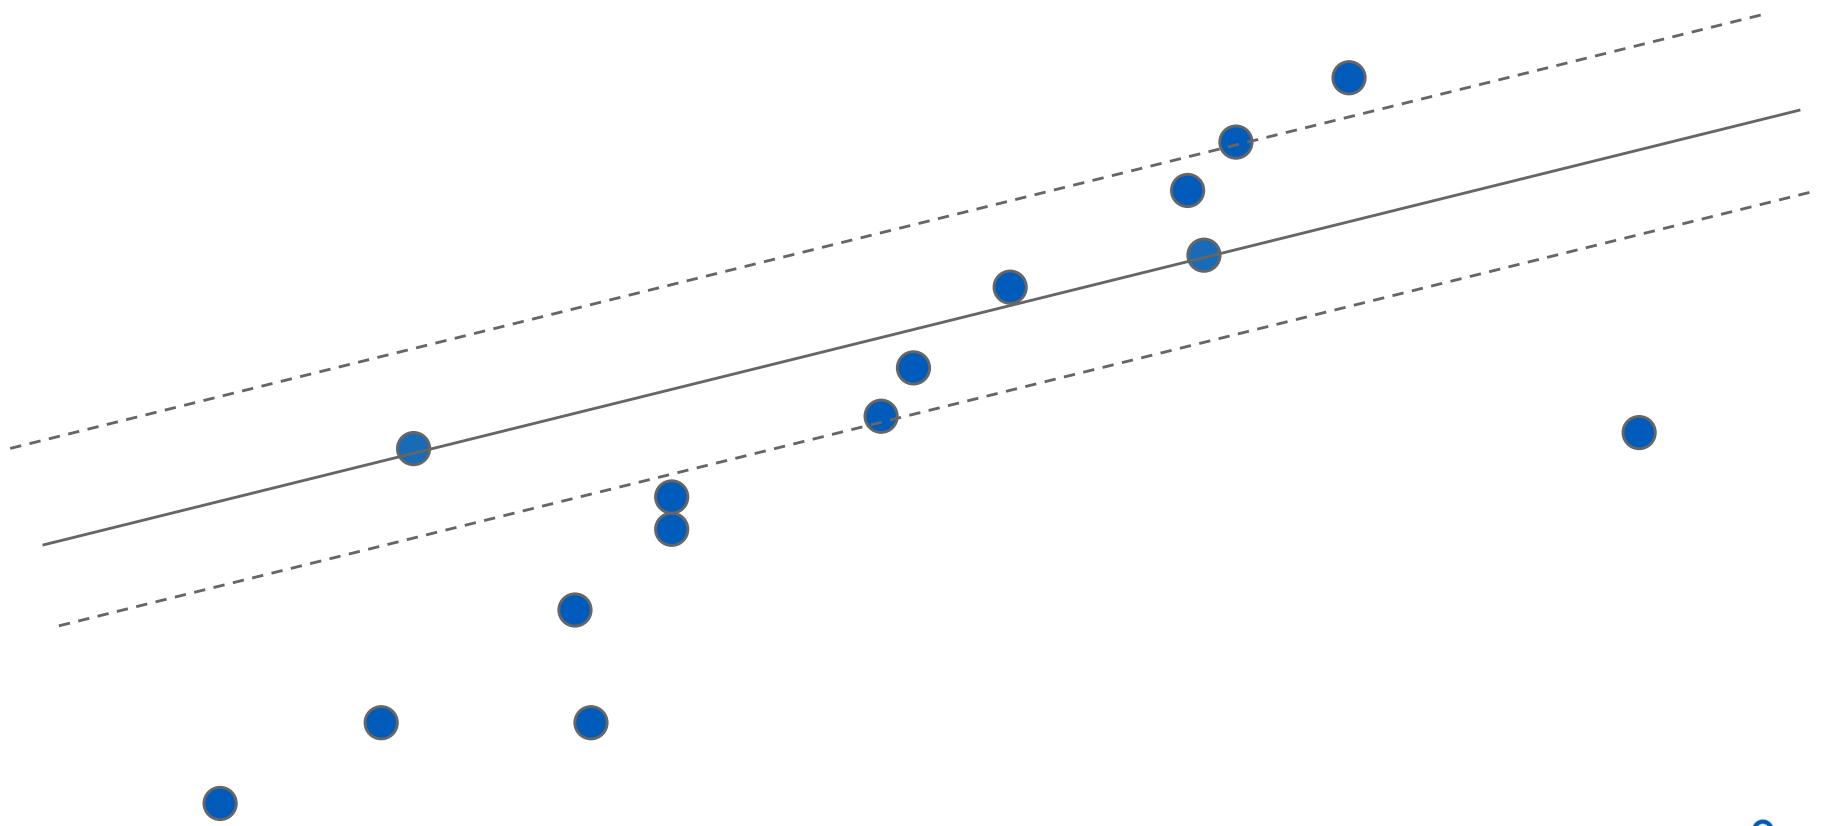
# RANSAC Line Fitting Example

Fit Line



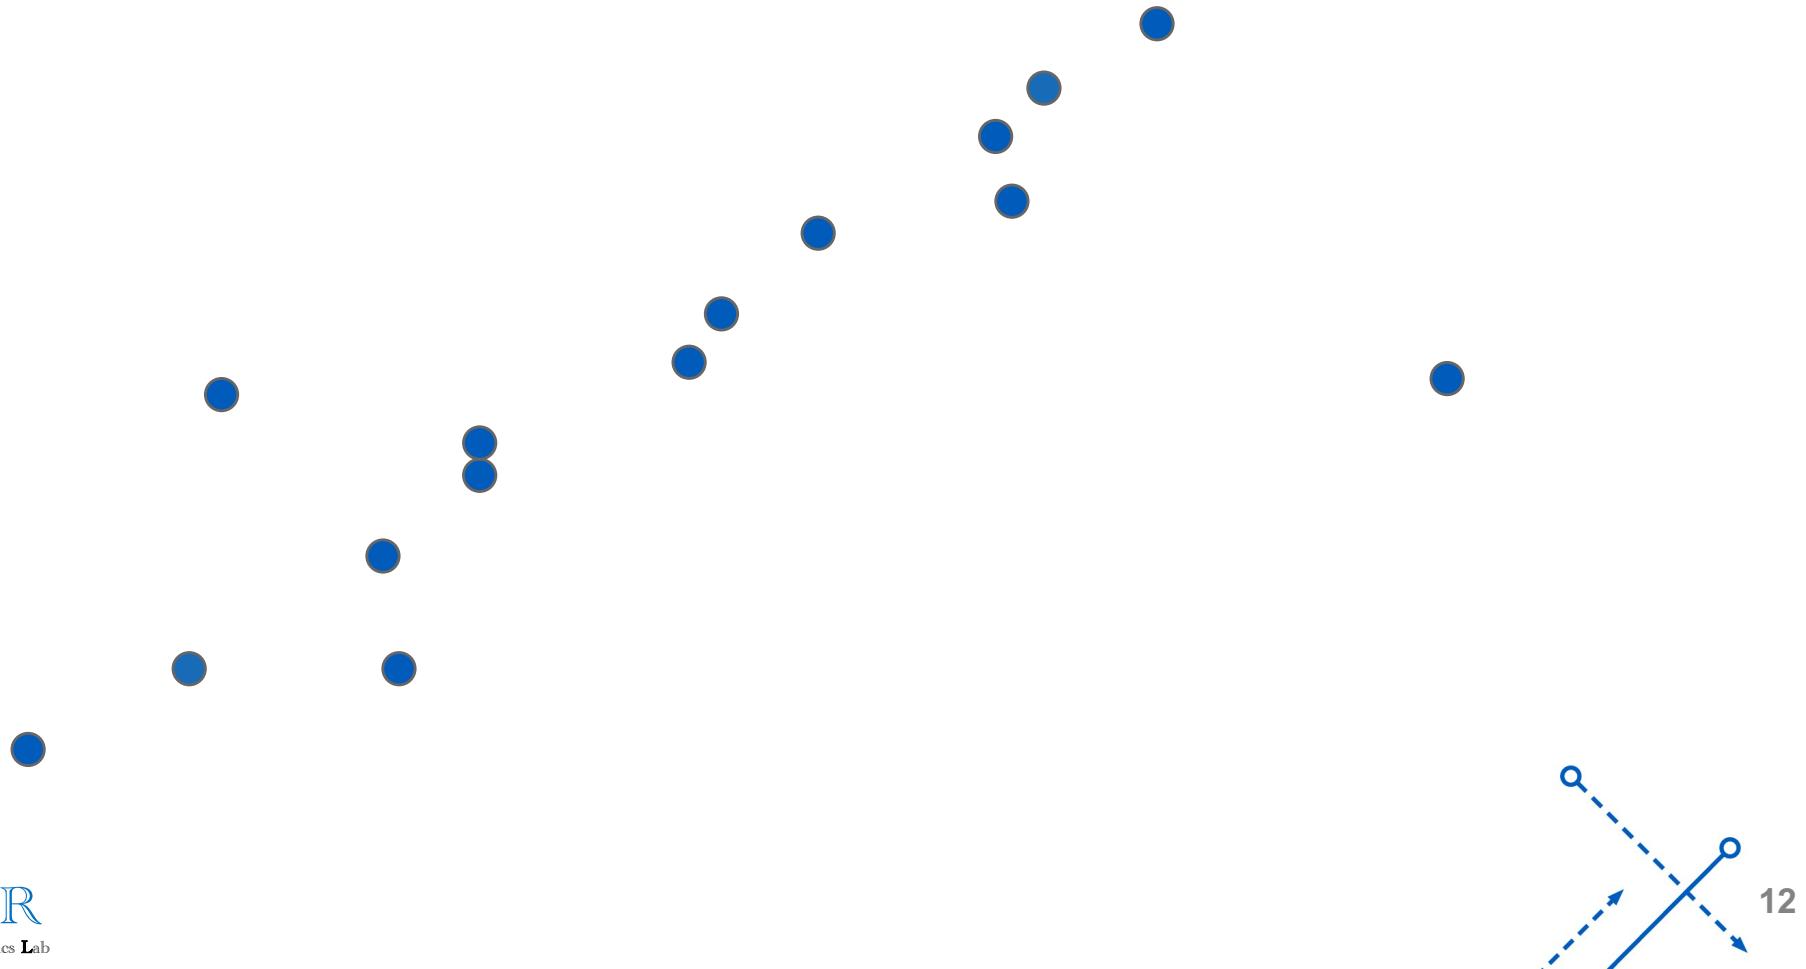
# RANSAC Line Fitting Example

Total number of points within a threshold of line.



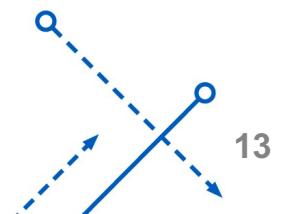
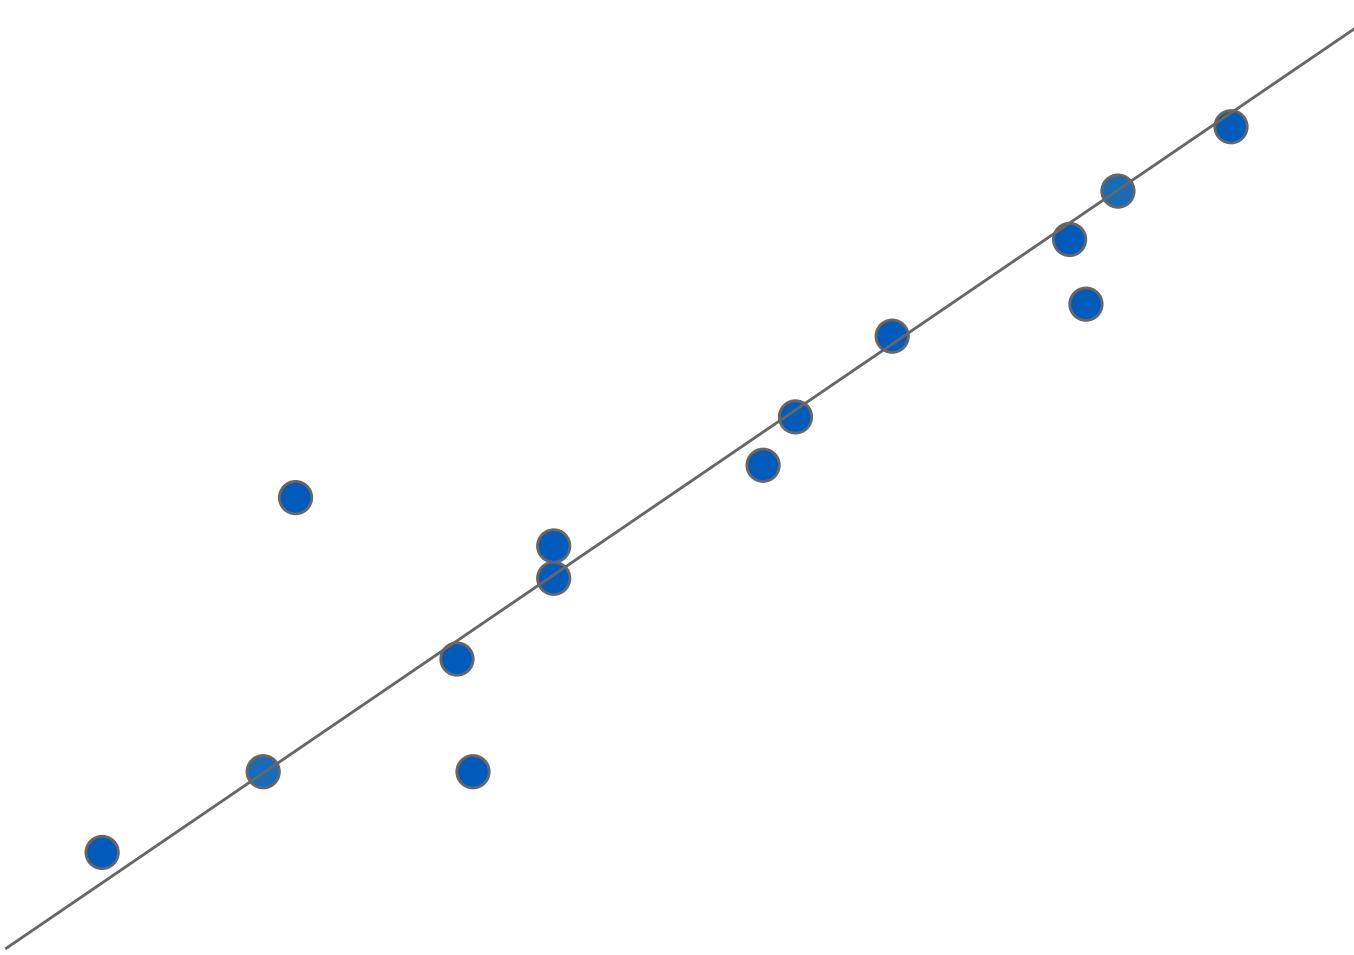
# RANSAC Line Fitting Example

Repeat, until get a good result



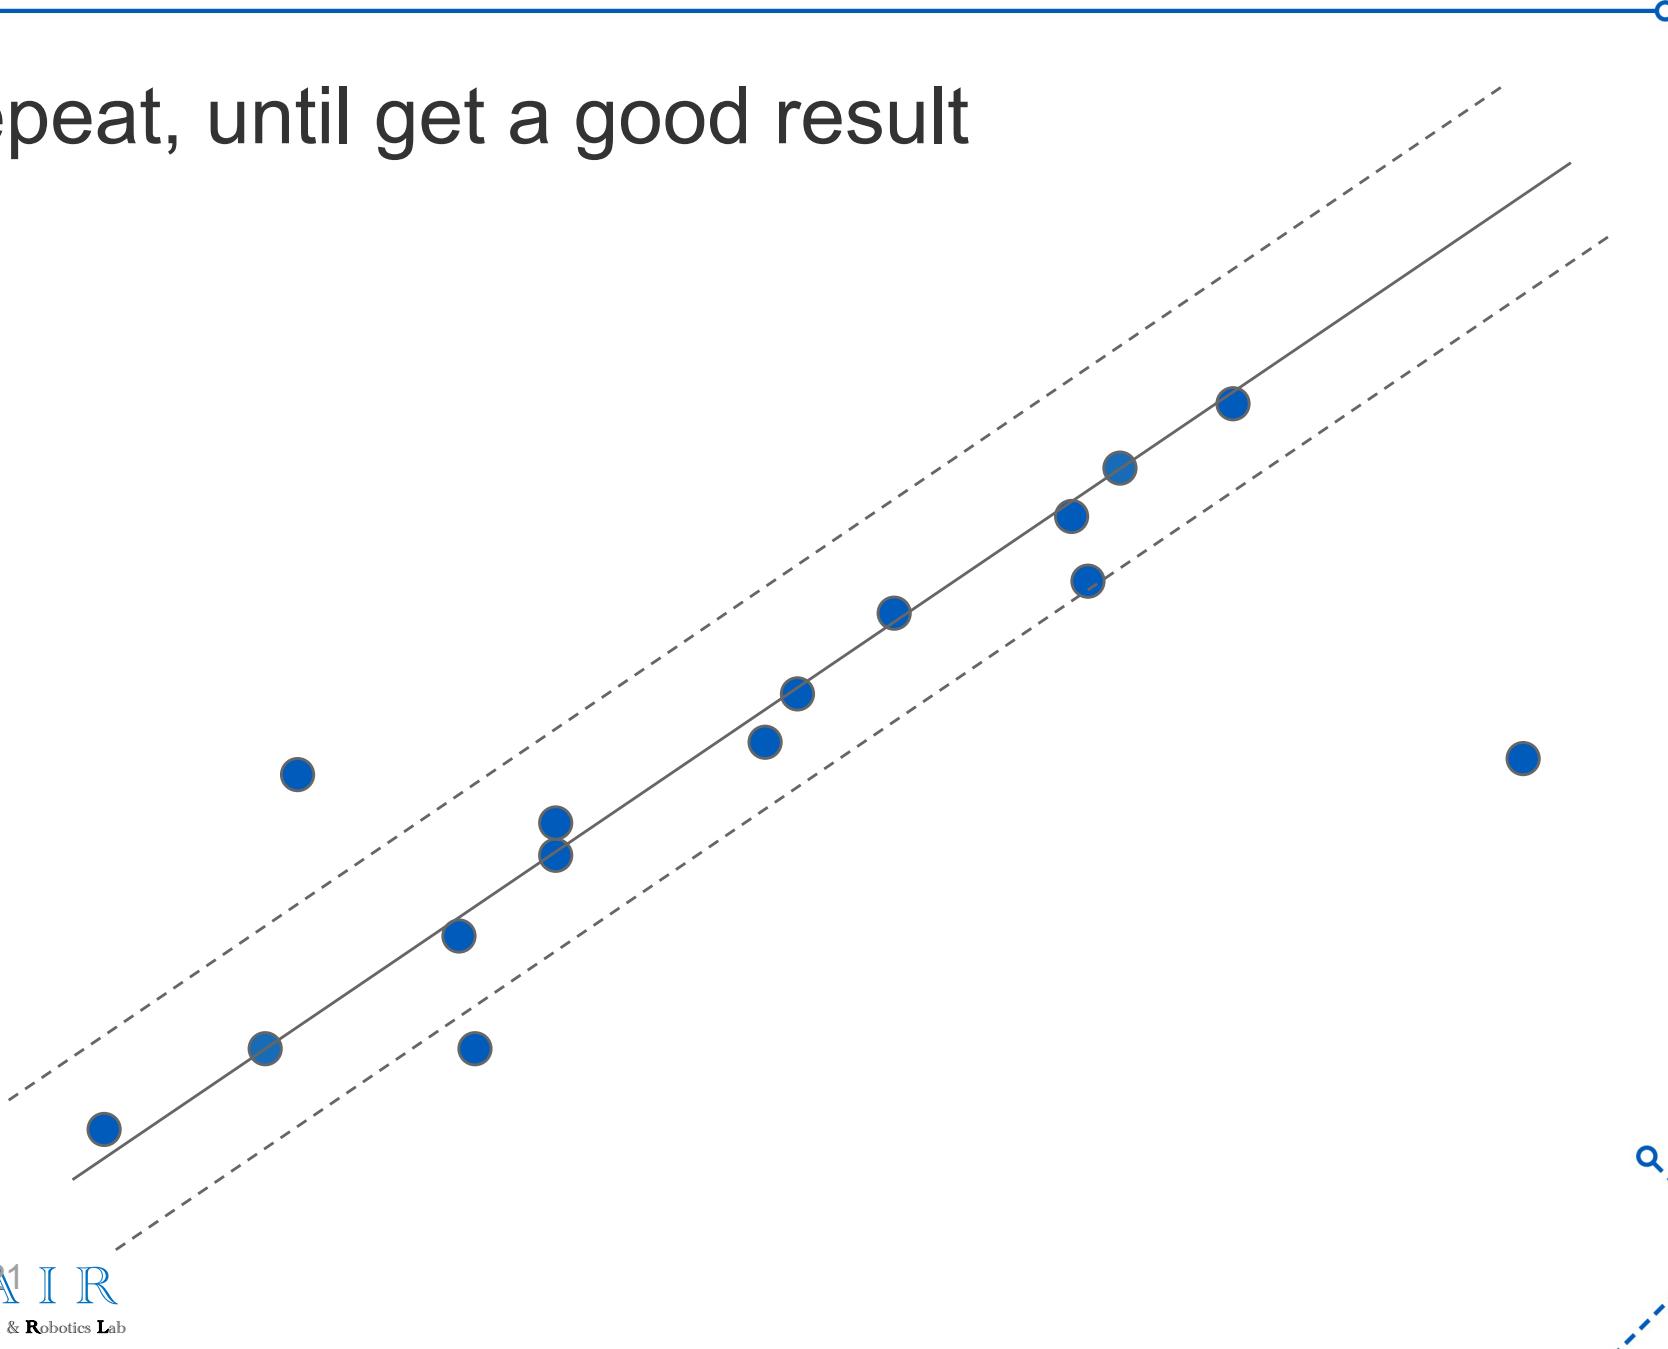
# RANSAC Line Fitting Example

Repeat, until get a good result



# RANSAC Line Fitting Example

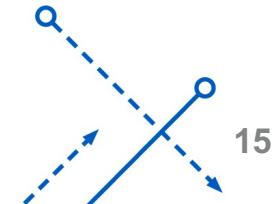
Repeat, until get a good result



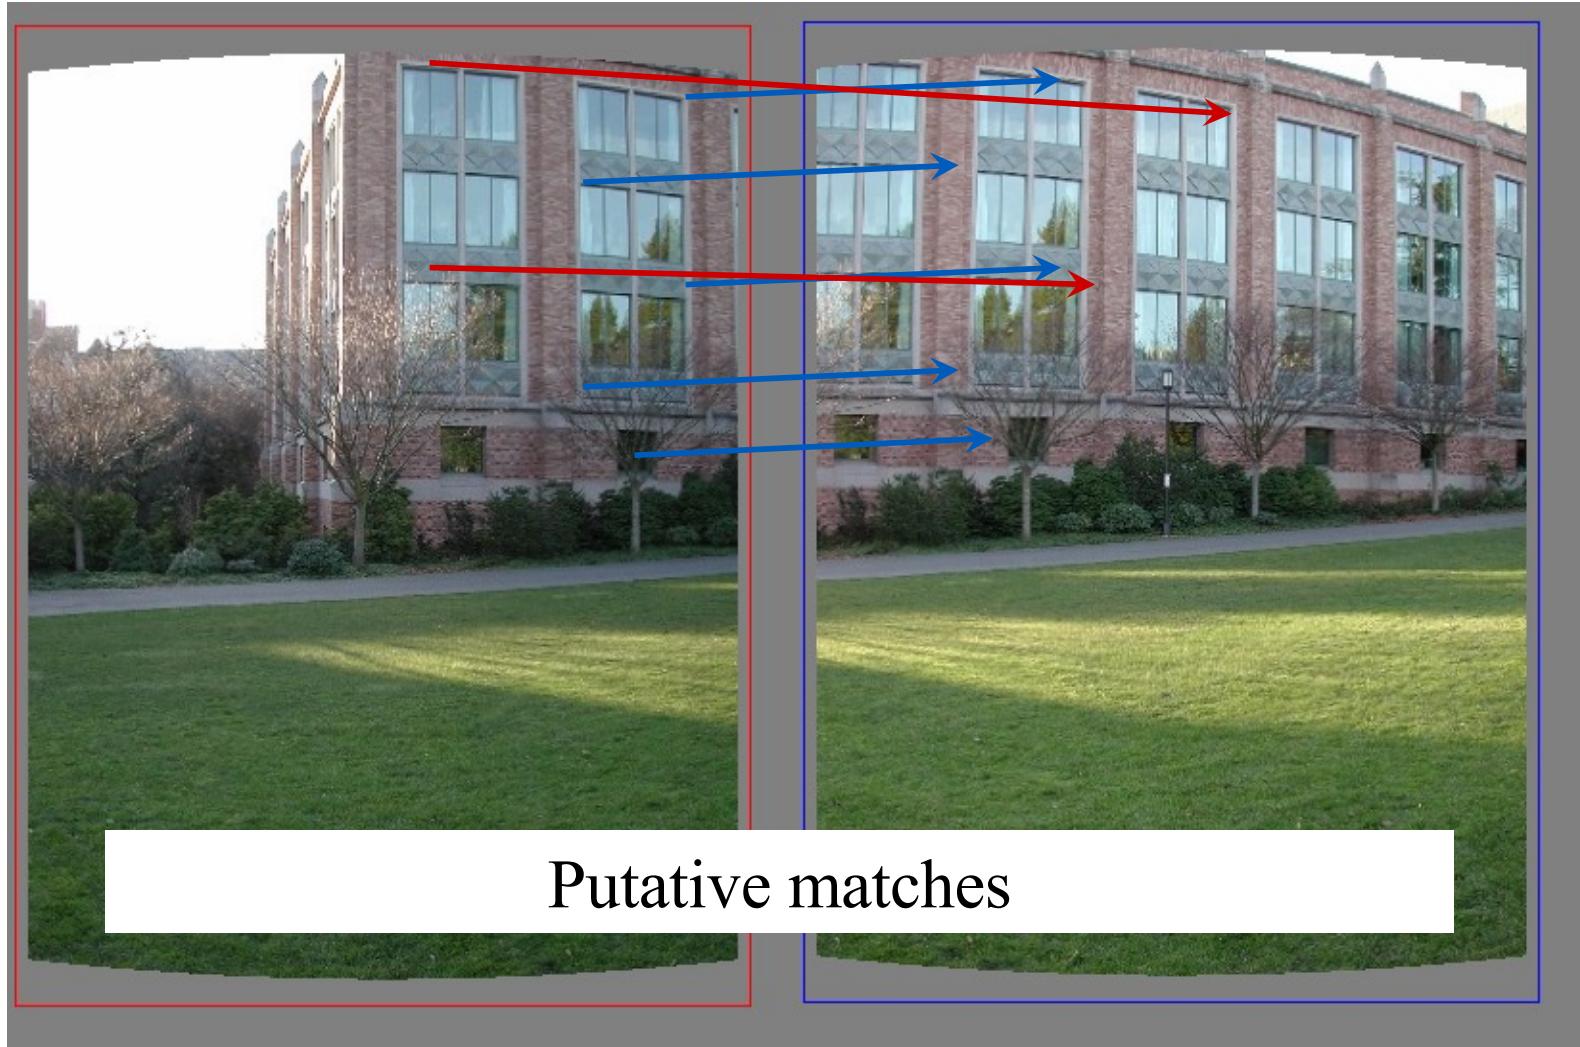
# How to choose parameters?

- Number of sampled points  $n$ : minimum points to fit a model.
- Inlier threshold  $\delta$ .
  - Choose  $\delta$  so that a good point with noise is likely (e.g., prob=0.95) within threshold.
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ , so that  $t^2 = 1.95^2 \sigma^2$
- To determine the number of iterations  $K$ .
  - Desired probability of success ( $p$ ): at least one useful result.
  - Let  $w$  be the probability of choosing an inlier when selecting a point.
    - $w = \text{number of inliers in data} / \text{number of points in data}$
    - $n$  points selected independently for estimating a model.
    - $w^n$ : the probability that all  $n$  points are inliers.
    - $1 - w^n$ : probability of at least one of the  $n$  points is an outlier.
    - $(1 - w^n)^k$ : after  $k$  iterations, never select a set of  $n$  inlier points.
  - $1 - p = (1 - w^n)^K$
  - $K = \frac{\log(1-p)}{\log(1-w^n)}$

n	$p = 0.99$ , proportion of outliers $(1 - w)$						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

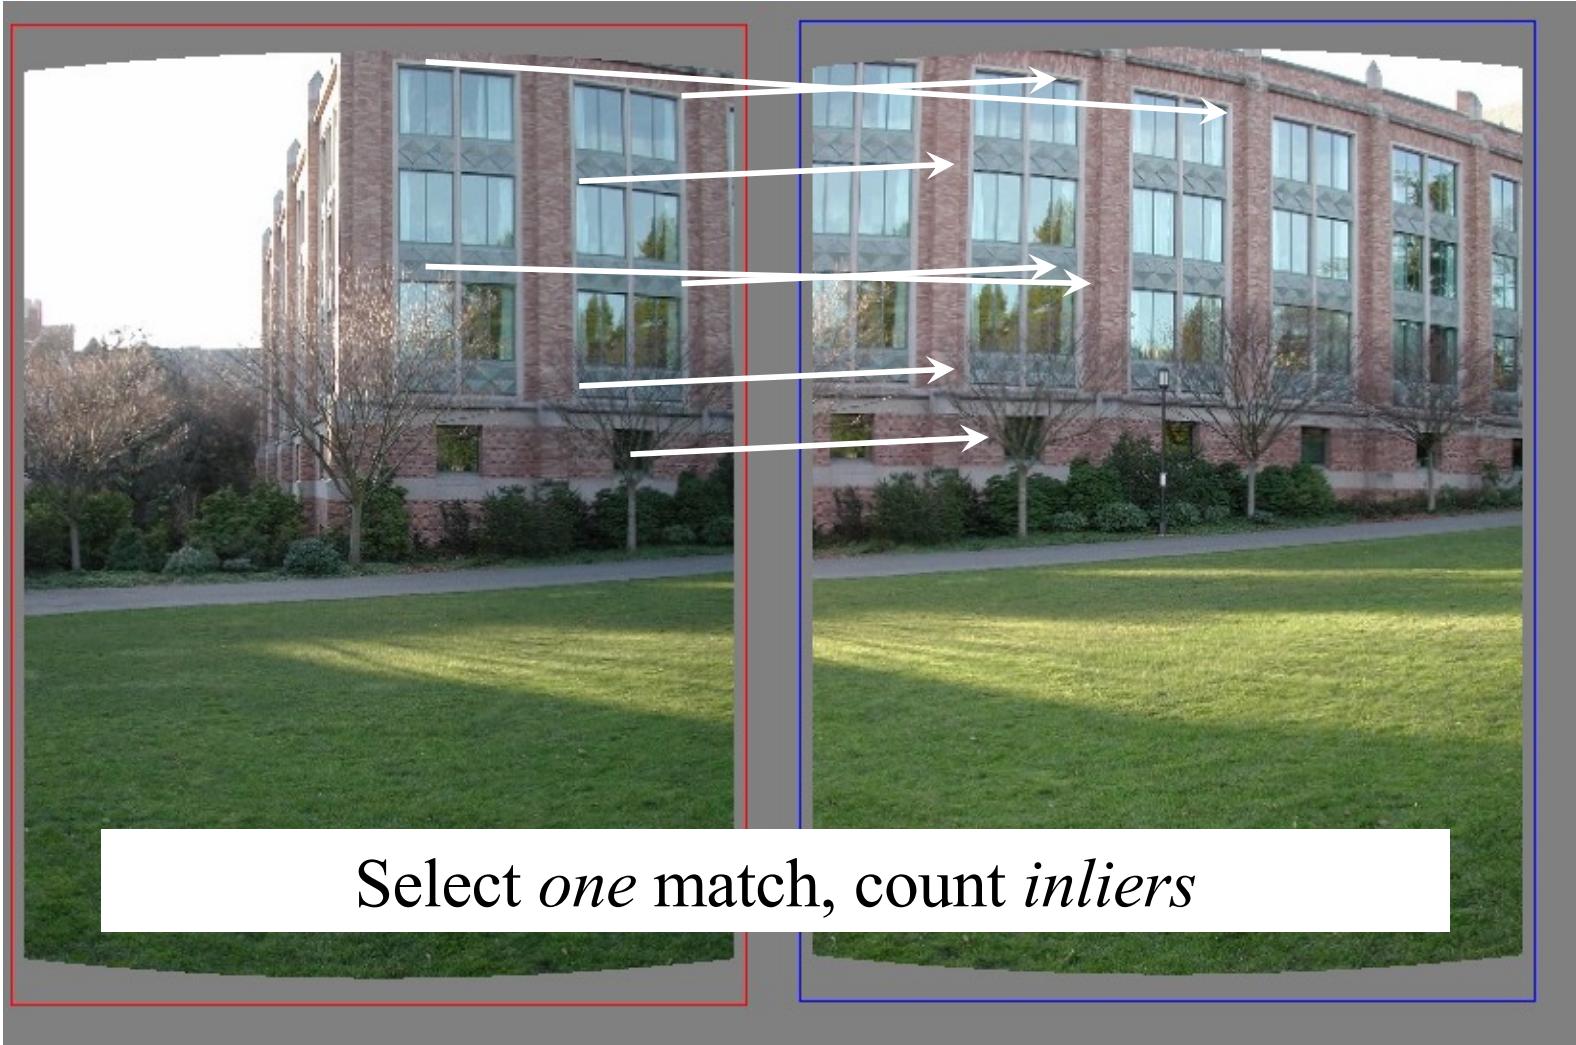


# RANSAC example: Translation

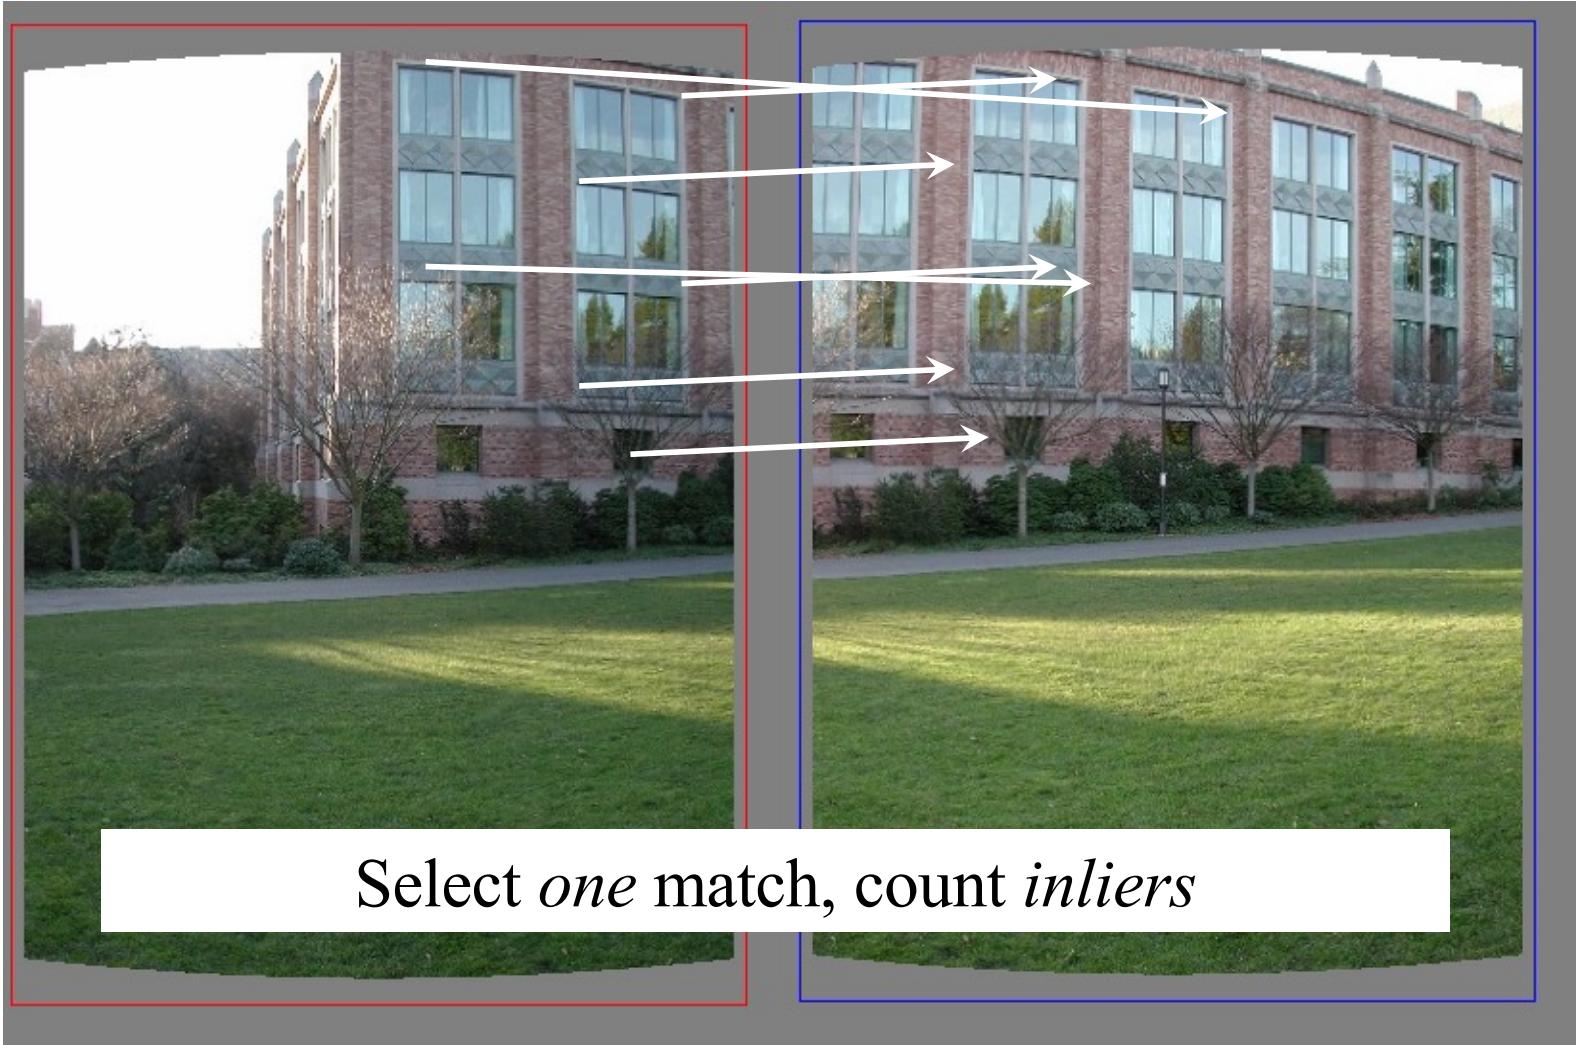


Source: Rick Szeliski

# RANSAC example: Translation

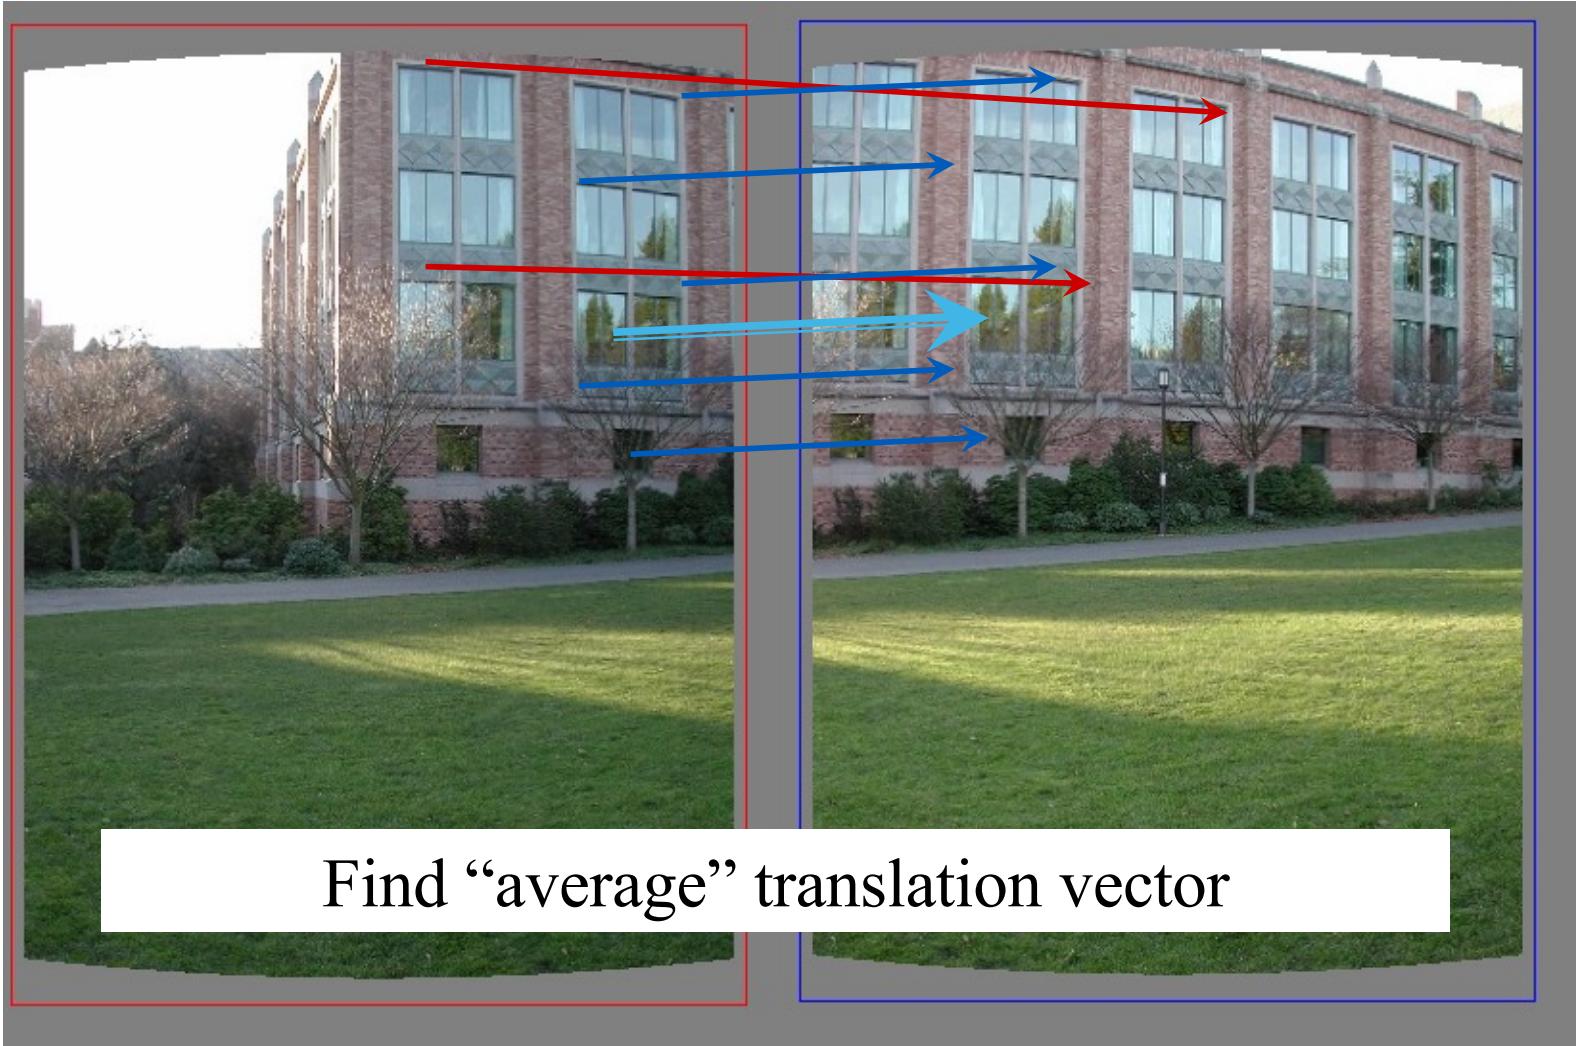


# RANSAC example: Translation



# RANSAC example: Translation

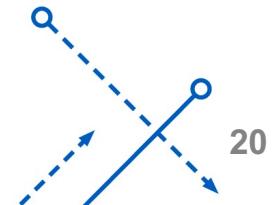
---



# Summary

---

- Choose:
  - Inlier threshold
    - Related to the amount of noise we expect
  - Number of rounds
    - Related to the percentage of outliers we expect
    - Related to the probability of success we are hoping for.



# RANSAC ALGORITHM

---

- Input:
  - Data: a set of observed data points
  - Model: a model to fit the data
  - Threshold: a threshold to determine inliers
- Output: BestModel: the model with the most inliers
- Repeat for a fixed number of iterations:
  - 1. Select a random subset of data
  - 2. Fit the model to the data points in the subset
  - 3. Determine the inliers by comparing the fitted model to data
  - 4. If the number of inliers exceeds the threshold
    - re-estimate the model using all the inliers
  - 5. Store the model if it has the most inliers seen so far
- Return BestModel



# RANSAC Algorithm

Given:

```
data - A set of observations.  
model - A model to explain the observed data points.  
n - The minimum number of data points required to estimate the model parameters.  
k - The maximum number of iterations allowed in the algorithm.  
t - A threshold value to determine data points that are fit well by the model (inlier).  
d - The number of close data points (inliers) required to assert that the model fits well to the data.
```

Return:

```
bestFit - The model parameters which may best fit the data (or null if no good model is found).
```

```
iterations = 0  
bestFit = null  
bestErr = something really large // This parameter is used to sharpen the model parameters to the best data fitting as iterations goes on.  
  
while iterations < k do  
    maybeInliers := n randomly selected values from data  
    maybeModel := model parameters fitted to maybeInliers  
    confirmedInliers := empty set  
    for every point in data do  
        if point fits maybeModel with an error smaller than t then  
            add point to confirmedInliers  
        end if  
    end for  
    if the number of elements in confirmedInliers is > d then  
        // This implies that we may have found a good model.  
        // Now test how good it is.  
        betterModel := model parameters fitted to all the points in confirmedInliers  
        thisErr := a measure of how well betterModel fits these points  
        if thisErr < bestErr then  
            bestFit := betterModel  
            bestErr := thisErr  
        end if  
    end if  
    increment iterations  
end while  
return bestFit
```



# RANSAC Properties

---

Good

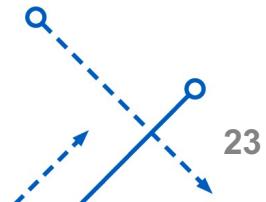
- Robust to outliers
- Applicable for larger number of model parameters than Hough transform.
- Optimization parameters are easier to choose than Hough transform.
  - Lines with normal form works for Hough, but slope-intercept form not.

Bad

- Computational time grows quickly with outliers and parameters
  - While Hough transform grows quickly with number of parameters.
- Not good for getting multiple fits
  - Hough transform can fit multiple lines simultaneously.

Common applications

- Computing a homography (e.g., image stitching)
- Estimating fundamental matrix (relating two views)





# IMAGE PROCESSING

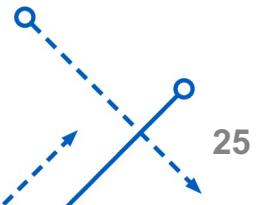
## Image Stitching



# Image Stitching

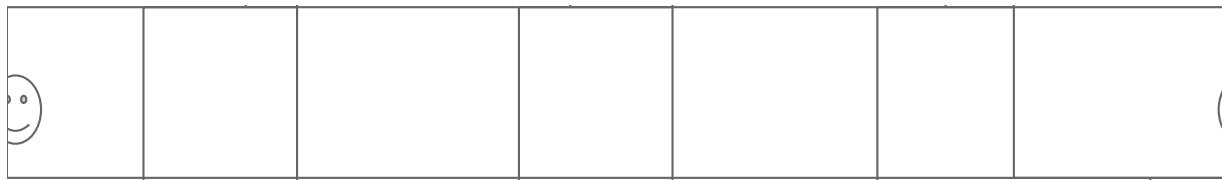
---

1. Align the images over each other
  - camera pan  $\leftrightarrow$  translation on cylinder
2. Blend the images together

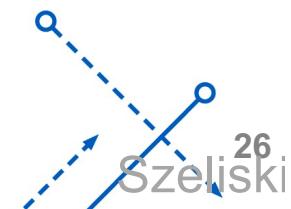


# Assembling the panorama

---

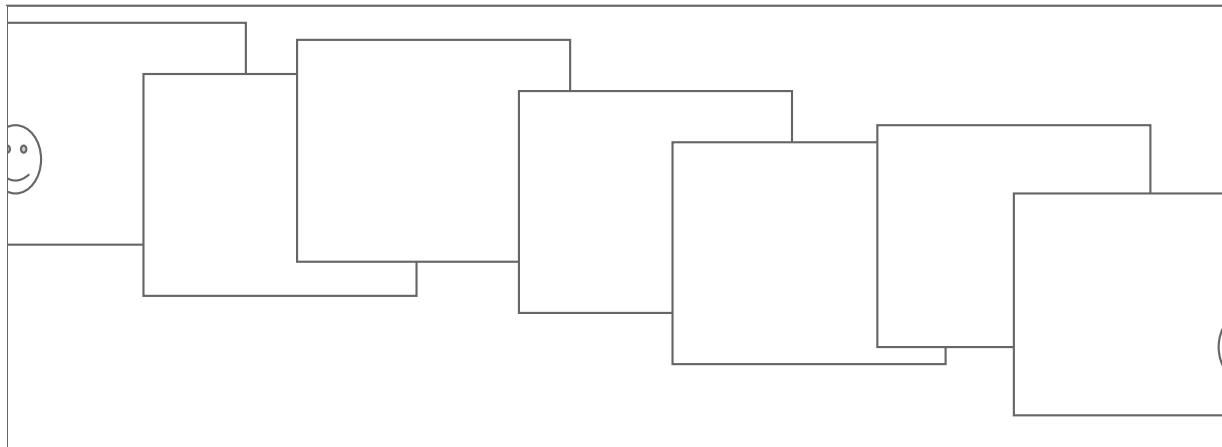


- Stitch pairs together, blend, then crop



# Problem: Drift

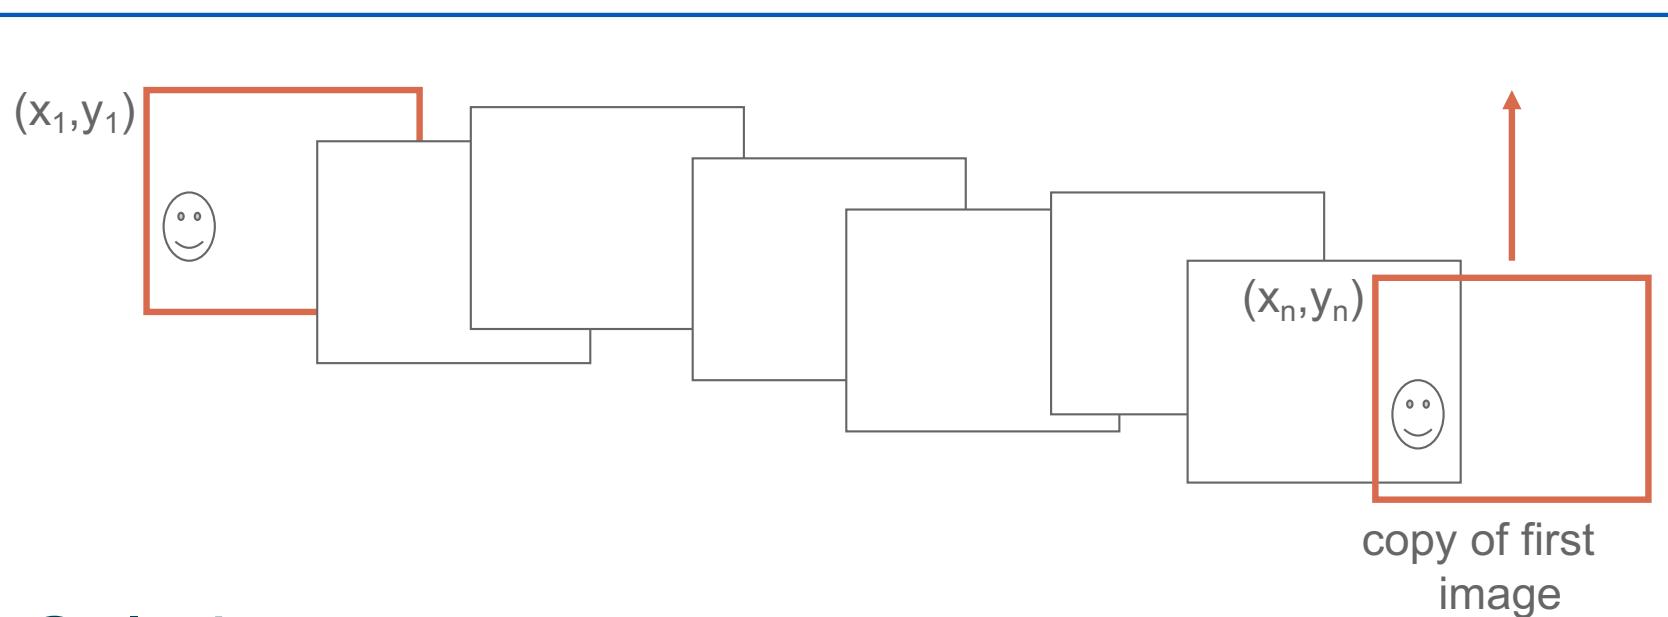
---



- Error accumulation
  - small (vertical) errors accumulate over time
  - apply correction so that sum = 0 (for 360° pan.)



# Problem: Drift

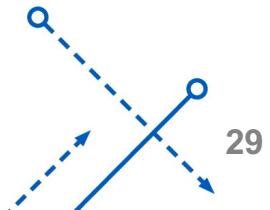


- **Solution**

- add another copy of first image at the end
- this gives a constraint:  $y_n = y_1$
- there are a bunch of ways to solve this problem
  - add displacement of  $(y_1 - y_n)/(n - 1)$  to each image after the first
  - compute a global warp:  $y' = y + ax$
  - run a big optimization problem, incorporating this constraint
    - best solution, but more complicated
    - known as “bundle adjustment”

# Full-view (360° spherical) panoramas

---

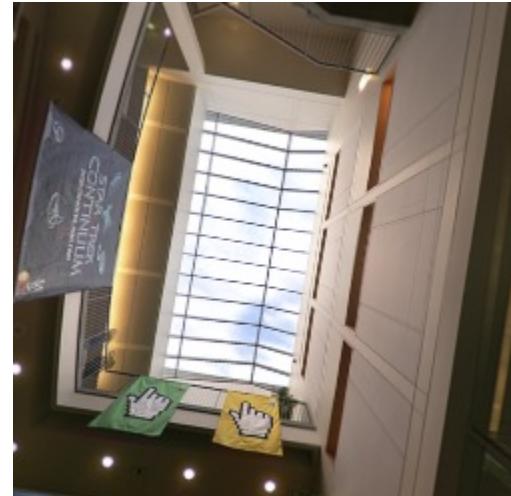


# Full-view Panorama



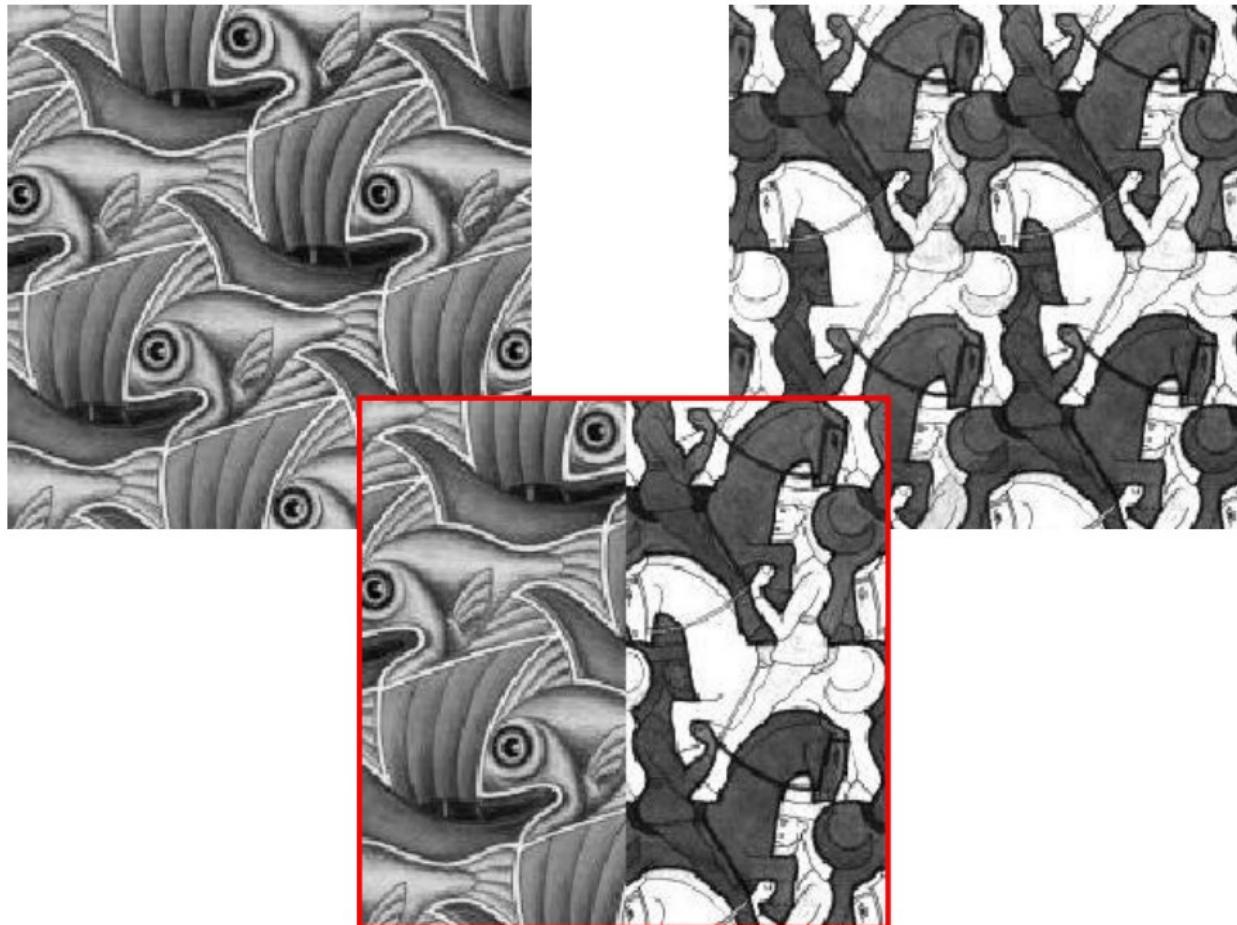
# Texture Mapped Model

---



# Image Blending

- How do we put images together so there is no obvious divisions?



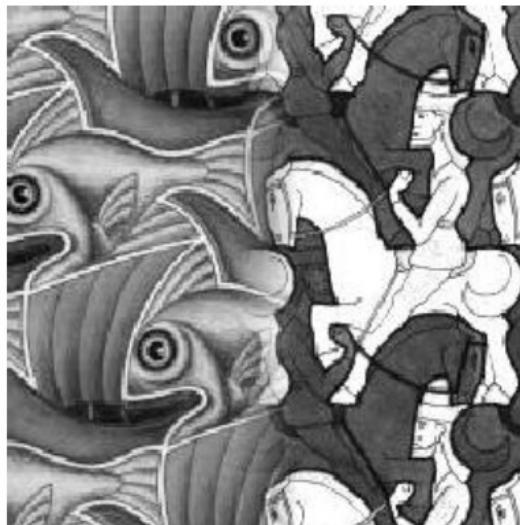
# Alpha Blending / Feathering



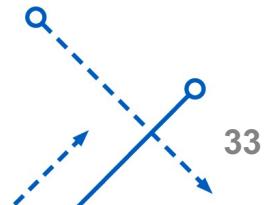
+



=

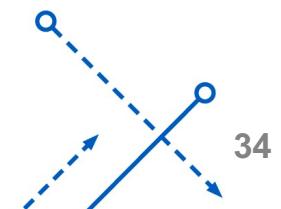
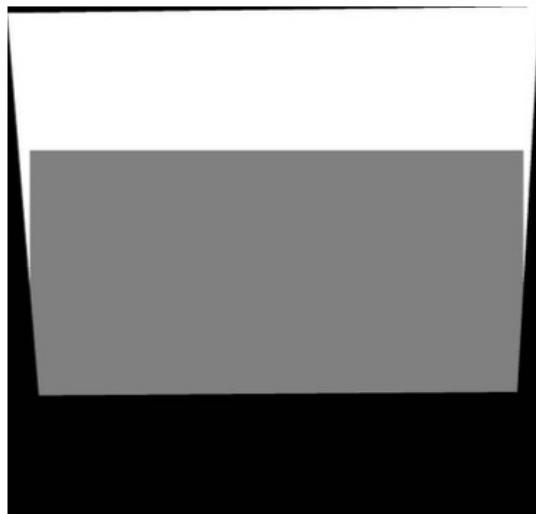


$$I_{\text{blend}} = \alpha I_{\text{left}} + (1-\alpha) I_{\text{right}}$$



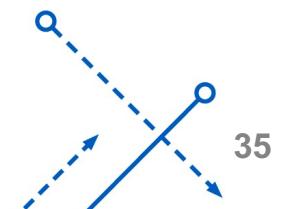
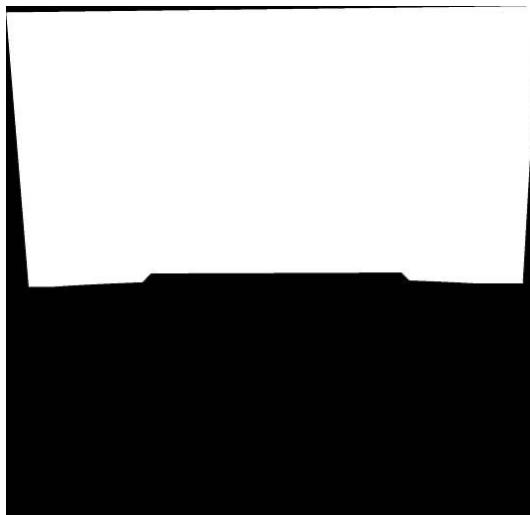
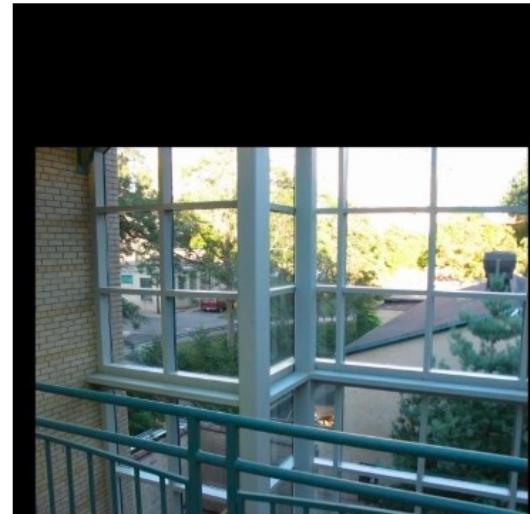
# Setting alpha: simple averaging

- Alpha = 0.5 in the overlap region



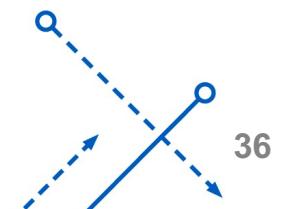
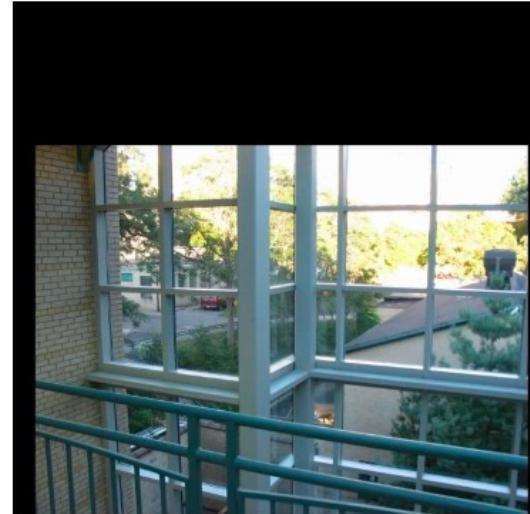
# Setting alpha: center seam

- Alpha = logical(dtrans1>dtrans2)



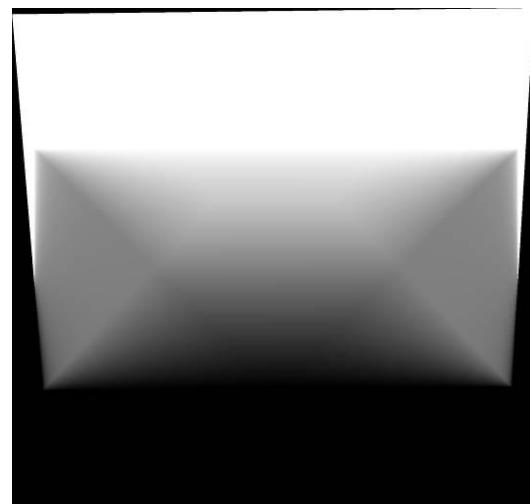
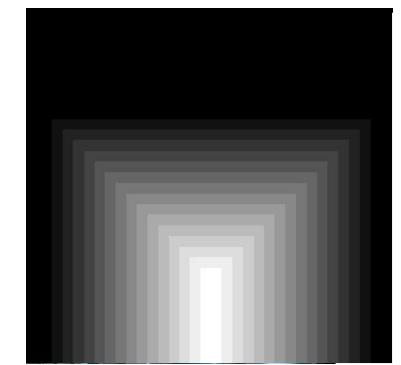
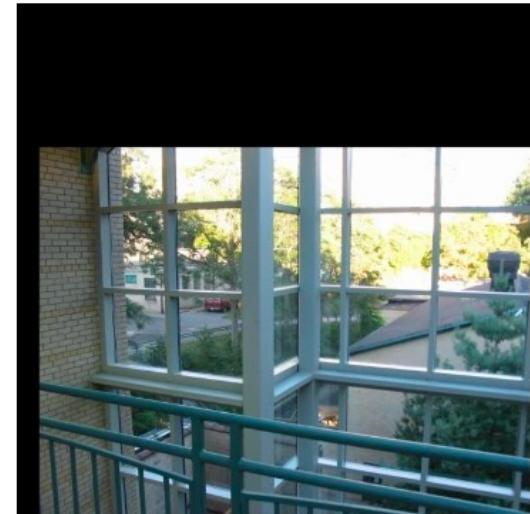
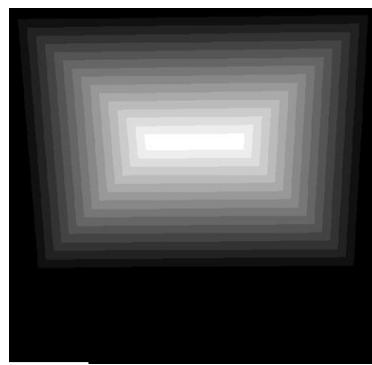
# Setting alpha: blurred seam

- Alpha = blurred



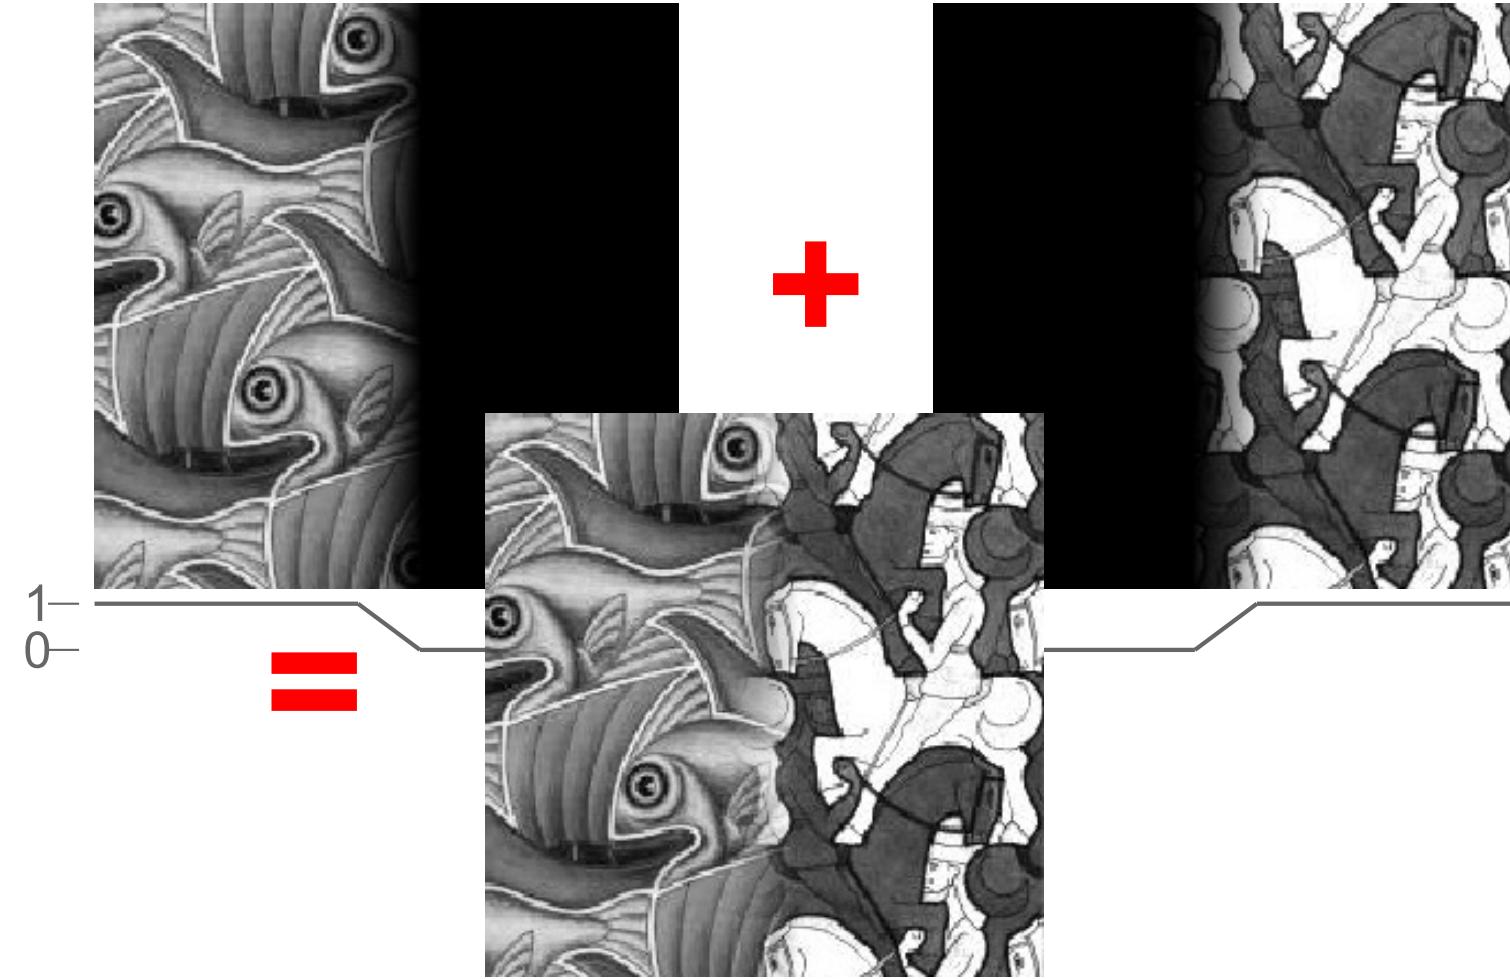
# Setting alpha: center weighting

- $\text{Alpha} = \text{dtrans1} / (\text{dtrans1} + \text{dtrans2})$

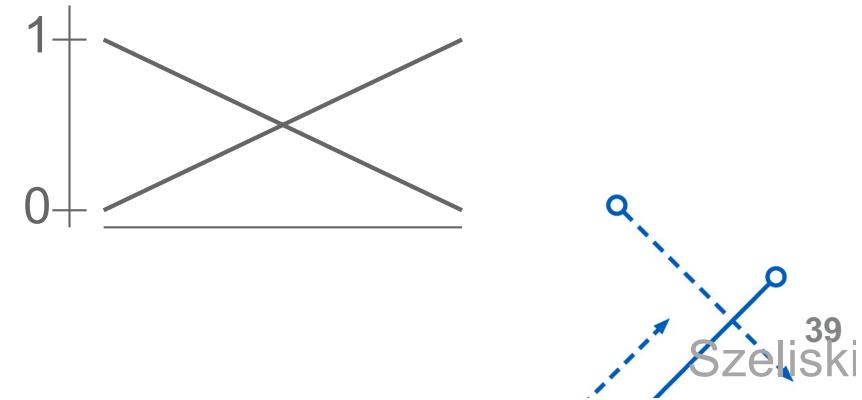
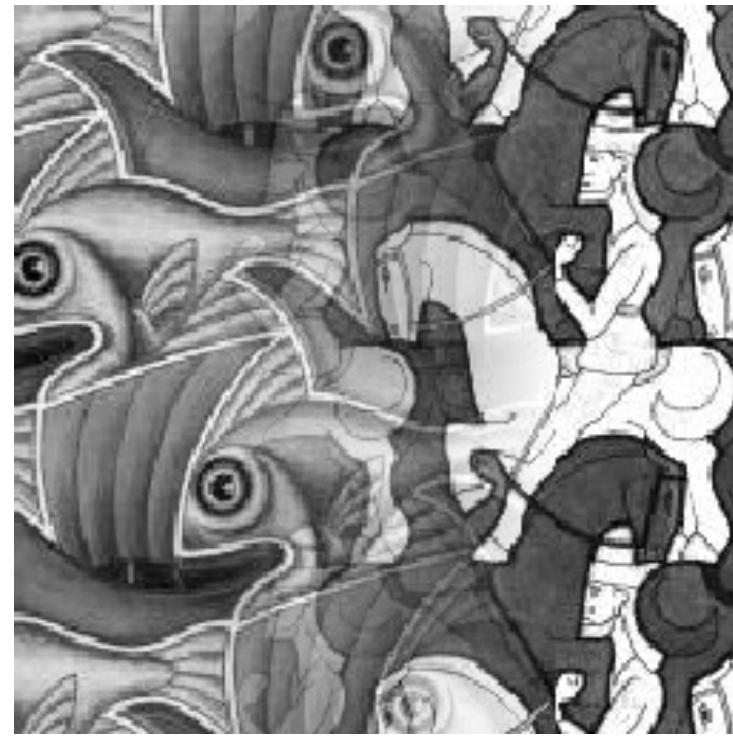
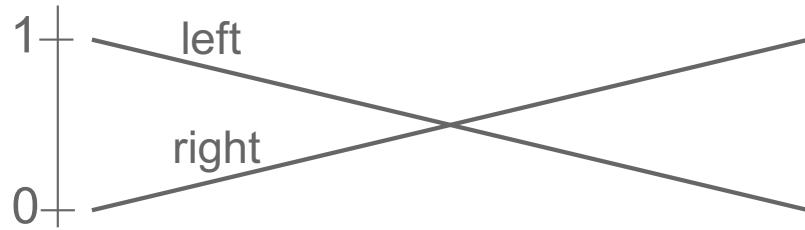


# Feathering

---

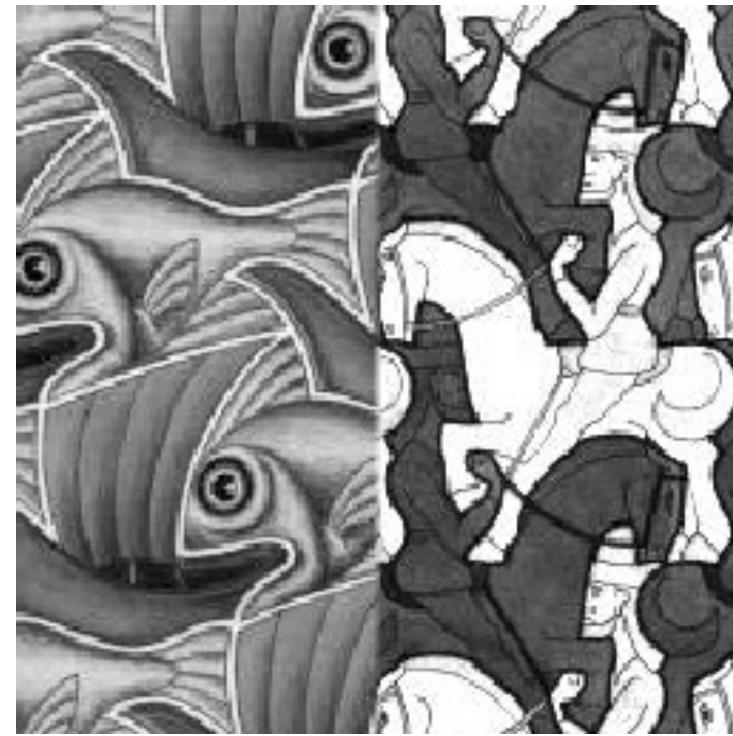


# Effect of window size



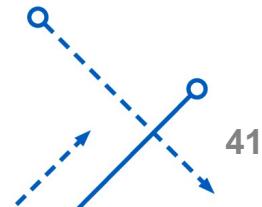
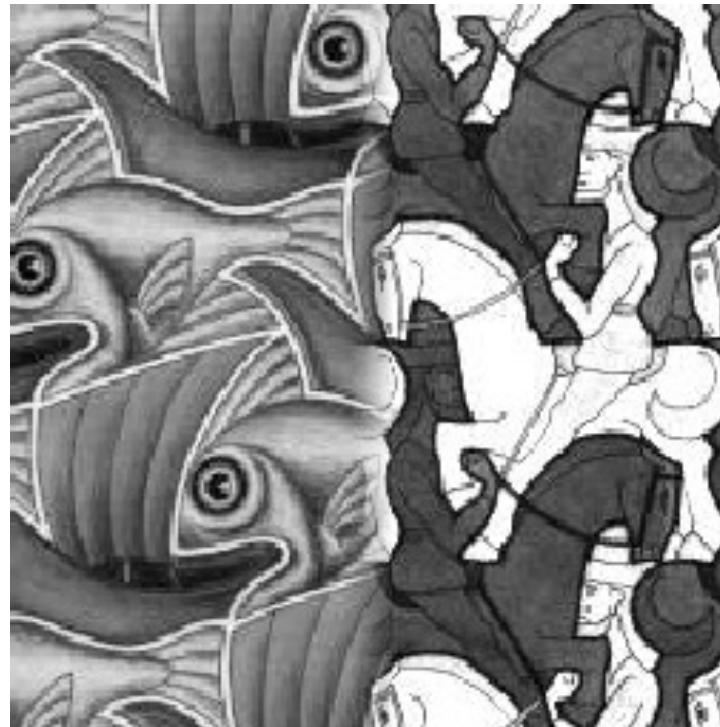
# Effect of window size

---



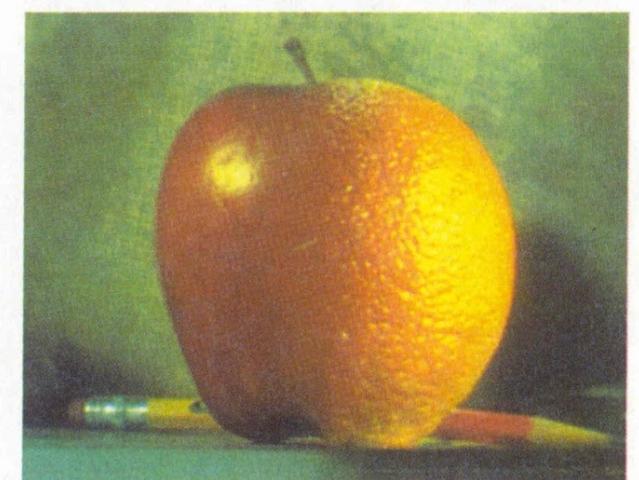
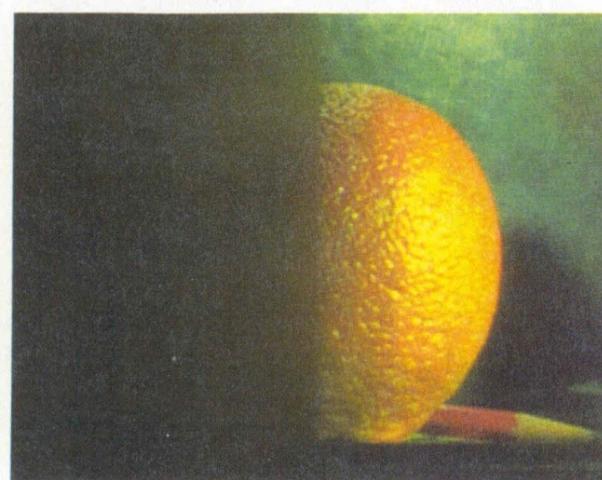
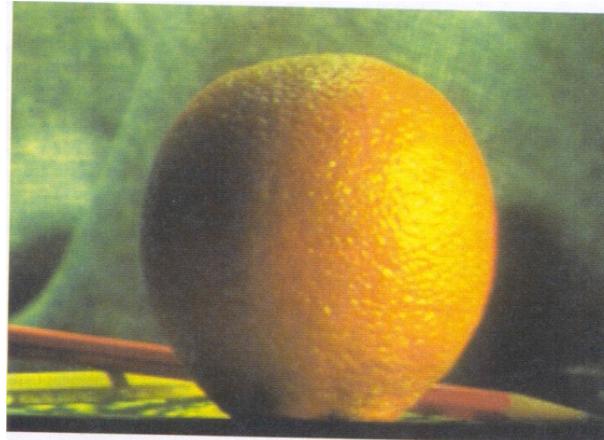
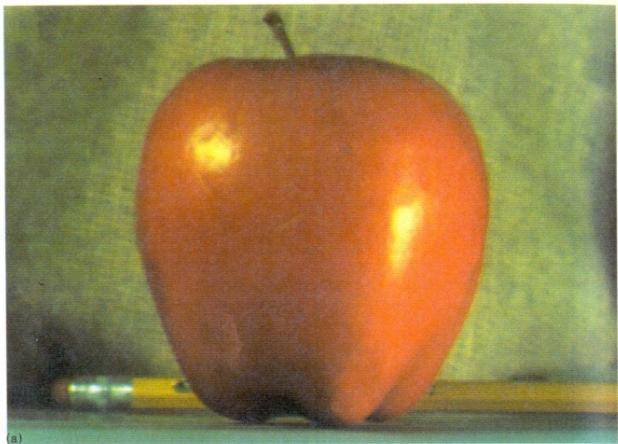
# Good Window Size

- “Optimal” Window: smooth but not ghosted.
- To avoid seams
  - window = size of largest prominent feature
- To avoid ghosting
  - window  $\leq 2 \times$  size of smallest prominent feature



# Pyramid Blending

---



# Laplacian Pyramid (Recap)

- Information captured at each level of a Gaussian (top) and Laplacian (bottom) pyramid
  - showing full resolution.

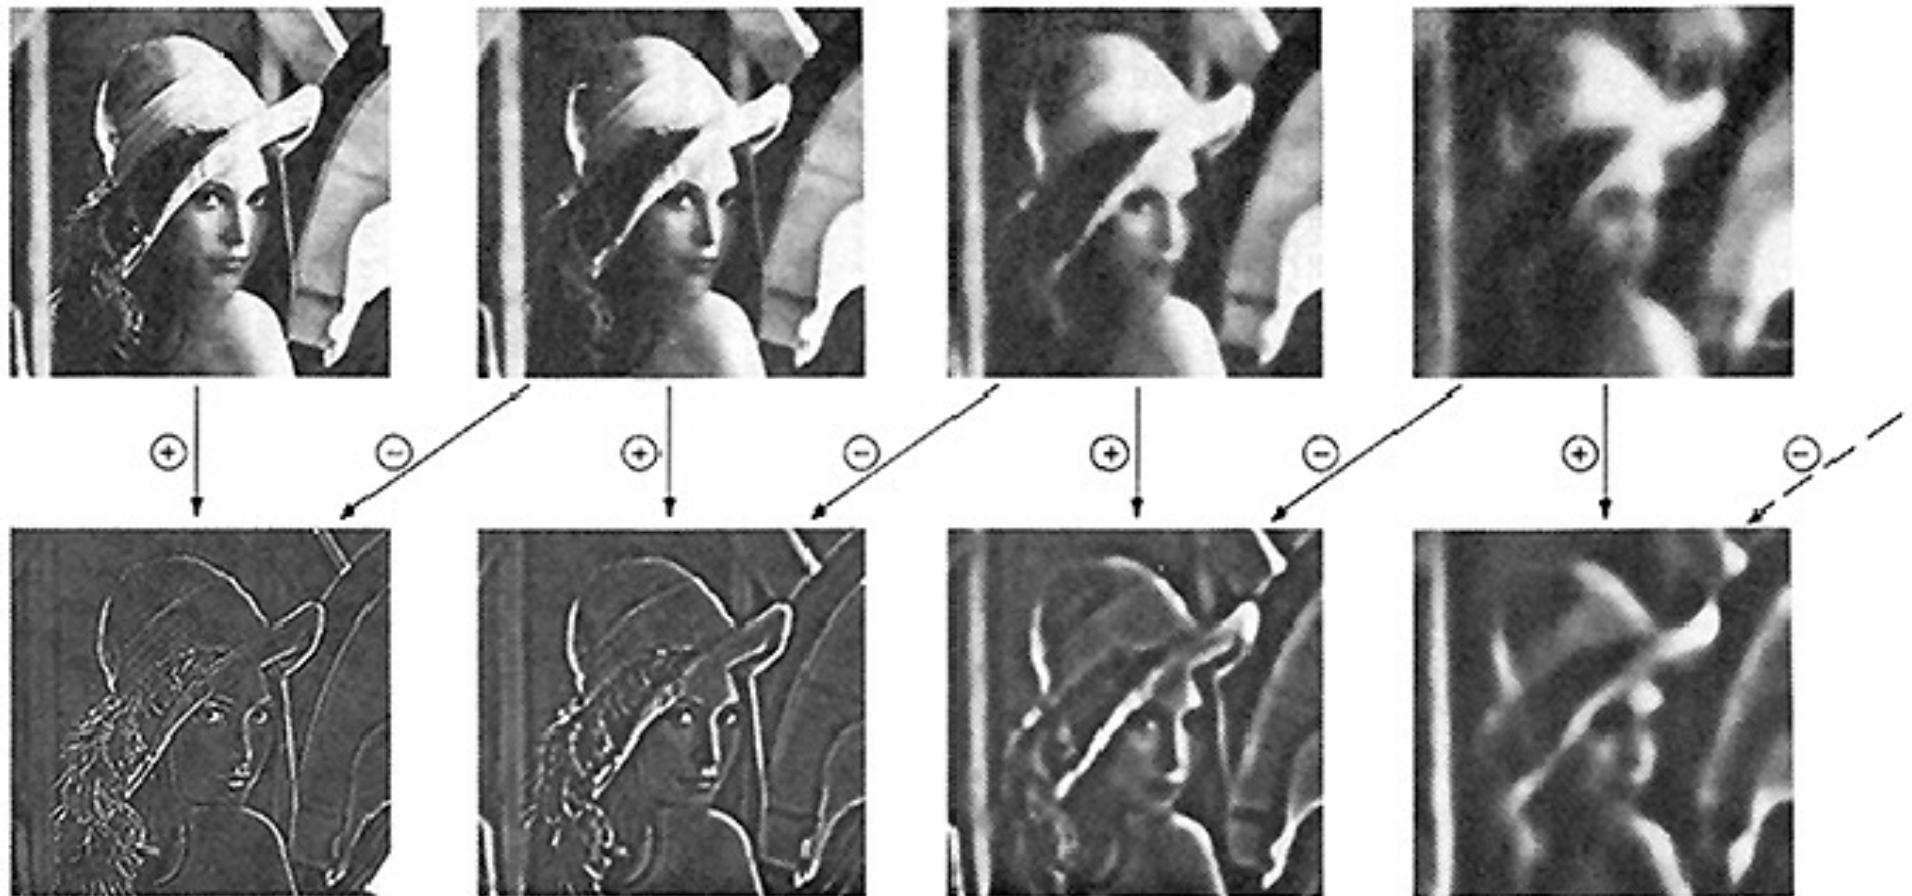
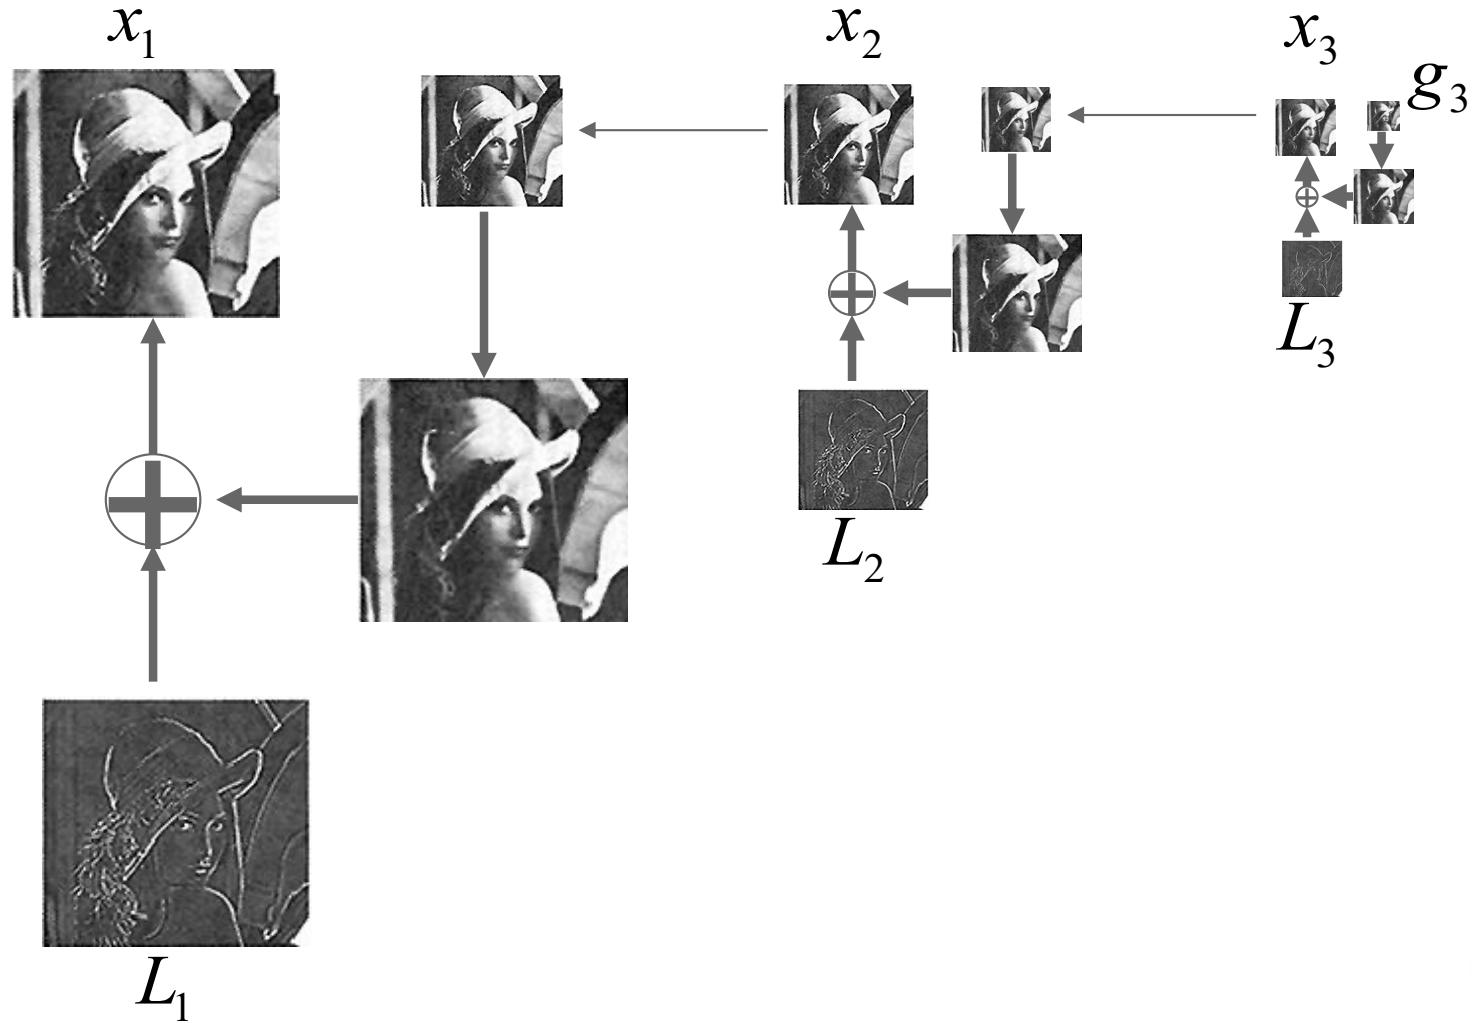


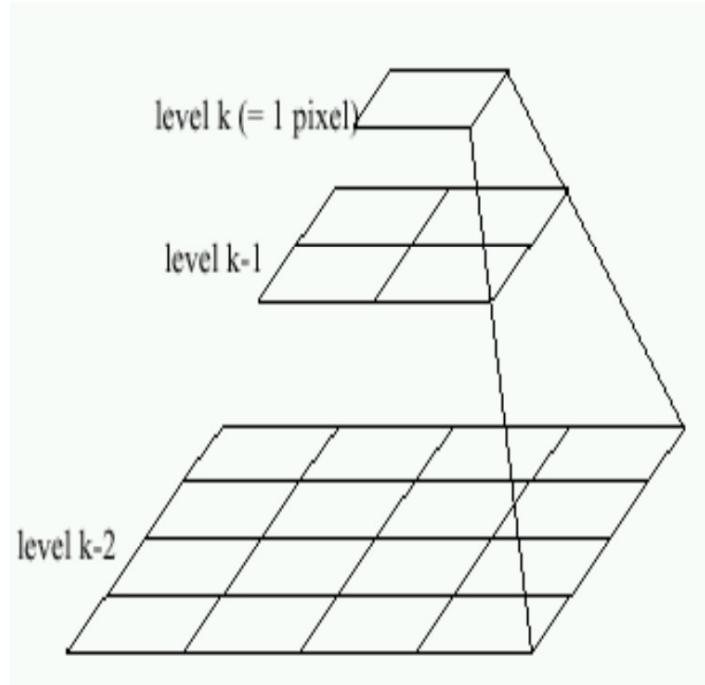
Fig. 5. First four levels of the Gaussian and Laplacian pyramid. Gaussian images, upper row, were obtained by expanding pyramid arrays (Fig. 4) through Gaussian interpolation. Each level of the Laplacian pyramid is the difference between the corresponding and next higher levels of the

# Laplacian Pyramid (Recap)

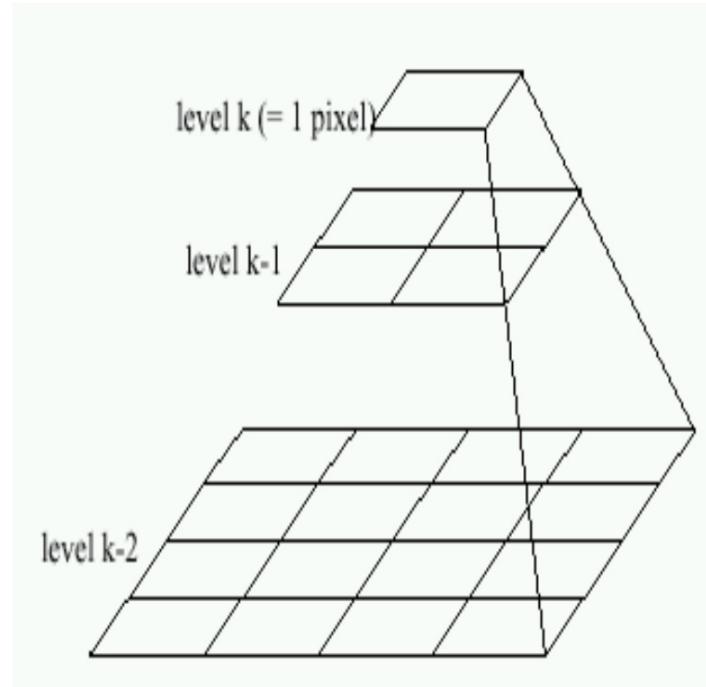
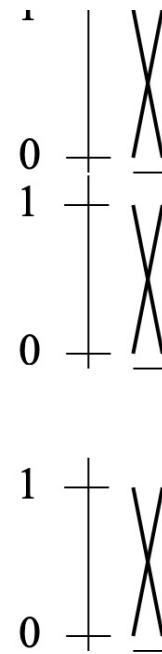
- Reconstruction: recover  $x_1$  from  $L_1$ ,  $L_2$ ,  $L_3$  and  $g_3$



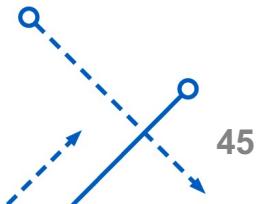
# Pyramid Blending



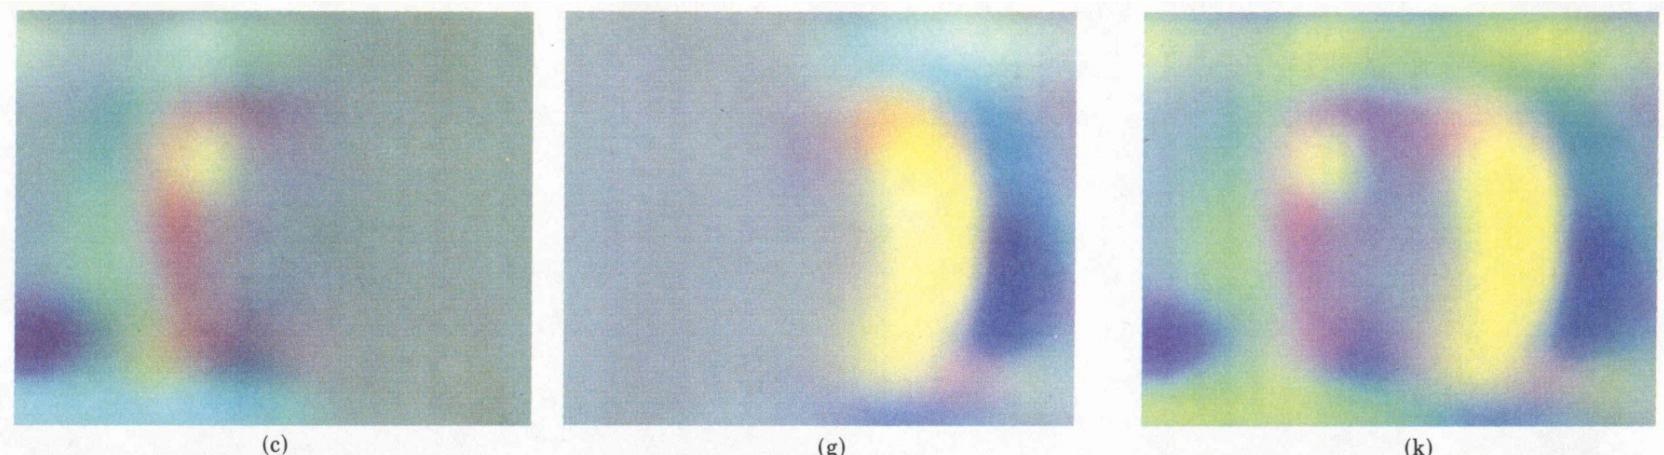
Left pyramid



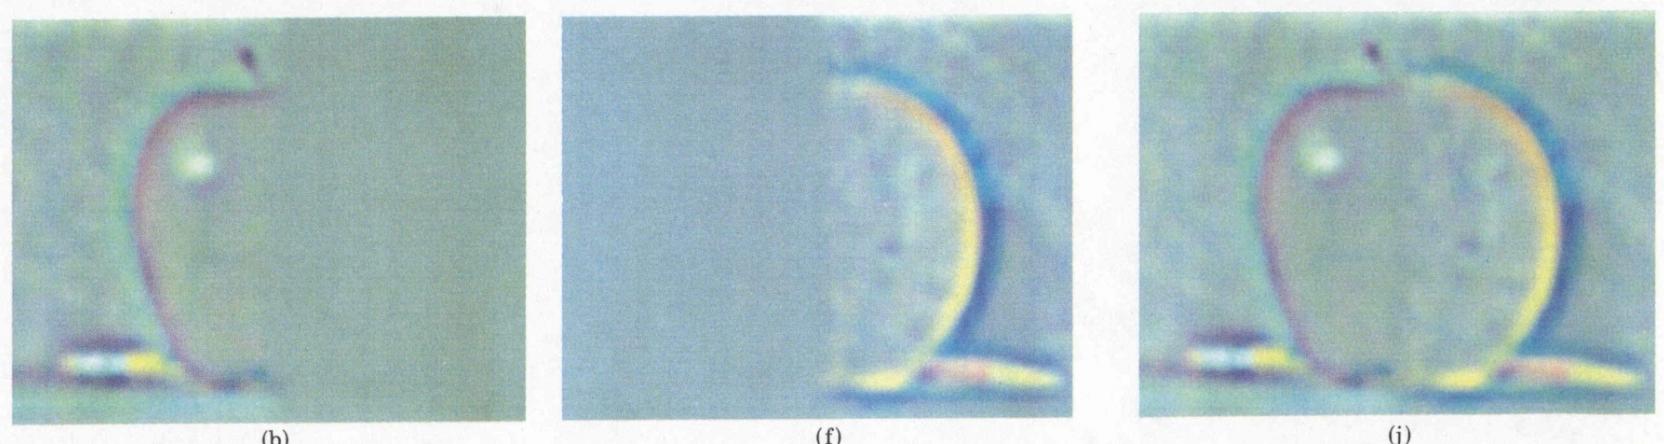
Right pyramid



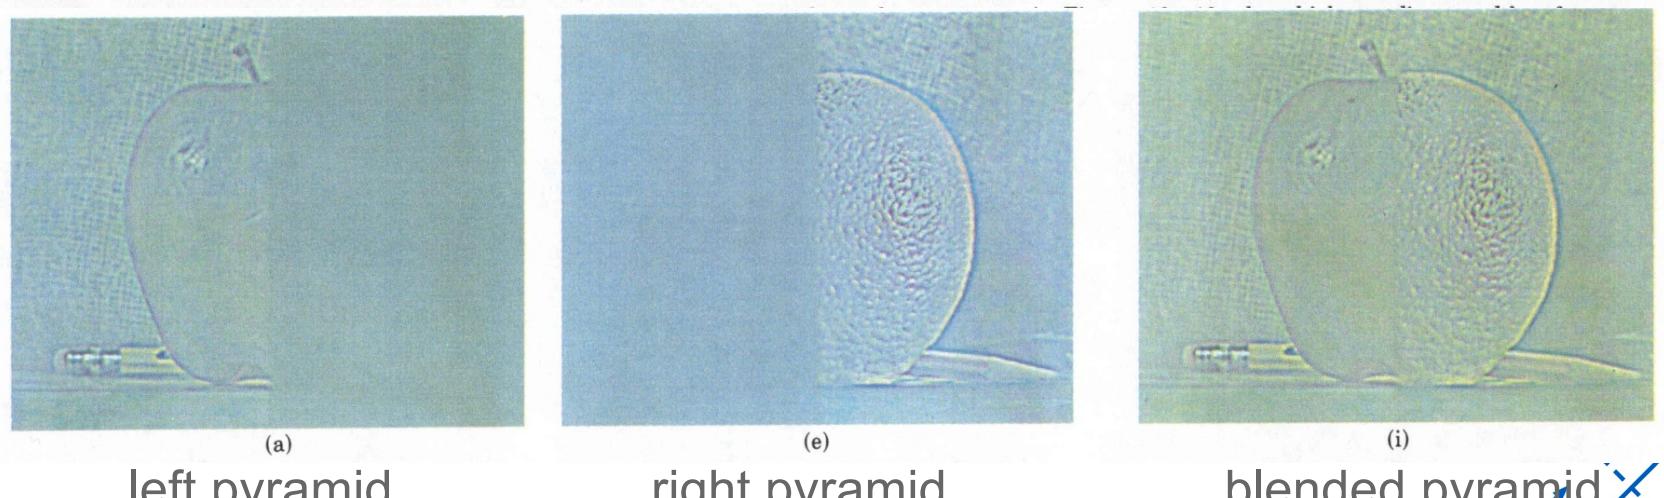
Laplacian  
level  
4



Laplacian  
level  
2



Laplacian  
level  
0



left pyramid

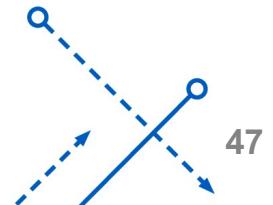
right pyramid

blended pyramid

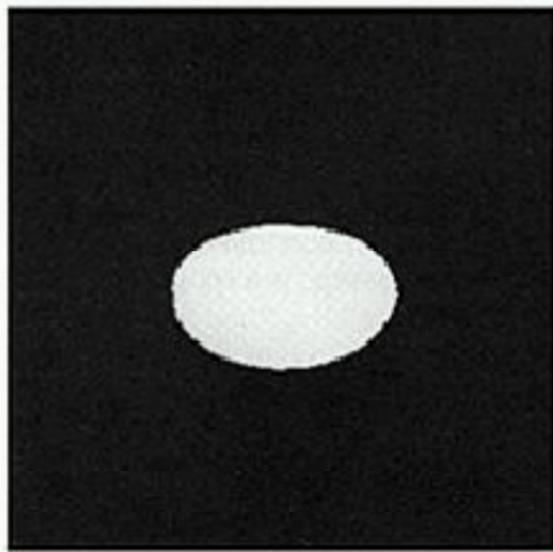
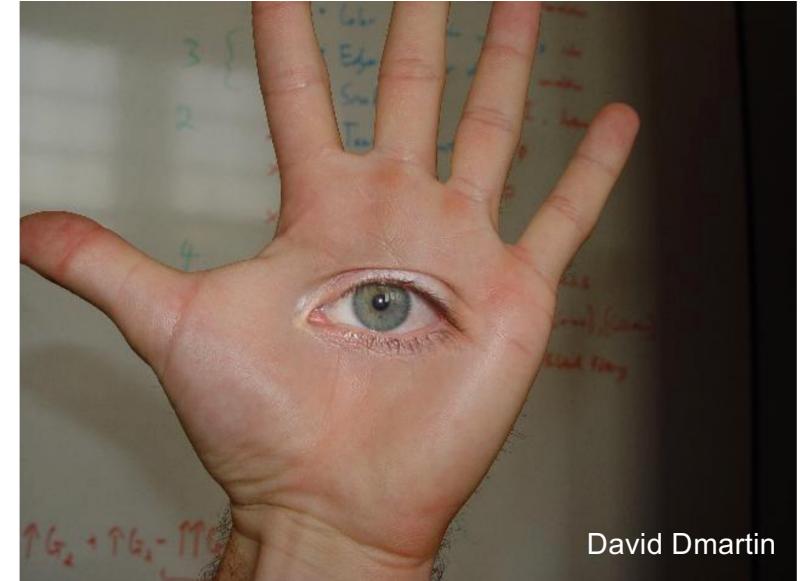
# Laplacian Pyramid: Blending

---

- 1. Build Laplacian pyramids LA and LB from images A and B
- 2. Build a Gaussian pyramid GR from selected region R
- 3. Combined pyramid LS from LA / LB using GR as weights:
  - $LS(i, j) = GR(i, j) * LA(i, j) + (1 - GR(i, j)) * LB(i, j)$
- 4. Collapse the LS pyramid to get the final blended image



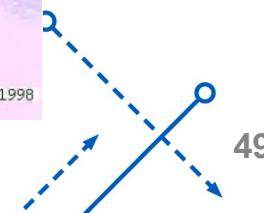
# Blending Regions



# Season Blending



10:24:59 15-NOV-1998



# Don't blend, CUT!

- So far we only tried to blend between two images.
- What about finding an optimal seam?



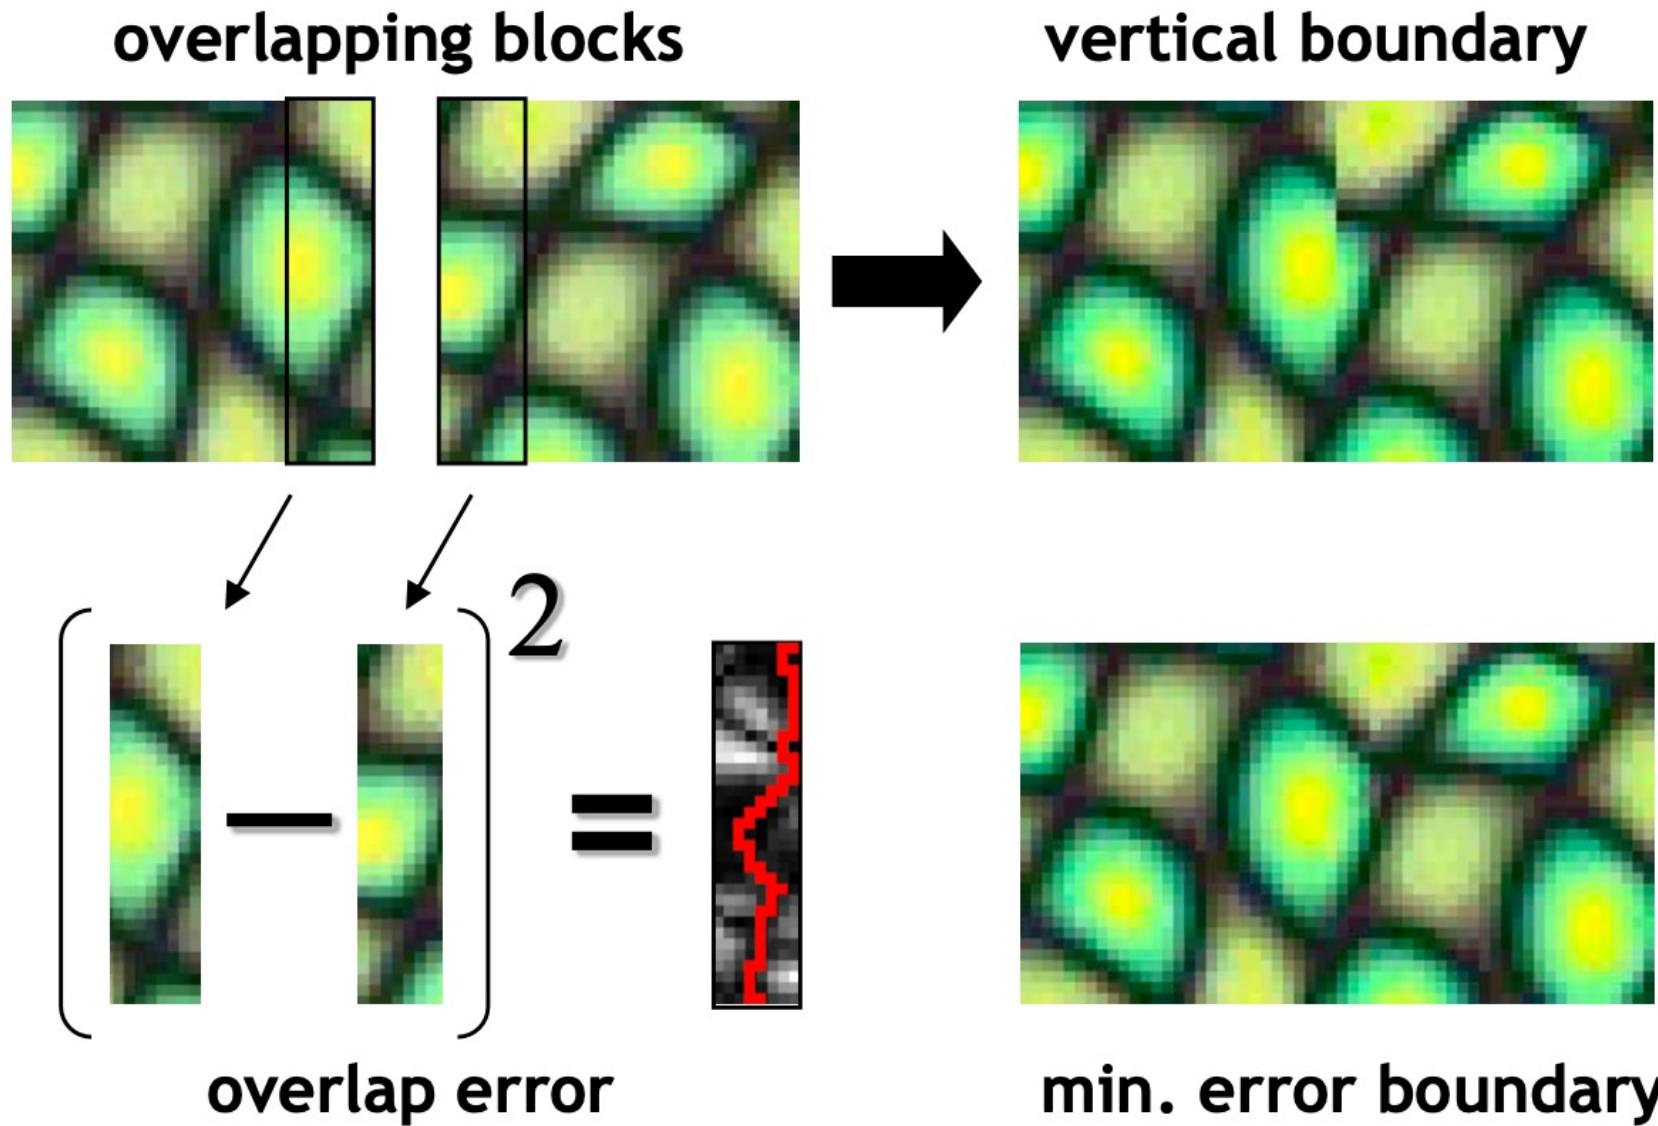
# Don't blend, CUT!

---

- Segment the mosaic
  - Single source image per segment
- Avoid artifacts along boundaries
  - Dijkstra's algorithm



# Minimal error boundary

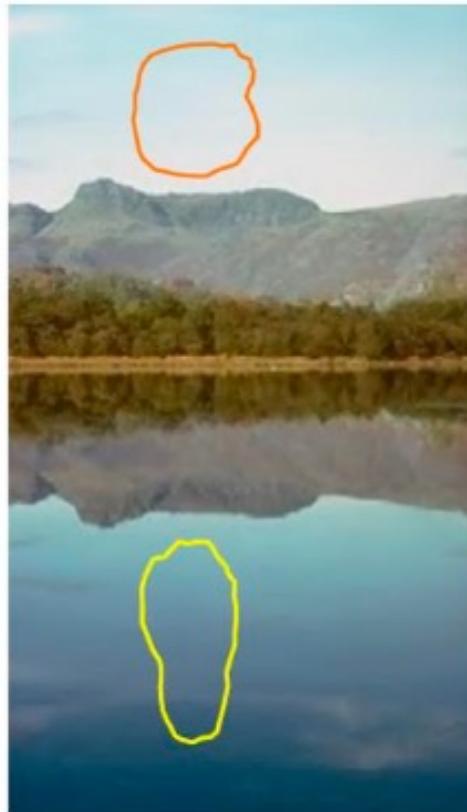


# Gradient Domain Blending

- Blend the gradients of the two images, then integrate.



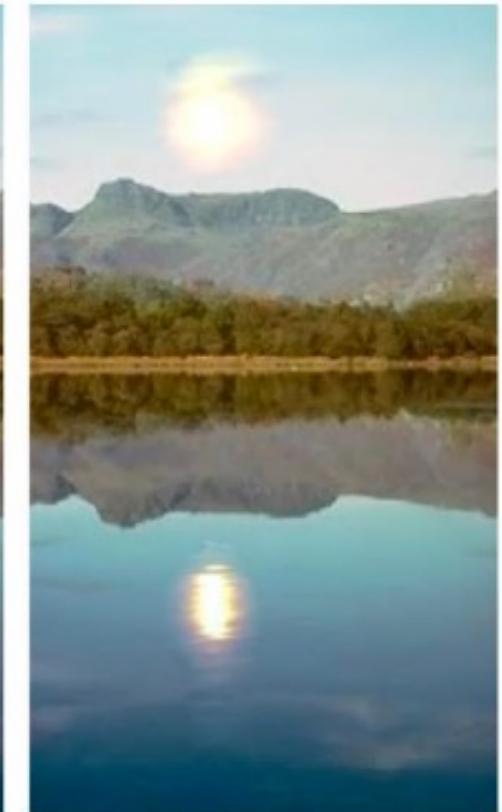
sources



destinations



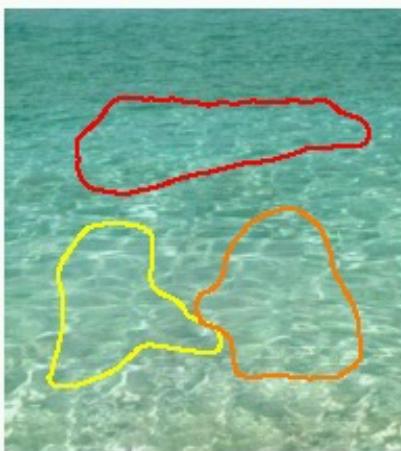
cloning



seamless cloning

# Gradient Domain Blending

- Blend the gradients of the two images, then integrate.



sources/destinations



cloning



seamless cloning

# Watch a video

## Interactive Digital Photomontage

Aseem Agarwala, Mira Dontcheva  
Maneesh Agrawala, Steven Drucker, Alex Colburn  
Brian Curless, David Salesin, Michael Cohen



# Final thought: What is a “panorama”?

- Tracking a subject
- Repeated (best) shots
- Multiple exposures
- “Infer” what photographer wants?
- Also referred to as  
“Computational Photography”

